

CS2200

Systems and Networks

Spring 2024

Lecture 9: Performance

Alexandros (Alex) Daglis
School of Computer Science
Georgia Institute of Technology
adaglis@gatech.edu

Announcements

- Homework 1 grades posted
 - Homework 3 due Wednesday
 - Lab 1 due Friday
 - Starting Chapter 5
-
- Will be using this logo on the slide *preceding* the first PS question
 - So you make sure you are already logged in
 - There will be a LOT of PS questions today!



Metrics

If want to try to make processors better, we have to take measurements

Common Metrics:

- Space → memory footprint
- Time → execution time
- Instruction frequency (or count)
 - Static
 - Dynamic
- Benchmarks

What determines execution time?

- CPI = “Cycles Per Instruction” → number of clock cycles each instruction takes
- $Execution\ time = (\sum CPI_j) * clock\ cycle\ time,$
where $1 \leq j \leq n$
- That’s a pretty tough sum to compute because modern computers can execute billions of instructions per second
- So, we approximate as
 $Execution\ time = n * CPI_{Avg} * clock\ cycle\ time,$
where n is the number of instructions
(executed—i.e., dynamic, not static—instruction count)
- This is known as the **Iron Law** of processor performance

The “Iron Law” of Processor Performance

$$\text{Processor Performance} = \frac{\text{Time}}{\text{Program}}$$

$$= \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Cycles}}{\text{Instruction}} \times \frac{\text{Seconds}}{\text{Cycle}}$$

ISA & Compiler

Microarchitecture

Circuit





What's the execution time?

1 GHz processor, 10K instructions, $CPI_{Avg} = 3$

- 0% A. 3 msec
- 0% B. 1 GHz
- 0% C. 0.003 msec
- 0% D. 30 μ sec
- 0% E. 30K
- 0% F. About 11, sir

$$10^{-9} \text{ sec/cy} \times 10^4 \text{ inst} \times 3 \text{ cy/inst}$$

$$= 3 \times 10^{-5} \text{ sec}$$

$$= .00003 \text{ sec}$$

$$= .03 \text{ msec}$$

$$= 30 \mu\text{sec}$$






Instruction Frequency

- *Static* instruction frequency refers to number of times a particular instruction occurs in compiled code.
 - Impacts memory footprint
 - If a particular instruction appears a lot in a program, can try to optimize amount of space it occupies by clever instruction encoding techniques in the instruction format.
- *Dynamic* instruction frequency refers to number of times a particular instruction is executed when program is run.
 - Impacts execution time of program
 - If dynamic frequency of an instruction is high then can try to make enhancements to datapath and control to ensure that CPI taken for its execution is minimized.



Static instruction frequency...

- A. Refers to the type of instructions in the instruction-set
- B. Refers to the frequency of occurrence of instructions in compiled code
- C. Refers to the frequency of occurrence of instructions that actually get executed ← **Dynamic instruction frequency**
- D. Refers to the clock frequency of the processor
- E. Is the basis for datapath design

-  Refers to the type of instructions in the instruction-set
-  Refers to the frequency of occurrence of instructions in compil...
-  Refers to the frequency of occurrence of instructions that actu...
-  Refers to the clock frequency of the processor
-  Is the basis for datapath design

What is the static frequency of ADD?

The ADD instruction occurs twice in a program that contains a total of 1000 instructions in the compiled code. All 1000 instructions get executed during a program run. One of the ADD instructions is in a 5-instruction loop that gets executed 1000 times.

$$2/1000 = 0.2\%$$



What about the **dynamic** frequency of ADD?

The ADD instruction occurs twice in a program that contains a total of 1000 instructions in the compiled code. All 1000 instructions get executed during a program run. One of the ADD instructions is in a 5-instruction loop that gets executed 1000 times.

- A. Two
- B. 0.2%
- C. $(1000/5995) * 100\%$
- D. $(1001/5995) * 100\%$
- E. $(1/5000) * 100\%$
- F. $(1001/5000) * 100\%$







ADDs executed:

$$1 + 1000 = 1001$$

Total executed:

$$(1000-5) \text{ inst} + 1000 * 5 = 5995$$

Dynamic add frequency: $1001/5995$

 Two	 0.2%
 $(1000/5995) * 100\%$	 $(1001/5995) * 100\%$
 $(1/5000) * 100\%$	 $(1001/5000) * 100\%$

I feel the need, the need for speed!

- How do we improve performance?
- Reduce execution time (of course)
- How?
- The Iron Law tells us:

$$\text{Execution time} = N \downarrow * \text{CPI} \downarrow * \text{Cycle Time} \downarrow$$

- How can we reduce the left-hand side?
- Reduce one or more of the right-hand factors

How can we measure improvement?

$$\text{speedup}_{\text{A over B}} = \frac{\text{Execution Time On Processor B}}{\text{Execution Time On Processor A}}$$

$$\text{speedup}_{\text{improved}} = \frac{\text{Execution Time Before Improvement}}{\text{Execution Time After Improvement}}$$

$$\text{improvement in execution time} = \frac{\text{old execution time} - \text{new execution time}}{\text{old execution time}}$$

Example

You improve your application's algorithm so that it runs in 29 seconds instead of 38 seconds.

What is the speedup?

$$E_{\text{Before}}/E_{\text{After}}$$

$$38 / 29 = 1.3103 \\ = 31\% \text{ speedup}$$

$$29 * 1.31 \sim 38$$

What is the improvement in execution time?

$$(E_{\text{old}} - E_{\text{new}}) / E_{\text{old}}$$

$$(38 - 29) / 38 = .2368... \\ = 24\%$$

$$38 - (24\% \text{ of } 38) \sim 29$$



Improvement in execution speed

It takes me 8 minutes to walk to class from my office. I can run twice as fast as I walk. If I walk half the distance and run the remaining half, how much time will I take to reach class from my office?

Rank

Responses

$$\begin{aligned}
 8 &= d/v \\
 n &= \frac{1}{2} d/v + \frac{1}{2} d/2v \\
 8v &= d \\
 n &= 4v/v + 4v/2v \\
 n &= 4 + 2 \\
 n &= 6
 \end{aligned}$$

Amdahl's Law

Amdahl's Law

$$\text{Time}_{\text{after}} = \frac{\text{Time}_{\text{affected}}}{\text{Amount of Improvement}} + \text{Time}_{\text{unaffected}}$$

My office walk: $6 = 4 / 2 + 4$

Improving an instruction

A processor spends 20% of its time on ADD instructions.
An engineer proposes to improve the speed of the ADD instruction by 4 times.
What is the speedup achieved by this modification?

The improvement only applies for the ADD instruction, so 80% of the execution time is unaffected by the improvement.

Original normalized execution time = 1

$$\begin{aligned}\text{New execution time} &= (\text{time spent in ADD}/4) + \text{remaining time} \\ &= 0.2 / 4 + .8 \\ &= 0.85\end{aligned}$$

$$\begin{aligned}\text{Speedup} &= \text{execution time before} / \text{execution time after} \\ &= 1 / 0.85 = 1.18 = 18\%\end{aligned}$$



Improving an instruction

An engineer is asked to improve the processor's overall performance by 2 times by optimizing the ADD instruction.

The processor spends 20% of its time on ADD instructions.

How much faster must the ADD instruction become?

- A. 2x
- B. 10x
- C. 100x
- D. That's impossible!!

Execution time?

CPI of Instruction Classes	Code 1	Code 2
R-type = 2	3	10
I-type = 10	3	1
J-type = 3	5	2
S-type = 4	2	3

13 inst
6+30
+15+8 =
59 cycles

16 inst
20+10+6
+12 =
48 cycles

- A. Code 1 is better than Code 2 since it has fewer instructions
- B. Code 2 is better than Code 1 since it has fewer instructions
- C. Code 1 is better than Code 2 since it takes fewer total clock cycles to execute
- D. Code 2 is better than Code 1 since it takes fewer total clock cycles to execute

Microarchitecture change?

- We have a computer with three types of instructions that have the following CPI:

Type	CPI
A	2
B	5
C	1

- An architect determines that she can reduce the CPI for B to 3 but will need to slow the clock speed of the processor. What is the maximum permissible slowing of the clock that will make this change worthwhile?
- Assume that all the workloads for this processor use 30% of A, 10% of B, and 60% of C types of instructions

How do we answer that?

Execution time of the old machine:

$$ET_o = N * (F_{Ao} * CPI_{Ao} + F_{Bo} * CPI_{Bo} + F_{Co} * CPI_{Co}) * C_o$$

(where F_x and CPI_x are the dynamic frequencies and CPIs of each type of instruction, respectively)

$$ET_o = N * (0.3 * 2 + 0.1 * 5 + 0.6 * 1) * C_o = N * 1.7 * C_o$$

Execution time for the new machine:

$$ET_n = N * (0.3 * 2 + 0.1 * 3 + 0.6 * 1) * C_n = N * 1.5 * C_n$$

For the design to be viable,

$$ET_n < ET_o$$

$$N * 1.5 * C_n < N * 1.7 * C_o$$

$$C_n < 1.7/1.5 * C_o$$

$$C_n < 1.13 * C_o$$

Maximum permissible increase in clock cycle time = 13%

Type	CPI
A	2
B	5
C	1

30% of A,
10% of B,
60% of C

Combining two instructions?

Instruction	CPI
Add	2
Shift	3
Others	2
Add/Shift	4

Assume a SHIFT instruction is always preceded by an ADD.

If SHIFT instructions represent 10% of the dynamic instruction frequency of a program, what is the speedup of the program with all {ADD, SHIFT} replaced by the combined instruction?

[HINT: For every 10 instructions in the original program, 2 instructions are the ADD/SHIFT combo. Thus the number of instructions in the new program shrinks to 90% of the original program.]

A solution

$$\begin{aligned} ET_o &= N * (F_{ADD} * 2 + F_{SHIFT} * 3 + F_{others} * 2) * C \\ &= N * (0.1 * 2 + 0.1 * 3 + 0.8 * 2) * C \\ &= 2.1 * N * C \end{aligned}$$

With the combo instruction replacing {ADD SHIFT}, the number of instructions in the new program shrinks to 0.9N in the new program. The frequency of the combo instruction is 1/9 and the other instructions are 8/9.

$$\begin{aligned} ET_n &= (0.9 * N) * (F_{COMBO} * 4 + F_{others} * 2) * C \\ &= (0.9 * N) * (1 / 9 * 4 + 8 / 9 * 2) * C \\ &= 2 * N * C \end{aligned}$$

$$\begin{aligned} \text{Speedup} &= \text{old execution time} / \text{new execution time} \\ &= 2.1NC / 2NC \\ &= 1.05 \end{aligned}$$

Instruction	CPI
Add	2
Shift	3
Others	2
Add/Shift	4

Benchmarks

- **Benchmarks** are a set of programs that are **representative** of the workload for a processor.
- The key difficulty is to be sure that the benchmark program selected **really** are representative of the prospective workload.
- Standard benchmark suites (SPEC, LINPACK, Whetstone, Dhrystone, and many more) are used to try to compare “apples” with “apples” by summarizing performance *across a set of programs*
 - E.g., SPEC uses perl, gcc, AI apps, compression, imaging apps, modeling apps, etc. to represent a common workload
- A radical new hardware design is hard to benchmark because there may not yet be a compiler or much code.

Using a benchmark

Some caveats:

- Testing a single system component (e.g., the processor) only gives a limited view:
e.g., memory organization and memory—processor—bus bandwidths are also key
- Some processors do well on certain benchmark programs and other do well on other programs
- A composite index can be useful when we want to compare two processors without knowing the exact kind of workload they are going to run, but we must be very cautious
 - More on this later

Reasons to be skeptical of a benchmark

- The vendor gave you the benchmark results (in polite company, we call this a conflict of interest)
- The vendor wrote the benchmark suite
- The benchmark suite doesn't seem to have any elements that represent your workload (e.g., you run web server farms and the benchmark represents only computationally intensive scientific calculations)
- The equipment being benchmarked is different from the equipment you want to evaluate (maybe a little different, maybe a lot different)
- The benchmark uses a different compiler suite than you plan to use

The cost of mistakenly choosing the wrong equipment is very high!



Comparing Multiple Programs

	Computer A	Computer B	Computer C
Program 1 (secs)	1	10	20
Program 2 (secs)	1000	100	20
Program 3 (secs)	1001	110	40

A is 10 times faster than B for program 1
B is 10 times faster than A for program 2
A is 20 times faster than C for program 1
C is 50 times faster than A for program 2
B is 2 times faster than C for program 1
C is 5 times faster than B for program 2

Each statement above is correct...

...but I just want to know which machine is the best?

Need a composite metric



Let's Try a Simpler Example

Two machines timed on two benchmarks

	<u>Machine A</u>	<u>Machine B</u>
Program 1	2 seconds	4 seconds
Program 2	12 seconds	8 seconds

How much faster is Machine A than Machine B?

Attempt 1: ratio of runtimes, normalized to Machine A runtimes

program1: $4/2$

program2 : $8/12$

- Machine A ran 2 times faster on program 1, $2/3$ times faster on program 2
- On *average*, Machine A is $(2 + 2/3) / 2 = 4/3$ times faster than Machine B

It turns this “averaging” stuff can fool us; watch...

Example (con't)

Two machines timed on two benchmarks

	Machine A	Machine B
Program 1	2 seconds	4 seconds
Program 2	12 seconds	8 seconds

How much faster is Machine A than B?

Attempt 2: ratio of runtimes, normalized to Machine B runtimes

program 1: $2/4$ program 2 : $12/8$

- Machine A ran program 1 in $1/2$ the time and program 2 in $3/2$ the time
- On average, $(1/2 + 3/2) / 2 = 1$
- Put another way, Machine A is 1.0 times faster than Machine B

Example (con't)

Two machines timed on two benchmarks

	Machine A	Machine B
Program 1	2 seconds	4 seconds
Program 2	12 seconds	8 seconds

How much faster is Machine A than B?

Attempt 3: ratio of aggregated runtimes, norm. to A

- Machine A took 14 seconds for both programs
- Machine B took 12 seconds for both programs
- Therefore, Machine A takes $14/12$ of the time of Machine B
- Put another way, Machine A is $6/7$ faster than Machine B

Which is Right?

Question:

- How can we get three different answers?

Answer:

- Because, while they are all reasonable calculations...

...each answers a different question

Need to be more precise in understanding and posing these performance & metric questions

Arithmetic and Harmonic Mean

Average of the execution time that tracks total execution time is the arithmetic mean

$$\frac{1}{n} \sum_{i=1}^n \textit{Time}_i$$

This is the definition for “average” you are most familiar with

If performance is expressed as a rate, then the average that tracks total execution time is the harmonic mean

$$\frac{n}{\sum_{i=1}^n \frac{1}{\textit{Rate}_i}}$$

This is a different definition for “average” you are probably less familiar with

Geometric Mean

- Used for relative rate (i.e., ratio) or normalized performance

$$\textit{Relative_Rate} = \frac{\textit{Rate}}{\textit{Rate}_{ref}} = \frac{\textit{Time}_{ref}}{\textit{Time}}$$

- Geometric mean

$$\sqrt[n]{\prod_{i=1}^n \textit{Relative_Rate}_i} = \frac{\sqrt[n]{\prod_{i=1}^n \textit{Rate}_i}}{\textit{Rate}_{ref}}$$

Why does the choice of the mean matter?

Benchmark	Ops (millions)	Computer 1	Computer 2	Speedup (C2 vs C1)
<i>Absolute performance (Time)</i>				
Program 1	100	1	20	
Program 2	100	1000	20	
Total time		1001	40	25
Avg (arith mean)		500	20	25

Why does the choice of the mean matter?

Benchmark	Ops (millions)	Computer 1	Computer 2	Speedup (C2 vs C1)
<i>Absolute performance (Time)</i>				
Program 1	100	1	20	
Program 2	100	1000	20	
Total time		1001	40	25
Avg (arith mean)		500	20	25
<i>Performance in MFLOPS (Rate)</i>				
Program 1		100	5	
Program 2		0.1	5	
Arith. mean		50.1	5	0.1
Geom. mean		3.2	5	1.6
Harm. mean		0.2	5	25

For rates use Harmonic Mean!

Problems with Arithmetic Mean

- Applications do not have the same probability of being run
- Longer programs weigh more heavily in the average

For example, two machines timed on two benchmarks

	<u>Machine A</u>	<u>Machine B</u>
Program 1	2 seconds (20%)	4 seconds (20%)
<u>Program 2</u>	<u>12 seconds (80%)</u>	<u>8 seconds (80%)</u>

- If we do arithmetic mean, Program 2 “counts more” than Program 1
An X% improvement in Program 2 changes the average more than an X% improvement in Program 1
- But perhaps Program 2 is 4 times more likely to run than Program 1

Weighted Execution Time

Often, one runs some programs more often than others. Therefore, we should *weight* the more frequently used programs' execution time

$$\sum_{i=1}^n \text{Weight}_i \times \text{Time}_i$$

Weighted Harmonic Mean

$$\frac{1}{\sum_{i=1}^n \frac{\text{Weight}_i}{\text{Rate}_i}}$$

Using a Weighted Sum (or weighted average)

	<u>Machine A</u>	<u>Machine B</u>
Program 1	2 seconds (20%)	4 seconds (20%)
Program 2	12 seconds (80%)	8 seconds (80%)
Total	10 seconds	7.2 seconds

Allows us to determine relative performance $10/7.2 = 1.39$

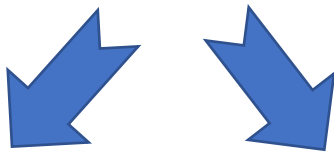
→ Machine B is 1.39 times faster than Machine A

(instead of $14/12 = 1.16x$ faster without weighting)

What if we only know normalized runtimes?

Normalize runtime of each program to a reference

	<u>Machine A (ref)</u>	<u>Machine B</u>
Program 1	2 seconds	4 seconds
Program 2	12 seconds	8 seconds



	<u>Machine A (norm to B)</u>	<u>Machine B (norm to A)</u>
Program 1	0.5	2.0
Program 2	1.5	0.666
Average?	1.0	1.333

- When we normalize A to B and average, it looks like A & B are the same.
- But when we normalize B to A and average, it looks like A is better!

Using Geometric Mean

	Machine A (norm to B)	Machine B (norm to A)
Program 1	0.5	2.0
Program 2	1.5	0.666
Geometric Mean	0.866	1.154

Note that $1.154 = 1/0.866$

Drawbacks:

- Does not reflect actual runtime because it normalizes
- Each application now counts equally

When is geomean useful?

Geometric mean of ratios is not proportional to total time

Use to compare machines when

- Relative performance on each program is known
- Relative runtime/weights of different programs is not known
- E.g., to aggregate speedups on set of programs

Rule of thumb: Use AM for times, HM for rates, GM for ratios

Summary of metrics

Name	Notation	Units	Comment
Memory footprint	-	Bytes	Total space occupied by the program in memory
Execution time	$(\sum_{1 \leq j \leq n} CPI_j) * \text{clock cycle time}$, where $1 \leq j \leq n$	Seconds	Running time of the program that executes n instructions
Arithmetic mean	$(E_1 + E_2 + \dots + E_p)/p$	Seconds	Average of execution times of constituent p benchmark programs
Weighted Arithmetic mean	$(f_1 * E_1 + f_2 * E_2 + \dots + f_p * E_p)$	Seconds	Weighted average of execution times of constituent p benchmark programs
Geometric mean	$p^{\text{th}} \text{ root } (E_1 * E_2 * \dots * E_p)$	Seconds	p^{th} root of the product of execution times of p programs that constitute the benchmark
Static instruction frequency		%	Occurrence of instruction i in compiled code
Dynamic instruction frequency		%	Occurrence of instruction i in executed code
Speedup (M_A over M_B)	E_B/E_A	Number	Speedup of Machine A over B
Speedup (improvement)	$E_{\text{Before}}/E_{\text{After}}$	Number	Speedup After improvement
Improvement in Exec time	$(E_{\text{old}} - E_{\text{new}})/E_{\text{old}}$	Number	New Vs. old
Iron Law	Time/Program = Insn/Prog * CPI * Time/Cycle	Seconds	Caution: optimization of one component may hurt another one, causing overall perf. loss!
Amdahl's law	$\text{Time}_{\text{after}} = \text{Time}_{\text{unaffected}} + \text{Time}_{\text{affected}}/x$	Seconds	x is amount of improvement