



CS2200 Systems and Networks Spring 2024

Lecture 10: Pipelining

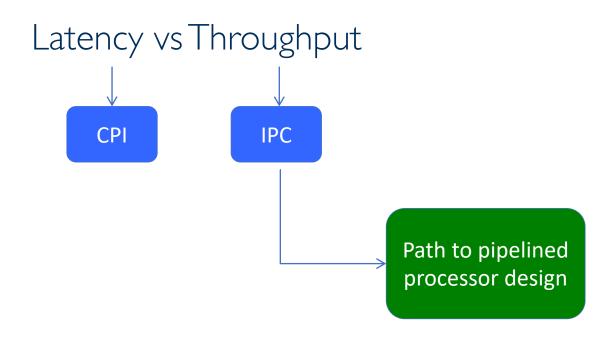
Alexandros (Alex) Daglis School of Computer Science Georgia Institute of Technology adaglis@gatech.edu

Lecture slides adapted from Bill Leahy and Charles Lively of Georgia Tech

Test | Logistics

- Syllabus
 - Chapters I-4: ISA, Processor Design, Interrupts
 - Chapter 5 (up to and including 5.6): Processor Performance
- Flow
 - 70% of questions released at 5pm on Sunday 2/11
 - Sunday 5pm to Monday 5pm
 - Clarification questions via Ed megathread (not email)
 - We will create ONE POST and post it on Tuesday, summarizing all clarifications as needed
 - If no clarification is made, assume that no clarification is needed
 - Test will be conducted in person during the lab on Wednesday 2/14
 - Questions released +30% hidden questions
 - Timed test (75 min) using honorlock and gradescope closed everything
 - Test taken without honorlock gets zero

Going faster...







station I
(place order)

station II (select bread)

station III (cheese)

station IV (meat)





station I
(place order)

station II (select bread)

station III (cheese)

station IV (meat)





station I (place order) station II
(select bread)

station III (cheese)

station IV (meat)





station I
(place order)

station II (select bread)

station III (cheese)

station IV (meat)





station I
(place order)

station II (select bread)

station III (cheese) station IV (meat)

station V
(veggies)





Net result: One complete sandwich every five cycles.



Bill's Mega-Sandwich Shop











station I
(place order)

station II (select bread)

station III (cheese)

station IV (meat)

station V
(veggies)

"What will You Have?"









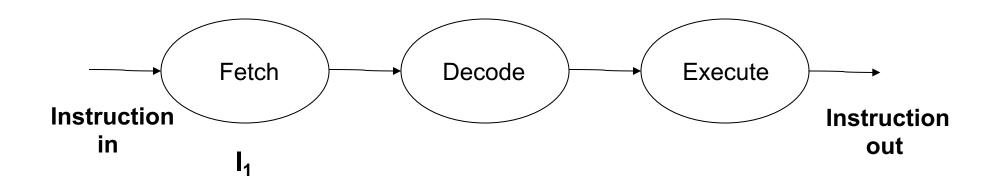


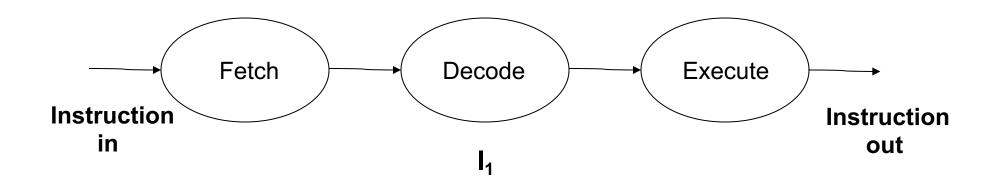


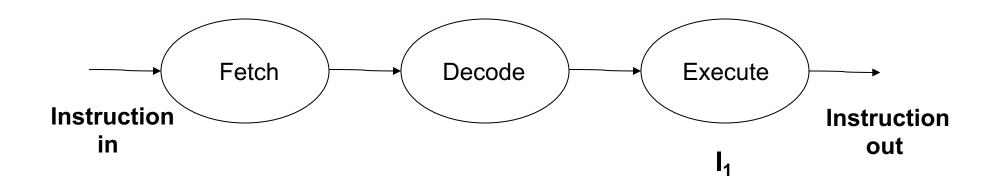


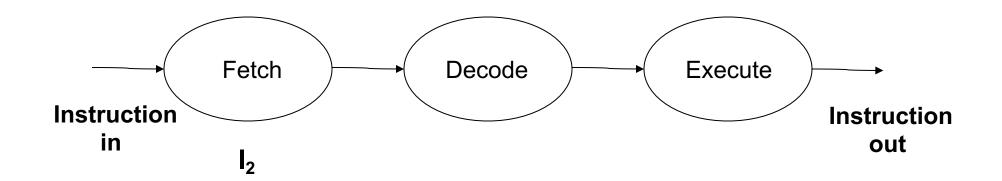
Net result: One complete sandwich every cycle! (Once you fill the pipeline.) A 5x speedup!

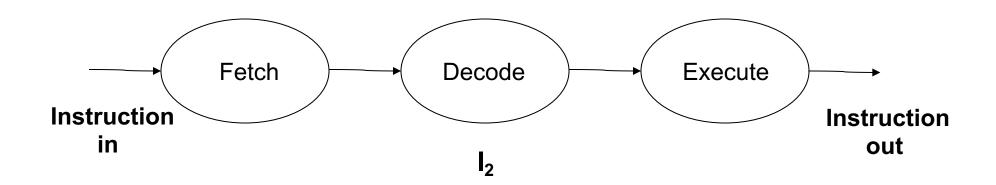


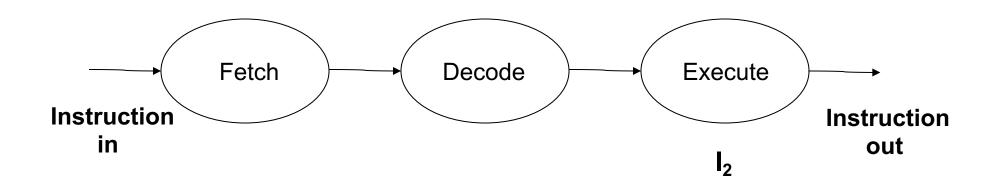




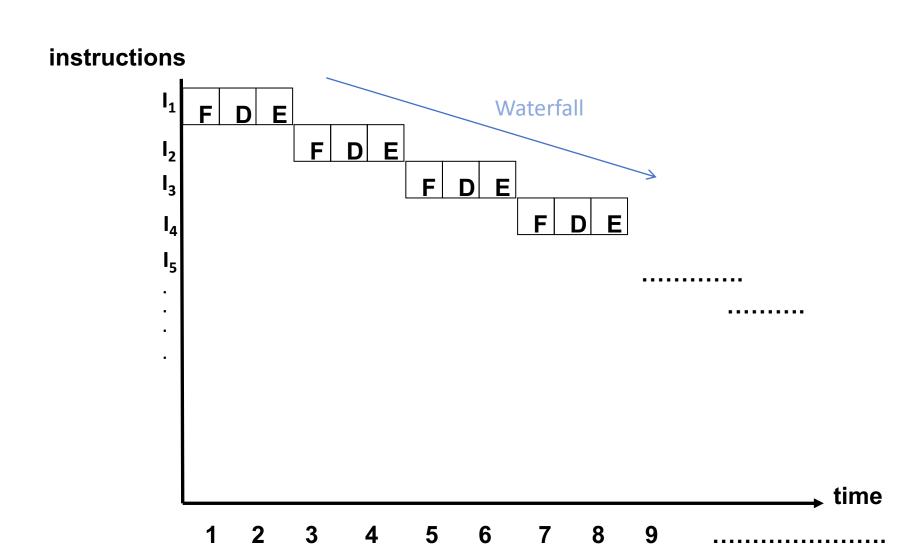




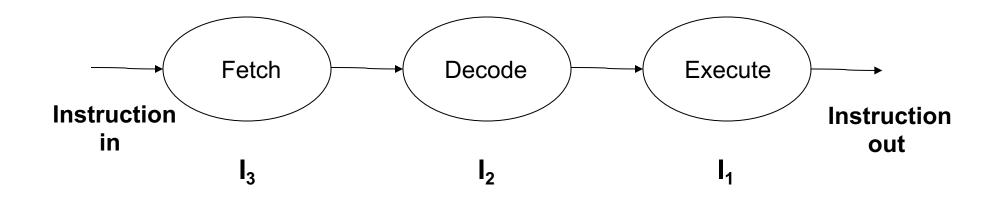




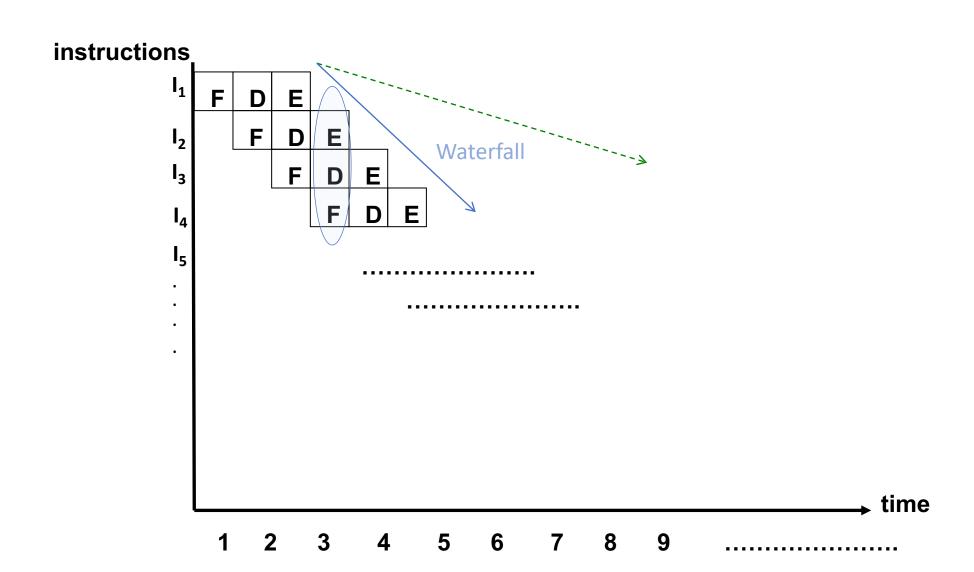
One at a time...



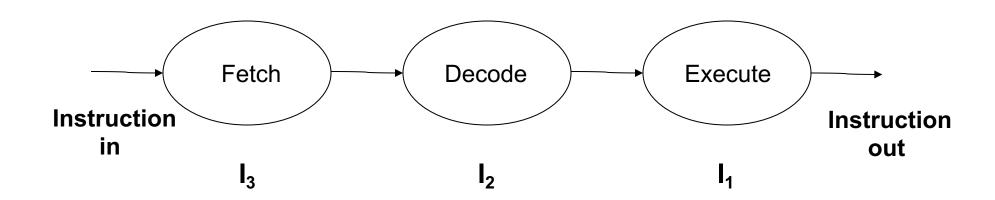
Three at a time?



Overlapping execution



What do we need in the datapath?



LC-2200 Datapath 32 LdPC -LdA LdB LdIR LdMAR PC MAR IR Addr IR[31..0] Din Din WrREG WrMEM registers memory 16x 1024x func ALU: IR[19..0] regno 00: ADD 32 bits 32 bits 01: NAND 20 10: A - B Dout Dout 11: A + 1 sign extend DrREG DrPC DrALU DrMEM DrOFF =0? IR[27..24] Rx: 4 -bit register number to control logic IR[23..20] Ry: 4 -bit register number to control logic LdZ IR[3..0] ▶ Rz: 4 -bit register number to control logic IR[31..28] ◆ OP: 4 -bit opcode to control logic

→ Z: 1 -bit boolean to control logic

What units are used?

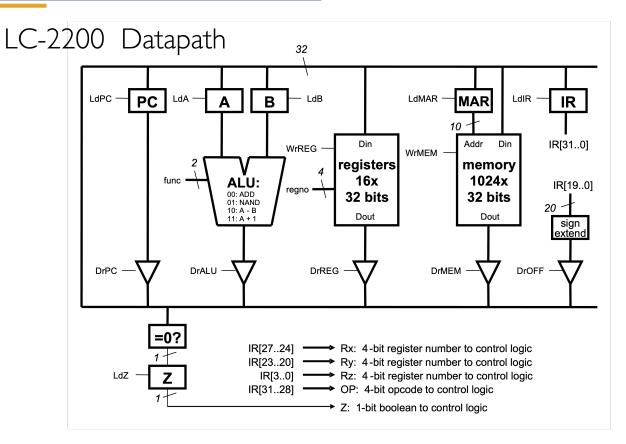
Macro State Units in Use

FETCH IR, ALU, PC, MEM

DECODE IR

EXECUTE (ADD) IR, ALU, Reg-file

EXECUTE (LW) IR, ALU, Reg-file, MEM, Sign extender



Imitation: The Sincerest Form of Flattery



station I (place order) New (5th order)



station II (select bread) 4th order



station III (cheese) 3rd order



station IV (meat) 2nd order

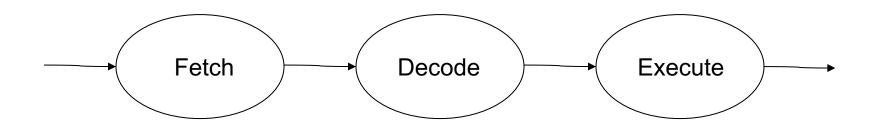


station V (veggies) 1st order

Points to note:

- Work done by each station is roughly equal
- Every sandwich goes through every station, regardless of what you want or don't want
- Order form and partially assembled sandwich is <u>passed from one station to the next</u>
- Each station does part of the work for assembling a sandwich

How do we mimic sandwich assembly?



- Unequal division of labor
- Violates first principle in sandwich assembly line

Bill's Cost-Cutting Measure









station I (place order) New (5th order)

station II (select bread) 4th order

















1st order



Bill's Cost-Cutting Measure







Station 1 (place order) 3rd order

station II (select bread) 2nd order

station III (cheese, meat, veggies) 1st order







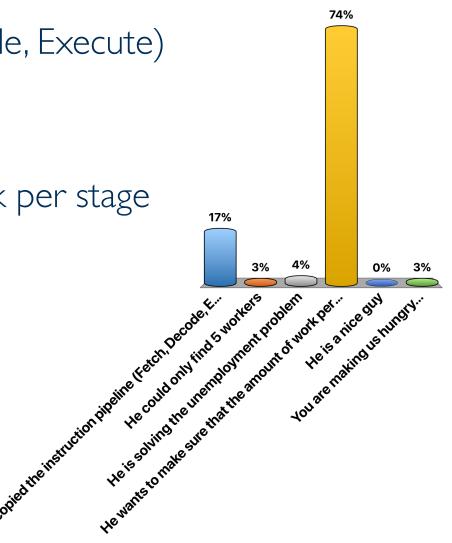
Disaster!

- Increased waiting times
- One sandwich every 3 cycles, not one every cycle!
- Two workers who only work about a third of the time

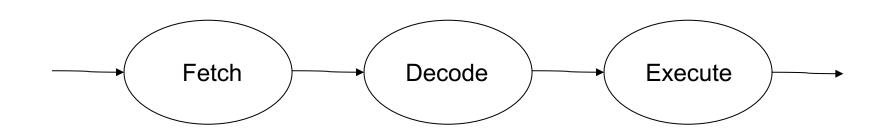


Bill originally used 5 stages rather than 3 in his sandwich pipeline because...

- A. He copied the instruction pipeline (Fetch, Decode, Execute)
- B. He could only find 5 workers
- C. He is solving the unemployment problem
- D. He wants to make sure that the amount of work per stage is roughly the same
- E. He is a nice guy
- F. You are making us hungry...



How do we mimic sandwich assembly?



- Get to the basics!
- What needs to happen for every instruction?
- Fetch into IR and increment PC
- Decode instruction and read register contents
- Perform arithmetic/logic (maybe)
 - Perform address computation (maybe)
- Fetch/store memory operand (maybe)
- Write to register (maybe)

Each step is roughly the same amount of work

How can we know what registers to read?

- What does the "bread guy" care about?
- How does the "bread guy" know what bread?
- It's the order form!



- What's the equivalent of the order form in the instruction pipeline?
- IR

LC-2200 Instruction set

bits 27-0:

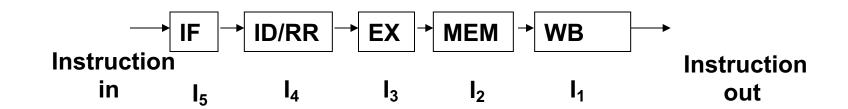
unused

R-type instructions (add, nand): bits 31-28: opcode; reg X; bits 27-24: reg Y; bits 19-4: unused (should be all 0s); bits 23-20: bits 3-0: reg Z I-type instructions (addi, lw, sw, beq): bits 31-28: opcode; bits 27-24: reg X bits 19-0: Imm. Offset bits 23-20: reg Y; J-type instructions (jalr): opcode; bits 31-28: bits 27-24: reg X bits 23-20: reg Y; bits 19-0: unused O-type instructions (halt):

bits 31-28:

opcode;

Every instruction goes through



- Some stages don't do anything for some instructions
- Each stage works on a different instruction

```
IF - fetch instruction into IR and increment PC
```

ID/RR - decode and read register contents

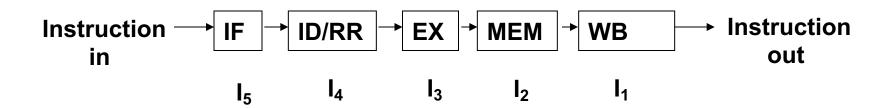
EX - perform arithmetic/logic (maybe)

- perform address computation (maybe)

MEM - fetch/store memory operand (maybe)

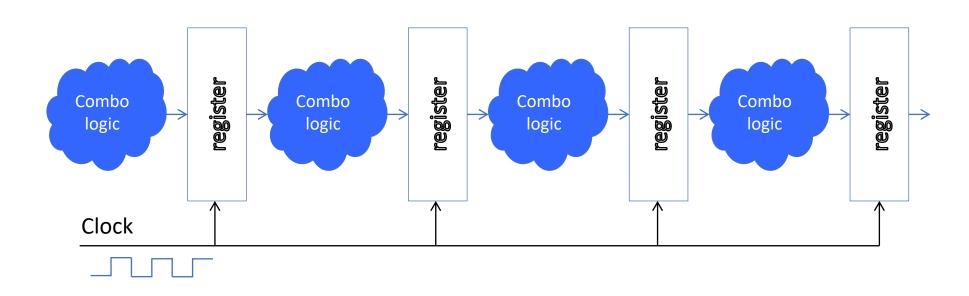
WB - write to register (maybe)

What moves between stages?



- Think of our sandwich assembly line
- Each station passes two things to the next
 - The order form
 - A partially assembled sandwich
- Analogies with our data path
 - Is the IR equivalent to the order form?
 - What is a "partially assembled data sandwich"?

Electrical isolation between stages



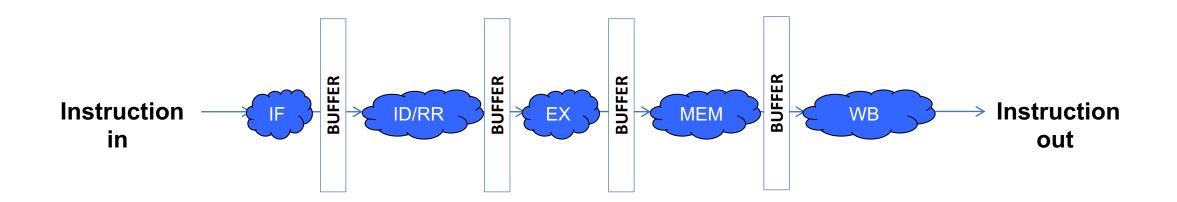
- Registers act as a wall to isolate the actions of each combinational logic cloud!
- On each clock tick, the registers send their output to drive the next stage of combinational logic

What's in these pipeline boxes?



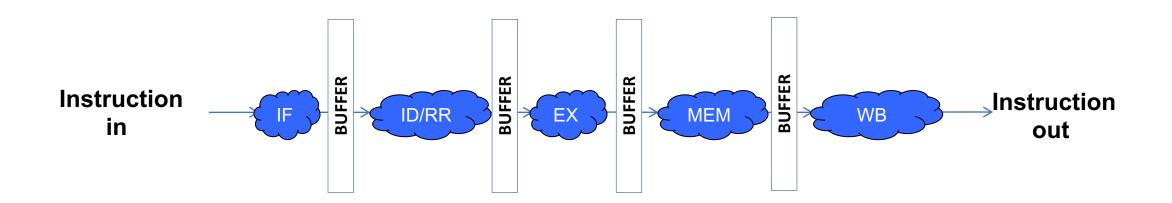
- How do we separate our pipeline boxes?
- The boxes are really just combinational logic clouds
- So we already know how to separate them...

This is where buffers come from...



- There is combinational logic in each cloud for performing the actions of that stage
- The buffers are merely clocked registers for passing the partially assembled results

Again: What does each stage do?



F - fetch instruction into IR and increment PC

ID/RR - read register contents

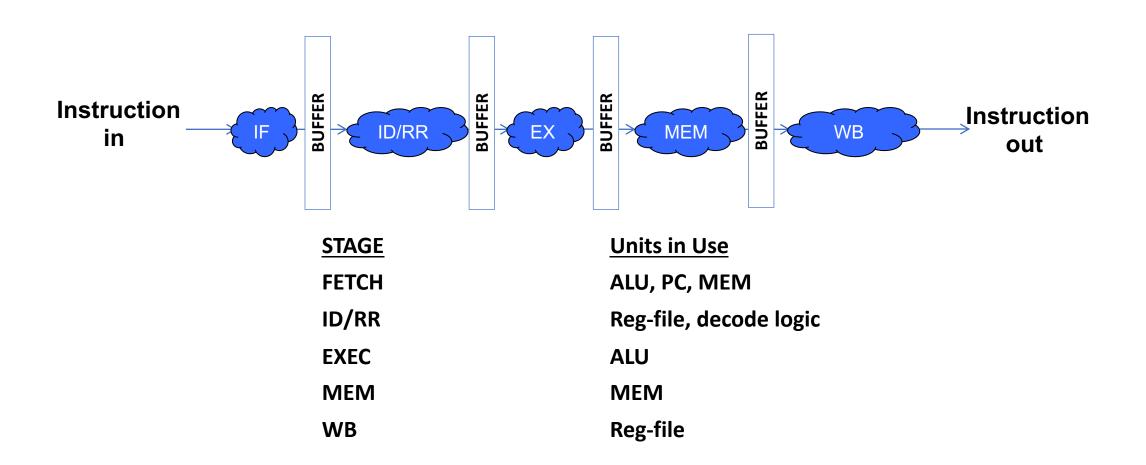
EX _ - perform arithmetic/logic (maybe)

- perform address computation (maybe)

MEM - fetch/store memory operand (maybe)

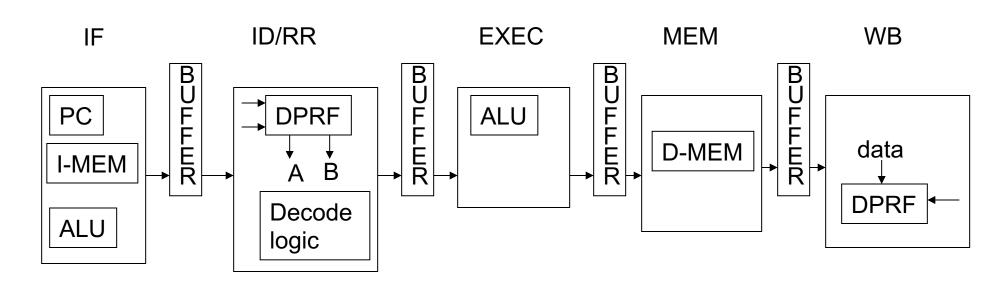
WB - write to register (maybe)

Again: What does each stage need?



Conclusion: we need a little datapath in each stage

A shiny new data path!



STAGE Units in Use

FETCH ALU, PC, MEM

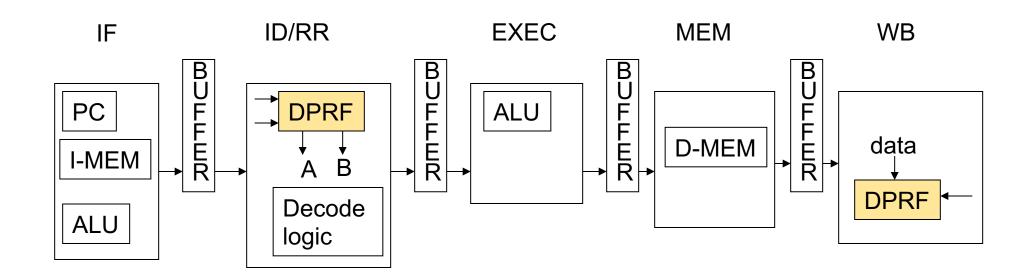
ID/RR Reg-file, decode logic

EXEC ALU

MEM MEM

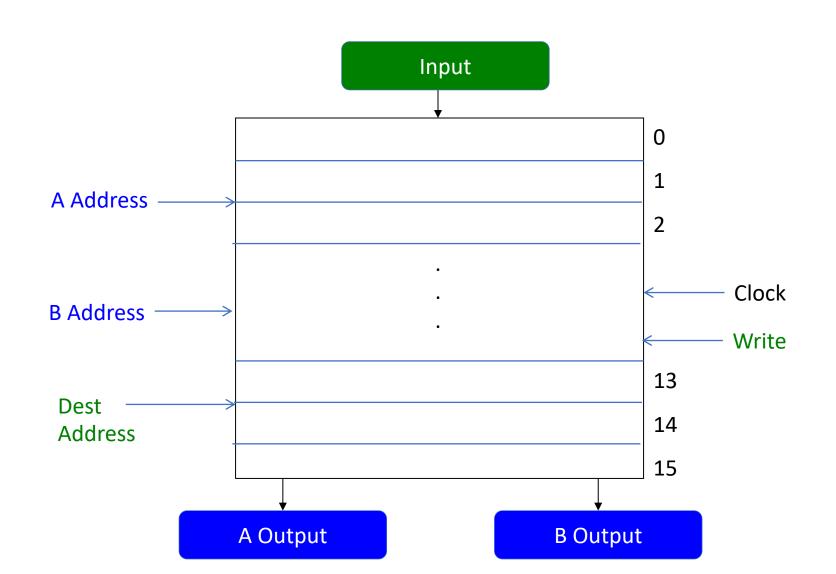
WB Reg-file

Two register files? Not.

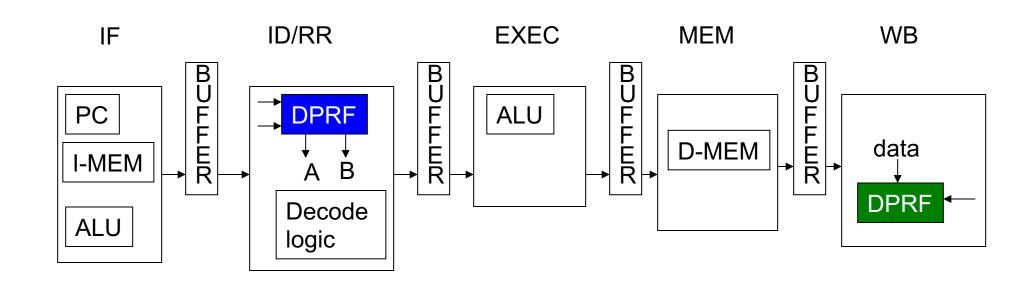


We use a register file in two stages

Do you remember this device?

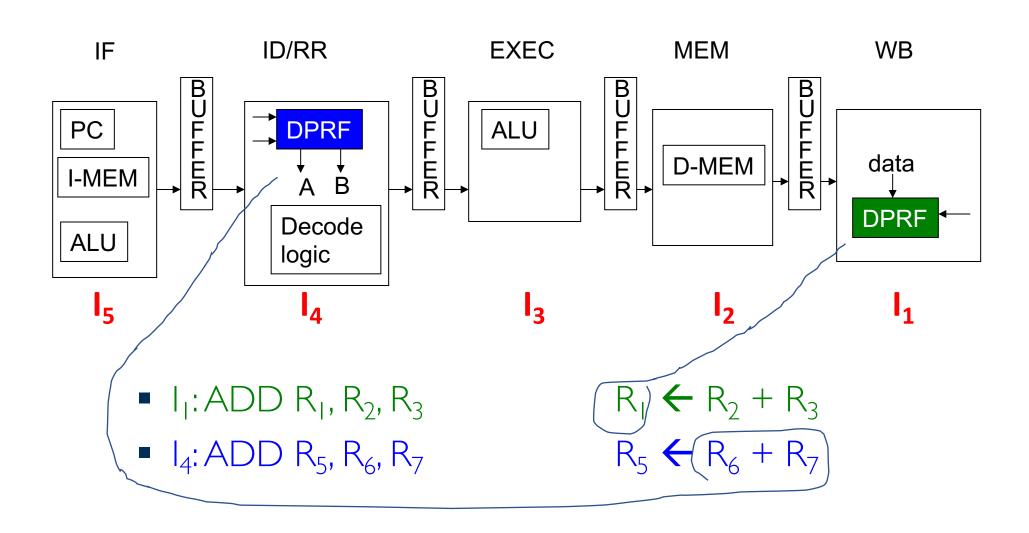


Two register files? Not.

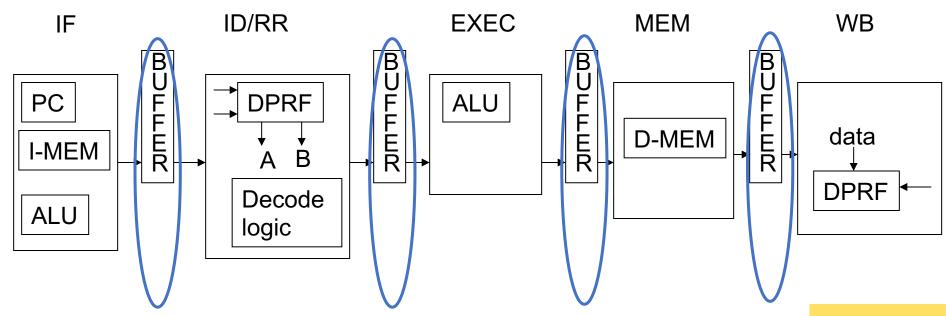


- We use a register file in two stages, but conveniently
 - ID/RR uses read ports (A addr, B addr B, A output, B output)
 - WB uses write port (Dest addr, Input)
- So we can happily share the same register file on each clock cycle

Example of sharing the register file



But what about those buffers?



- What goes in those buffers? Same? Different?
- The different stages seem to need different information
- What was that sandwich analogy?

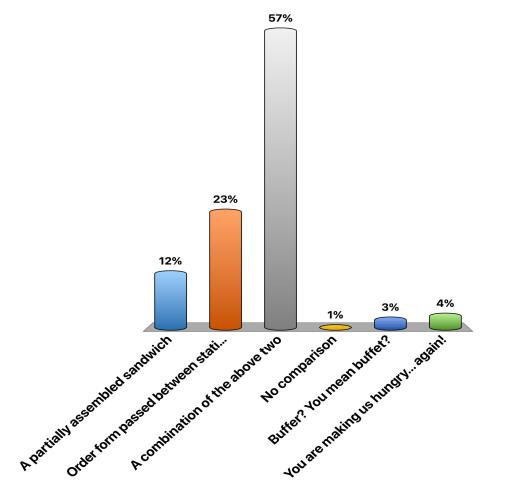




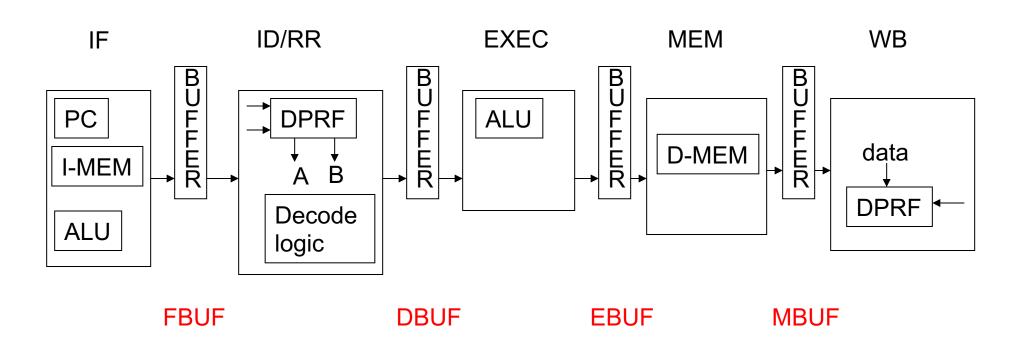
Comparing the instruction pipeline...

...to Bill's sandwich pipeline, the buffers between stages serve the same function as

- A. A partially assembled sandwich
- B. Order form passed between stations
- C. A combination of the above two
- D. No comparison
- E. Buffer? You mean buffet?
- F. You are making us hungry... again!

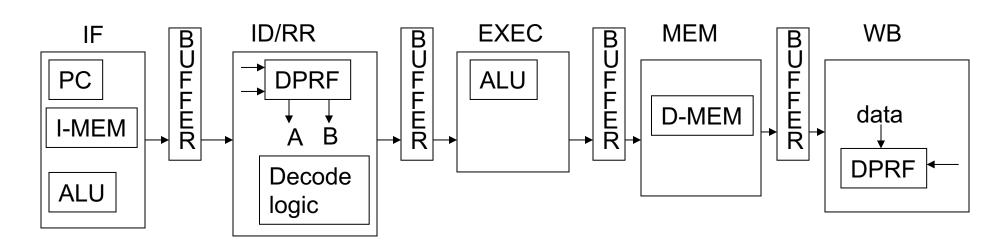


So what makes up our partially assembled data sandwich?



- The IR (or parts of it) is the order form
- The data sandwich is the partial execution results
- Do the buffers have to all carry the same information? Or can they be different?

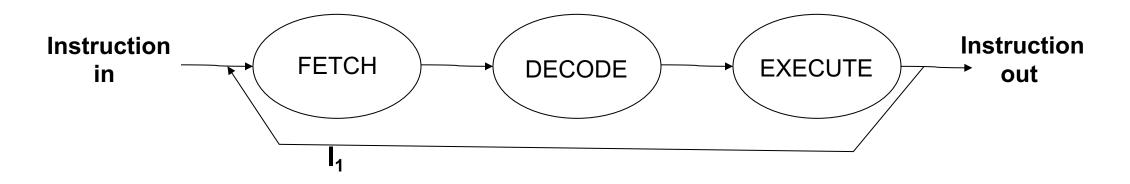
What does each stage need in the buffers?



FBUF DBUF EBUF MBUF

	Vame	Output of Stage	Contents	
F	BUF	F	Primarily contains instruction read from memory	
	DBUF	ID/RR	Decoded IR and values read from register file	
E	BUF	EX	Primarily contains result of ALU operation plus other parts of the	
			instruction depending on the instruction specifics	
1	1BUF	MEM	Same as EBUF if instruction is not LW or SW; If instruction is LW, then	
			buffer contains the contents of memory location read	

Our non-pipelined LC-2200

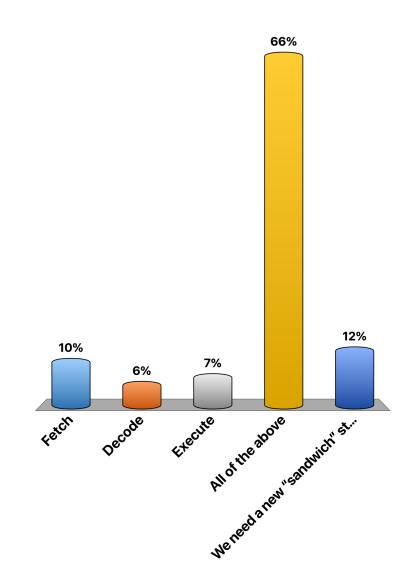


At each point, the processor is in exactly one macro state



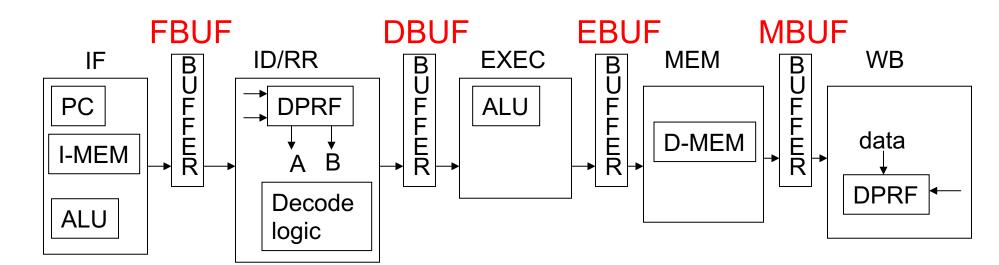
What "macro state" is a pipelined processor in at any point in time?

- A. Fetch
- B. Decode
- C. Execute
- D. All of the above
- E. We need a new "sandwich" state



Anatomy of an ADD instruction

- IF stage (cycle 1)
- ID/RR stage (cycle 2)
- EX stage (cycle 3)
- MEM stage (cycle 4)
- WB stage (cycle 5)



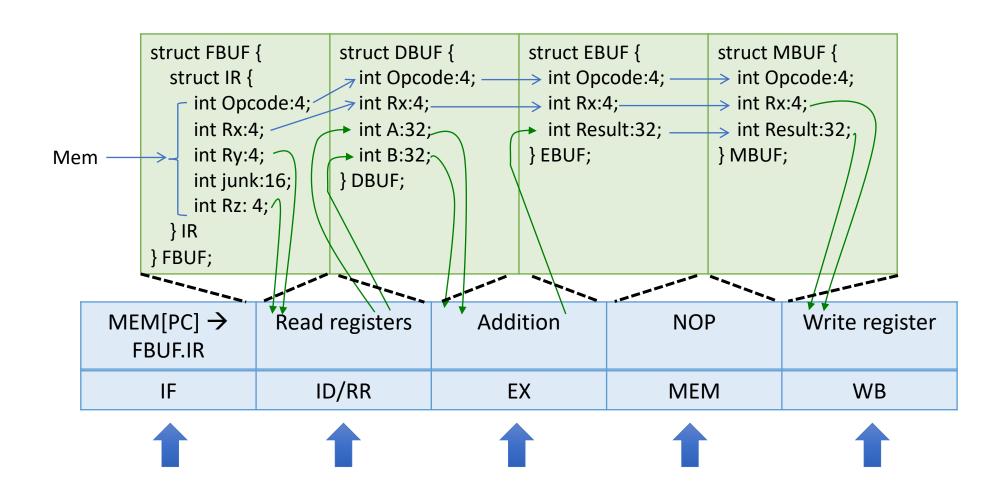
Actions for an ADD instruction

Stage	Actions	Comments
IF	MEM[PC] → FBUF.IR PC + I → PC	Instruction is fetched from memory and placed in FBUF (which is essentially the IR)
ID/RR	DPRF[FBUF.IR.Ry] → DBUF.A DPRF[FBUF.IR.Rz] → DBUF.B FBUF.IR.Opcode → DBUF.Opcode FBUF.IR.Rx → DBUF.Rx	Read register Ry into DBUF.A Read register Rz into DBUF.B Copy opcode from FBUF to DBUF Copy Rx from FBUF to DBUF
EX	DBUF.A + DBUF.B → EBUF.Result DBUF.OPCODE → EBUF.Opcode DBUF.Rx → EBUF.Rx	Perform addition Copy opcode from DBUF to EBUF Copy Rx from DBUF to EBUF
MEM	EBUF → MBUF	Copy EBUF to MBUF
WB	MBUF.Result → DPRF[MBUF.Rx]	Write the result of addition into register Rx

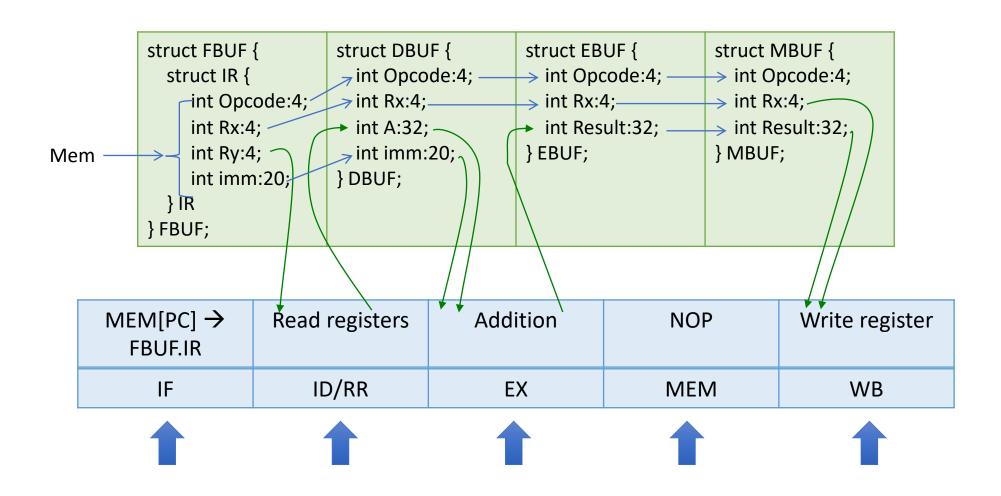
Actions for an ADD instruction

Stage	Actions	Comments	
IF	MEM[PC] → FBUF.IR	Instruction is fetched from memory and placed	
	PC + I → PC	Considering only the Add	
ID/RR	DPRF[FBUF.IR.Ry] → DBUF.A	instruction, quantify the sizes of	
	DPRF[FBUF.IR.Rz] → DBUF.B FBUF.IR.Opcode → DBUF.Opcode FBUF.IR.Rx → DBUF.Rx	the various buffers between the stages of the pipeline.	
		EDUE	
EX	DBUF.A + DBUF.B -> EBUF.Result	FBUF? lition	
	DBUF.OPCODE → EBUF.Opcode DBUF.R× → EBUF.R×	DBUF? le from DBUF to EBUF FBUF? m DBUF to EBUF	
	DDOI.IX / LDOI.IX		
MEM	EBUF → MBUF	MBUF? to MBUF	

How does this pipeline work for ADD?



How does this pipeline work for ADDI?



How to generalize design of pipeline register?

- Look at each instruction in the ISA
- Decide what needs to be carried over from one stage to the next

Opcode other fields

Design the DBUF register for LC-2200

- Don't optimize by overloading the different fields of the register
- Decide what fields need to be in DBUF for each instruction in the ISA

ADD/NAND	ADDI	BEQ	And so on
struct DBUF { int Opcode:4; int Rx:4; int A:32; int B:32; } DBUF;	struct DBUF { int Opcode:4; int Rx:4; int A:32; int imm:20; } DBUF;	struct DBUF { int Opcode:4; int A:32; int B:32; int imm:20; int PC:32; } DBUF;	

Now take the UNION of the requirements

ADD/NAND	ADDI	BEQ	And so on
struct DBUF { int Opcode:4; int Rx:4; int A:32; int B:32; } DBUF;	struct DBUF { int Opcode:4; int Rx:4; int A:32; int imm:20; } DBUF;	struct DBUF { int Opcode:4; int Rx:4; int imm:20; int PC:32; } DBUF;	



```
struct DBUF {
  int Opcode:4;
  int Rx:4;
  int A:32;
  int B:32;
  int imm:20;
  int PC:32;
} DBUF;
```

Now take the UNION of the requirements

```
struct DBUF {
  int Opcode:4;
  int Rx:4;
  int A:32;
  int B:32;
  int imm:20;
  int PC:32;
} DBUF;
```

- In the ID/RR stage, fill in fields depending on the instruction in FBUF.IR
- Leave unneeded fields unfilled
- What fields would we fill in for
 JALR \$at,\$ra # \$ra ← PC; PC ← \$at

For the BEQ instruction, the fields in DBUF not filled in...

```
BEQ Rx, Ry, offset ; If (Rx == Ry) PC \leftarrow PC + offset
```

- A. A (contents of Ry)
- B. B (contents of Rx)
- C. PC
- D. Immediate value
- E. Rx specifier

```
struct DBUF {
  int Opcode:4;
  int Rx:4;
  int A:32;
  int B:32;
  int imm:20;
  int PC:32;
} DBUF;
```

