



CS2200
Systems and Networks
Spring 2023

Lecture 12: Pipeline Hazards, pt 2

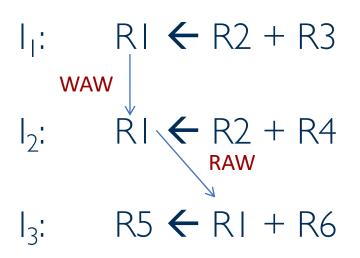
Alexandros (Alex) Daglis School of Computer Science Georgia Institute of Technology adaglis@gatech.edu

Lecture slides adapted from Bill Leahy and Charles Lively of Georgia Tech

Pipeline Hazards

- Structural hazards
 - Datapath/resource limitations
- Data hazards
 - Program limitations
- Control hazards
 - Program limitations

What hazards?



Register file

В	RP
В	RP
В	RP
В	RP
	RP
В	RP
В	RP
В	RP

Must ensure that I_3 gets the correct value of R1 (from I_2 instead of I_1)

- Without forwarding, need to make sure I_1 does not clear the busy bit $-I_2$ should do that
 - i.e., 13 cannot move past ID/RR before 12 is also done with WB
- With forwarding, need to make sure R1 value is forwarded from EX stage (where I_2 is), not from MEM (where I_1 is)

Control hazard

IF* ID/RR EX MEM WB

- 1.BEQ
- 2.ADD
- 3.NAND
- 4.LW

* We do not actually know what the instruction is in the Fetch Stage before we decode it

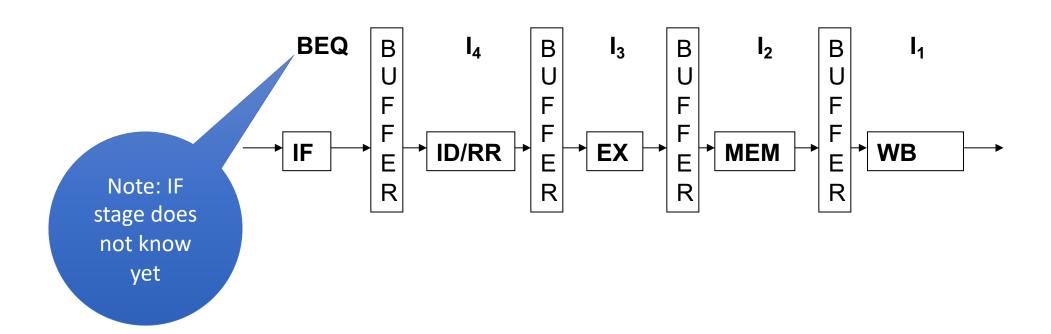
Is there any analog in the sandwich pipeline?

Conservative handling of branches

- Wait until branch outcome is known
 - Stall the pipeline → send NOPs from IF state
- Resume normal execution when branch outcome is known
 - Result of comparison (A-B) is known
 - If there is need to branch, PC gets the target address of the branch
- Two cases
 - Branch TAKEN (i.e. A==B)
 - Branch NOT TAKEN (i.e. A!=B)

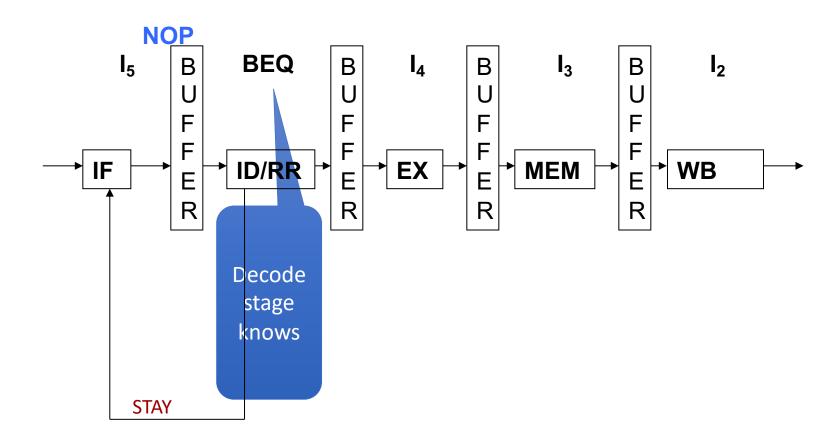
Conservative branch handling

Cycle 1



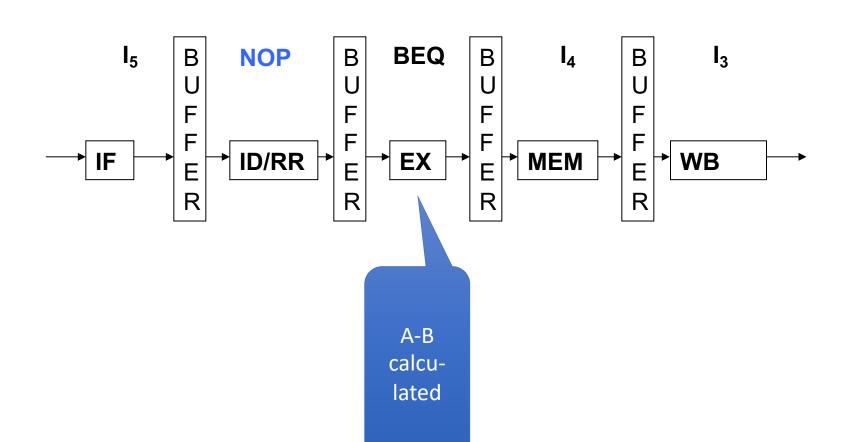
Branch handling

Cycle 2

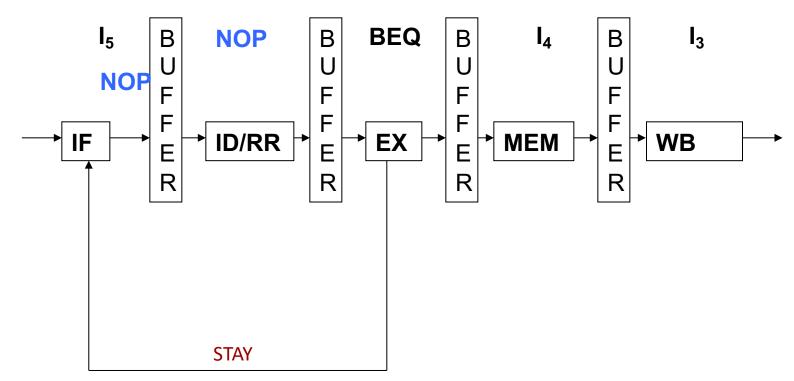


Branch handling

Cycle 3

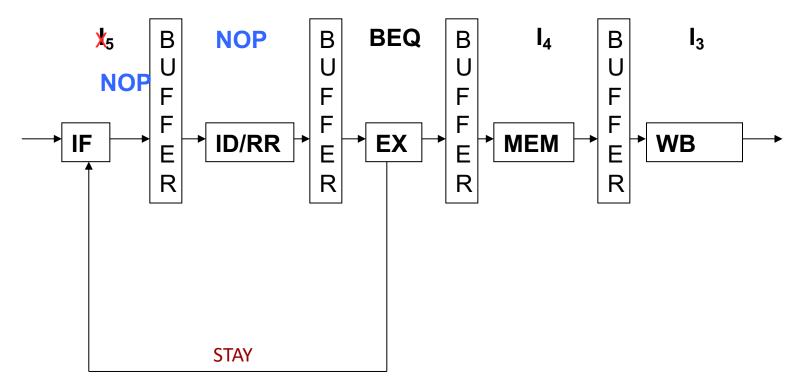


Cycle 3



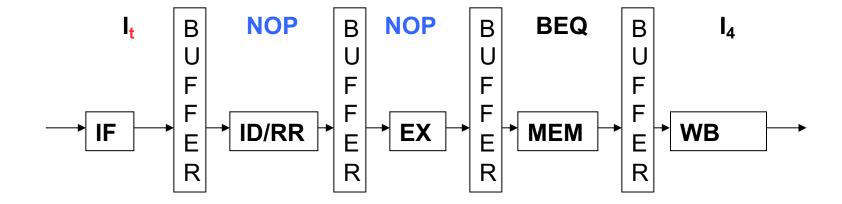
A-B is calculated as zero and branch is going to be taken PC+offset is being calculated

Cycle 3



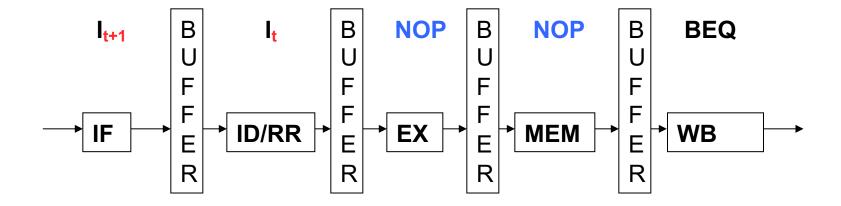
A-B is calculated as zero and branch is going to be taken PC+offset is being calculated > value is clocked into PC at end of EX

Cycle 4

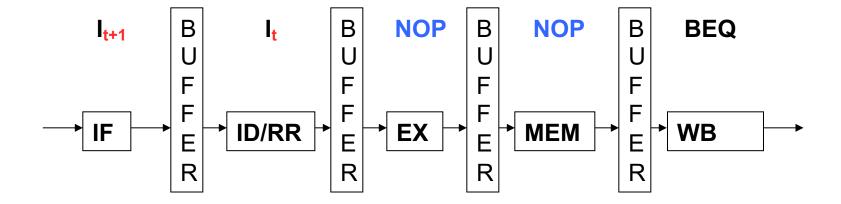


We used updated PC to fetch the target of branch

Cycle 5



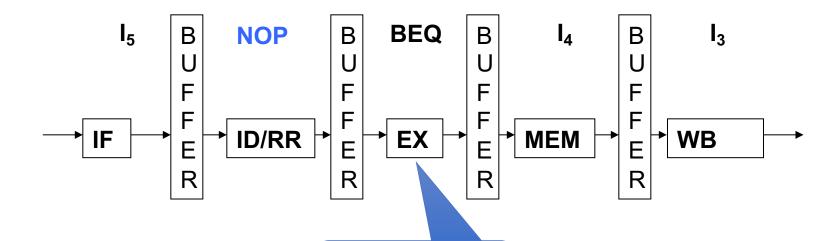
Cycle 5



Outcome: Two pipeline bubbles

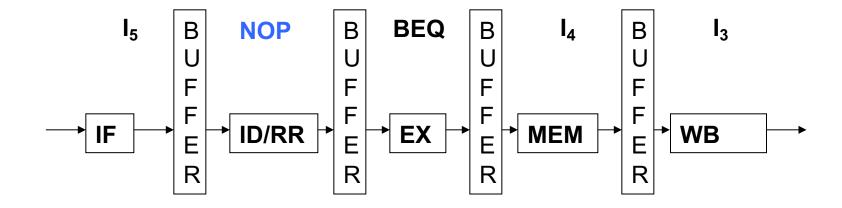
Branch handling

Cycle 3



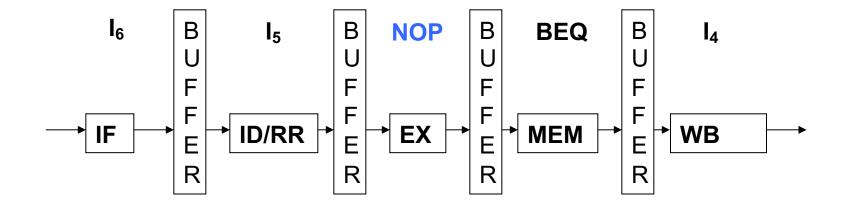
A-B calculated

Cycle 3



A-B calculated as non-zero; branch not taken PC is already pointing to the sequential path

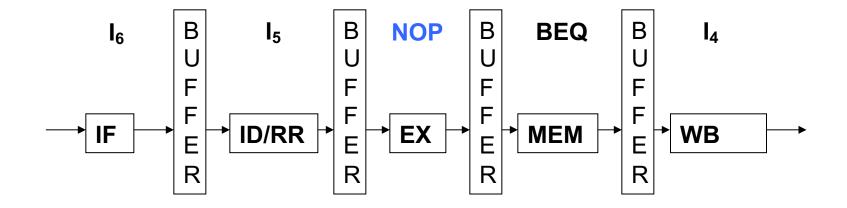
Cycle 4



A-B calculated as non-zero; branch not taken Normal pipeline operation resumed



Cycle 4



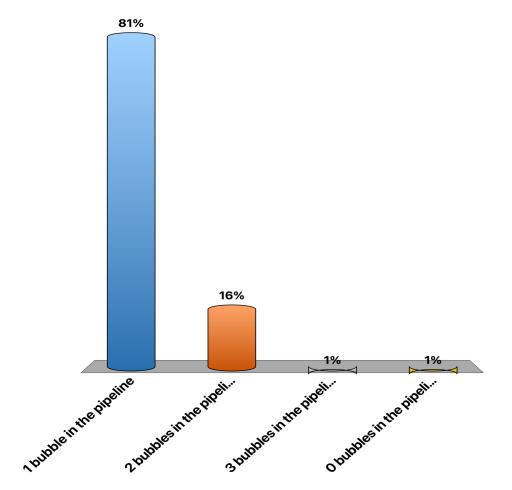
Outcome: one pipeline bubble



Question

Conservative handling of branches introduces at least...

- A. I bubble in the pipeline
- B. 2 bubbles in the pipeline
- C. 3 bubbles in the pipeline
- D. 0 bubbles in the pipeline



Code for our examples

Assume the following sequence of instructions:

BEQ LI ADD LW LI NAND SW

The instruction executed after BEQ will be ADD or NAND, depending on BEQ's outcome

Perfect pipeline

Cycle	IF	ID/RR	EX	MEM	WB
Number					
1	BEQ	-	-	-	-
2	ADD	BEQ		-	-
3	NAND	ADD	BEQ	-	-
4	SW	NAND	ADD	BEQ	-
5	-	SW	NAND	ADD	BEQ
6	-	-	SW	NAND	ADD
7			-	SW	NAND
8				-	SW
9					

Unfortunately stalls ruin a perfect pipeline

Conservative Branch Handling

Cycle Number	IF	ID/RR	EX	MEM	WB
1	BEQ	-	-	-	-
2	ADD	BEQ	-	-	-
3	(ADD)	NOP	BEQ	-	-
4	NAND	NOP	NOP	BEQ	-
5	SW	NAND	NOP	NOP	BEQ
6	-	SW	NAND	NOP	NOP
7	-	-	SW	NAND	NOP
8			-	SW	NAND
9				-	SW

Cycle 2: Still Don't Know If We Branch

Only at the end of EX stage is the target address of BEQ known

Cycle Number	IF	ID/RR	EX	MEM	WB
1	BEQ	-	-	-	-
2	ADD	BEQ	-	-	-
3	(ADD)	NOP	BEQ	-	-
4	NAND	NOP	NOP	BEQ	-
5	SW	NAND	NOP	NOP	BEQ
6	-	SW	NAND	NOP	NOP
7	-	-	SW	NAND	NOP
8			-	SW	NAND
9				-	SW

What should cycle 2 do?

- → ID/RR tells IF to stall after fetching
- → IF sends NOP to ID/RR at end of cycle

Cycle 3: Branch taken...

IF sends another NOP into ID/RR at the end of cycle 3
- (ADD is on the wrong path and gets squashed)
EX sets PC to PC + offset, so NAND is fetched

Cycle Number	IF	ID/RR	EX	MEM	WB
1	BEQ	-	-	-	-
2	ADD	BEQ	-	-	-
3	(ADD)	NOP	BEQ	-	-
4	NAND	NOP	NOP	BEQ	-
5	SW	NAND	NOP	NOP	BEQ
6	-	SW	NAND	NOP	NOP
7	-	-	SW	NAND	NOP
8			-	SW	NAND
9				-	SW

Assume BR is taken

A - BPC \leftarrow PC + offset

Cycle 4: Branch Taken

Only at the end of the EX stage was the target addr of BEQ known

Cycle	IF	ID/RR	EX	MEM	WB
Number					
1	BEQ	-	-	-	-
2	ADD	BEQ	-	-	-
3	(ADD)	NOP	BEQ	-	-
4	NAND	NOP	NOP	BEQ	-
5	SW	NAND	NOP	NOP	BEQ
6	-	SW	NAND	NOP	NOP
7	-	-	SW	NAND	NOP
8			-	SW	NAND
9				-	SW

We have these two NOPs introduced by BEQ stalling the pipeline



Cycle 9: Finished

Cycles 5-9, the pipeline runs normally (i.e. no new stalls)

Cycle	IF	ID/RR	EX	MEM	WB
Number					
1	BEQ	-	-	-	-
2	ADD	BEQ	-	-	-
3	(ADD)	NOP	BEQ	-	-
4	NAND	NOP	NOP	BEQ	-
5	SW	NAND	NOP	NOP	BEQ
6	-	SW	NAND	NOP	NOP
7	-	-	SW	NAND	NOP
8			-	SW	NAND
9				-	SW



Cycle 4: Branch Not Taken

EX on cycle 3 allows IF and ID/RR to proceed normally on cycle 4

	Cycle	IF	ID/RR	EX	MEM	WB
	Number					
	1	BEQ	_	-	-	-
İ	2	ADD	BEQ	-	-	-
	3	(ADD)	NOP	BEQ	-	-
	4	LW	ADD	NOP	BEQ	-
ĺ	5		LW	ADD	NOP	BEQ
	6	-		LW	ADD	NOP
	7	-	-		LW	ADD
	8			-		LW
	9				-	

Finish here with stalls

What's the CPI difference?

Cycle	IF	ID/RR	EX	MEM	WB
Number					
1	BEQ	-	-	_	-
2	ADD	BEQ	-	-	-
3	(ADD)	NOP	BEQ	-	-
4	NAND	NOP	NOP	BEQ	-
5	SW	NAND	NOP	NOP	BEQ
6	-	SW	NAND	NOP	NOP
7	-	-	SW	NAND	NOP
8			-	SW	NAND
9				-	SW



Branch taken: 9 total cycles / 3 instructions = 3 CPI

Not taken: 8 total cycles / 3 instructions = 2.66 CPI

Strategies for improving branches

- Branch prediction
- Delayed branch

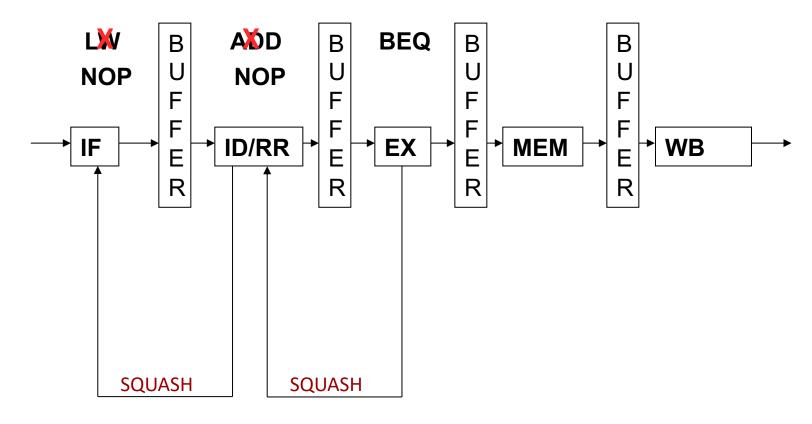
Branch Prediction

Assume branch is not taken

```
Cycle IF* ID/RR EX MEM WB
1 BEQ
2 ADD BEQ
3 LW ADD BEQ
```

^{*} We do not actually know what the instruction is in the Fetch Stage, before it's decoded

Branch Prediction



Our prediction was wrong! Branch will be taken Squash the instructions at IF and ID/RR

Static Branch "not taken" prediction

Given the following sequence of instructions:

BEQ L'

ADD

LW

. . . .

L1 NAND

SW

Using the branch prediction approach, what is the observed CPI for the 3 instructions (BEQ, ADD, LW) or (BEQ, NAND, SW) in each case?

Happy path [⊙] – prediction correct!

Cycle	IF	ID/RR	EX	MEM	WB
Number					
1	BEQ	-	-	-	-
2	ADD	BEQ	-	-	-
3	LW	ADD	BEQ	-	-
4	-	LW	ADD	BEQ	-
5	-	-	LW	ADD	BEQ
6	-	-	-	LW	ADD
7	_	-	-	_	LW

7 cycles / 3 instructions – 2.33 CPI

What if prediction turns out false?

Sad path 😊 — Branch prediction failed

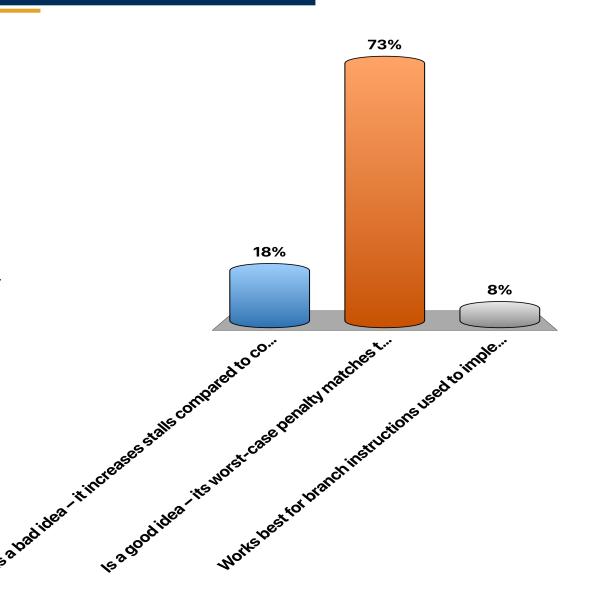
Squash ADD and LW and replace with NOPs; basically we degrade to the conservative approach

Cycle	IF	ID/RR	EX	MEM	WB
Number					
1	BEQ	-	-	-	-
2	ADD	BEQ	-	-	-
3	LW	ADD	BEQ	-	-
4	NAND	NOP	NOP	BEQ	-
5	SW	NAND	NOP	NOP	BEQ
6	-	SW	NAND	NOP	NOP
7	-	-	SW	NAND	NOP
8			-	SW	NAND
9				-	SW



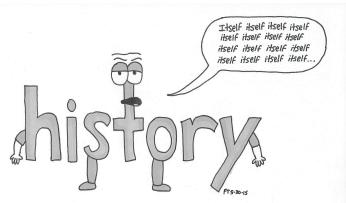
Static "not taken" prediction...

- A. Is a bad idea it increases stalls compared to conservative branch handling when most branches are actually taken
- B. Is a good idea its worst-case penalty matches that of conservative branch handling
- C. Works best for branch instructions used to implement loops

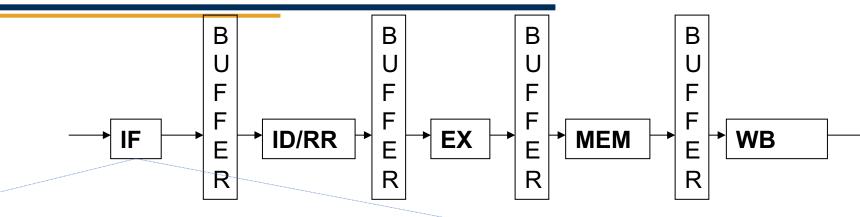


Branch prediction heuristics

- How do we (quickly) guess which way a branch will go?
- Can often tell by comparing target with current PC
 - Loops usually branch backwards (target < PC)
 - Guess the branch will be taken
 - Can construct conditionals to usually branch forwards (target > PC)
 - Guess the branch won't be taken
- Some ISAs have 2 different families of branch instructions so that the compiler can signal whether it thinks the branch will be taken or not taken
- Or we could keep a history...



Branch History Table and Branch Target Buffer



PC Instr 1000 BEQ

•••

1300 BEQ

. .

1500 BEQ

••

1600 BEQ

..

1800 BEQ

PC of branch instruction	Taken/Not Taken	PC of branch target
1000	Т	1200
1300	N	xxxx
1500	Т	1100
1600	N	xxxx
1800	N	xxxx

•••



Delayed branching

- Ignore the problem in hardware (i.e. no automatic generation of NOPs in fetch stage)
- Give the NOP instructions to the compiler
 - E.g. the instruction after a branch is always executed
- Let the compiler worry about it
 - Compile in a NOP instruction after every branch instruction
- Why could this be a good idea?
 - The compiler may be able to optimize...

Example using MIPS-style branch delays

```
; Add 7 to each element of a ten-element array
; whose address is in a0
     addi $t1, $a0, 40; 40 = 10 elements * 4 bytes
 loop:
         $a0, $t1, done; When a0=t1 we are done
     beq
                       ; branch delay slot
     nop
     lw $t0, 0($a0)
     addi $t0, t0, 7
     sw $t0, 0($a0)
     addi $a0, $a0, 4
     beg $zero, $zero, loop
                       ; branch delay slot
     nop
 done: halt
```

Let's assume one delay slot

Assume pipeline determines if branch is taken in its second stage

Delayed branch – before optimization

Delayed branch – during optimization

```
; Add 7 to each element of a ten-element array
; whose address is in a0
     addi $t1, $a0, 40
 loop:
          $a0, $t1, done
     beq
                         ; branch delay slot
     nop
    _ lw $t0, 0($a0)
                                   OK if LW is
     addi $t0, $t0, 7
                                   redundantly
      sw $t0, 0($a0)
     addi $a0, $a0, 4
                                executed on loop exit
     beq $zero, $zero, loop
                         ; branch delay slot
     nop
 done: halt
```

Delayed branch – during optimization

Delayed branch – after opt

```
; Add 7 to each element of a ten-element array
; whose address is in a0
    addi t1, a0, 40
loop:
    beq $a0, $t1, done
    lw $t0, 0($a0) ; branch delay slot

    addi $t0, $t0, 7
    sw $t0, 0($a0)

    beq $zero, $zero, loop
    addi $a0, $a0, 4 ; branch delay slot
    done: halt
```

Delayed branching

- Ignore the problem in hardware (i.e. no automatic generation of NOPs in fetch stage)
- Give the NOP instructions to the compiler
 - E.g. the instruction after a branch is always executed
- Let the compiler worry about it
 - Compile in a NOP instruction after every branch instruction
- Why could this be a good idea?
 - The compiler may be able to optimize...
- Why is this mostly a bad idea?
 - Exposes microarchitectural detail to programmer breaks abstraction!
 - Dynamic branch prediction proved to be very effective

Name	Pros	Cons	Use cases
Stall the pipeline	Simple strategy, no hardware needed for squashing instructions	Loss of performance	Early pipelined machines such as IBM 360 series
Branch Prediction (branch not taken)	Results in good performance with virtually no additional hardware since the instruction is anyhow being fetched from the sequential path already in IF stage	Needs ability to squash instructions in partial execution in the pipeline	Processors such as Intel Pentium, AMD Athlon, and PowerPC use this technique; typically they also employ sophisticated branch target buffers; MIPS R4000 uses a combination 1-delay slot plus a 2-cycle branch-not-taken prediction
Dynamic History- based Branch Prediction	Results in good performance but requires slightly more elaborate hardware design	Since the new PC value that points to the target of the branch is not available until the branch instruction is in EX stage, this technique requires more elaborate hardware assist to be practical	-
Delayed Branch	No need for any additional hardware for either stalling or squashing instructions; It involves the compiler by exposing the pipeline delay slots and takes its help to achieve good performance	With increase in depth of pipelines of modern processors, it becomes increasingly difficult to fill the delay slots by the compiler Exposes microarchitectural details to ISA level Assembly harder to understand	Older RISC architectures such as MIPS, PA-RISC, SPARC

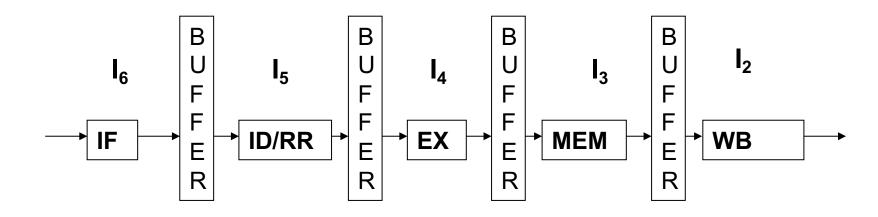
Summary of hazards

Instruction	Type of Hazard	Potential Stalls	With Data forwarding	With branch prediction (branch not taken)
ADD, NAND	Data	0, 1, 2, or 3	0	Not Applicable
LW	Data	0, 1, 2, or 3	0 or 1	Not Applicable
BEQ	Control	1 or 2	Not Applicable	0 (success) or
				2 (mispredict)

^{*} Structural hazards can always be eliminated by replicating bottleneck hardware resource

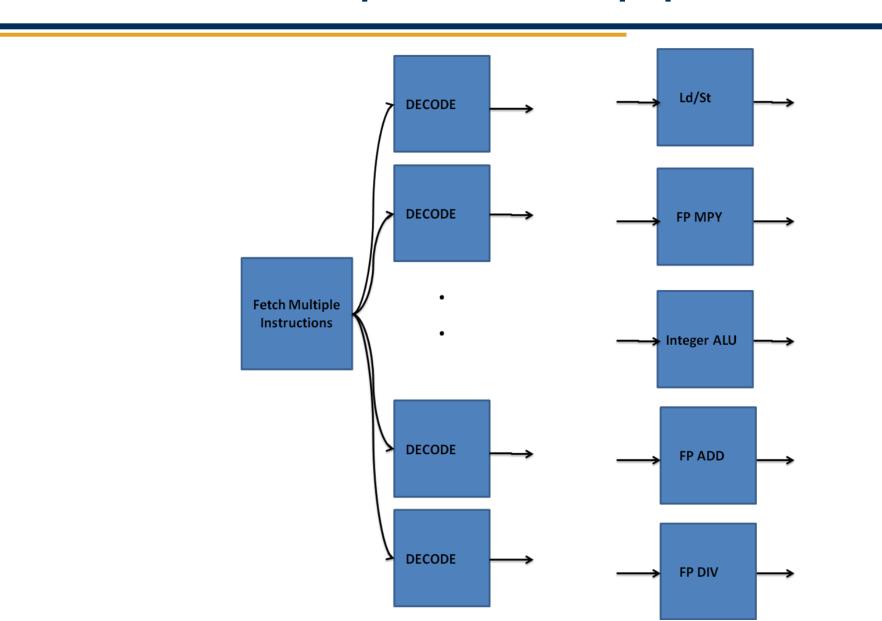
^{**} These penalties refer to the specific 5-stage pipeline we saw. Different penalties for different pipelines.

Program discontinuities



- External interrupt can occur at any time
- When do we take the interrupt?
- Any one of the 5 instructions may cause an exception or trap
- As soon as an instruction raising trap/exception reaches WB: STOP, FLUSH, go to INT state
 - Slow to respond to external events

Superscalar pipelines



Different Pipeline Depths

