

CS2200

Systems and Networks

Spring 2024

Lecture 23:

Programmed IO & DMA

Alexandros (Alex) Daglis
School of Computer Science
Georgia Institute of Technology
adaglis@gatech.edu

Roadmap

- Wrap up parallel systems
 - Hardware support for threads
- Programmed IO and DMA
 - Chapter 10 (10.1 – 10.7)
- Networking
 - Chapter 13

Synchronization support in a uniprocessor

Synchronization support in a uniprocessor


- Thread creation/termination

Synchronization support in a uniprocessor

- Thread creation/termination
- Communication among threads

Synchronization support in a uniprocessor

- Thread creation/termination
- Communication among threads




Nothing special needed...

PT, TLB, Cache all the same

Synchronization support in a uniprocessor

- Thread creation/termination
- Communication among threads
- Synchronization among threads
 - How do we implement mutex_lock?



Nothing special needed...

PT, TLB, Cache all the same

Proposed implementation of mutex

Proposed implementation of mutex

- Lock():
 while (mem_lock != 0)
 block the thread
 mem_lock = 1;

Unlock():
 mem_lock = 0

Proposed implementation of mutex

- Lock():
 while (mem_lock != 0)
 block the thread
 mem_lock = 1;

Unlock():
 mem_lock = 0

- What could go wrong?

Proposed implementation of mutex

- Lock():
 while (mem_lock != 0)
 block the thread
 mem_lock = 1;

Oops! These
instructions aren't
atomic

Unlock():
 mem_lock = 0

- What could go wrong?

Proposed implementation of mutex

- Lock():
 while (mem_lock != 0)
 block the thread
 mem_lock = 1;

Oops! These instructions aren't atomic

Unlock():
 mem_lock = 0

How did we deal with that earlier?

- What could go wrong?

Proposed implementation of mutex

- Lock():
 while (mem_lock != 0)
 block the thread
 mem_lock = 1;



Oops! These instructions aren't atomic

Unlock():
 mem_lock = 0

How did we deal with that earlier?

- What could go wrong?

We used OS mutex calls to make instructions inseparable.

Proposed implementation of mutex

- Lock():
 while (mem_lock != 0)
 block the thread
 mem_lock = 1;



Oops! These instructions aren't atomic

Unlock():
 mem_lock = 0

How did we deal with that earlier?

We used OS mutex calls to make instructions inseparable.

But now we ARE the OS!

- What could go wrong?

Lock() using machine instructions

T1

➡ Lock()
START LD R1,mem_lock
BZ SET
JSR block_thread
B START
SET ADDI R1,#1
ST R1,mem_lock
RET

T2

➡ Lock()
START LD R1,mem_lock
BZ SET
JSR block_thread
B START
SET ADDI R1,#1
ST R1,mem_lock
RET

R1	mem_lock	R1
x	0	x

Lock() using machine instructions

T1

T2

Lock()
START

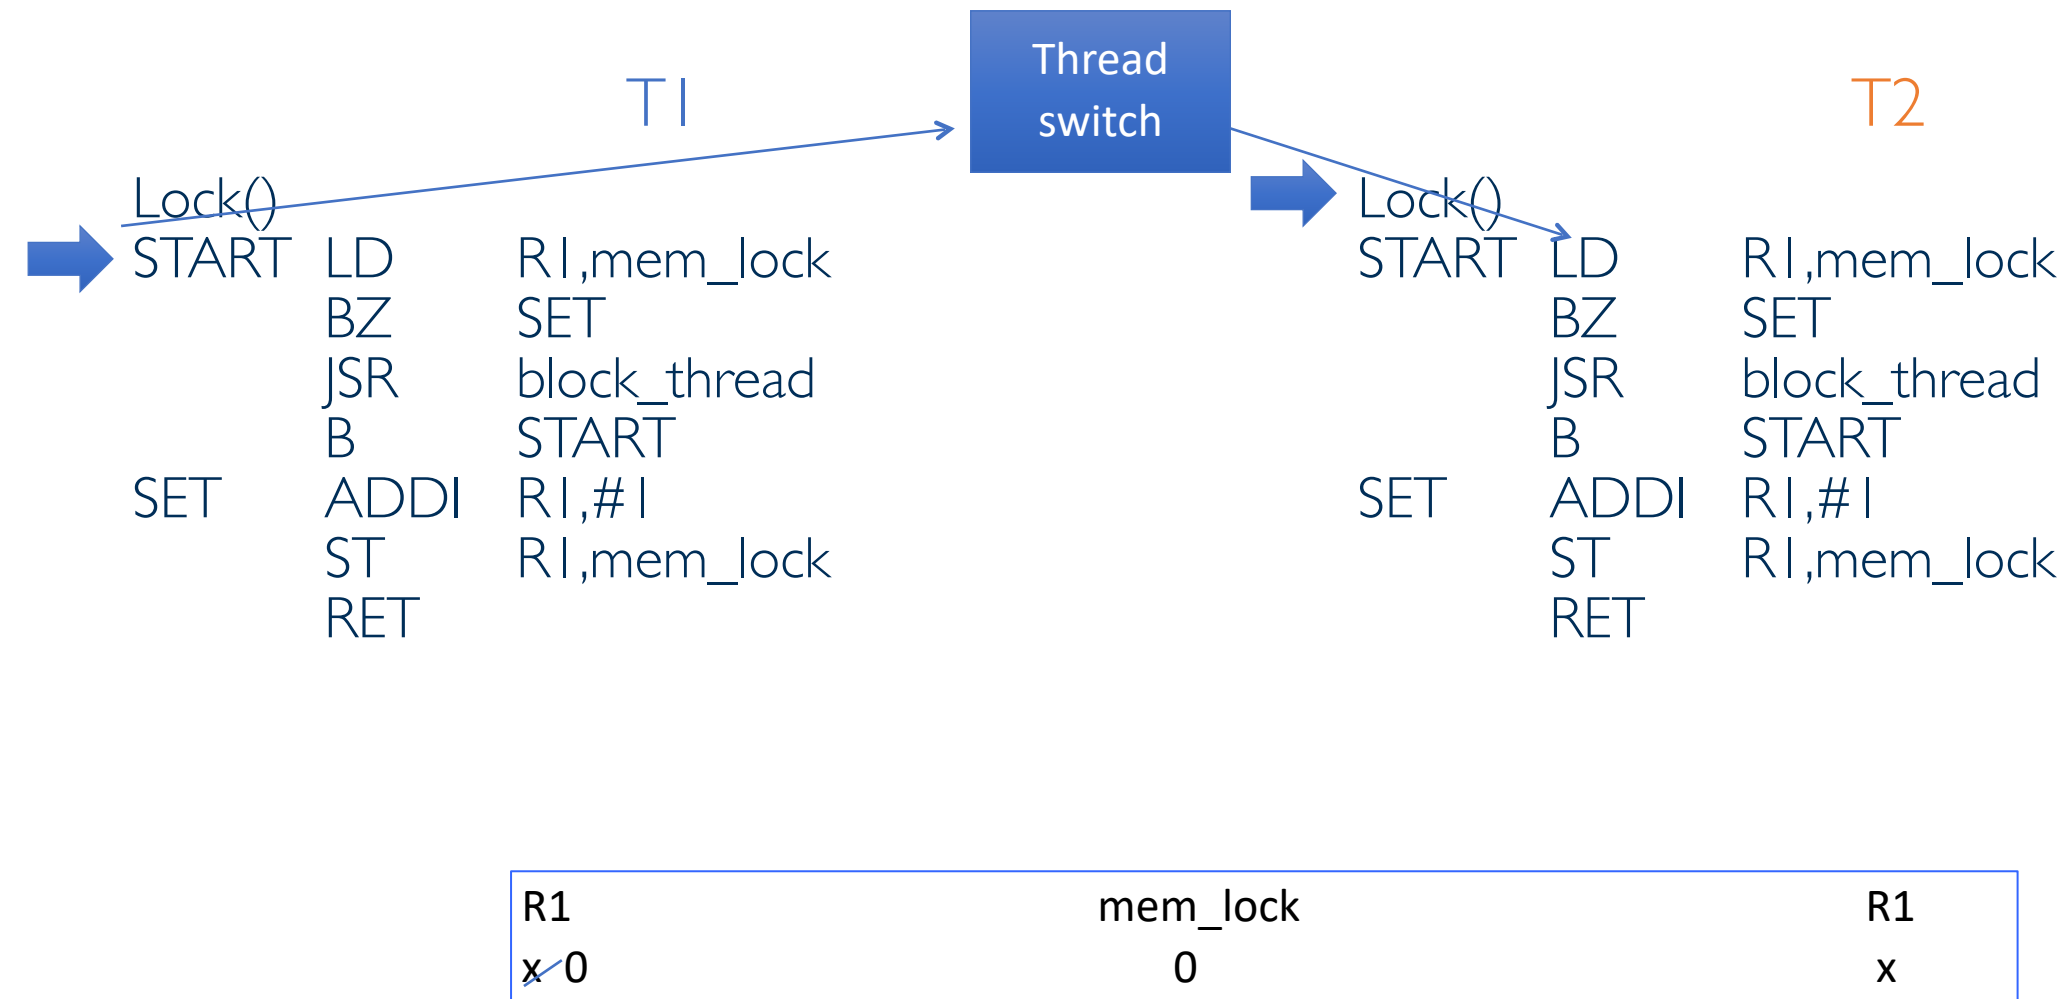
LD R1,mem_lock
BZ SET
JSR block_thread
B START
SET ADDI R1,#1
ST R1,mem_lock
RET

Lock()
START

LD R1,mem_lock
BZ SET
JSR block_thread
B START
SET ADDI R1,#1
ST R1,mem_lock
RET

R1	mem_lock	R1
x 0	0	x

Lock() using machine instructions



Lock() using machine instructions

T1

T2

Lock()
START

LD R1,mem_lock
BZ SET
JSR block_thread
B START
SET ADDI R1,#1
ST R1,mem_lock
RET

Lock()
START

LD R1,mem_lock
BZ SET
JSR block_thread
B START
SET ADDI R1,#1
ST R1,mem_lock
RET

R1	mem_lock	R1
x 0	0	x 0

Lock() using machine instructions

T1

→ Lock()
START LD R1,mem_lock
BZ SET
JSR block_thread
B START
SET ADDI R1,#1
ST R1,mem_lock
RET

T2

→ Lock()
START LD R1,mem_lock
BZ SET
JSR block_thread
B START
SET ADDI R1,#1
ST R1,mem_lock
RET

R1	mem_lock	R1
x 0	0	x 0

Lock() using machine instructions

T1

Lock()
START
LD R1,mem_lock
BZ SET
JSR block_thread
B START
SET ADDI R1,#1
ST R1,mem_lock
RET

T2

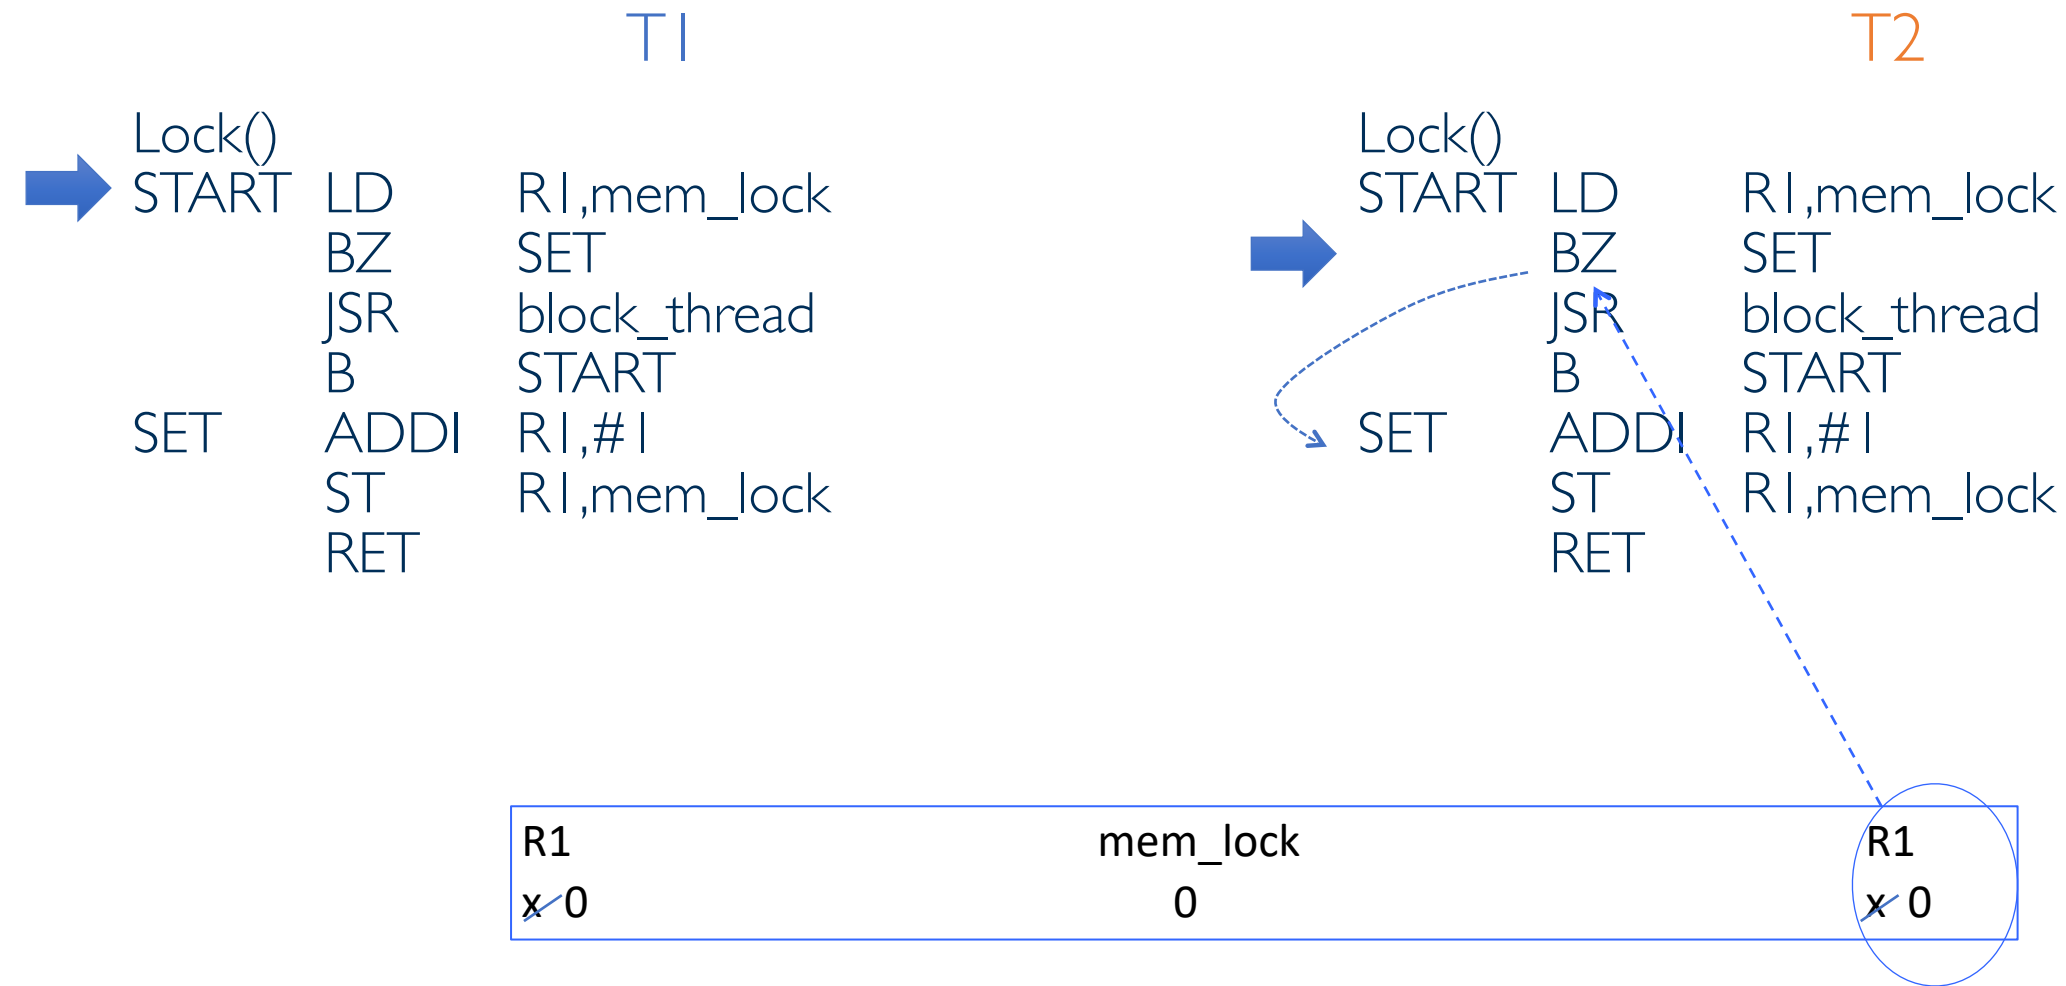
Lock()
START
LD R1,mem_lock
BZ SET
JSR block_thread
B START
SET ADDI R1,#1
ST R1,mem_lock
RET

R1
~~x~~ 0

mem_lock
0

R1
~~x~~ 0

Lock() using machine instructions



Lock() using machine instructions



R1	mem_lock	R1
x 0	0	x 0 1

Lock() using machine instructions



R1	mem_lock	R1
x 0	0 1	x 0 1

Lock() using machine instructions

T1

T2

Lock()
START

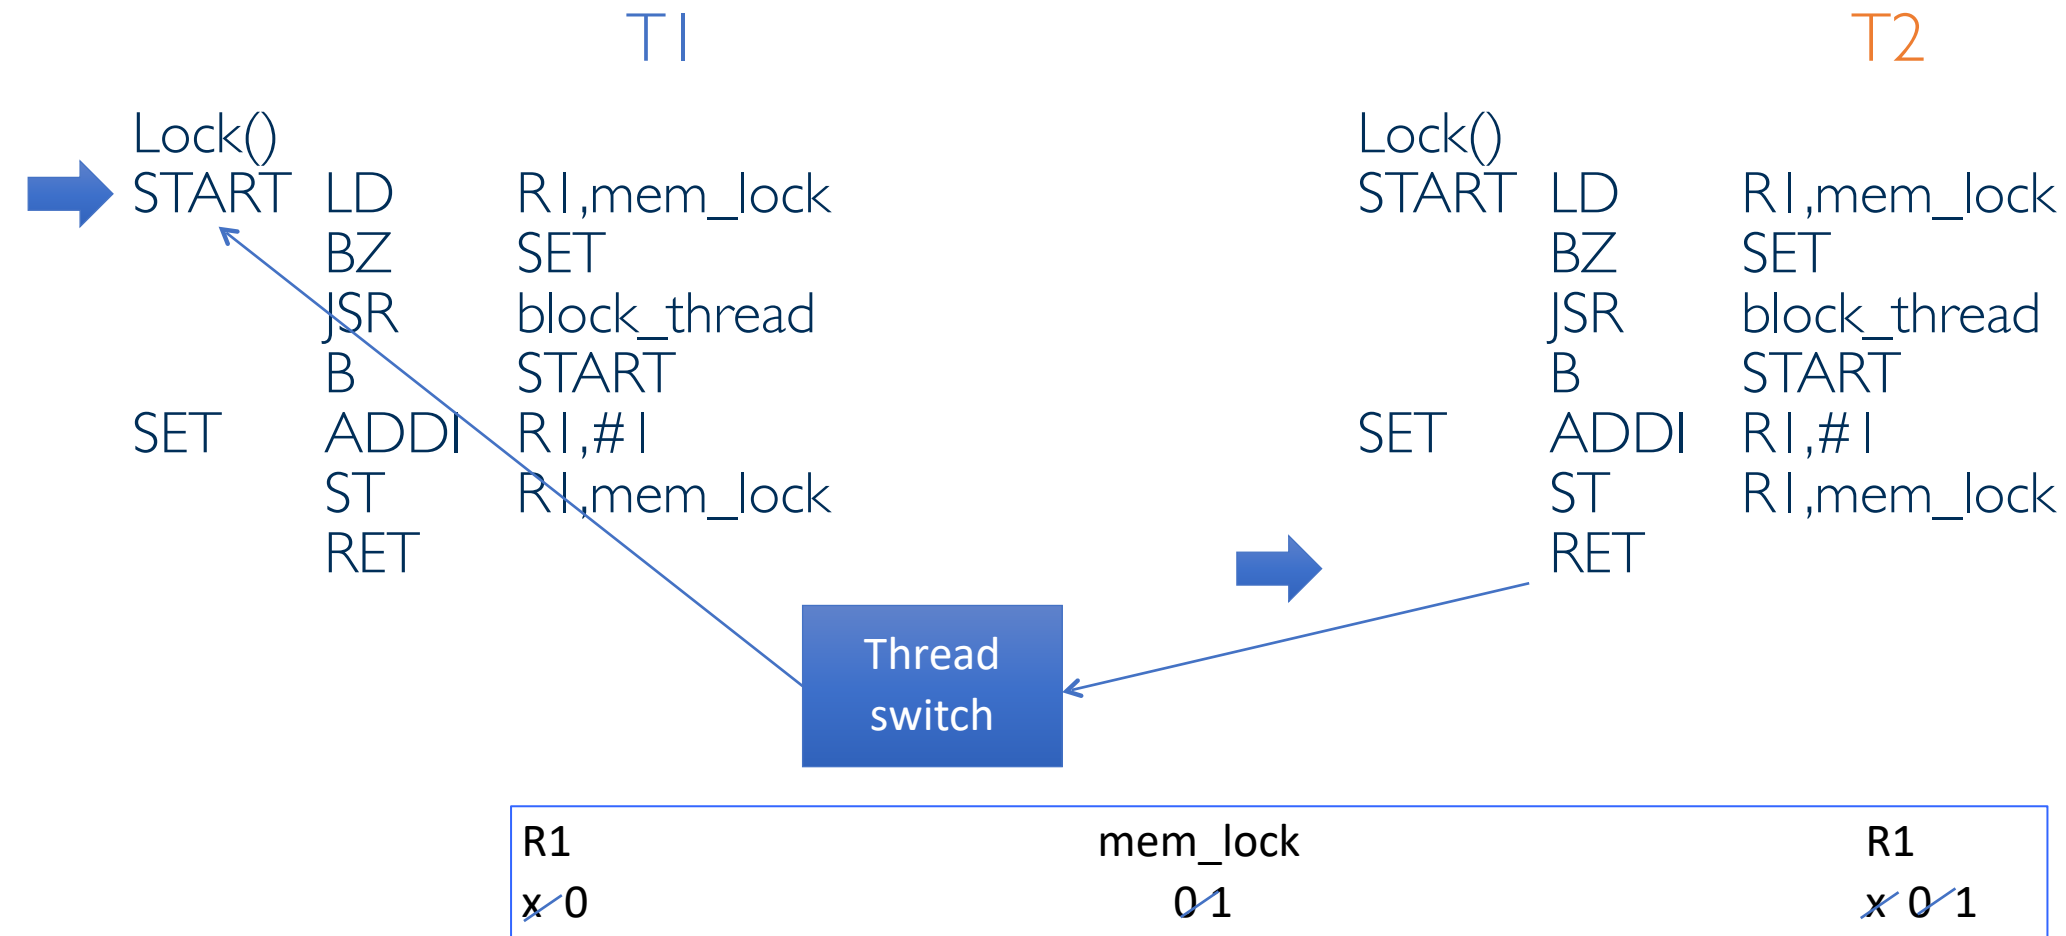
LD R1,mem_lock
BZ SET
JSR block_thread
B START
SET ADDI R1,#1
ST R1,mem_lock
RET

Lock()
START

LD R1,mem_lock
BZ SET
JSR block_thread
B START
SET ADDI R1,#1
ST R1,mem_lock
RET

R1	mem_lock	R1
x 0	0 1	x 0 1

Lock() using machine instructions

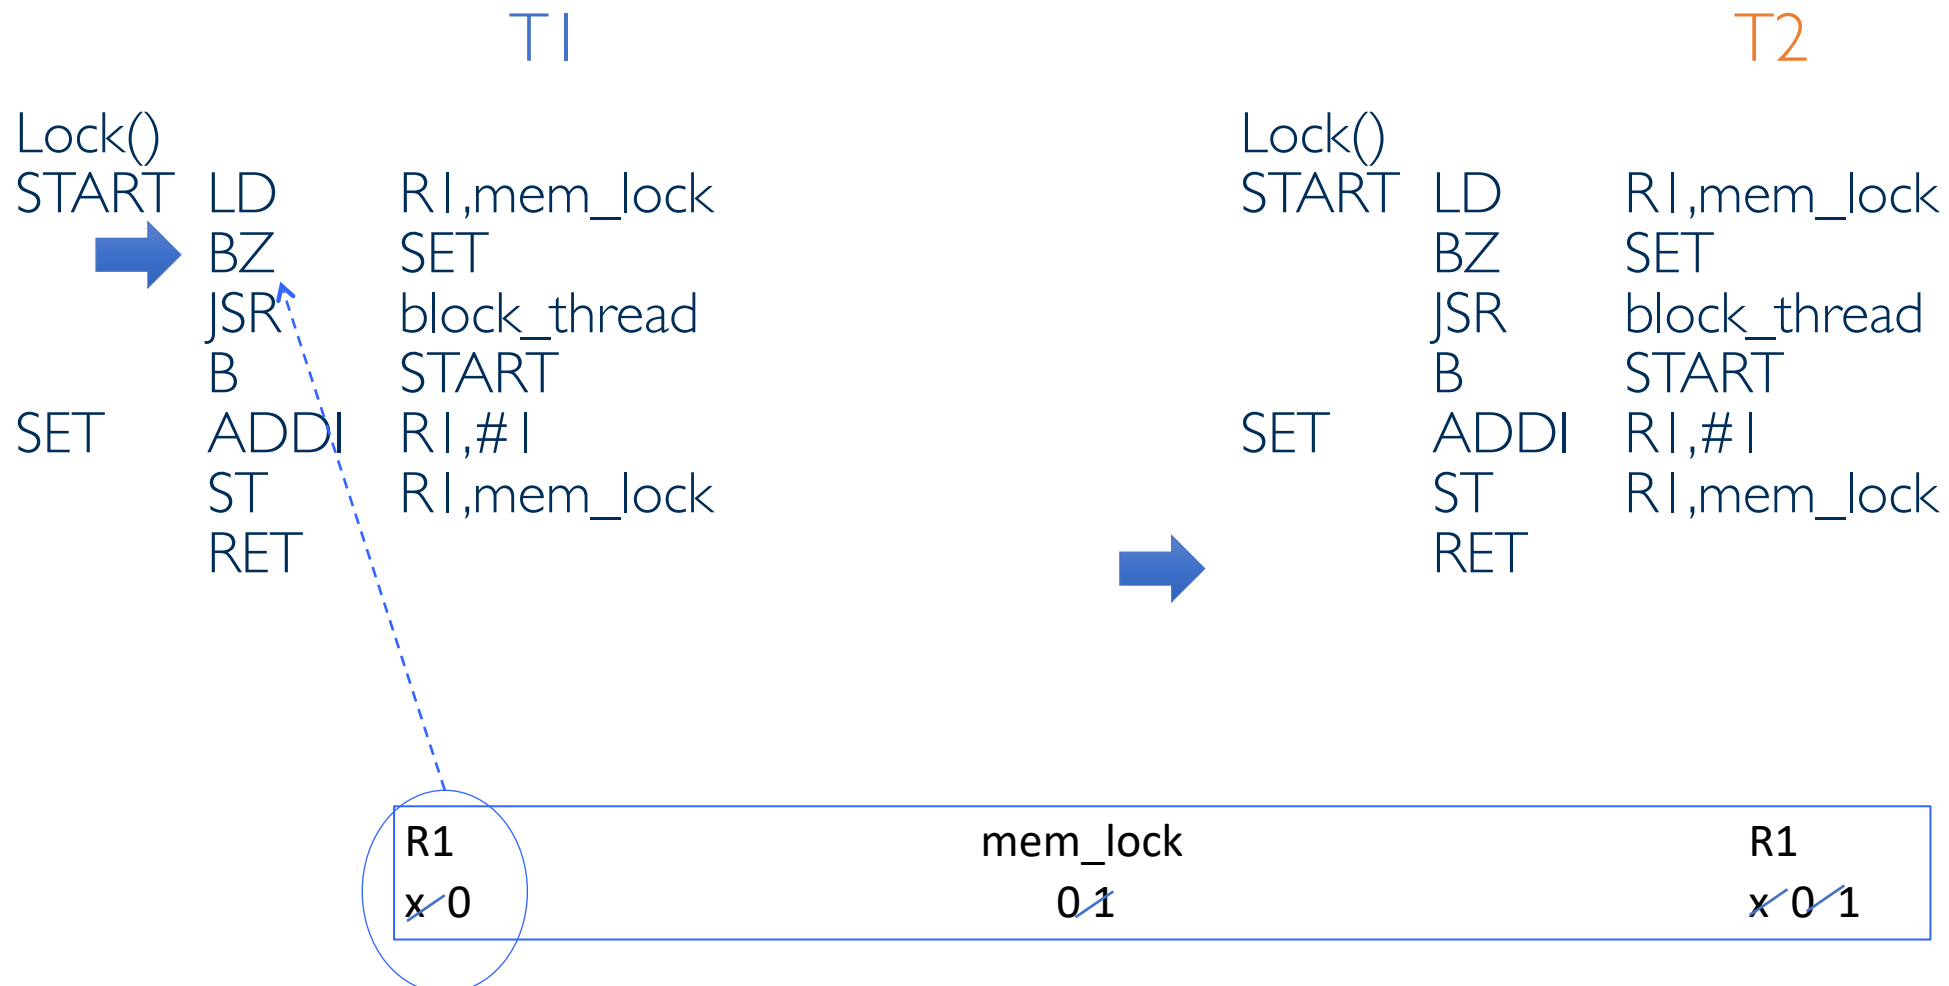


Lock() using machine instructions

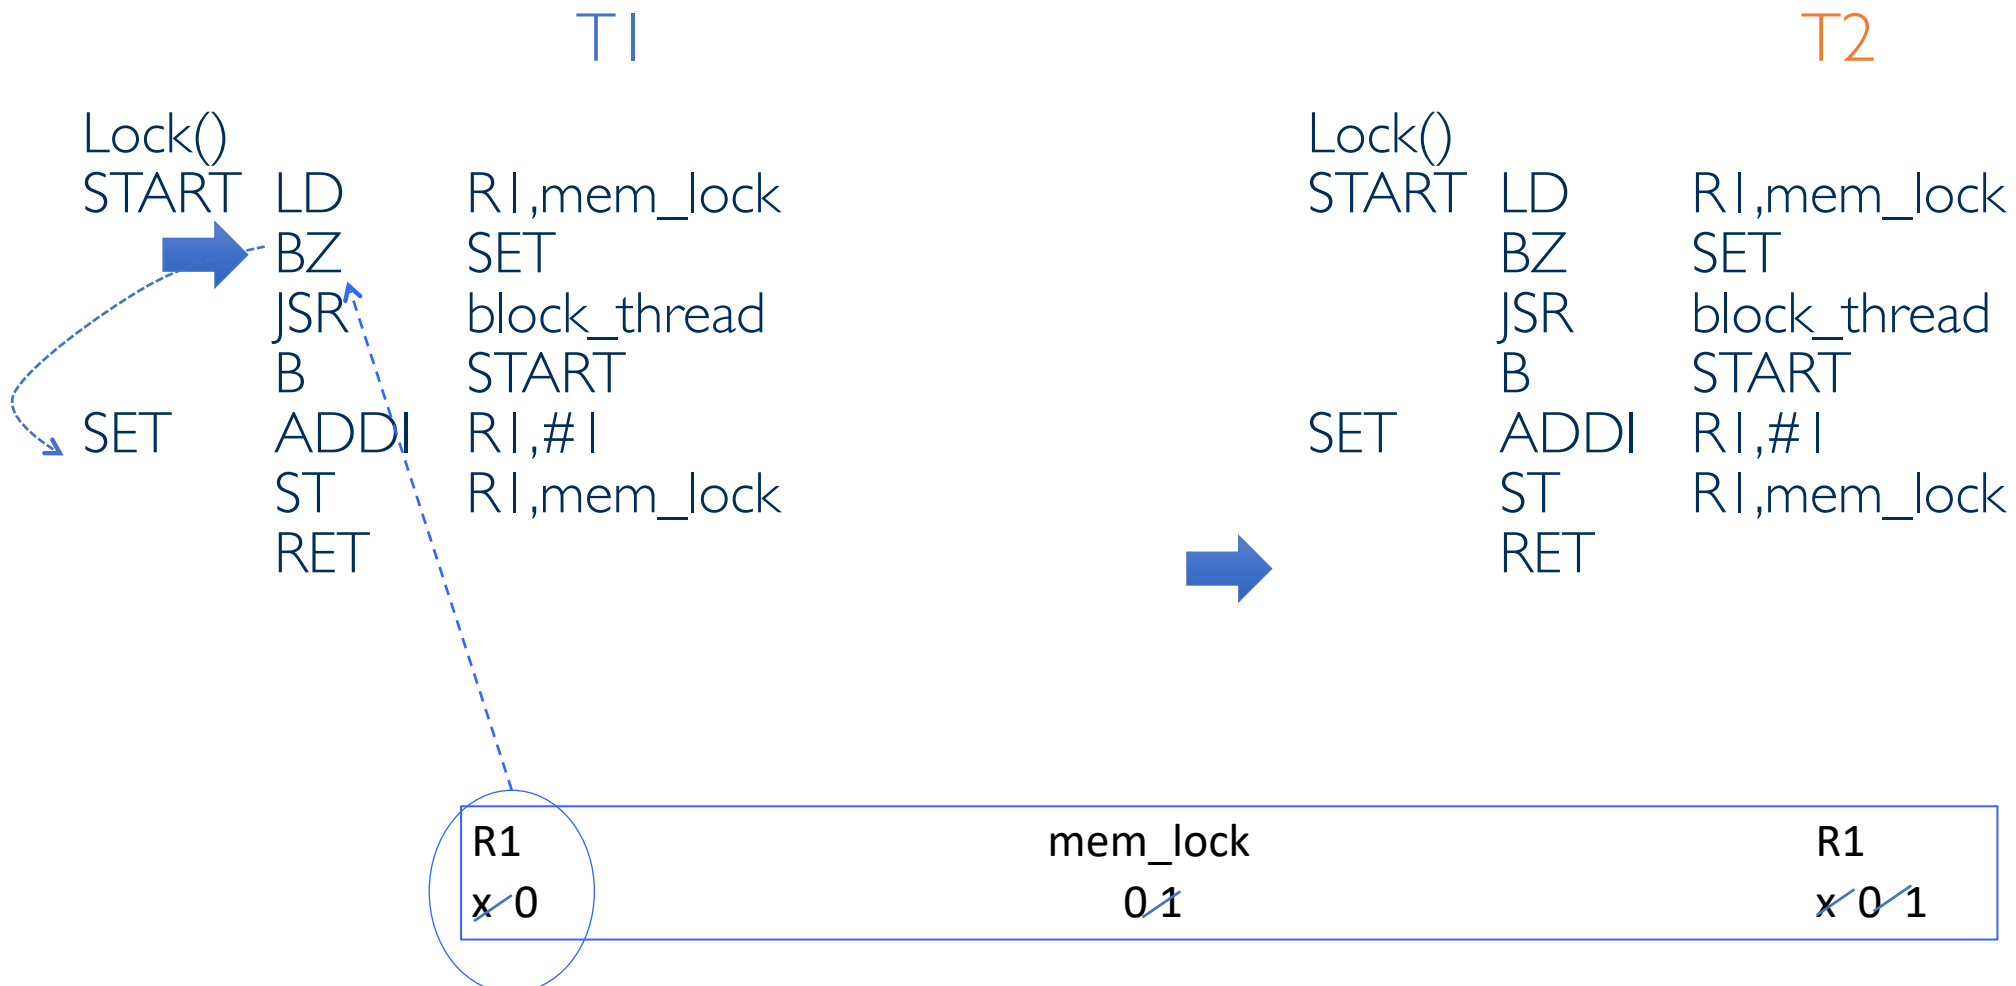


R1	mem_lock	R1
x 0	0 1	x 0 1

Lock() using machine instructions



Lock() using machine instructions



Lock() using machine instructions

T1

Lock()
START

LD R1,mem_lock
BZ SET
JSR block_thread
B START
SET ADDI R1,#1
ST R1,mem_lock
RET

T2

Lock()
START

LD R1,mem_lock
BZ SET
JSR block_thread
B START
SET ADDI R1,#1
ST R1,mem_lock
RET

R1
~~x~~ 0 1

mem_lock
~~0~~ 1

R1
~~x~~ 0 1

Lock() using machine instructions


Lock()
START

→

	T1
LD	R1,mem_lock
BZ	SET
JSR	block_thread
B	START
SET	ADDI R1,#1
	ST R1,mem_lock
	RET

Lock()
START

	T2
LD	R1,mem_lock
BZ	SET
JSR	block_thread
B	START
SET	ADDI R1,#1
	ST R1,mem_lock
	RET



Lock() using machine instructions

Lock()
START


→

	T1
LD	R1,mem_lock
BZ	SET
JSR	block_thread
B	START
SET	ADDI R1,#1
	ST R1,mem_lock
	RET

Lock()
START

→

	T2
LD	R1,mem_lock
BZ	SET
JSR	block_thread
B	START
SET	ADDI R1,#1
	ST R1,mem_lock
	RET



What happened?!?

What happened?!?

- We weren't in a critical section (We're implementing the lock for a critical section – how could we be in one??)

What happened?!?

- We weren't in a critical section (We're implementing the lock for a critical section – how could we be in one??)
- Interrupts *can* happen between instructions

What happened?!?

- We weren't in a critical section (We're implementing the lock for a critical section – how could we be in one??)
- Interrupts **can** happen between instructions

Lock():

```
while (mem_lock != 0)
    block the thread
mem_lock = 1;
```

Unlock():

```
mem_lock = 0
```

What happened?!?

- We weren't in a critical section (We're implementing the lock for a critical section – how could we be in one??)
- Interrupts **can** happen between instructions

Lock():

```
while (mem_lock != 0)
    block the thread
mem_lock = 1;
```

Unlock():

```
mem_lock = 0
```

- We need the test and assignment to be atomic/indivisible!

Hardware to the rescue!

Hardware to the rescue!

- We need an atomic Read-Modify-Write instruction

Hardware to the rescue!

- We need an atomic Read-Modify-Write instruction
- TEST-AND-SET <memory-location>
 - load current value in <memory-location>
 - store 1 in <memory-location>

Hardware to the rescue!

- We need an atomic Read-Modify-Write instruction
- TEST-AND-SET <memory-location>
 - load current value in <memory-location>
 - store 1 in <memory-location>
- Atomically:
 - Test L and set it to 1
 - If L tested originally as 0, we've claimed the lock
 - If L tested as 1, we need to try again

Hardware to the rescue!

- We need an atomic Read-Modify-Write instruction
- TEST-AND-SET <memory-location>
 load current value in <memory-location>
 store 1 in <memory-location>
- Atomically:
 - Test L and set it to 1
 - If L tested originally as 0, we've claimed the lock
 - If L tested as 1, we need to try again
- Work-alikes
 - Compare-and-swap (IBM 370)
 - Fetch-and-add (Intel x86)
 - Load-linked/store-conditional (MIPS, ARM, ...)

Our new implementation of mutex

Our new implementation of mutex

- Lock():
 while (test-and-set (&mem_lock))
 block the thread

Unlock():
 mem_lock = 0

Example: Implementing Mutex Lock

```
static int shared-lock = 0; /* global variable to
                             both T1 and T2 */
/* shared procedure for T1 and T2 */
int binary-semaphore(int *L)
{
    int X;

    X = test-and-set(L);

    /* X = 0 for successful return */
    return(X);
}
```

Example: Implementing Mutex Lock

```
static int shared-lock = 0; /* global variable to
                             both T1 and T2 */
/* shared procedure for T1 and T2 */
int binary-semaphore(int *L)
{
    int X;

    X = test-and-set(L);

    /* X = 0 for successful return */
    return(X);
}
```

Two threads T1 and T2 execute the following statement simultaneously:
 MyX = binary_semaphore(&shared-lock);
where MyX is a local variable in each of T1 and T2.

Example: Implementing Mutex Lock

```
static int shared-lock = 0; /* global variable to
                             both T1 and T2 */
/* shared procedure for T1 and T2 */
int binary-semaphore(int *L)
{
    int X;

    X = test-and-set(L);

    /* X = 0 for successful return */
    return(X);
}
```

Two threads T1 and T2 execute the following statement simultaneously:
MyX = binary_semaphore(&shared-lock);
where MyX is a local variable in each of T1 and T2.

What are the possible values returned to T1 and T2?

Example: Implementing Mutex Lock

```
static int shared-lock = 0; /* global variable to
                             both T1 and T2 */
/* shared procedure for T1 and T2 */
int binary-semaphore(int *L)
{
    int X;


    X = test-and-set(L);

    /* X = 0 for successful return */
    return(X);
}
```


Two threads T1 and T2 execute the following statement simultaneously:
MyX = binary_semaphore(&shared-lock);
where MyX is a local variable in each of T1 and T2.

What are the possible values returned to T1 and T2?
Getting 0 0 isn't possible!

How do we implement Symmetric Multi-Processing (SMP)?




How do we implement Symmetric Multi-Processing (SMP)?



The System (hardware+OS) has to ensure 3 things:


How do we implement Symmetric Multi-Processing (SMP)?



The System (hardware+OS) has to ensure 3 things:

1. Threads of the same process share the same PT


How do we implement Symmetric Multi-Processing (SMP)?



The System (hardware+OS) has to ensure 3 things:

1. Threads of the same process share the same PT
2. Threads have synchronization atomicity

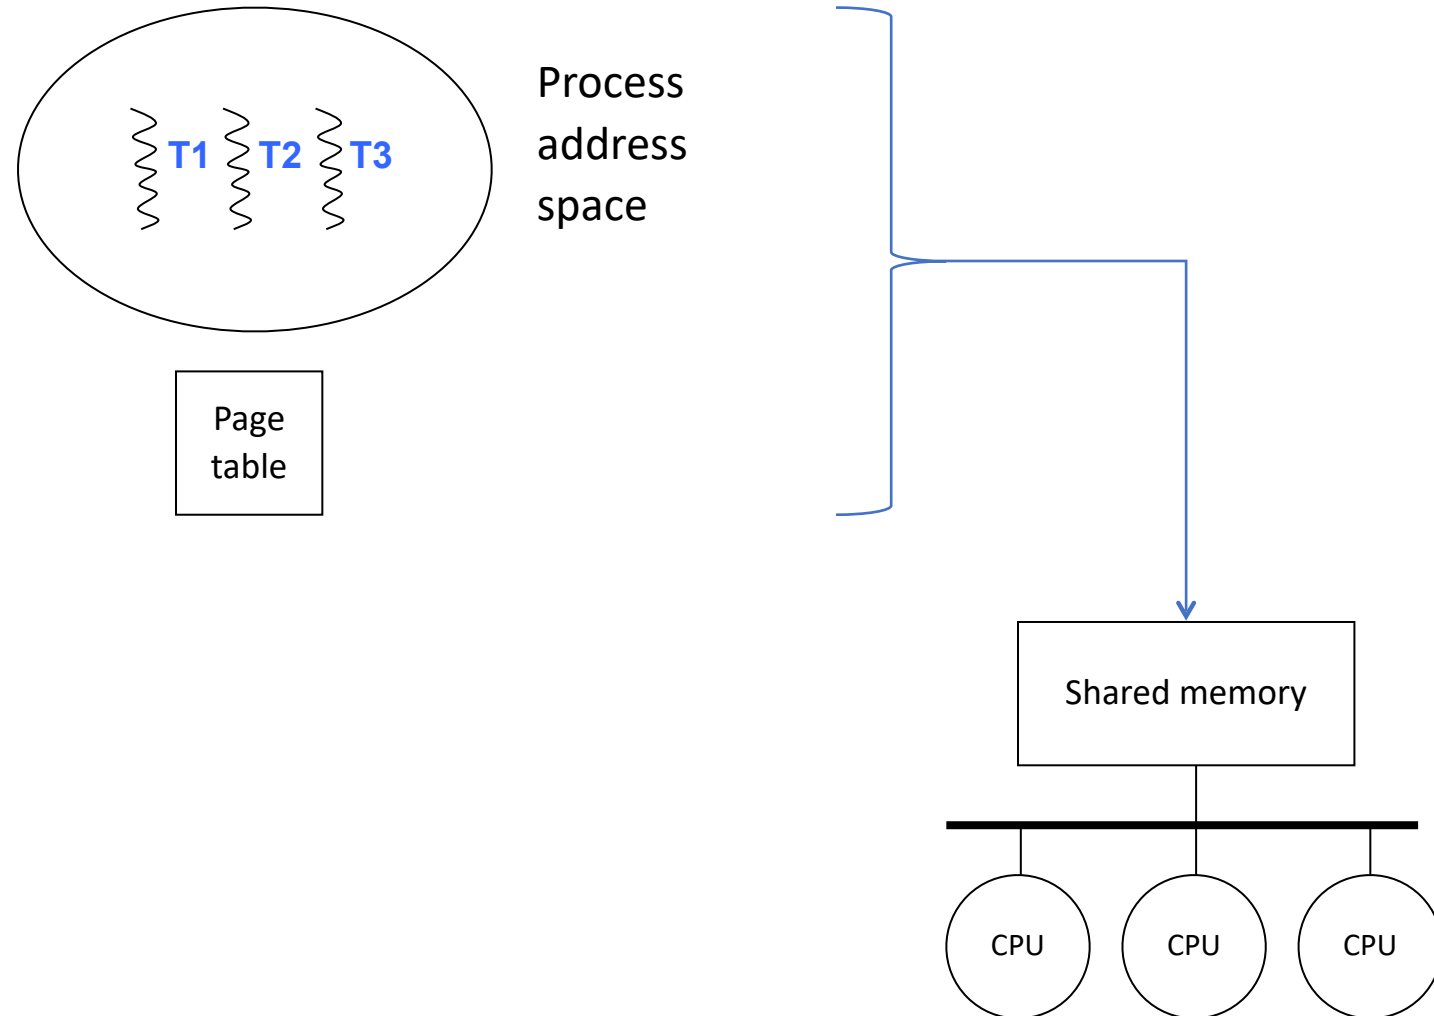
How do we implement Symmetric Multi-Processing (SMP)?



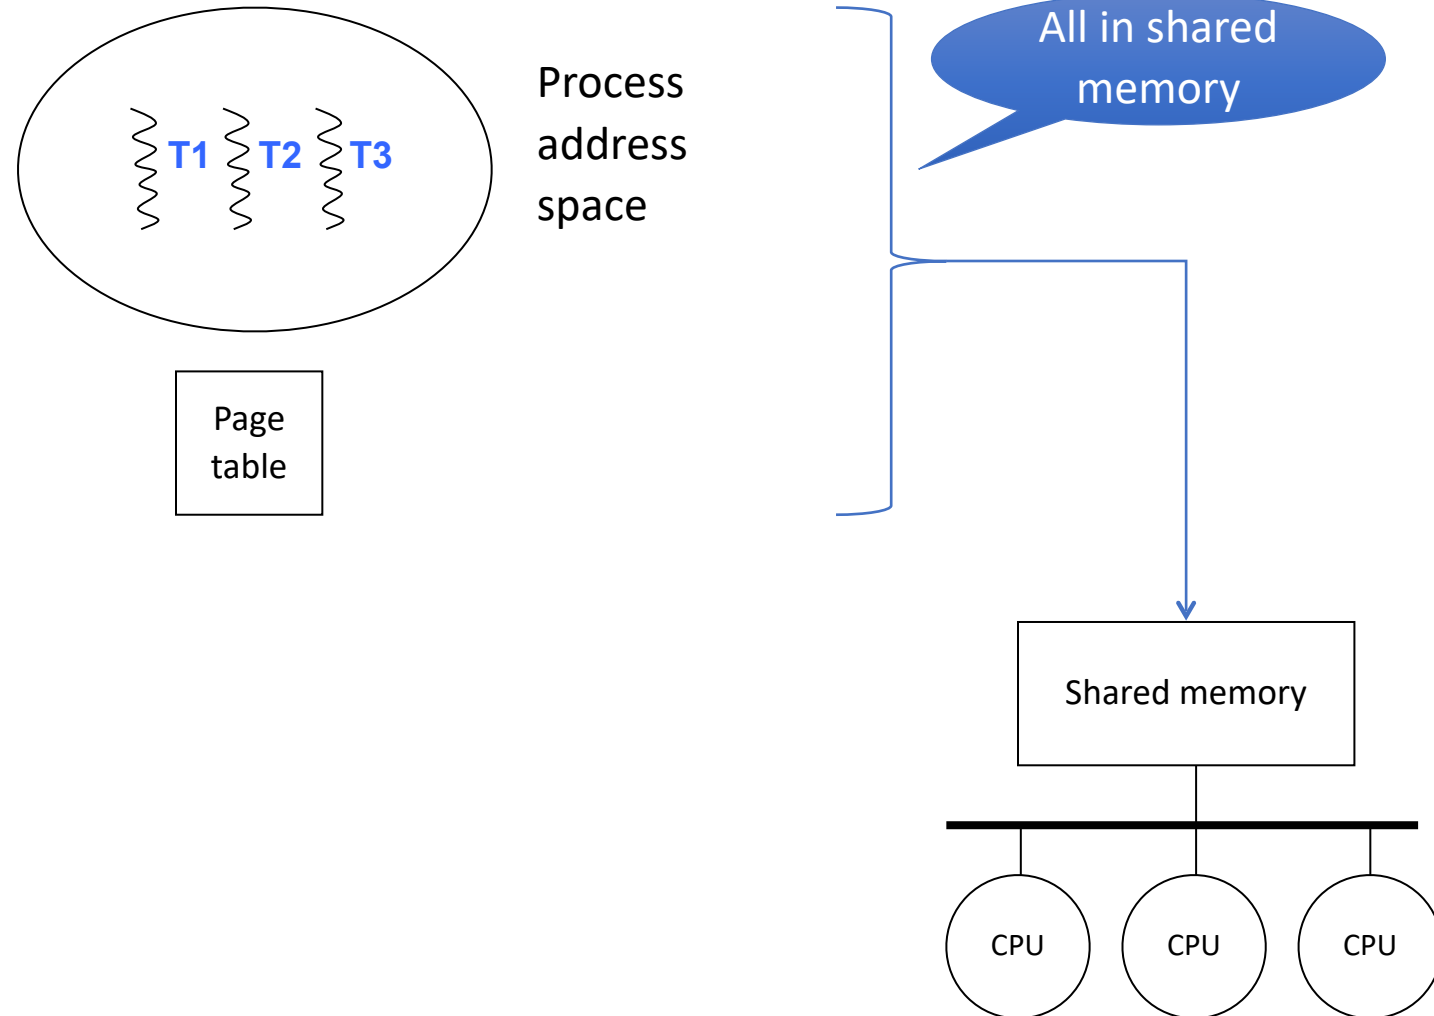
The System (hardware+OS) has to ensure 3 things:

1. Threads of the same process share the same PT
2. Threads have synchronization atomicity
3. Threads have identical views of memory

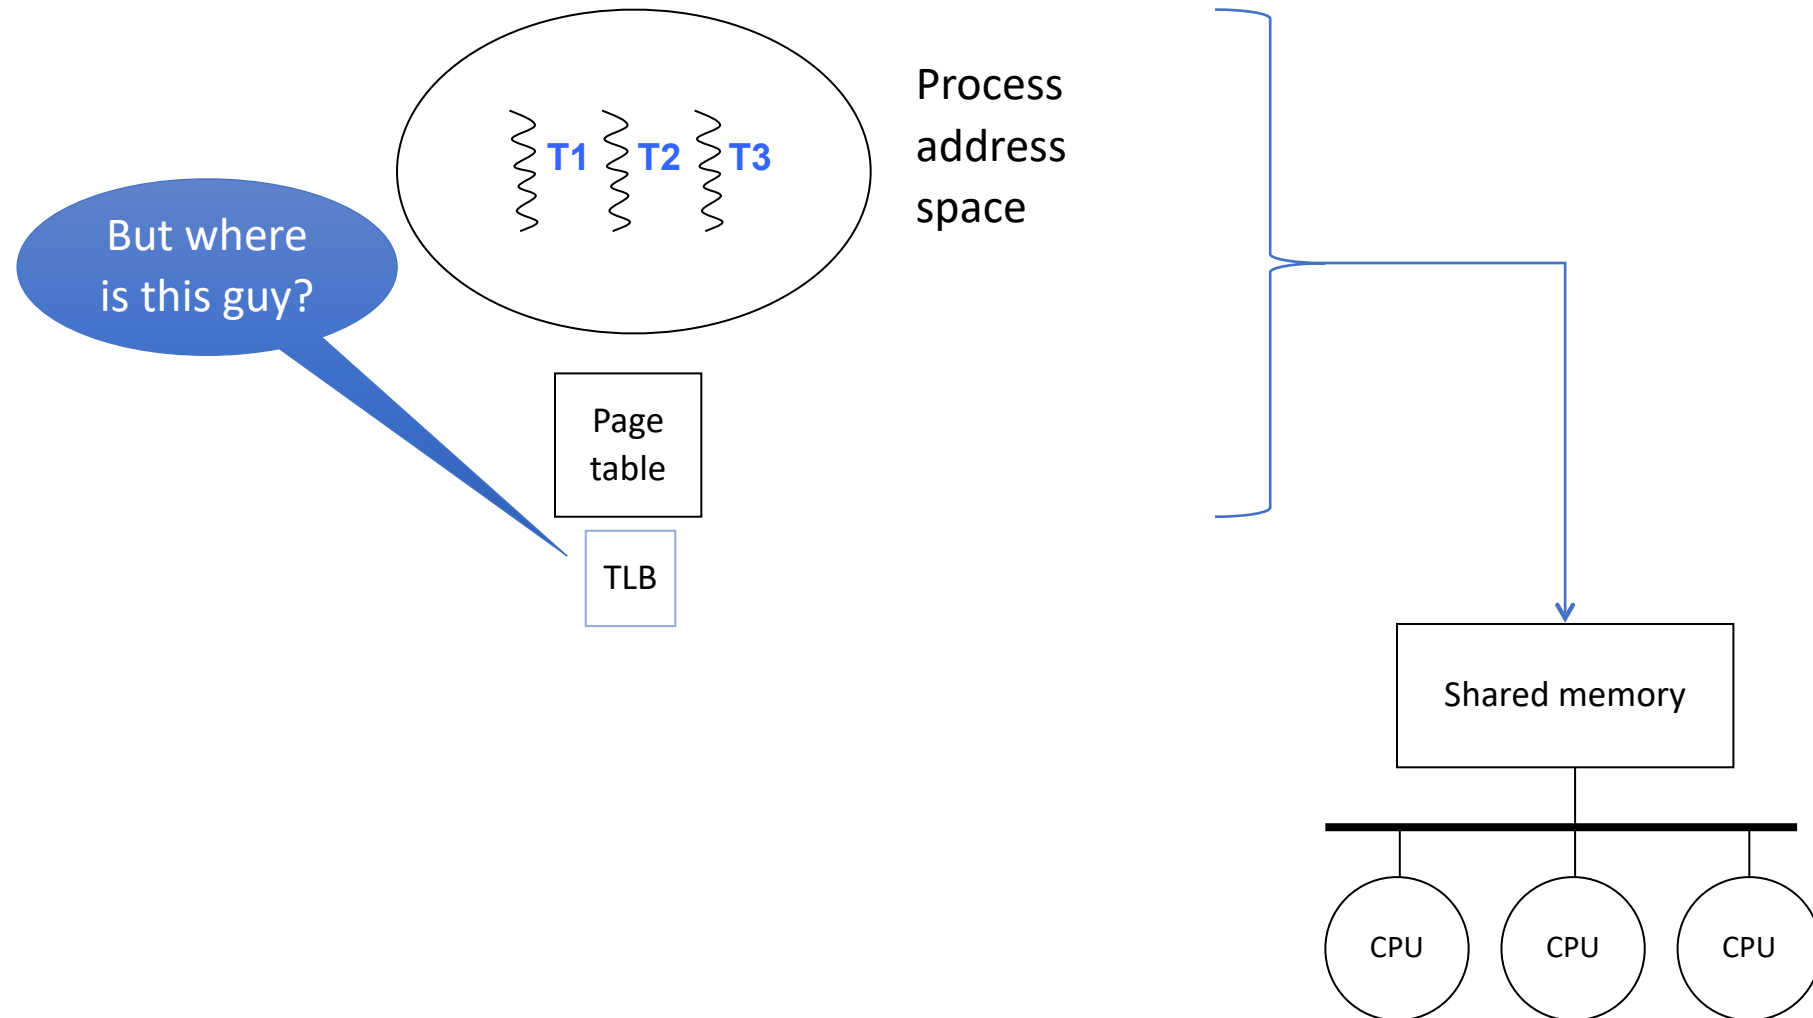
I) All threads share the same page table



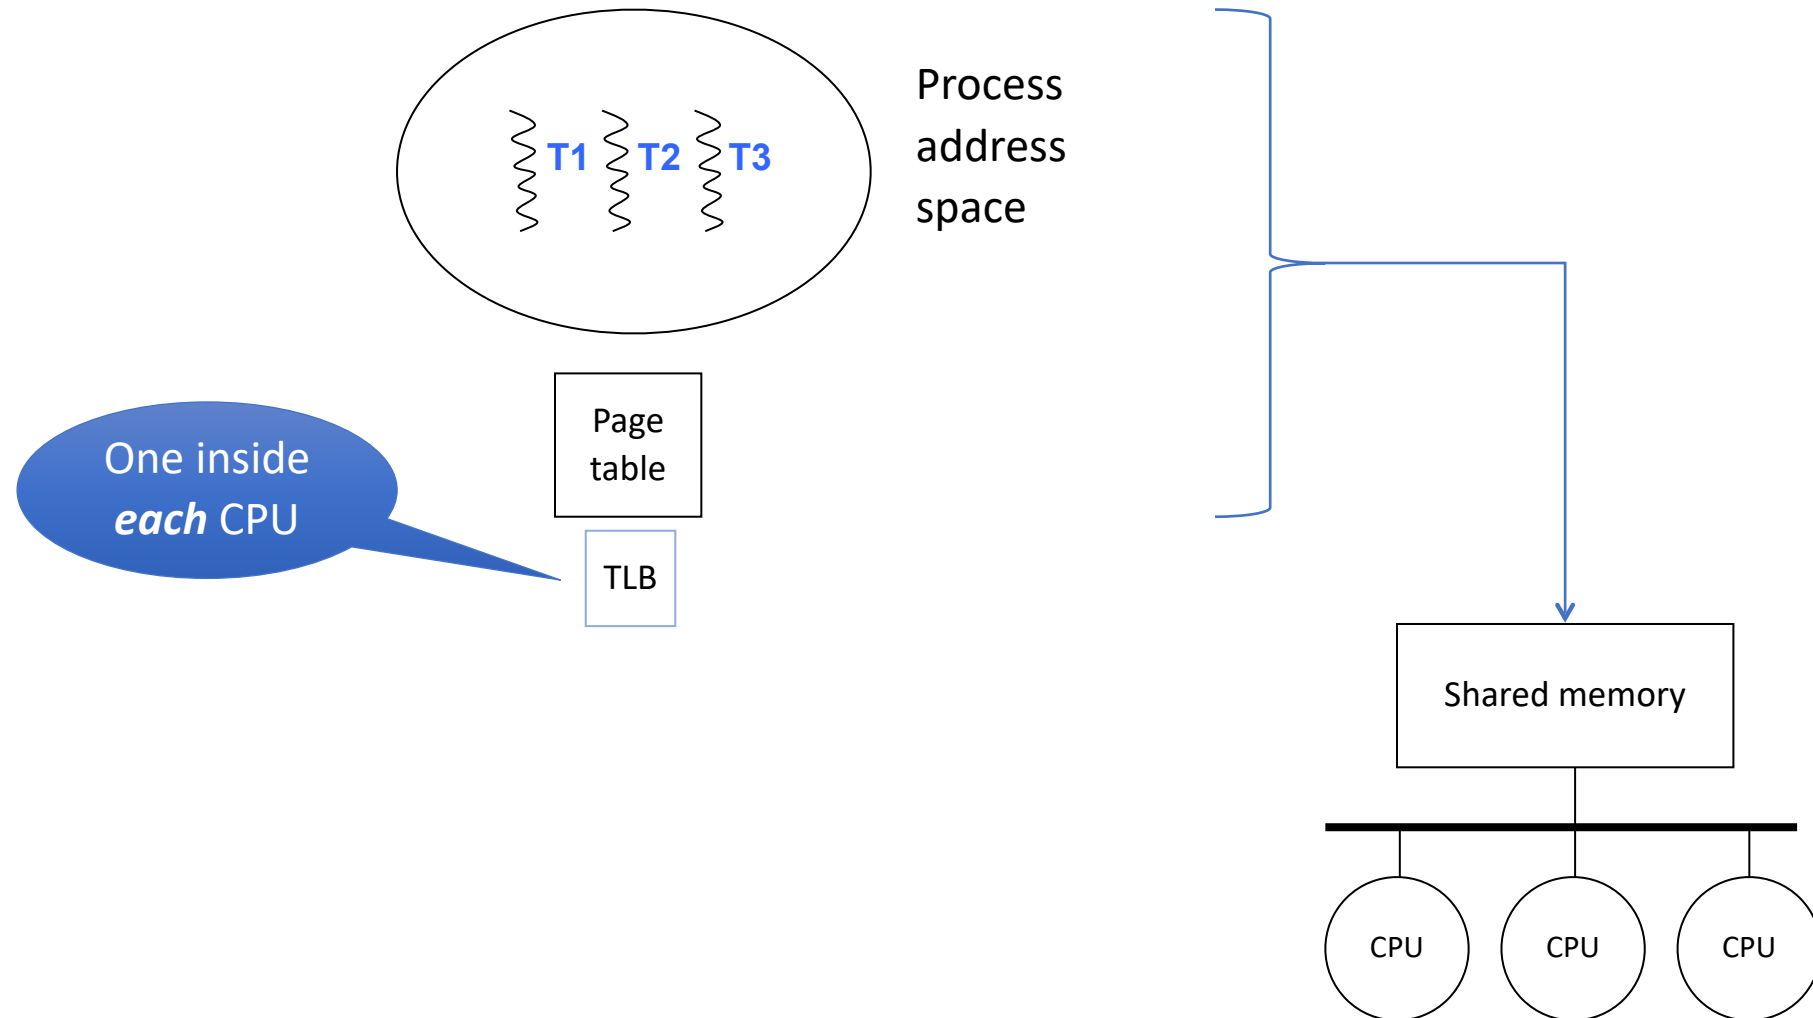
I) All threads share the same page table



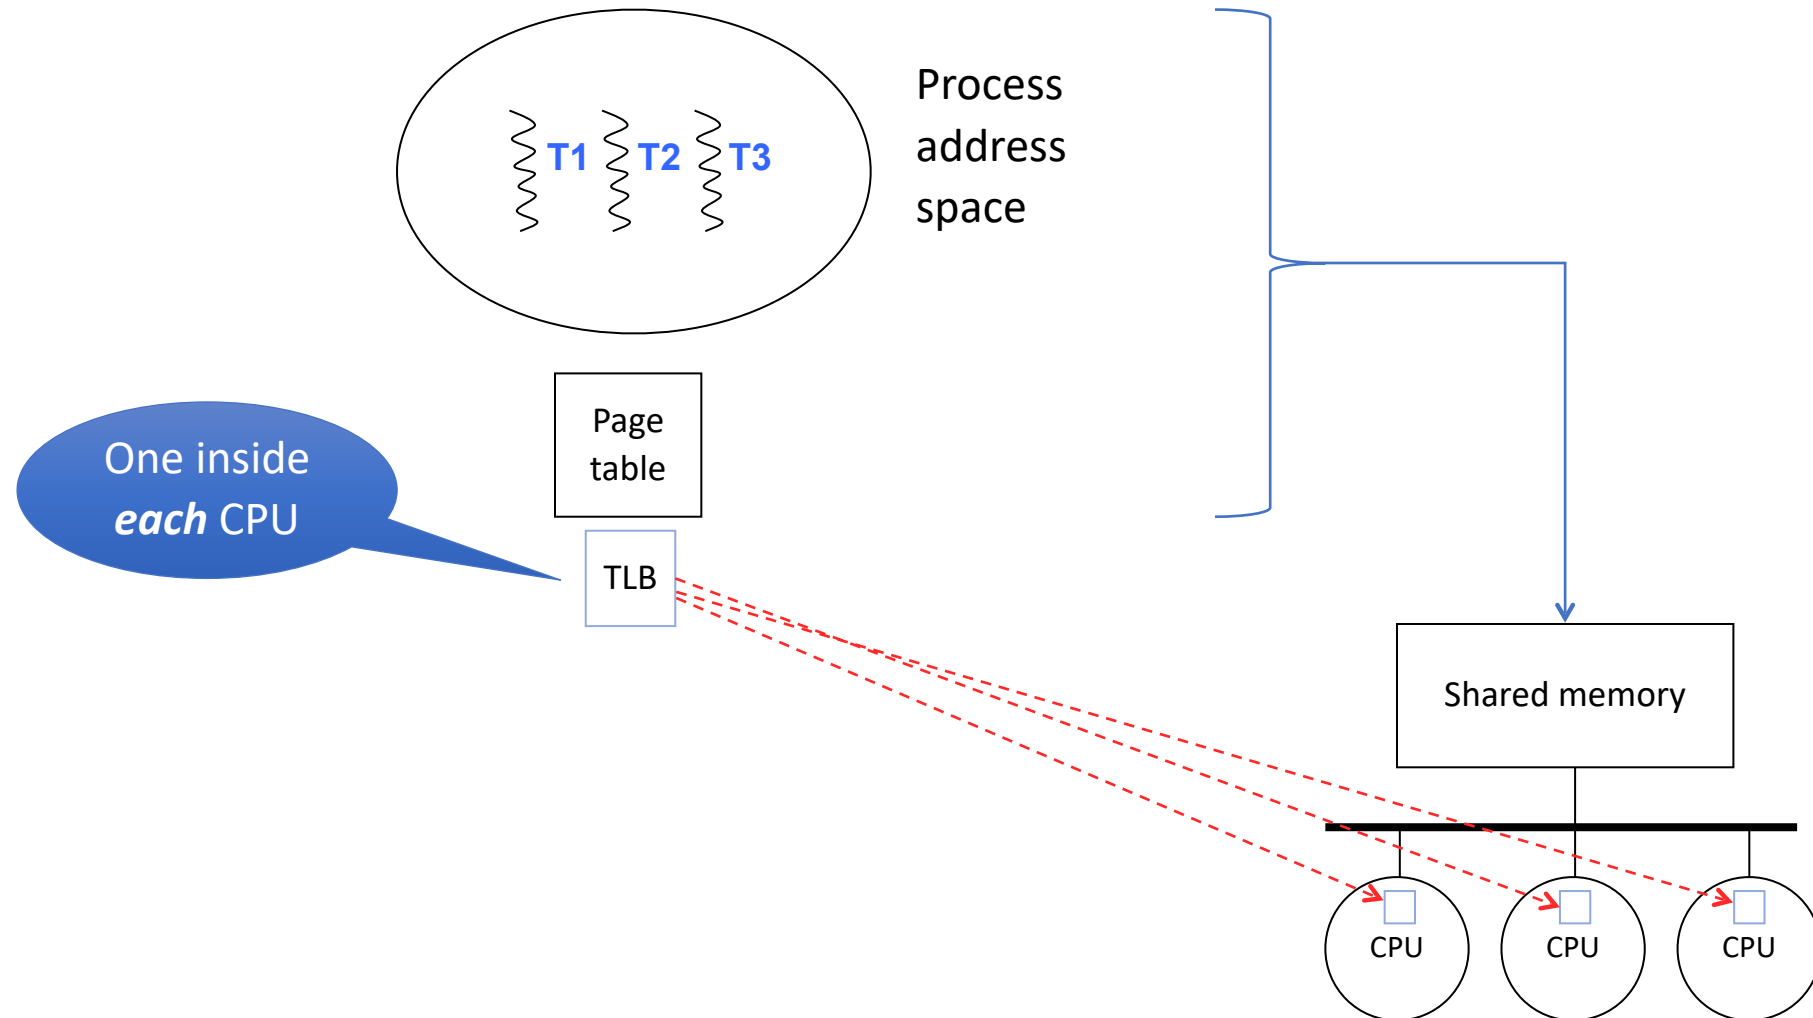
I) All threads share the same page table



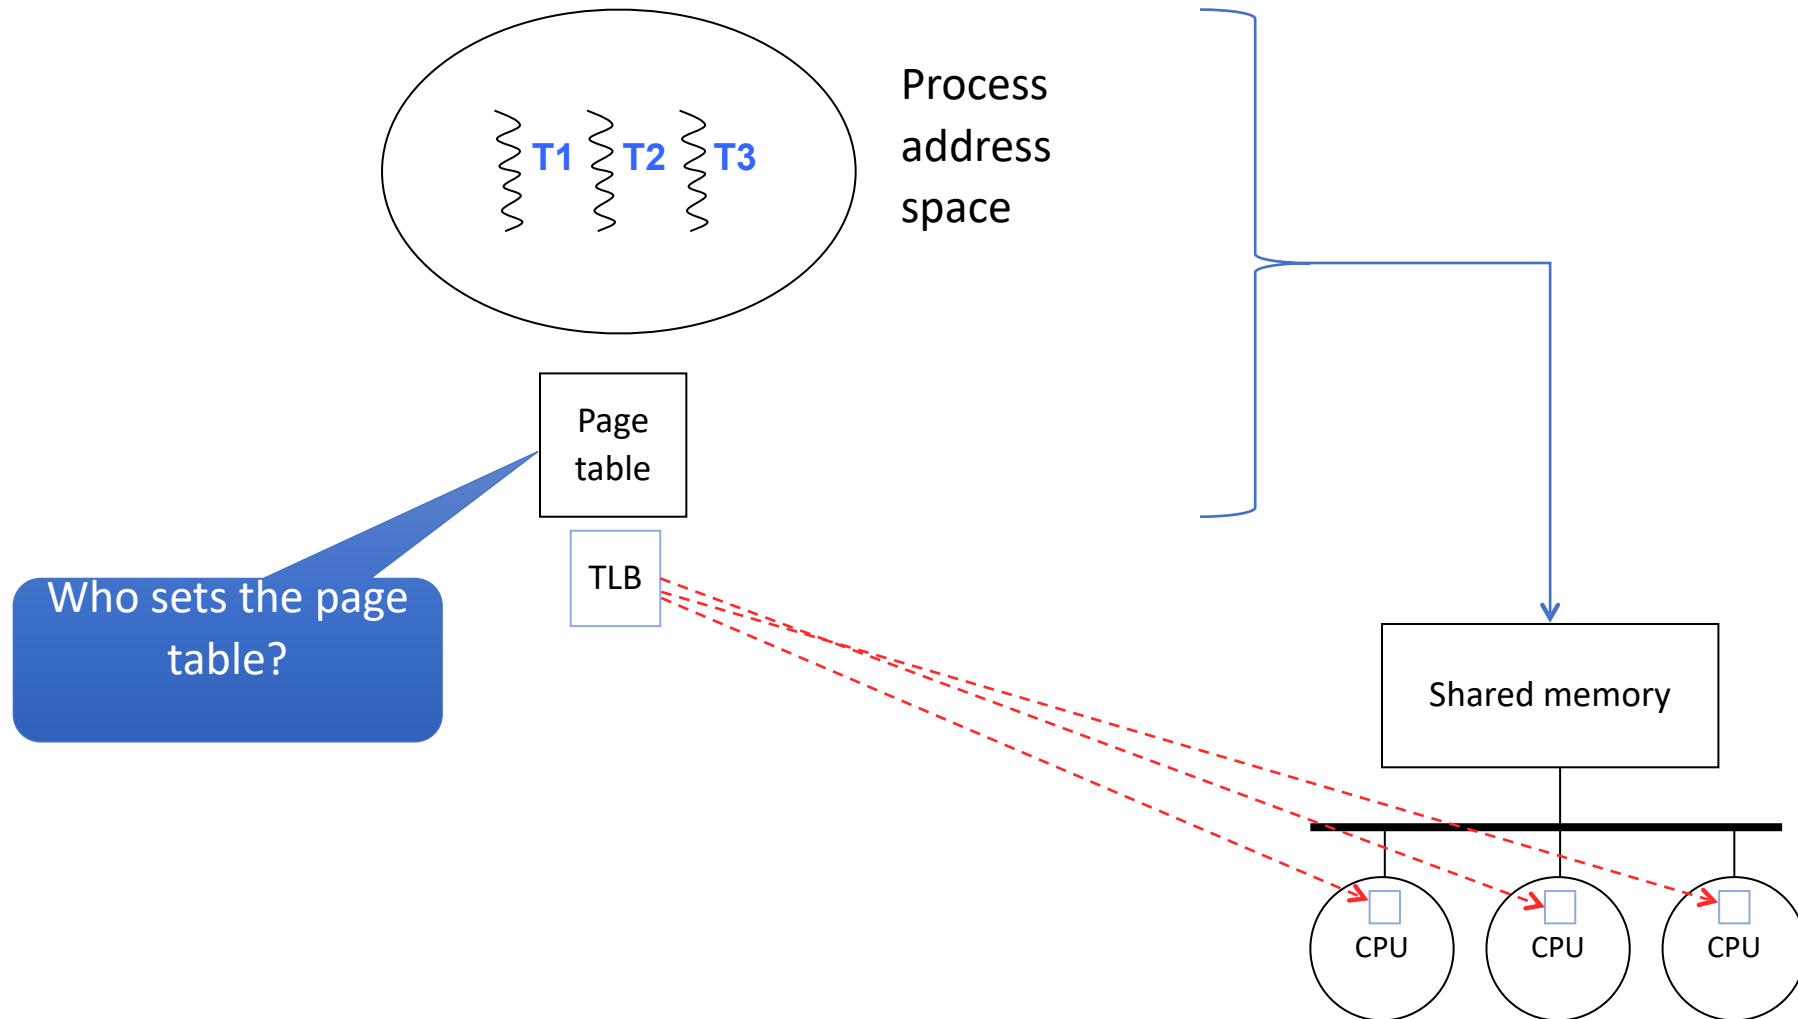
I) All threads share the same page table



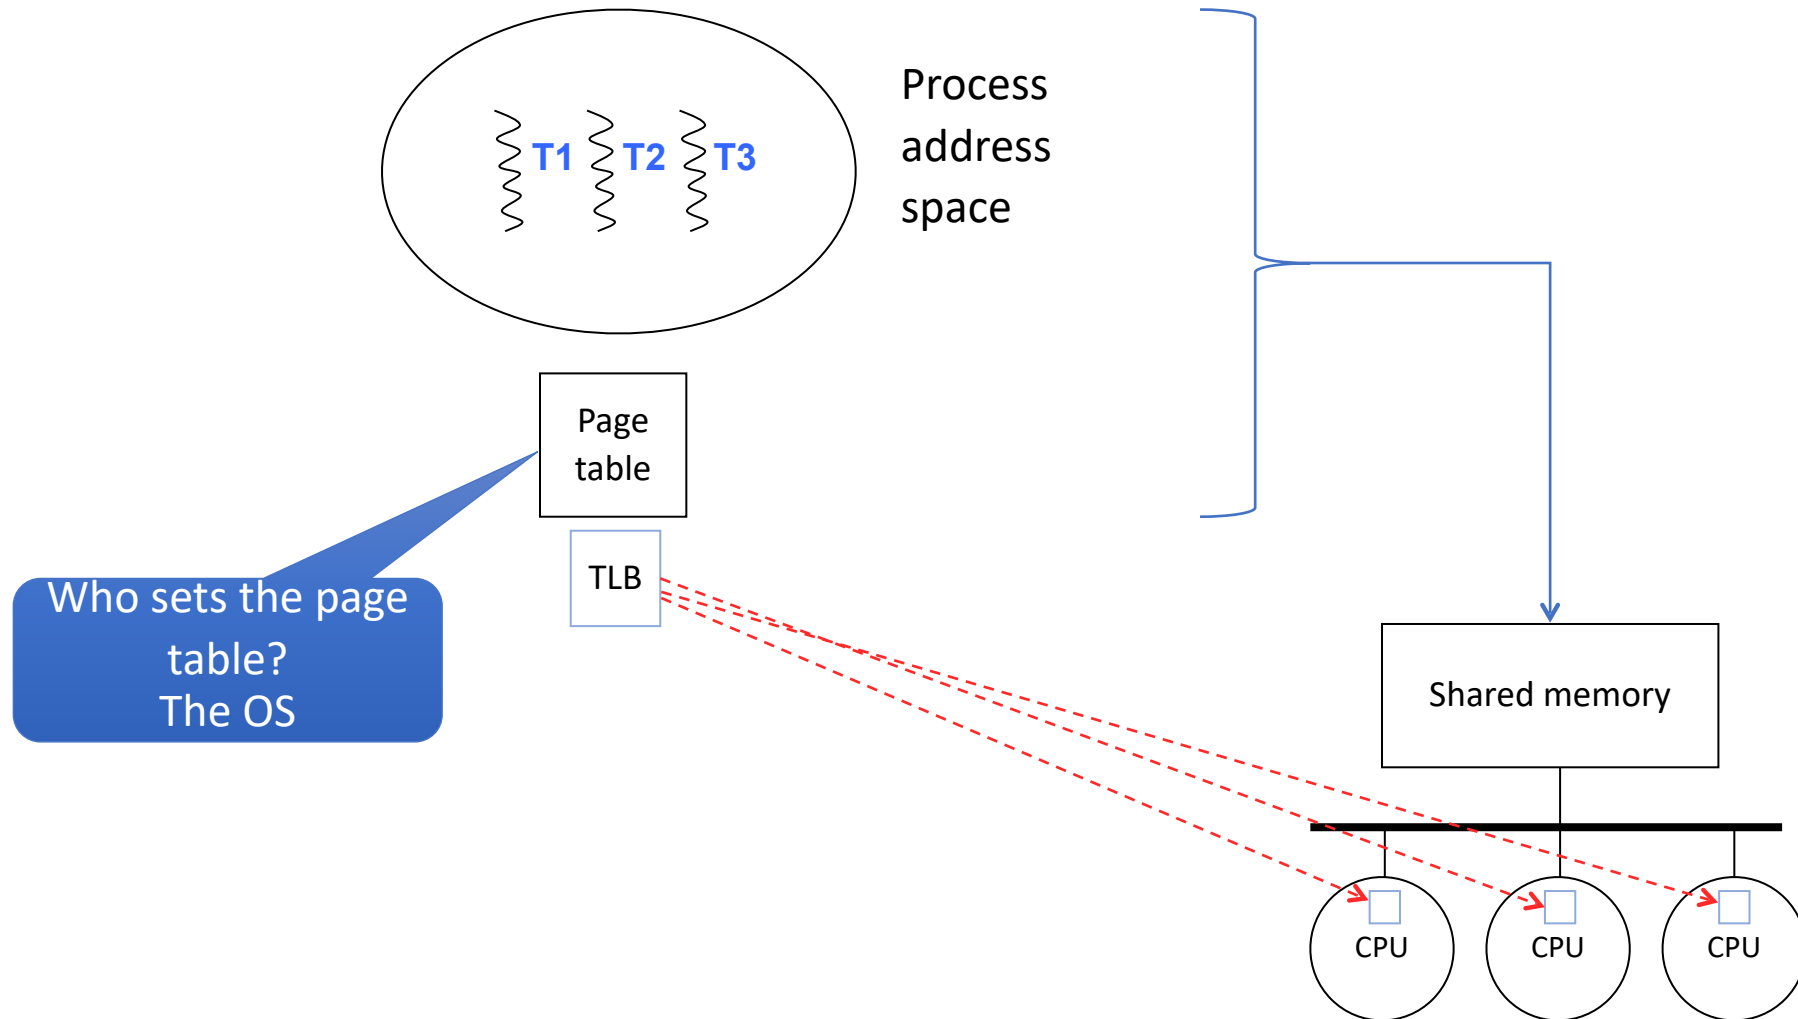
1) All threads share the same page table



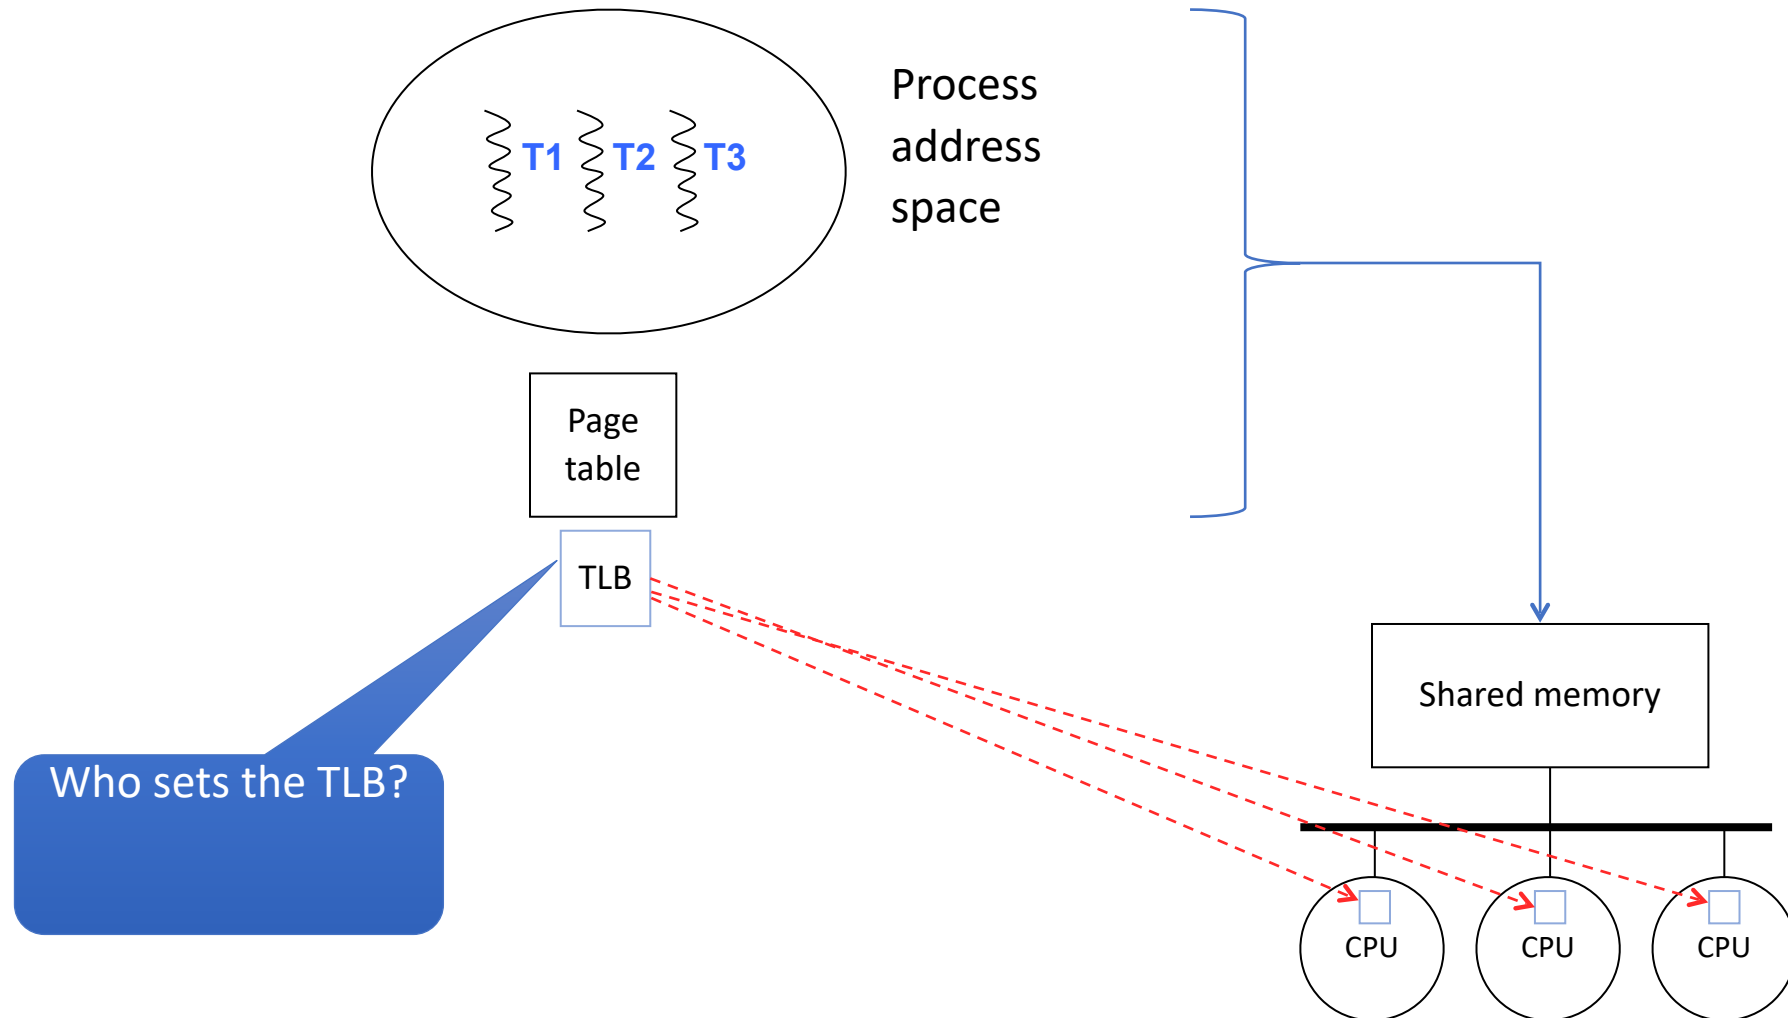
I) All threads share the same page table



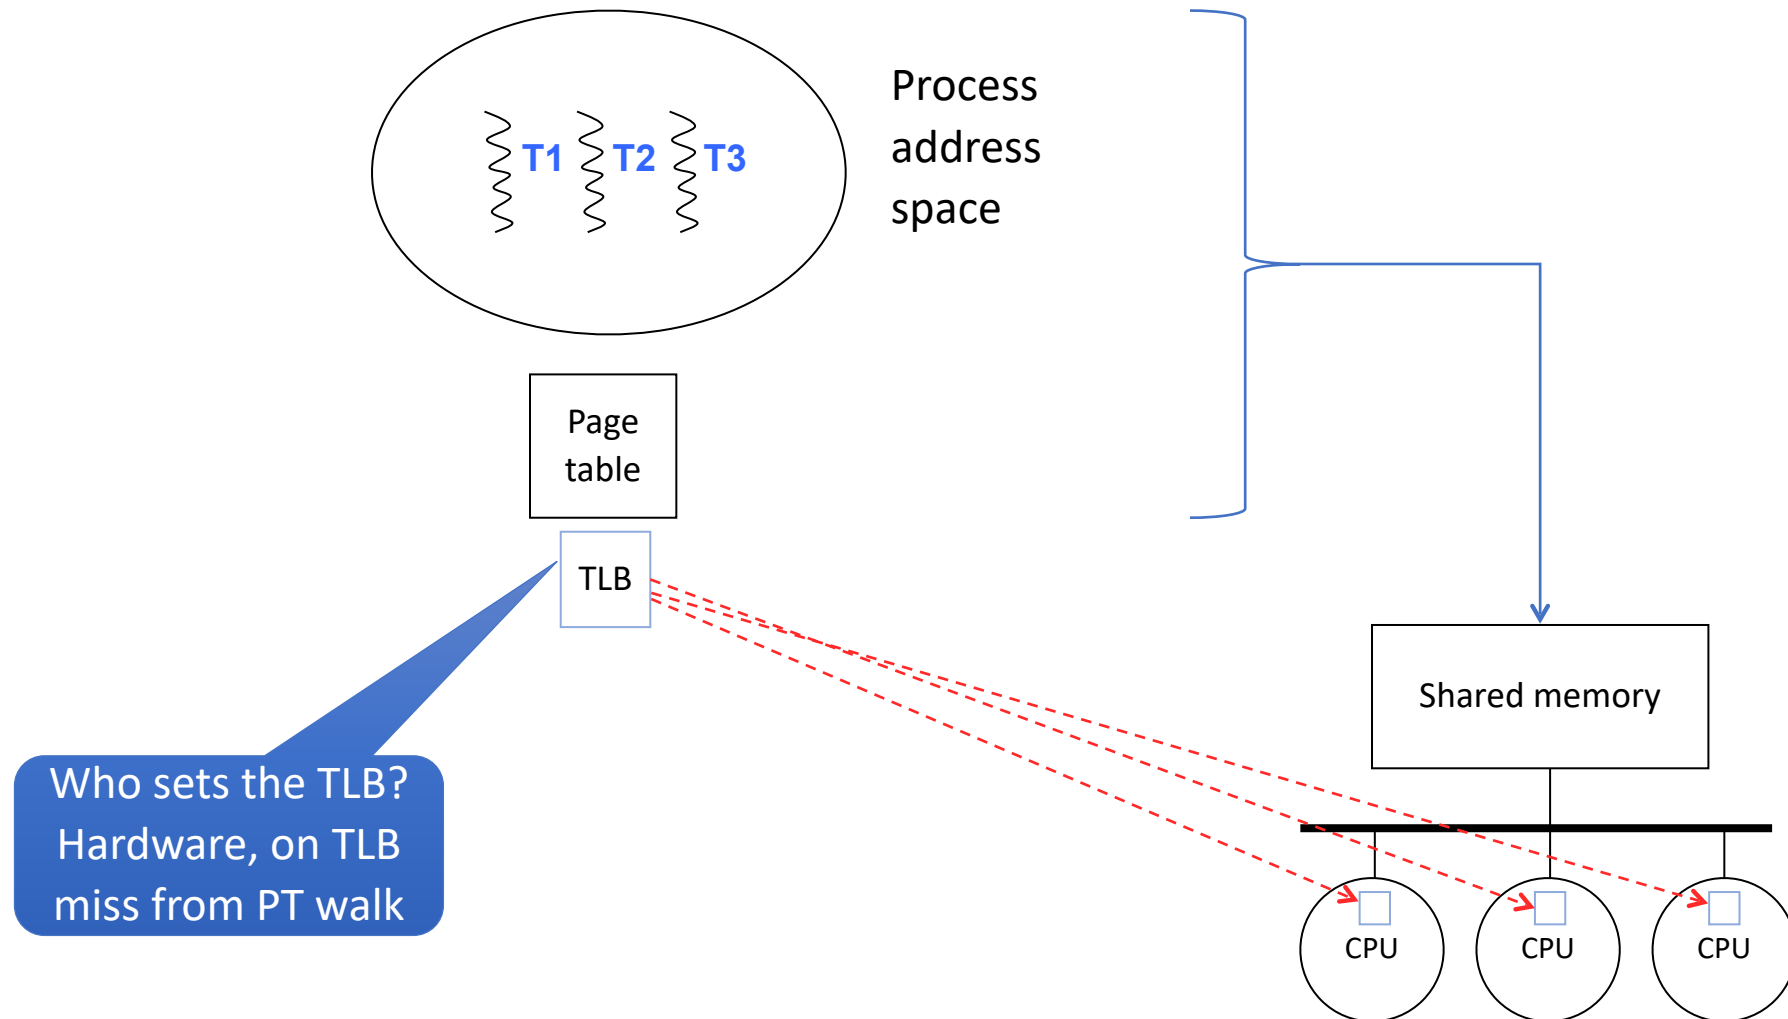
I) All threads share the same page table



1) All threads share the same page table



1) All threads share the same page table



SMP context switch handling

- As in the single-CPU case, the TLB must be flushed of user-space addresses on context switch

SMP context switch handling

- As in the single-CPU case, the TLB must be flushed of user-space addresses on context switch
- How do multiple processors complicate this?

SMP context switch handling

- As in the single-CPU case, the TLB must be flushed of user-space addresses on context switch
- How do multiple processors complicate this?
 - Basically, they don't
 - Any time a CPU is switched to a new thread, the OS flushes user entries from that CPU's TLB
 - There's no need to affect other TLBs

SMP page replacement handling

- On page replacement by the OS (which can happen on any of the CPUs)

SMP page replacement handling

- On page replacement by the OS (which can happen on any of the CPUs)
- OS must
 - Evict the TLB entry for that page
(must happen even on a uniprocessor)

SMP page replacement handling

- On page replacement by the OS (which can happen on any of the CPUs)
- OS must
 - Evict the TLB entry for that page
(must happen even on a uniprocessor)
 - Tell the OS on other CPUs to evict the corresponding TLB entry (if present) using software interrupts

SMP page replacement handling

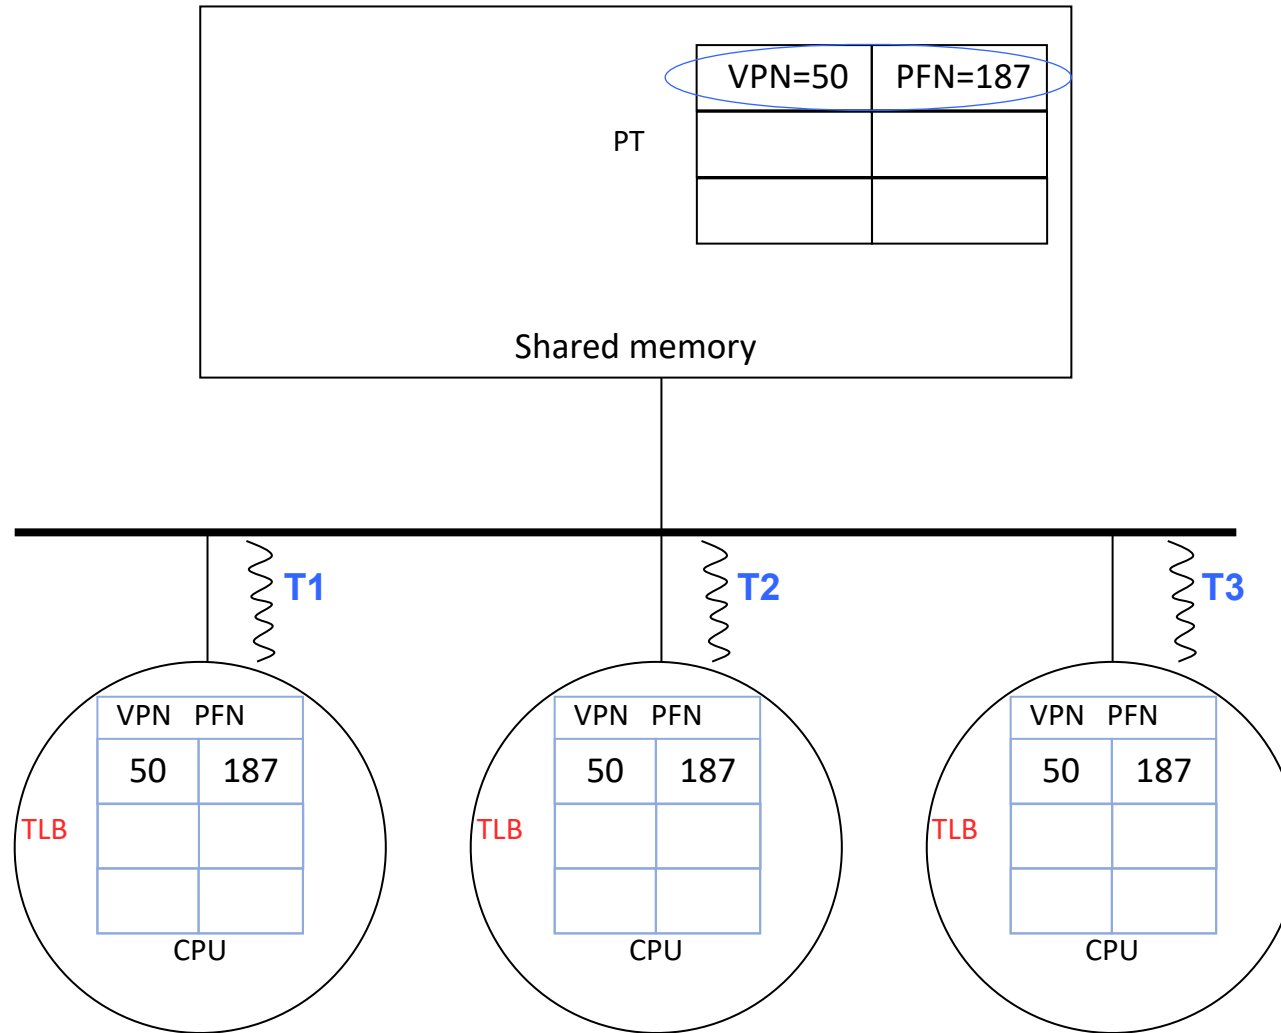
- On page replacement by the OS (which can happen on any of the CPUs)
- OS must
 - Evict the TLB entry for that page
(must happen even on a uniprocessor)
 - Tell the OS on other CPUs to evict the corresponding TLB entry (if present) using software interrupts
 - This is called **TLB Shutdown** and is expensive but necessary

SMP page replacement handling

- On page replacement by the OS (which can happen on any of the CPUs)
- OS must
 - Evict the TLB entry for that page
(must happen even on a uniprocessor)
 - Tell the OS on other CPUs to evict the corresponding TLB entry (if present) using software interrupts
 - This is called **TLB Shutdown** and is expensive but necessary
- All of this happens in software by the OS
 - ➔ Another partnership of hardware and software

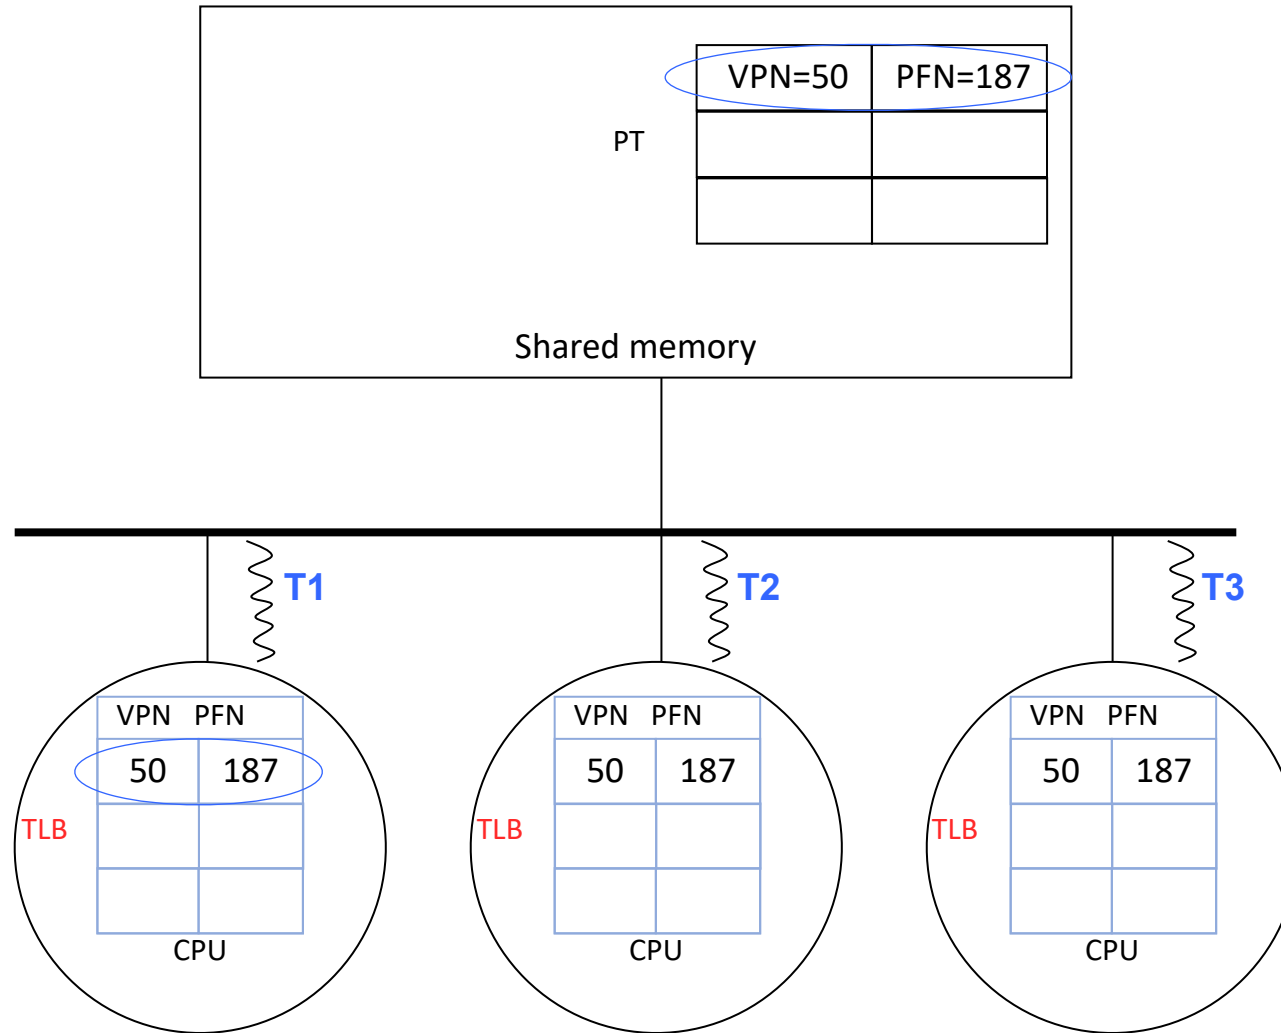
Example: Handling an SMP page fault

Note that the TLBs have each pulled in VPN=50 because each of the threads referenced that page



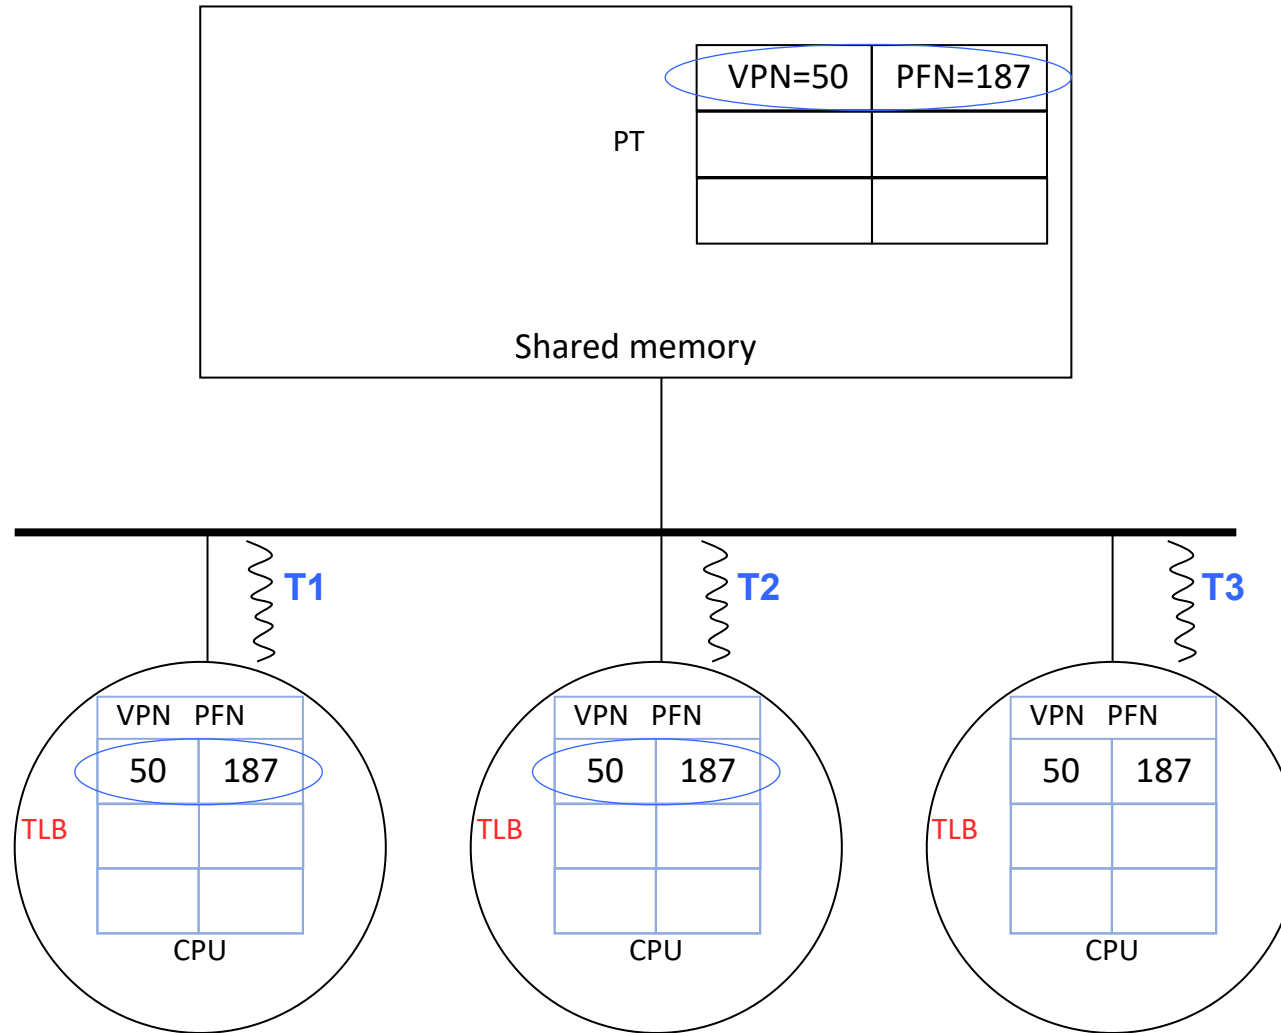
Example: Handling an SMP page fault

Note that the TLBs have each pulled in VPN=50 because each of the threads referenced that page



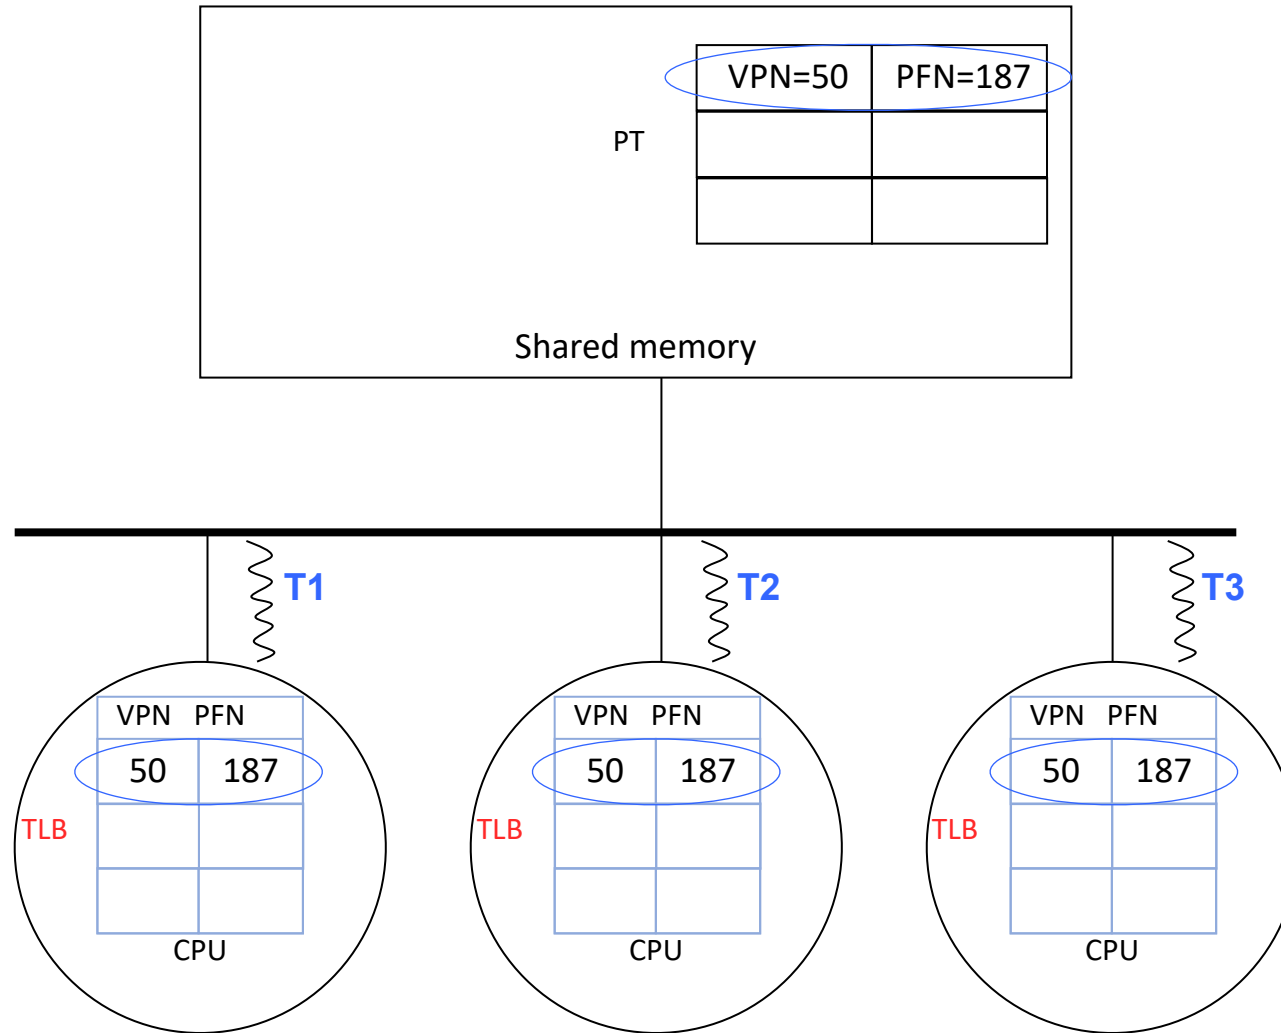
Example: Handling an SMP page fault

Note that the TLBs have each pulled in VPN=50 because each of the threads referenced that page



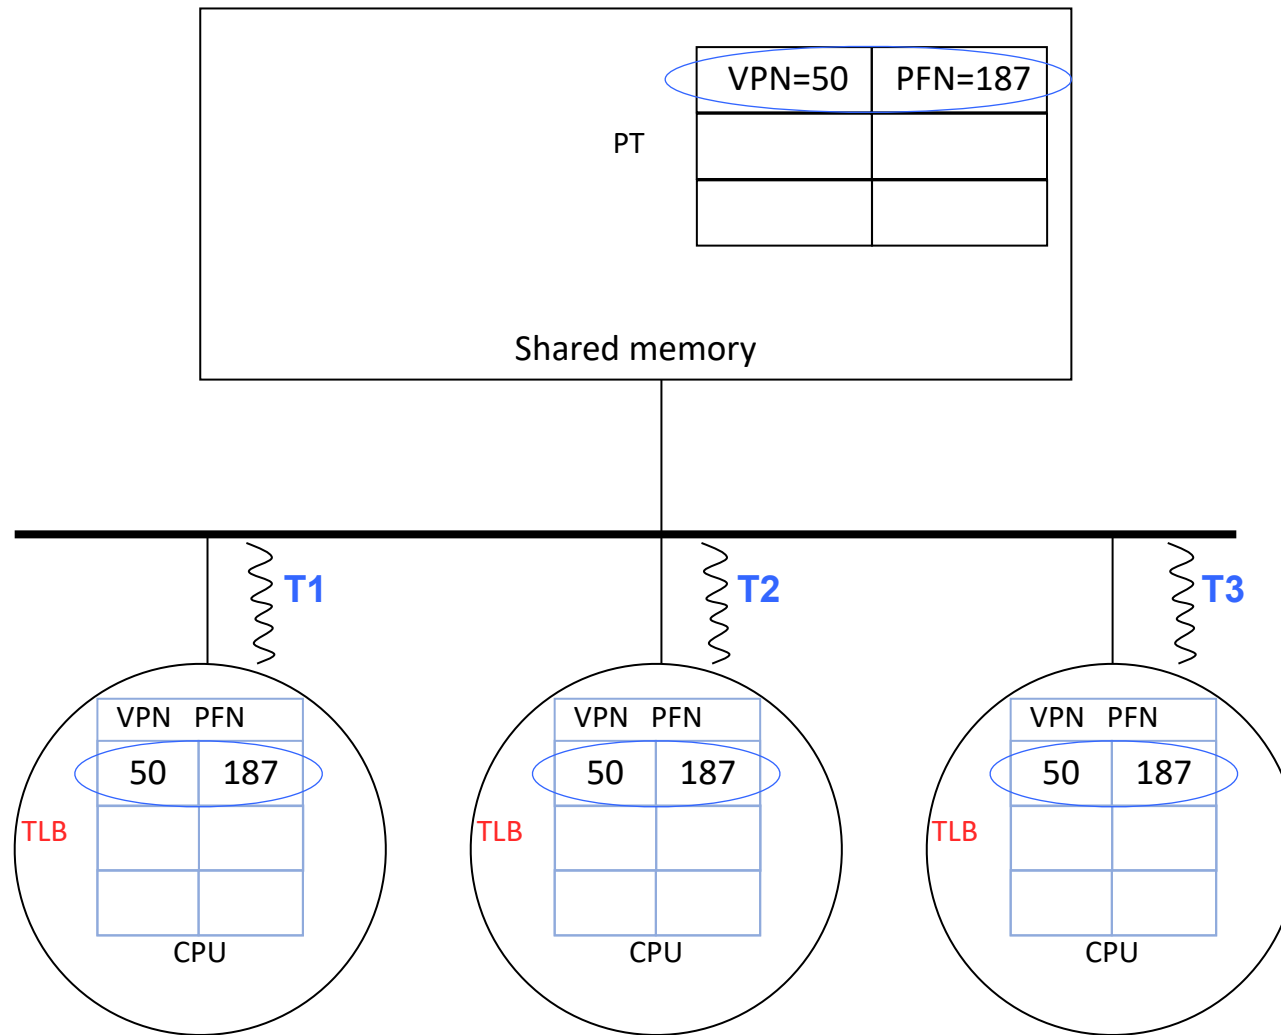
Example: Handling an SMP page fault

Note that the TLBs have each pulled in VPN=50 because each of the threads referenced that page



Example: Handling an SMP page fault

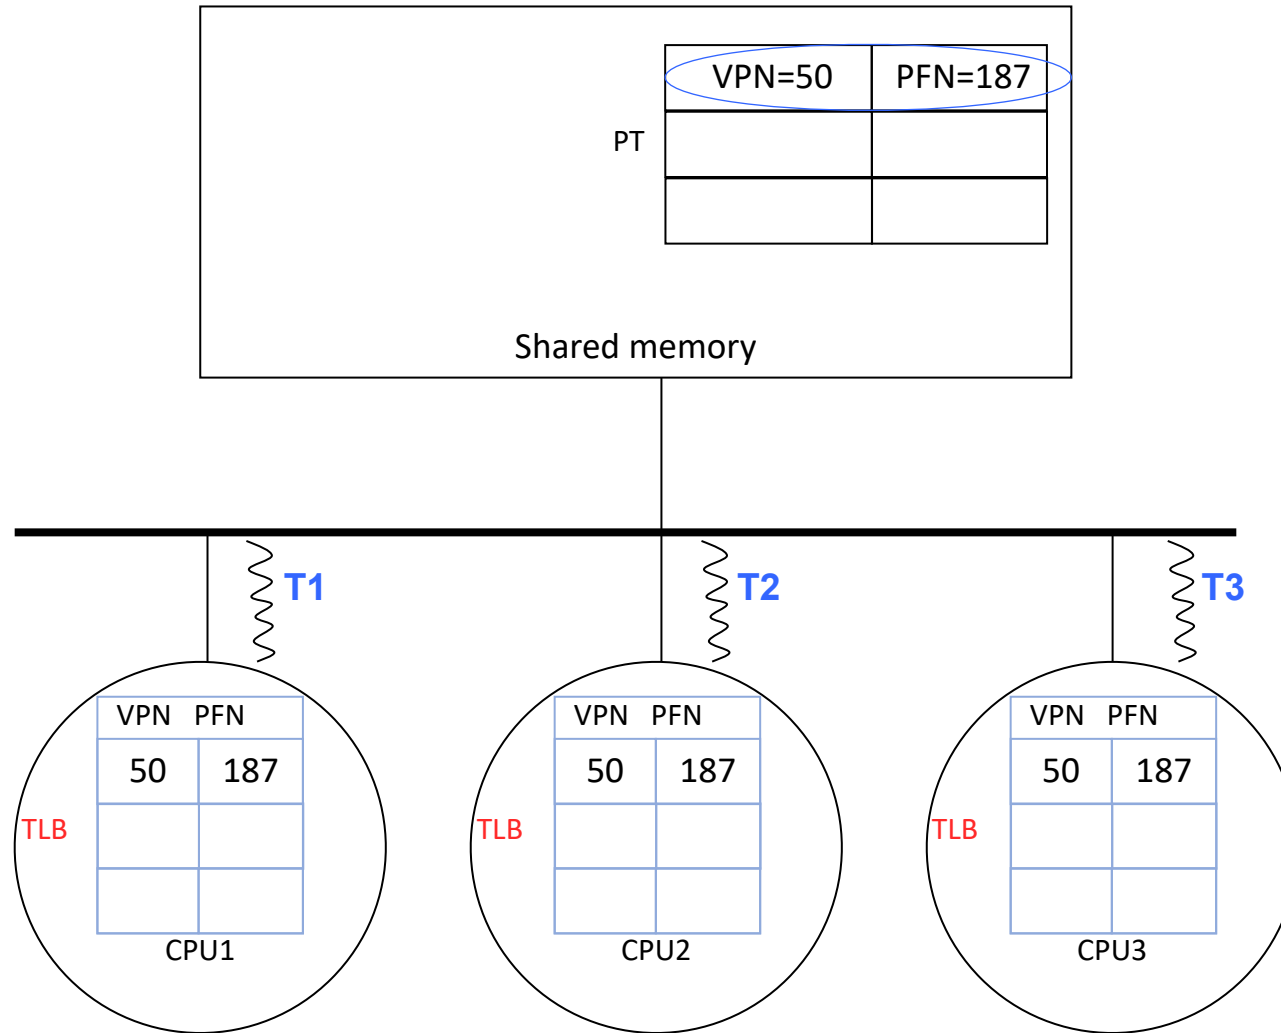
Note that the TLBs have each pulled in VPN=50 because each of the threads referenced that page



- Assume
 - T1 encounters a page fault on VPN=48
 - OS decides to evict VPN=50
 - And use PFN=187 for hosting VPN=48

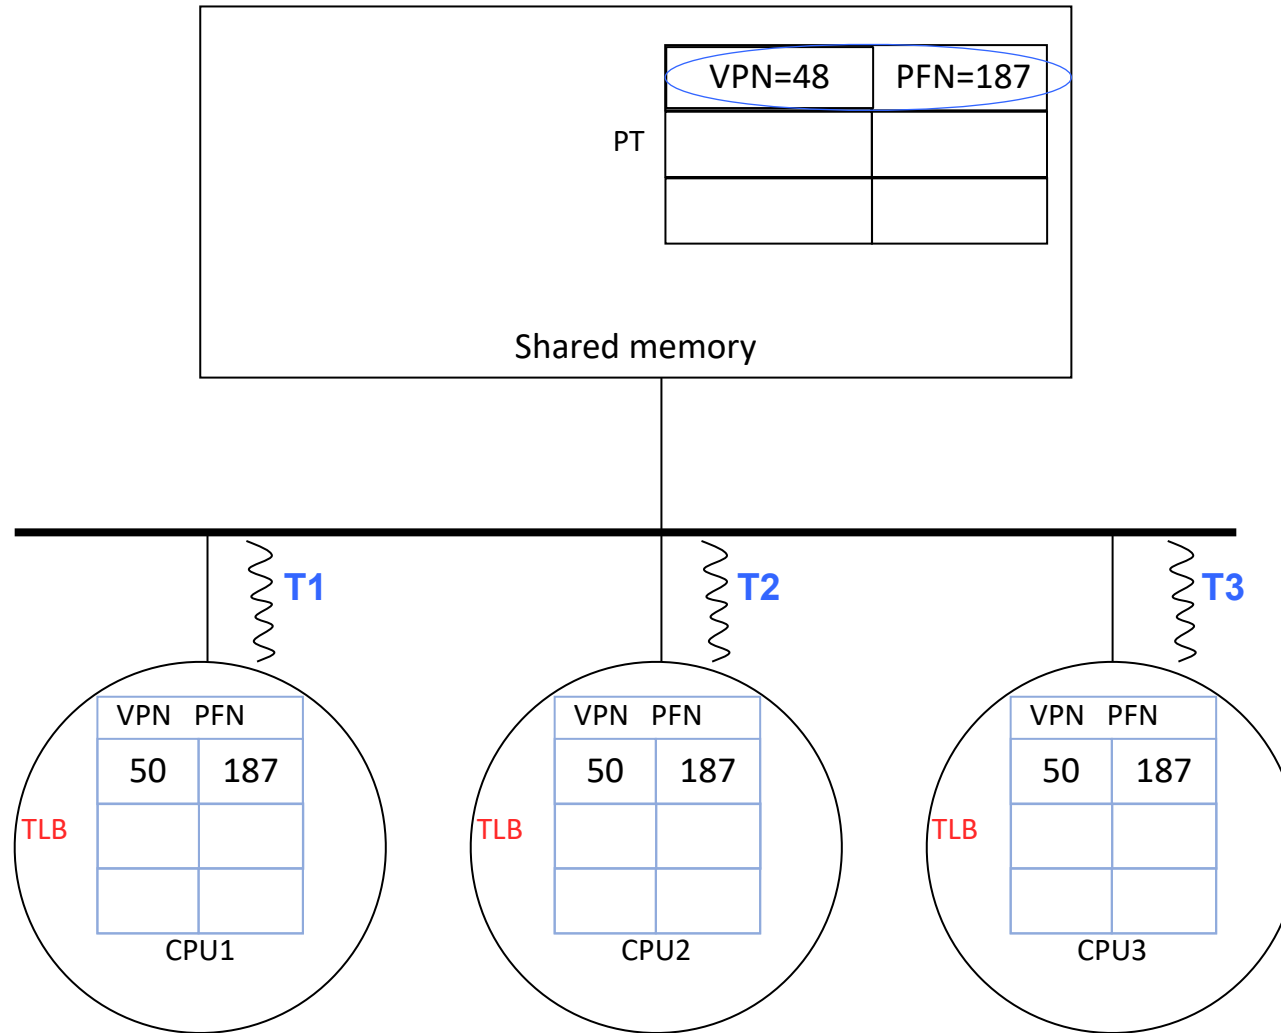
Example: Handling an SMP page fault

OS changes the page table entry for VPN=50



Example: Handling an SMP page fault

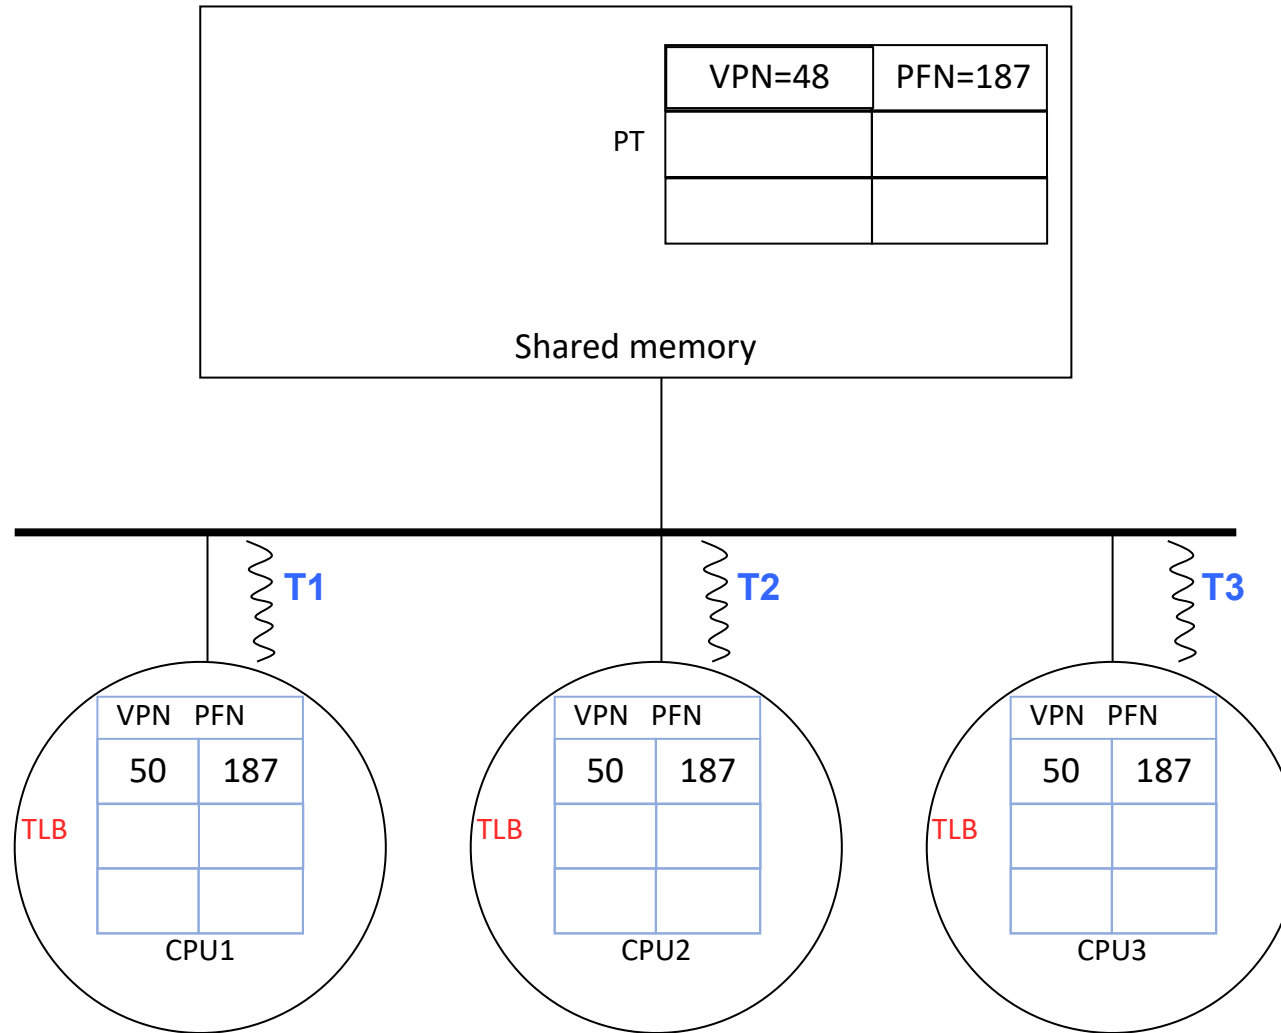
OS changes the page table entry for VPN=50



Example: Handling an SMP page fault

OS changes the page table entry for VPN=50

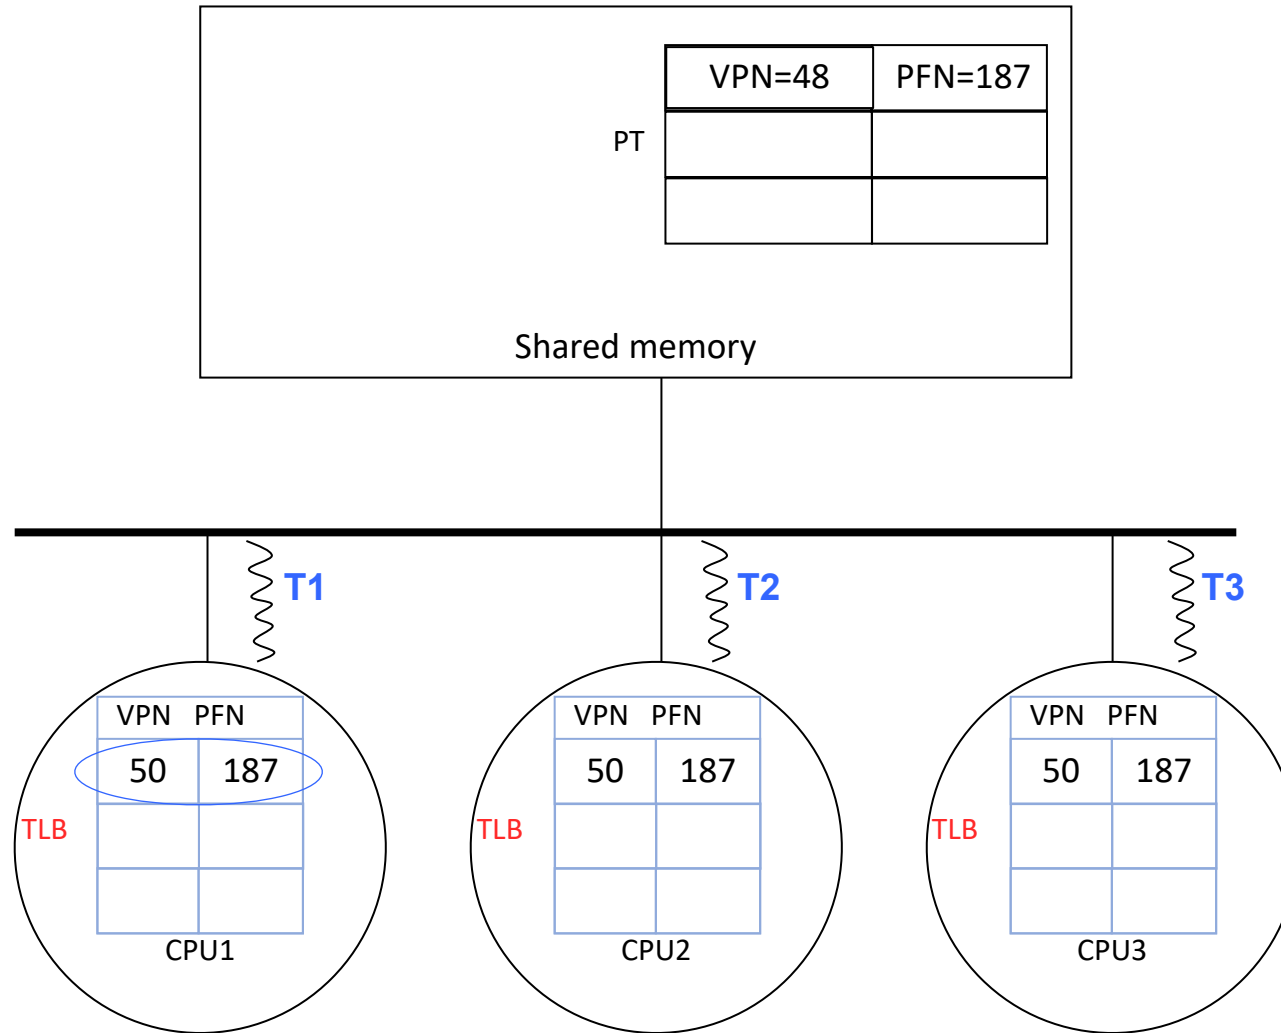
Then because it's running on CPU1, it evicts the TLB entry for VPN 50



Example: Handling an SMP page fault

OS changes the page table entry for VPN=50

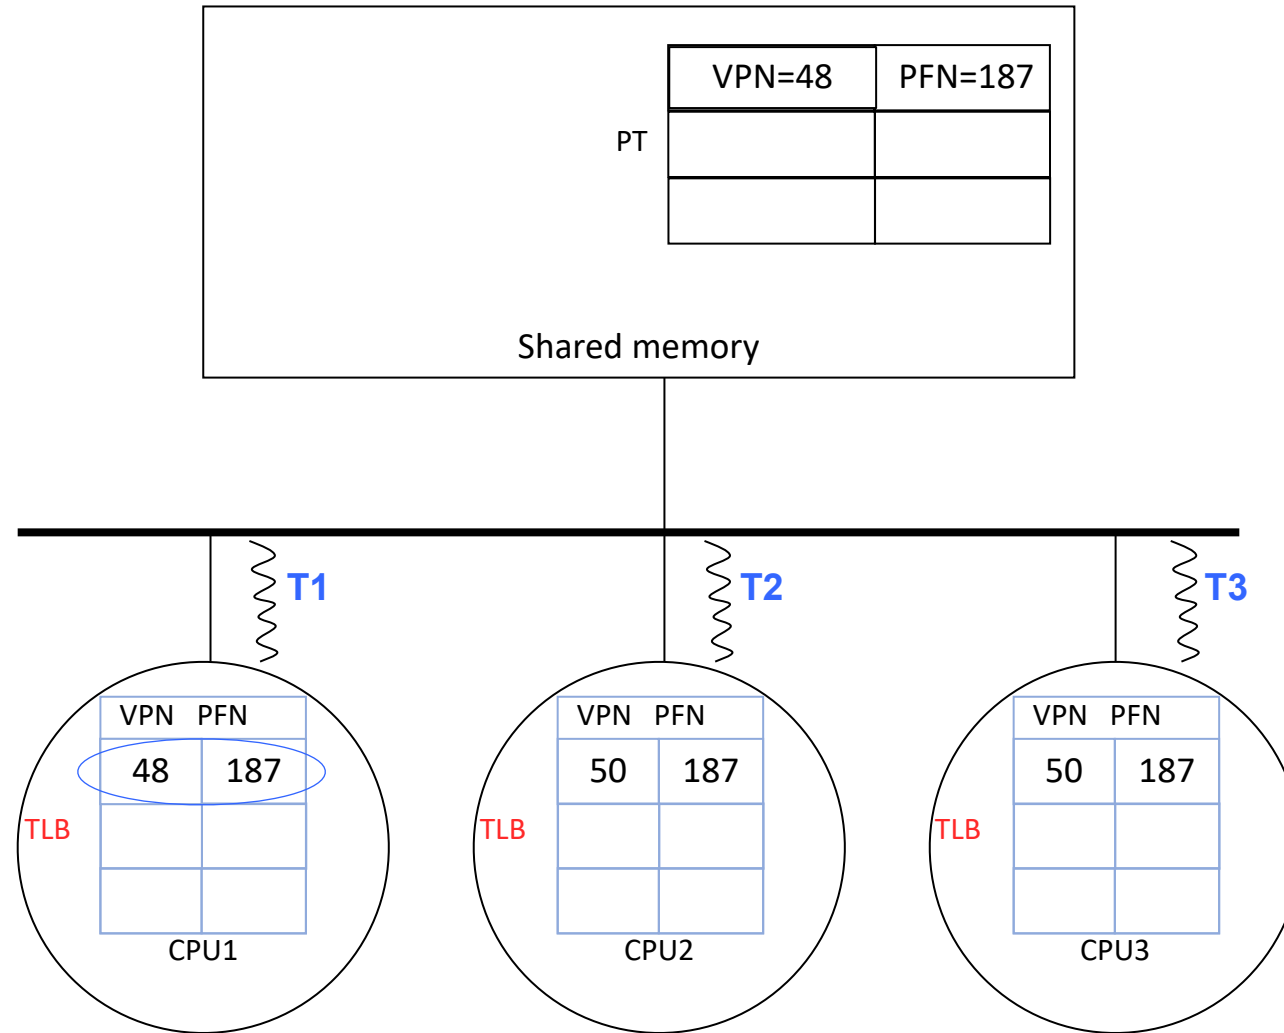
Then because it's running on CPU1, it evicts the TLB entry for VPN 50



Example: Handling an SMP page fault

OS changes the page table entry for VPN=50

Then because it's running on CPU1, it evicts the TLB entry for VPN 50

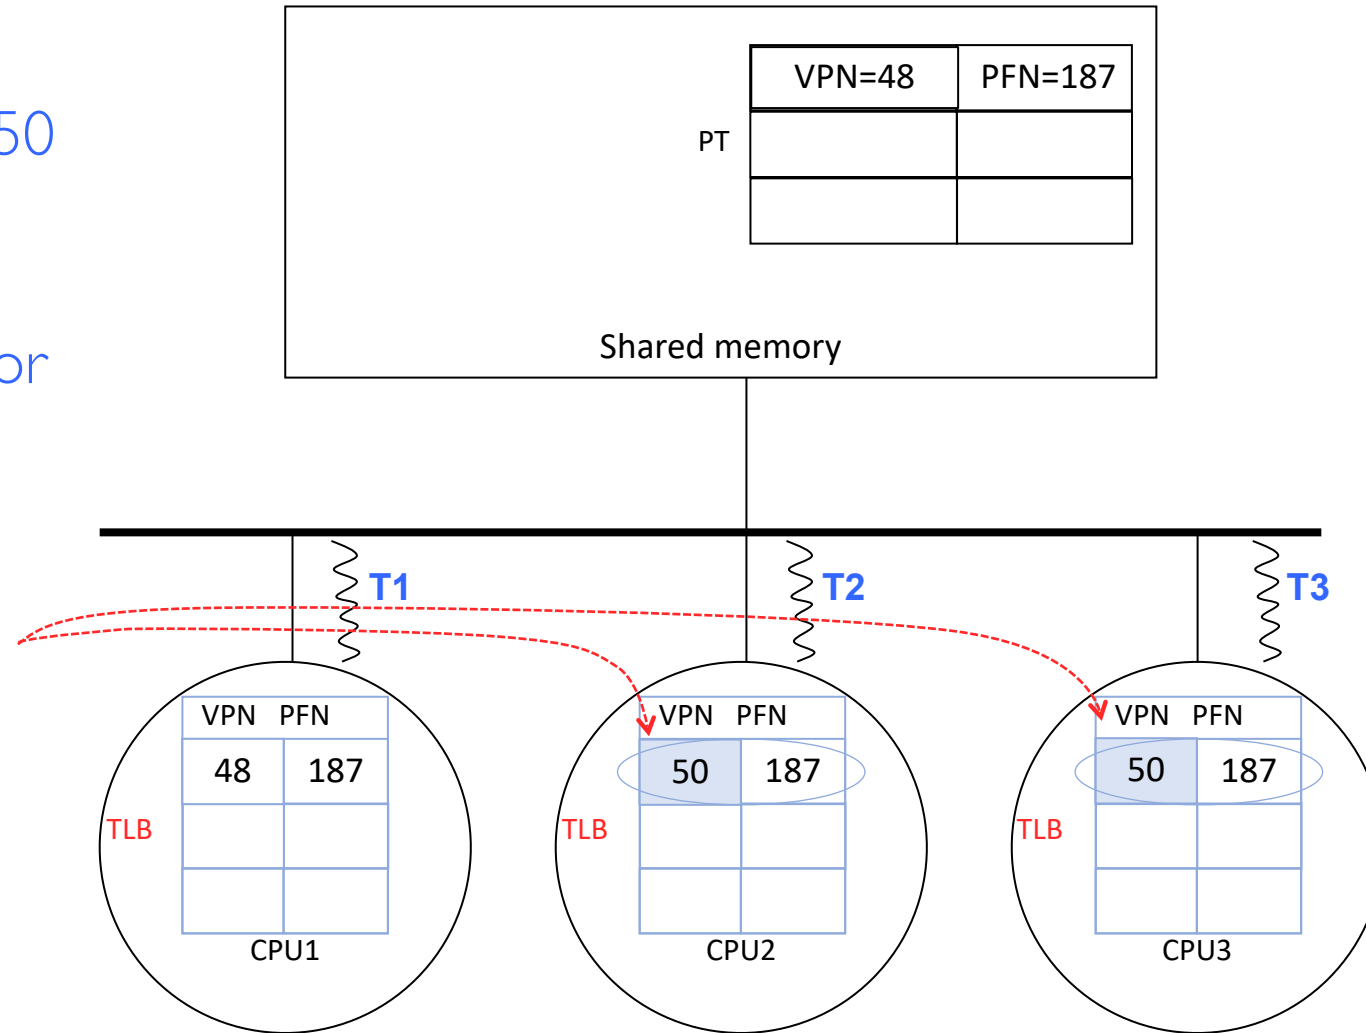


Example: Handling an SMP page fault

OS changes the page table entry for VPN=50

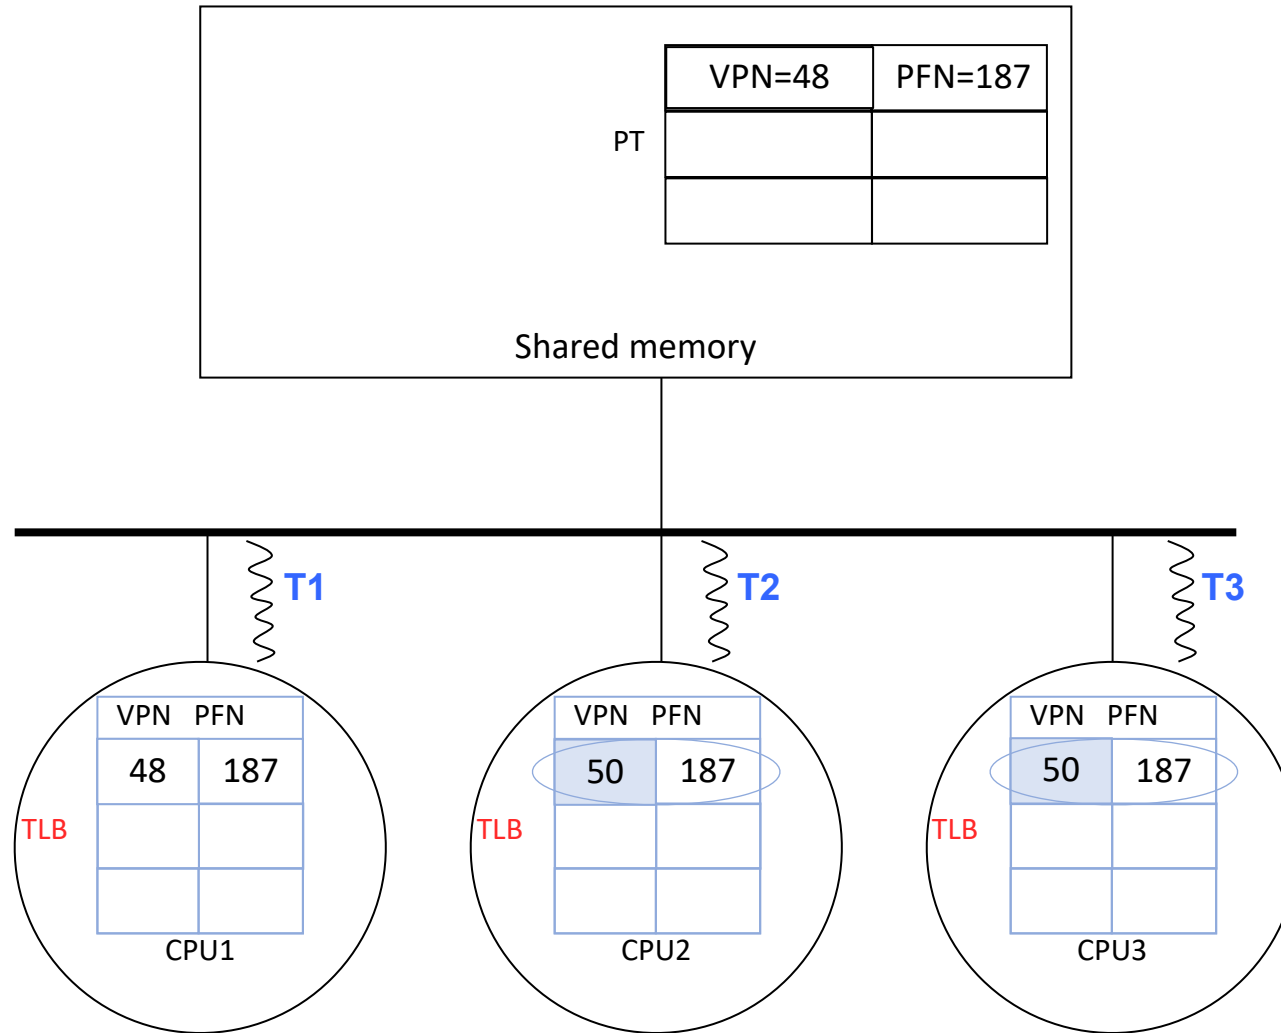
Then because it's running on CPU1, it evicts the TLB entry for VPN 50

Now we've got stale TLB entries in CPU2 and CPU3



Example: Handling an SMP page fault

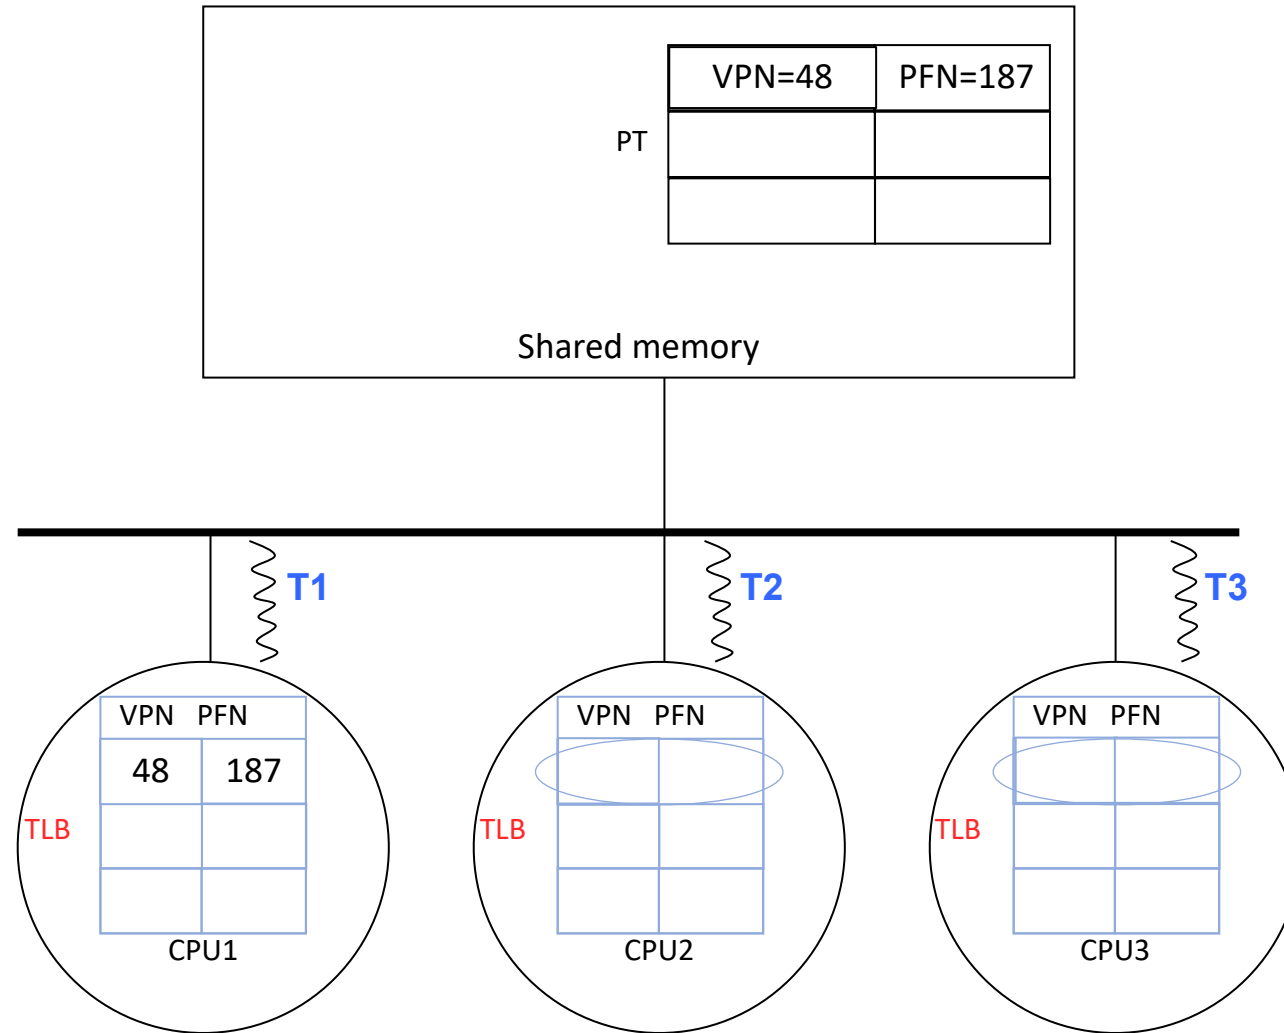
Then we have the TLB
Shutdown



Example: Handling an SMP page fault

Then we have the TLB
Shutdown

The OS arranges to
invalidate the
corresponding TLB
entries on the other
CPUs

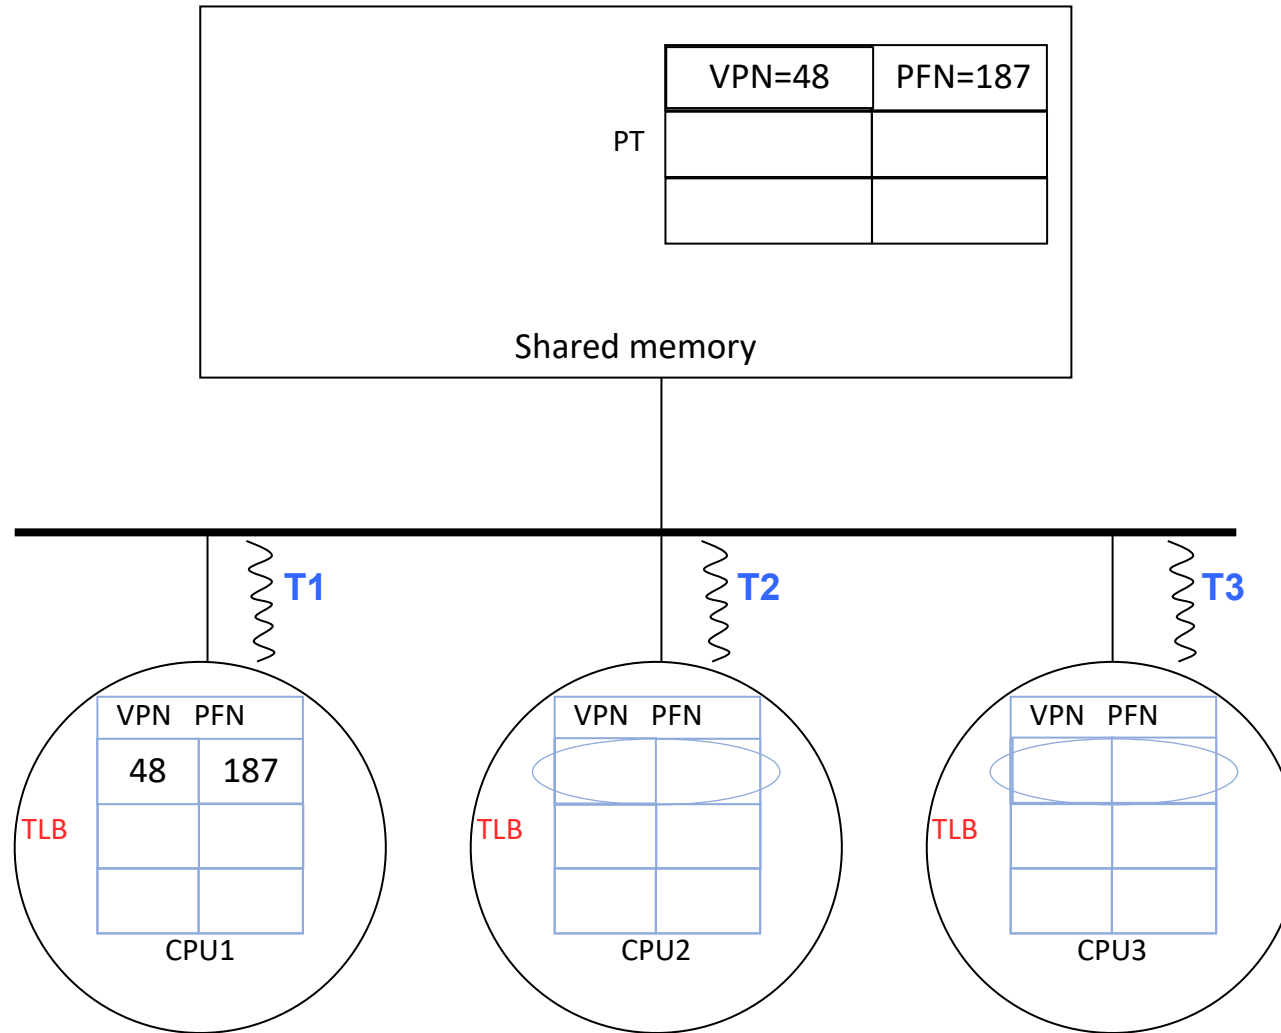


Example: Handling an SMP page fault

Then we have the TLB
Shutdown

The OS arranges to
invalidate the
corresponding TLB
entries on the other
CPUs

And the CPUs can pull
in the updated PTE
when (and if) they
next reference
VPN=48





Ensuring that all threads of a process share an address space in an SMP is


- A. Impossible
- B. Trivially achieved since the page table resides in shared memory
- C. Achieved by careful replication of the page table by the operating system for each thread
- D. Achieved by special-purpose hardware that no one has told us about yet



Keeping the TLBs consistent in an SMP

- A. Is the responsibility of the programmer
- B. Is the responsibility of the hardware
- C. Is the responsibility of the operating system
- D. Is not possible

How do we implement Symmetric Multi Processing (SMP)?



The System (hardware+OS) has to ensure 3 things:

1. Threads of the same process share the same PT
2. Threads have synchronization atomicity
3. Threads have identical views of memory

2) Threads have synchronization atomicity

2) Threads have synchronization atomicity

- We already introduced the TEST-AND-SET instruction

2) Threads have synchronization atomicity

- We already introduced the TEST-AND-SET instruction
- It should be easy on a multiprocessor, right?

2) Threads have synchronization atomicity

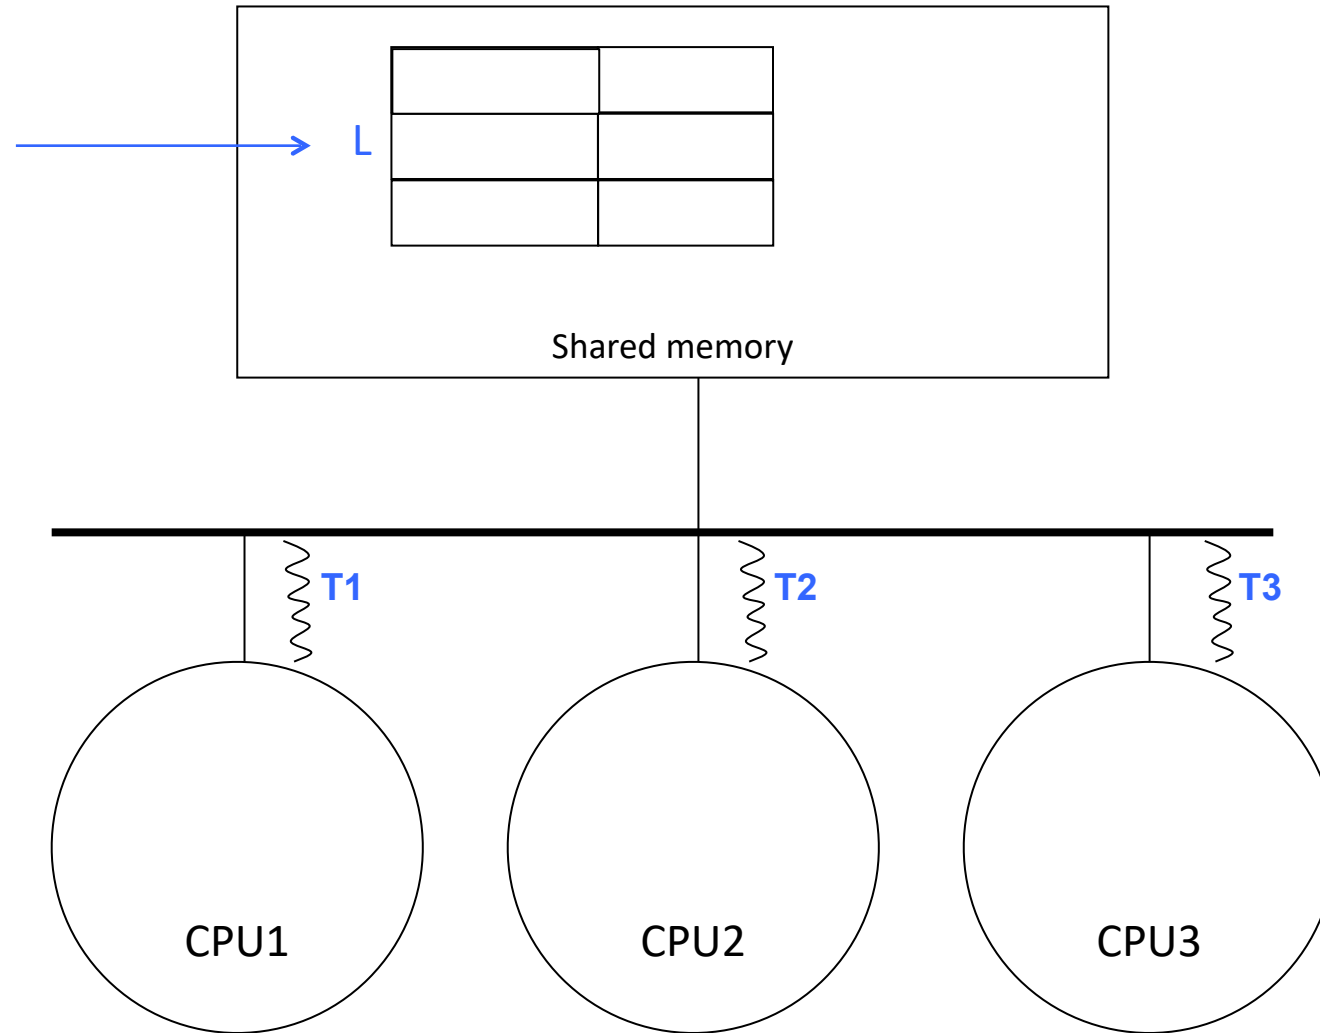
- We already introduced the TEST-AND-SET instruction
- It should be easy on a multiprocessor, right?
- The location we use for synchronization is in shared memory, so no sweat.

2) Threads have synchronization atomicity

- We already introduced the TEST-AND-SET instruction
- It should be easy on a multiprocessor, right?
- The location we use for synchronization is in shared memory, so no sweat.
- What could go wrong?

2) Threads have synchronization atomicity

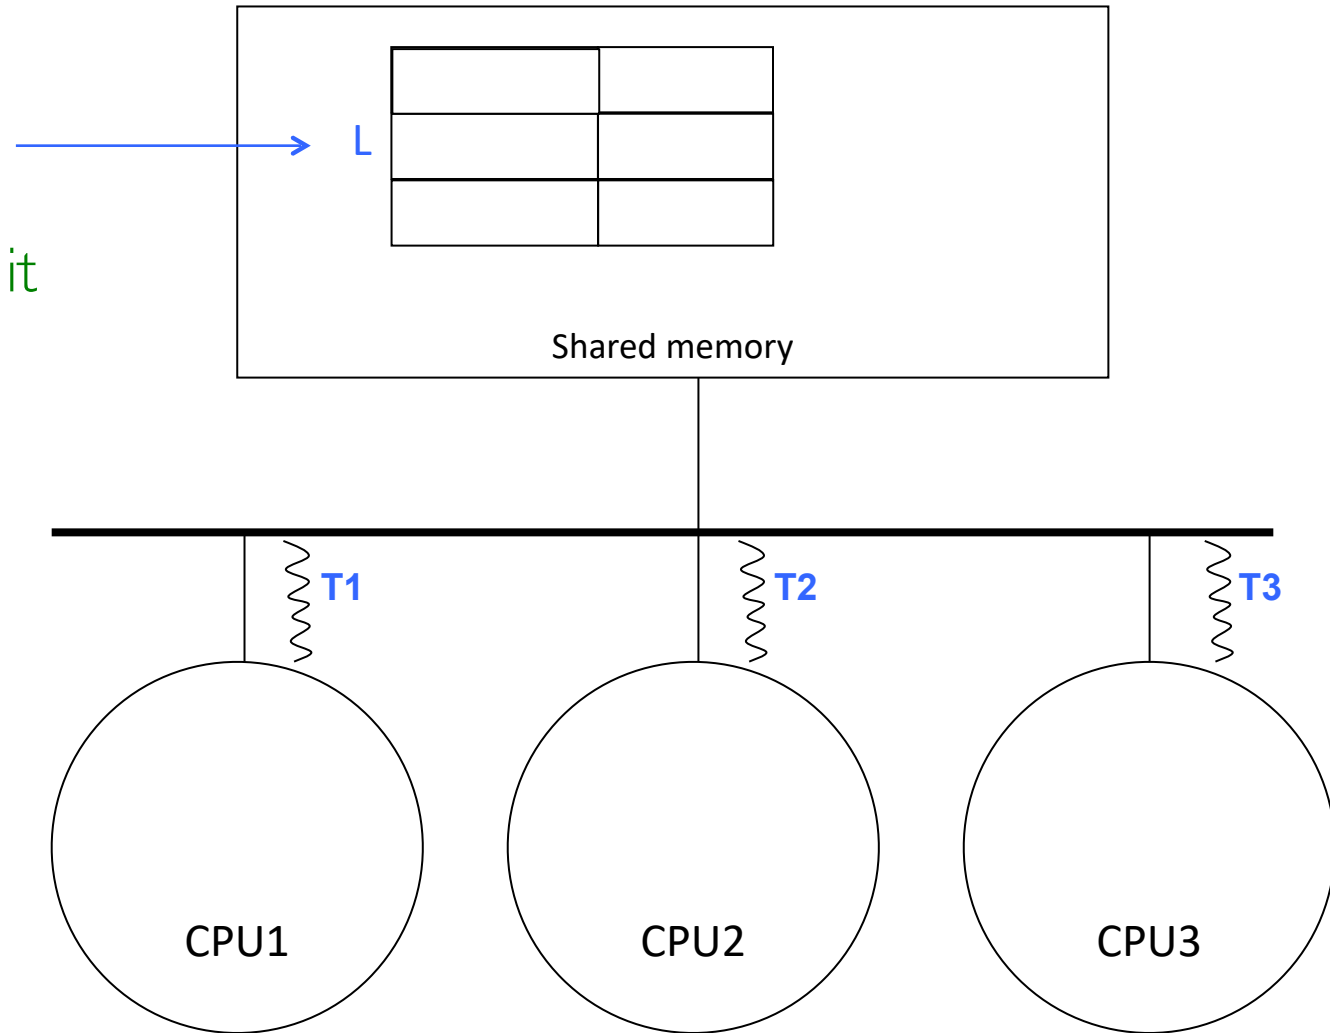
We have memory location L for TEST-AND-SET



2) Threads have synchronization atomicity

We have memory location L for TEST-AND-SET

Each CPU can access it

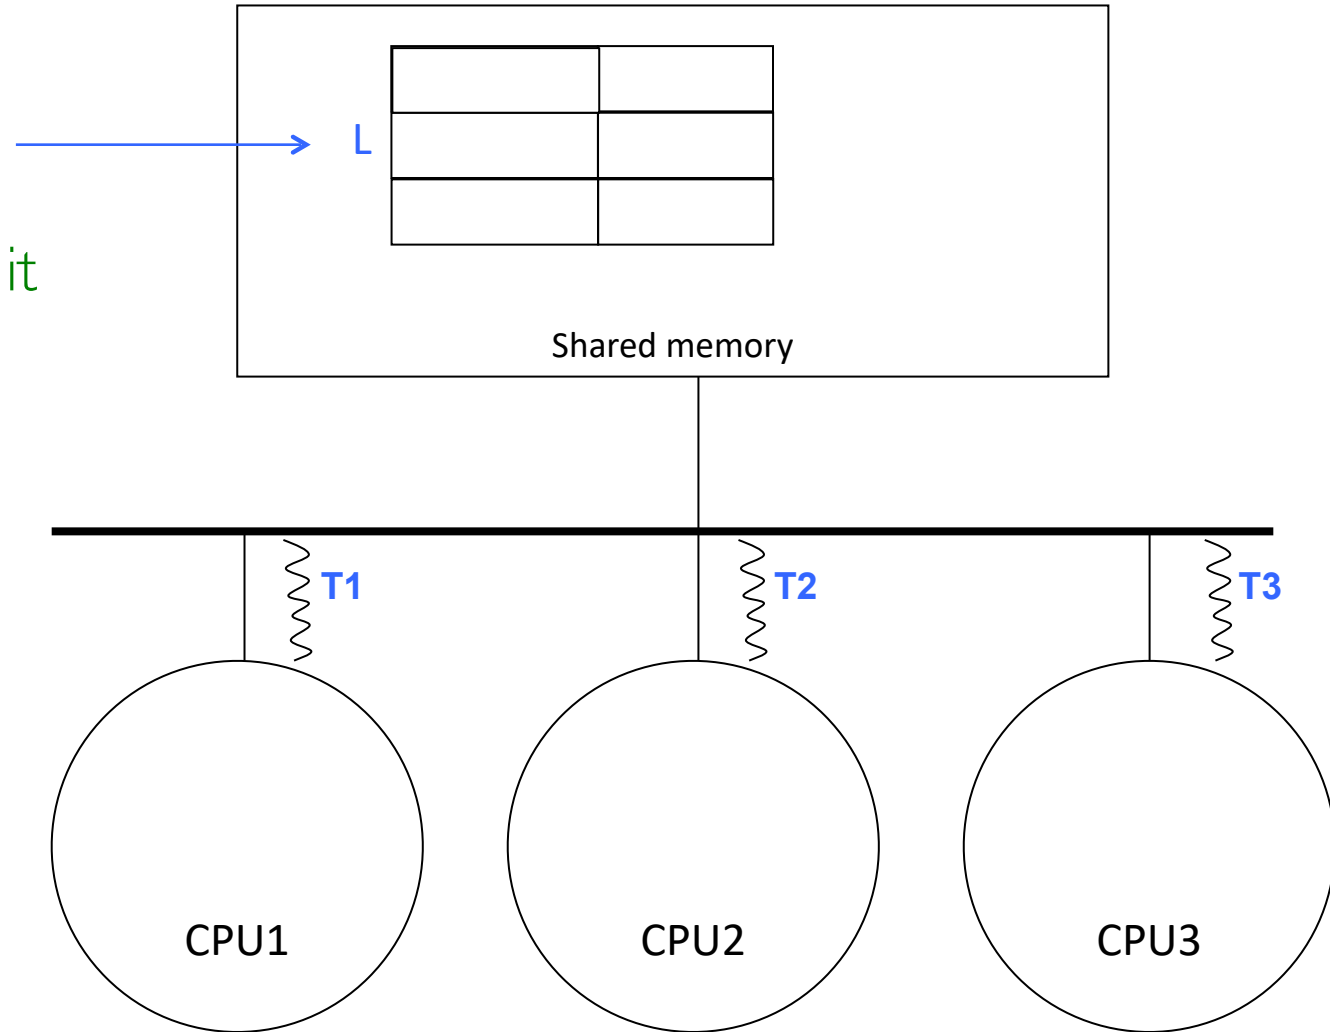


2) Threads have synchronization atomicity

We have memory location L for TEST-AND-SET

Each CPU can access it

But...what did we forget?



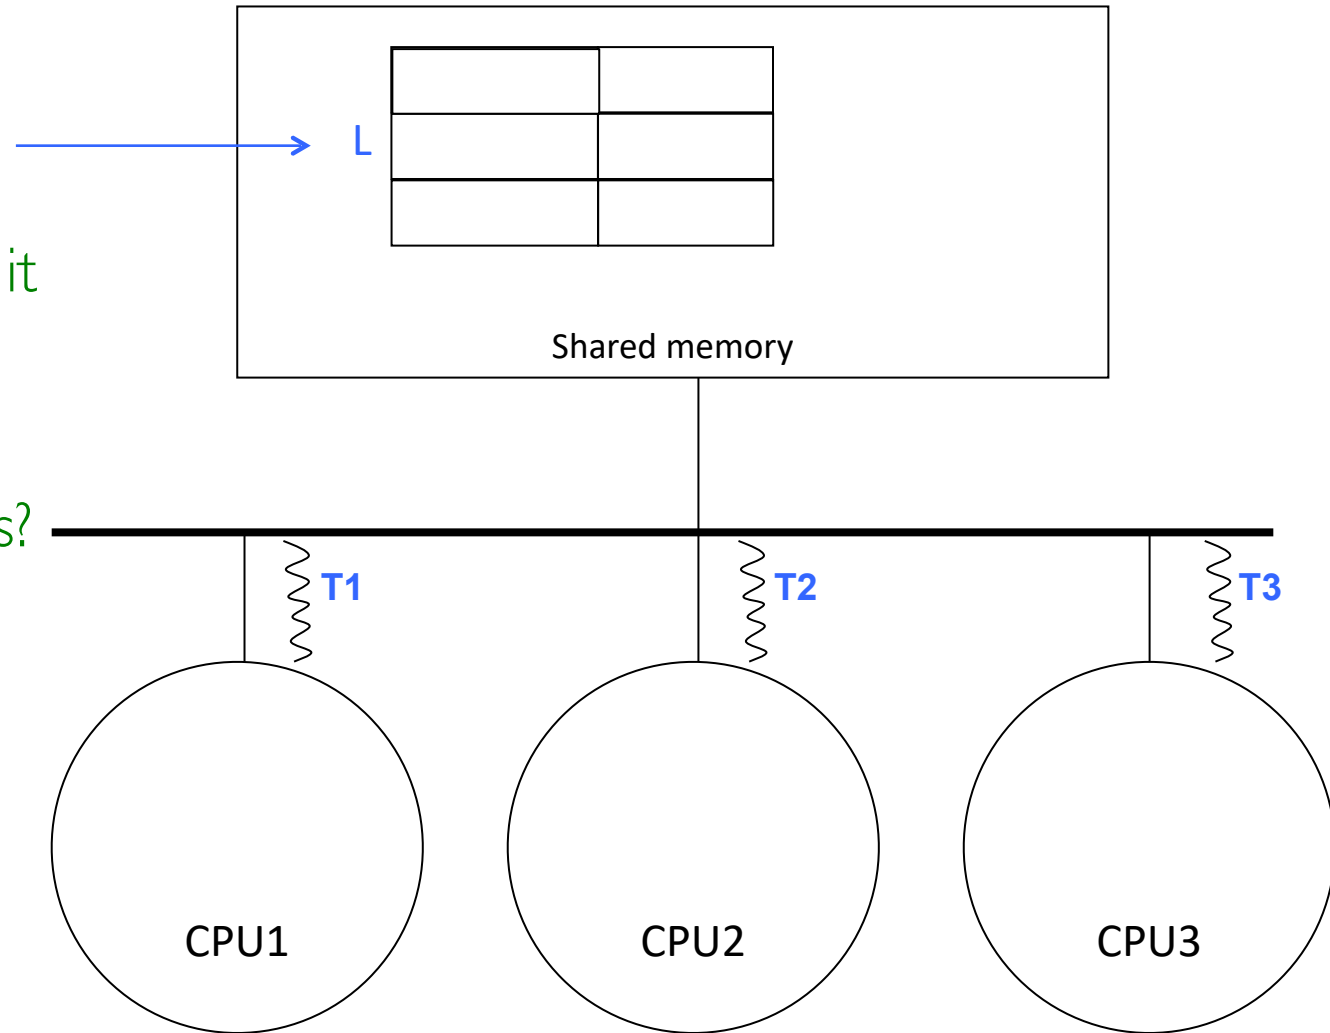
2) Threads have synchronization atomicity

We have memory location L for TEST-AND-SET

Each CPU can access it

But...what did we forget?

Where are the caches?



2) Threads have synchronization atomicity

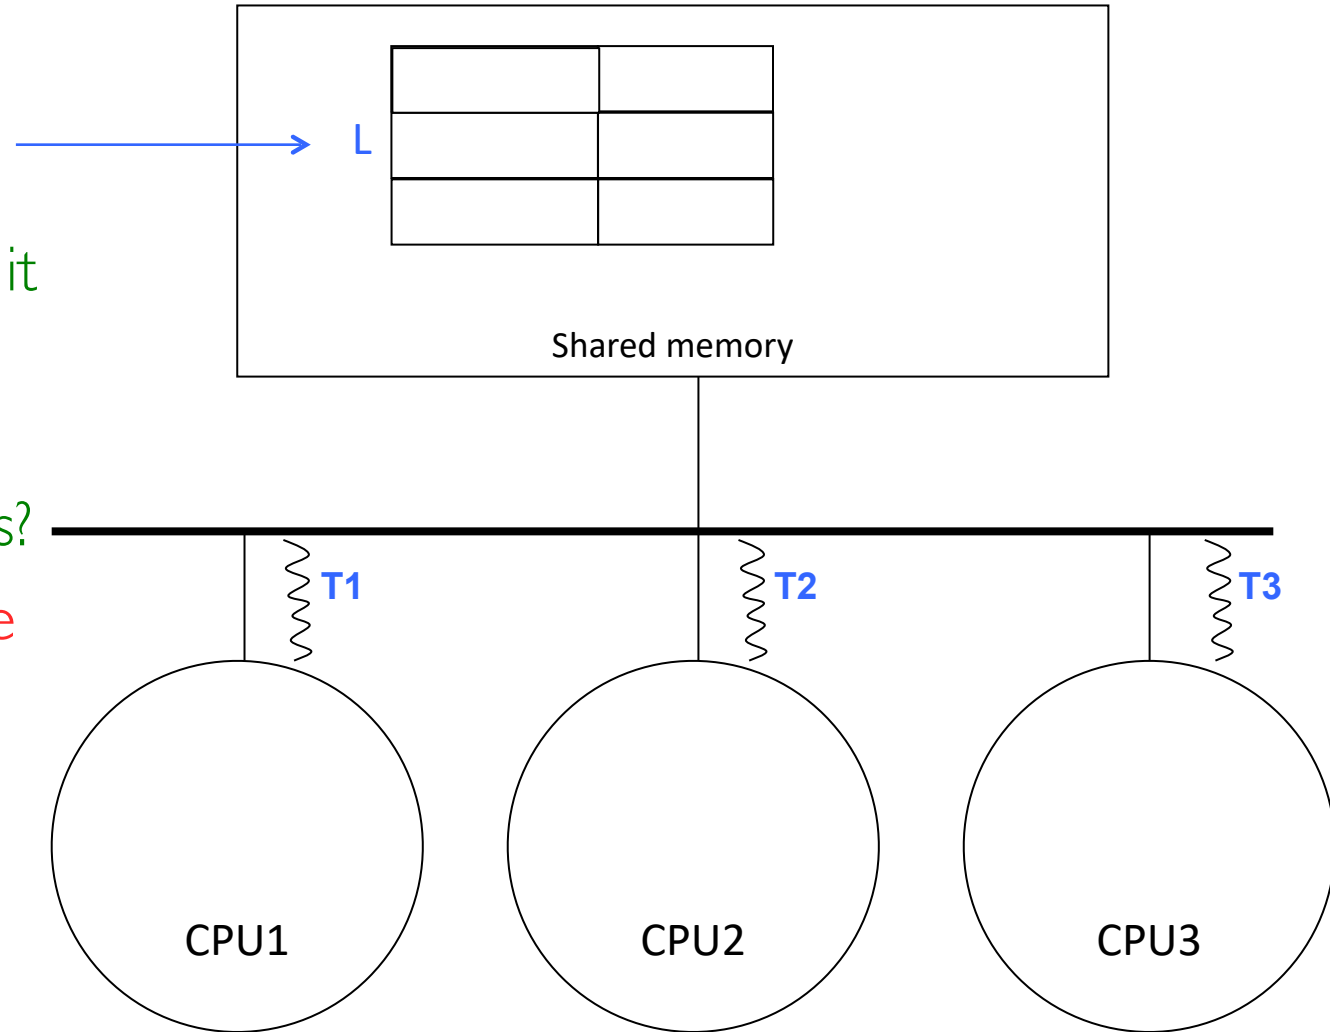
We have memory location L for TEST-AND-SET

Each CPU can access it

But...what did we forget?

Where are the caches?

In the CPUs of course



2) Threads have synchronization atomicity

We have memory location L for TEST-AND-SET

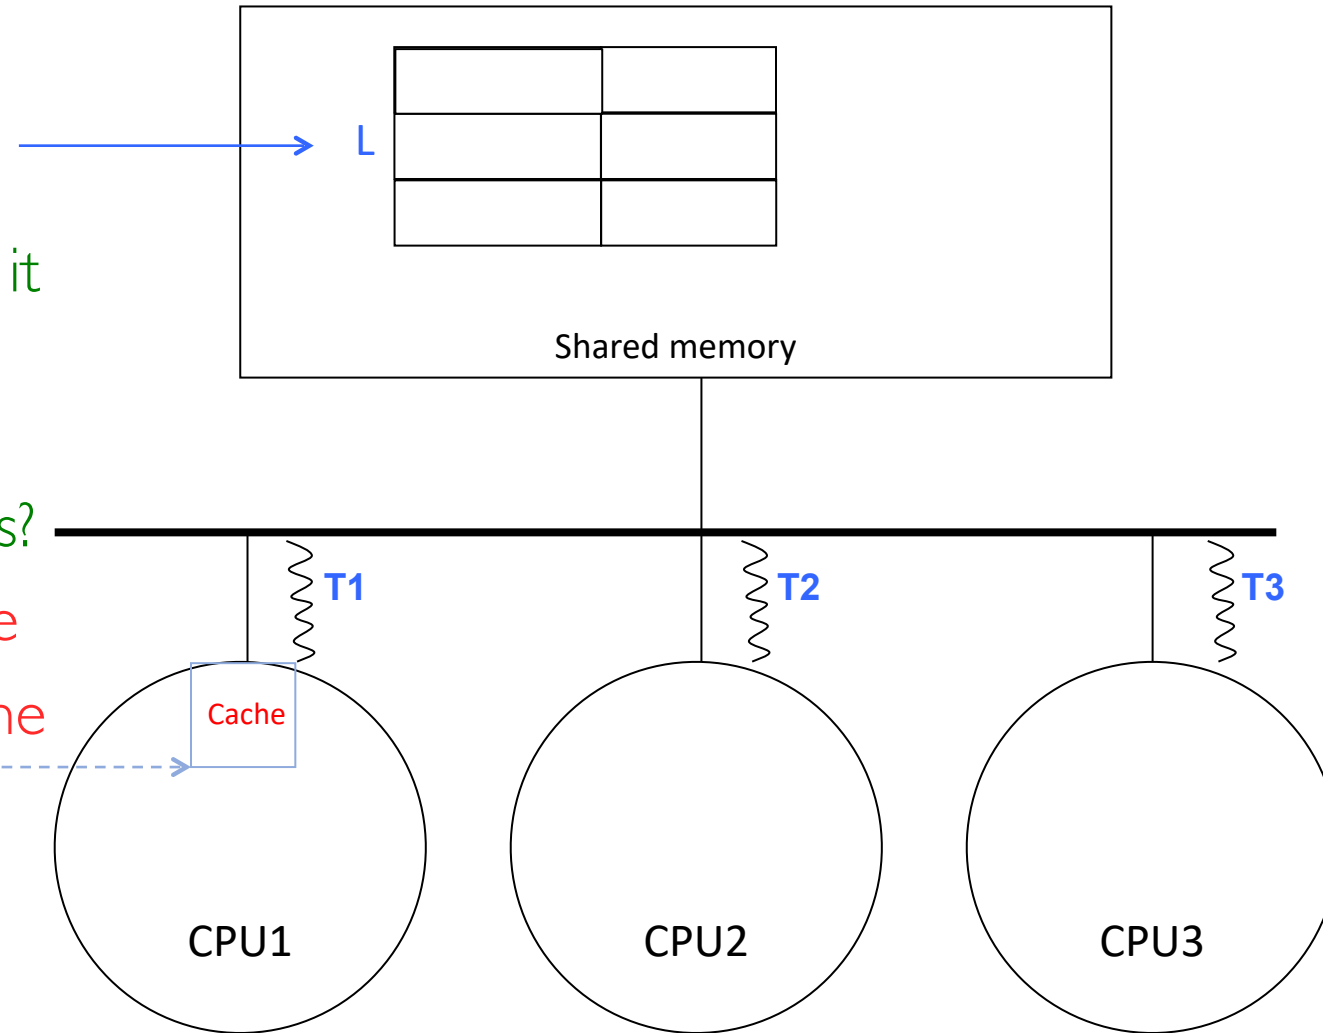
Each CPU can access it

But...what did we forget?

Where are the caches?

In the CPUs of course

If L is cached, the cache is wrong!



2) Threads have synchronization atomicity

We have memory location L for TEST-AND-SET

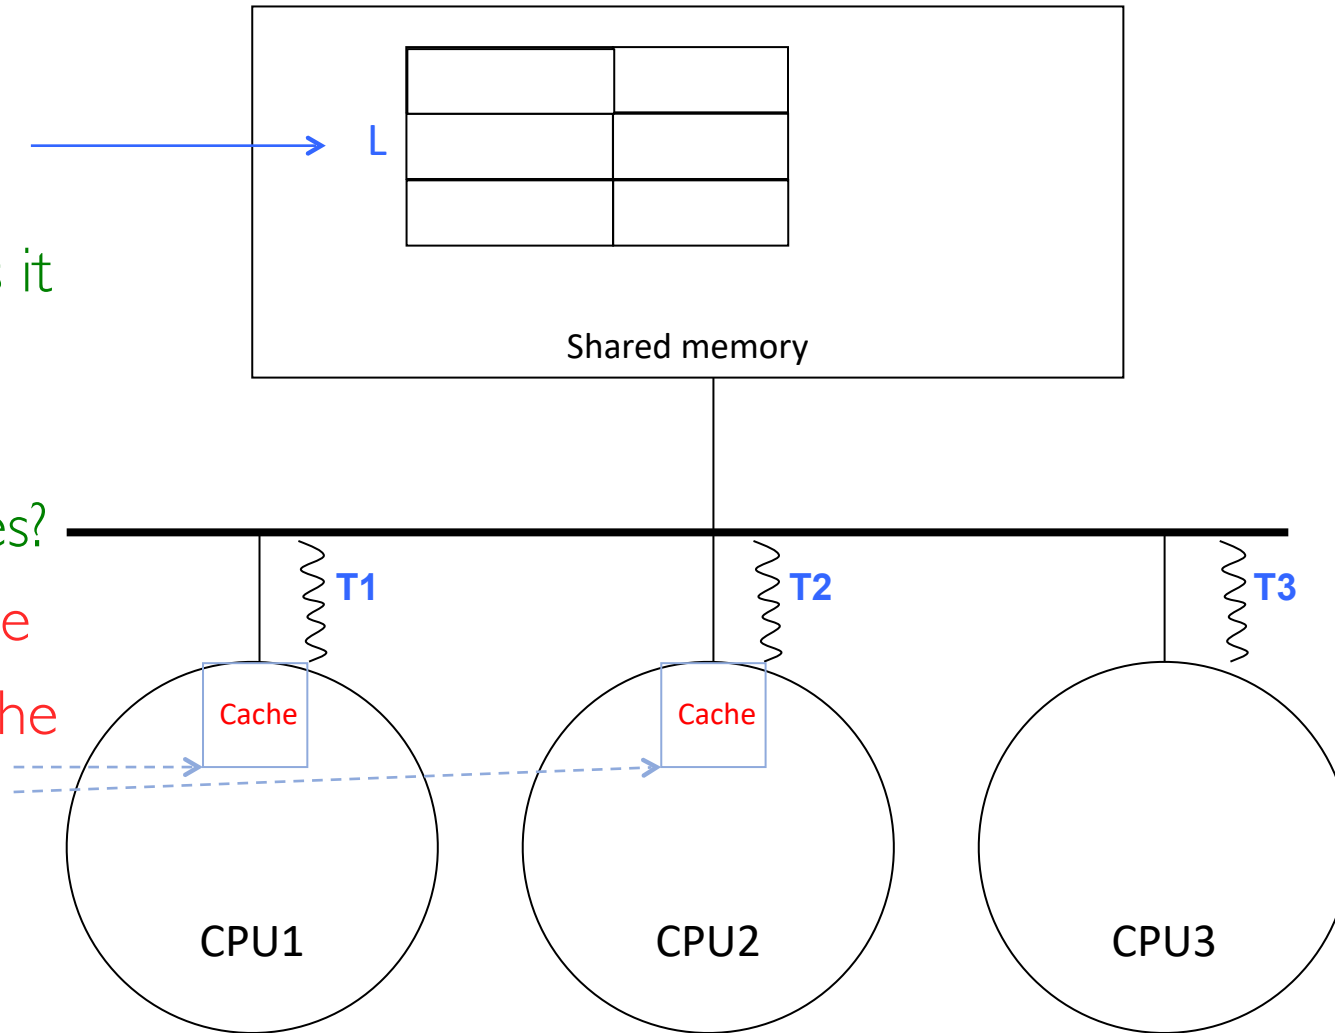
Each CPU can access it

But...what did we forget?

Where are the caches?

In the CPUs of course

If L is cached, the cache is wrong!



2) Threads have synchronization atomicity

We have memory location L for TEST-AND-SET

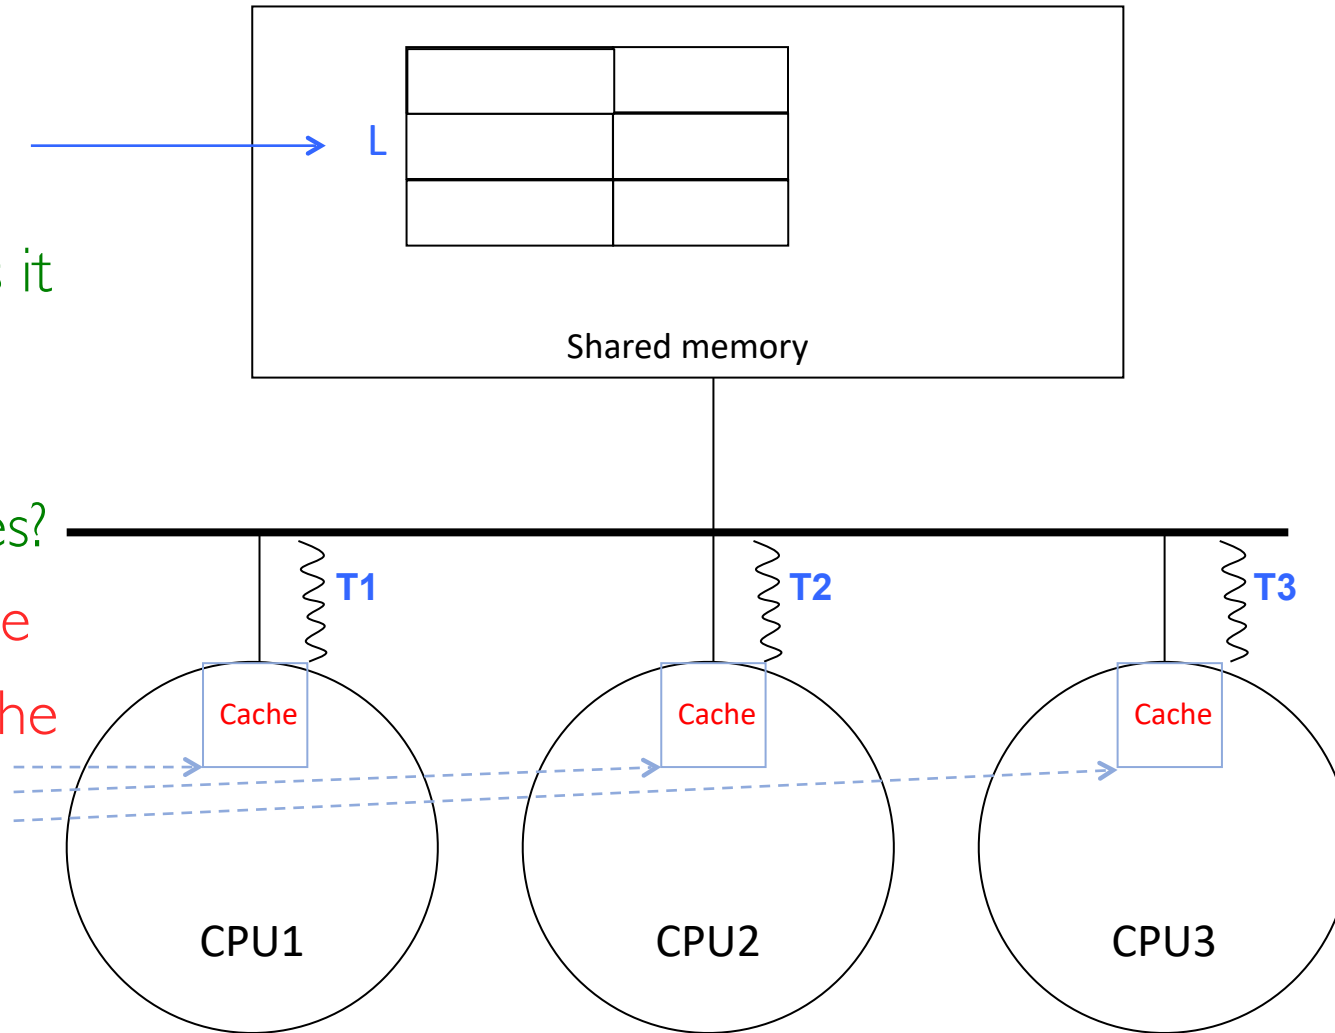
Each CPU can access it

But...what did we forget?


Where are the caches?

In the CPUs of course

If L is cached, the cache is wrong!



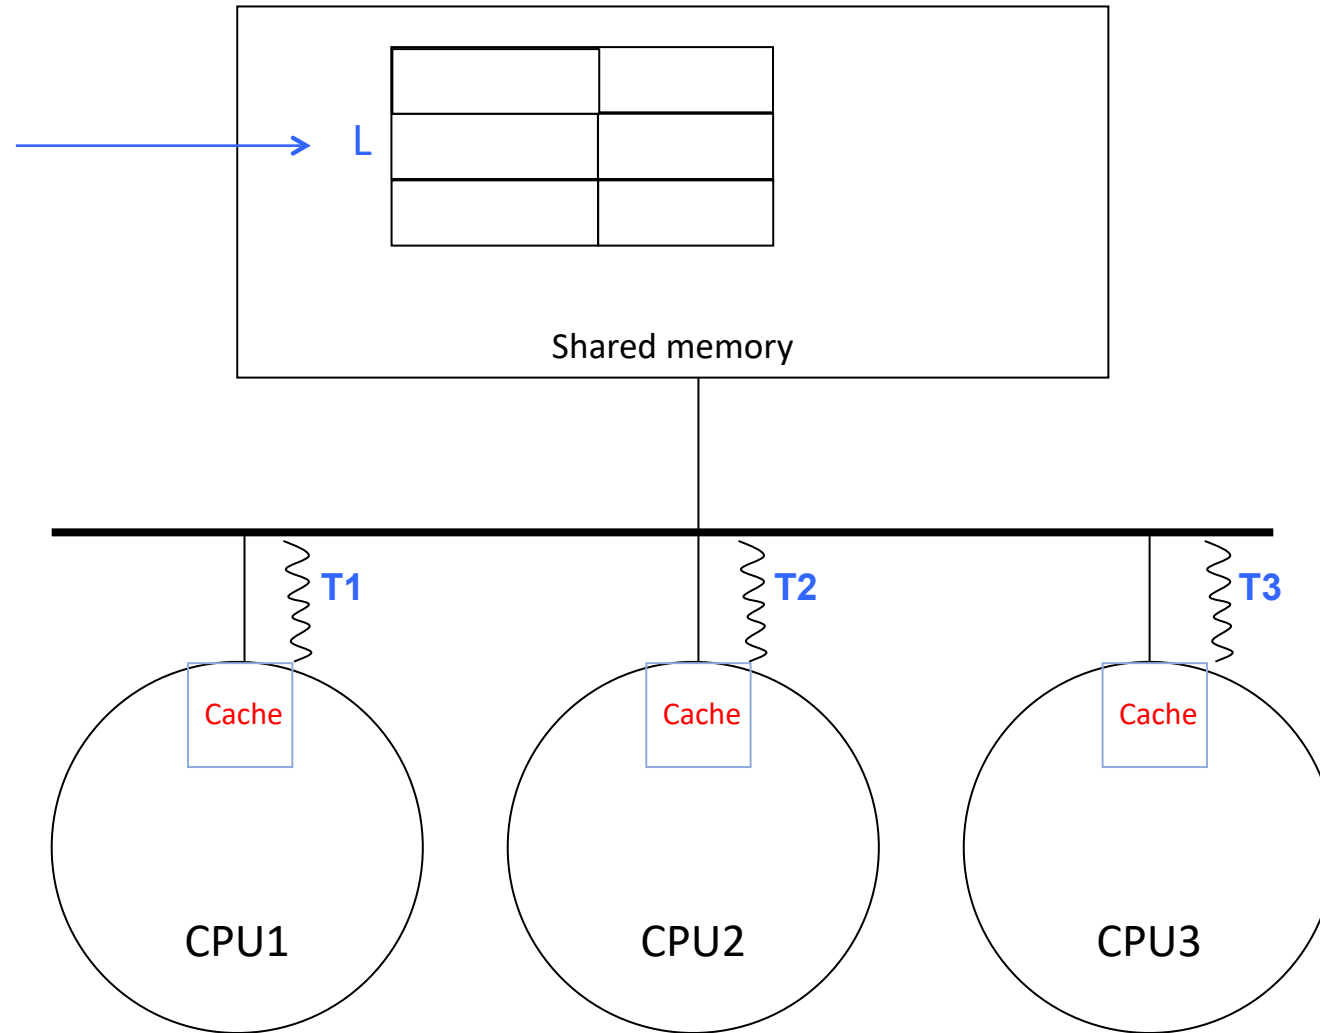
2) Threads have synchronization atomicity



2) Threads have synchronization atomicity

What shall we do?

One solution is to
bypass the cache for
the T&S instruction



Requirements for SMP

1. ~~Threads of the same process share the same PT~~
2. ~~Threads have synchronization atomicity~~
3. Threads have identical views of memory

Requirements for SMP

1. ~~Threads of the same process share the same PT~~
2. ~~Threads have synchronization atomicity~~
3. Threads have identical views of memory
 - This implies that access to a memory location returns the same value on all CPUs

Requirements for SMP

1. ~~Threads of the same process share the same PT~~
2. ~~Threads have synchronization atomicity~~
3. Threads have identical views of memory
 - This implies that access to a memory location returns the same value on all CPUs
 - But caches make copies of data

Requirements for SMP

1. ~~Threads of the same process share the same PT~~
2. ~~Threads have synchronization atomicity~~
3. Threads have identical views of memory
 - This implies that access to a memory location returns the same value on all CPUs
 - But caches make copies of data
 - We'll refer to the method of keeping all the copies of the same data across caches as a cache coherence protocol

3) Threads have identical views of memory

- Two possible solutions, in hardware

3) Threads have identical views of memory

- Two possible solutions, *in hardware*
- In both cases, cache becomes active and monitors or *snoops* the bus

3) Threads have identical views of memory

- Two possible solutions, *in hardware*
- In both cases, cache becomes active and monitors or *snoops* the bus
- Let's watch a memory location change value from *X* to *X'*

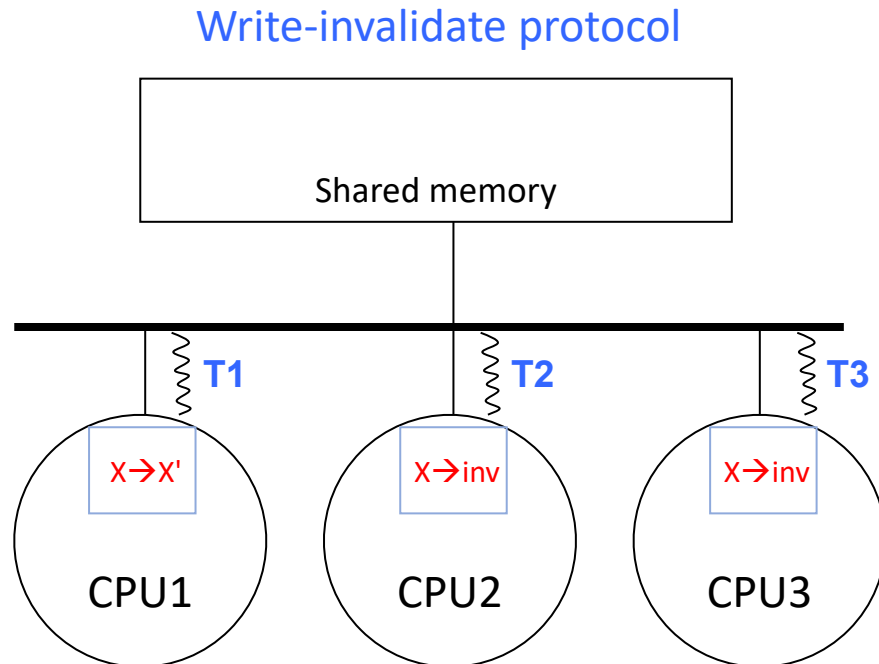
3) Threads have identical views of memory

- Two possible solutions, *in hardware*
- In both cases, cache becomes active and monitors or *snoops* the bus
- Let's watch a memory location change value from *X* to *X'*

Write-invalidate protocol

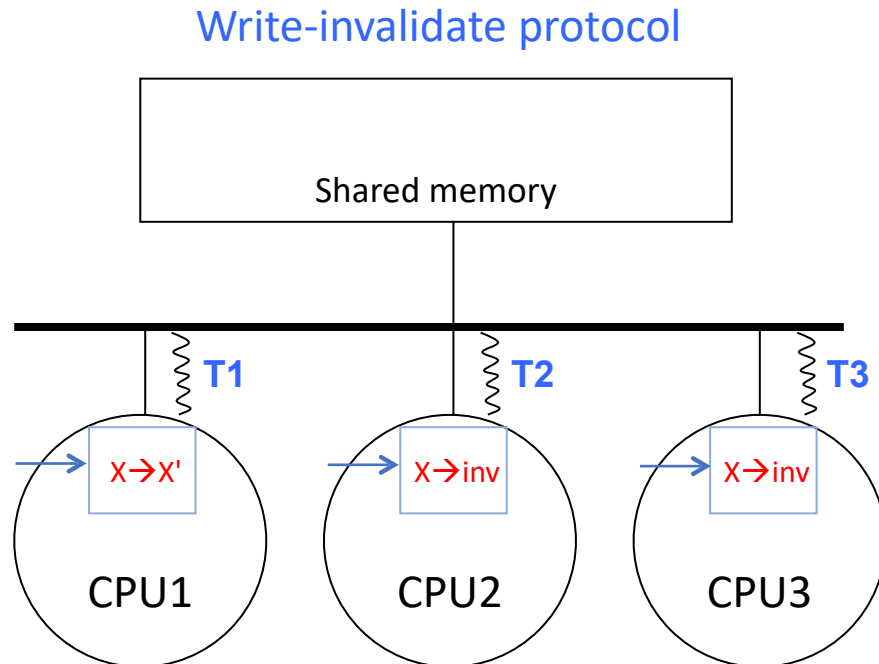
3) Threads have identical views of memory

- Two possible solutions, **in hardware**
- In both cases, cache becomes active and monitors or *snoops* the bus
- Let's watch a memory location change value from **X** to **X'**



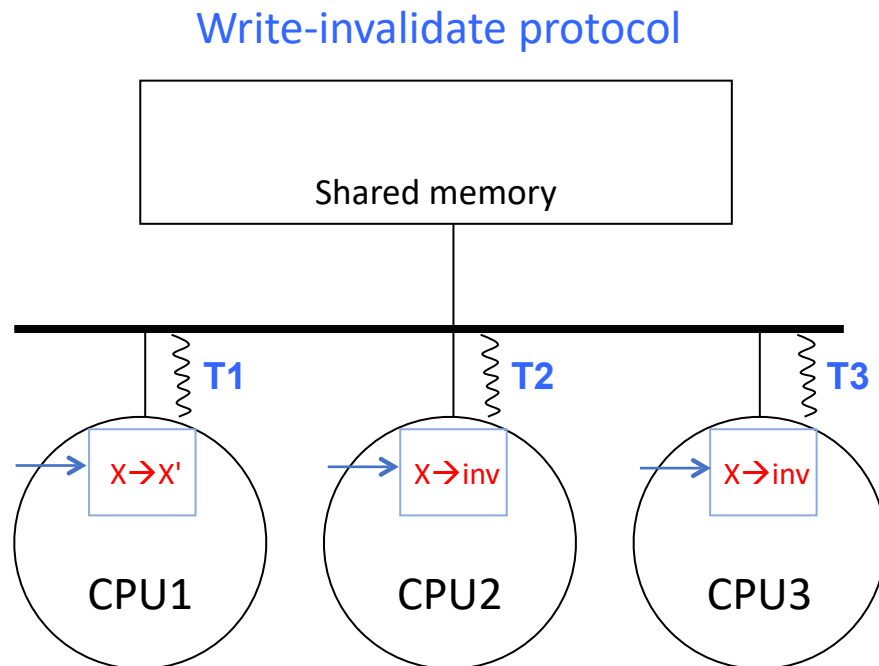
3) Threads have identical views of memory

- Two possible solutions, **in hardware**
- In both cases, cache becomes active and monitors or *snoops* the bus
- Let's watch a memory location change value from **X** to **X'**



3) Threads have identical views of memory

- Two possible solutions, **in hardware**
- In both cases, cache becomes active and monitors or *snoops* the bus
- Let's watch a memory location change value from **X** to **X'**




Write-update protocol


3) Threads have identical views of memory

- Two possible solutions, **in hardware**
- In both cases, cache becomes active and monitors or *snoops* the bus
- Let's watch a memory location change value from **X** to **X'**

Write-invalidate protocol




Write-update protocol




3) Threads have identical views of memory

- Two possible solutions, **in hardware**
- In both cases, cache becomes active and monitors or *snoops* the bus
- Let's watch a memory location change value from **X** to **X'**

Write-invalidate protocol



Write-update protocol






Keeping the caches coherent in an SMP

- A. ...is the responsibility of the user program
- B. ...is the responsibility of the hardware
- C. ...is the responsibility of the operating system
- D. ...is impossible
- E. ...is why we don't allow caches in SMP systems



Keeping the caches coherent in an SMP

- A. ...is the responsibility of the user program
- B. ...is the responsibility of the hardware
- C. ...is the responsibility of the operating system
- D. ...is impossible
- E. ...is why we don't allow caches in SMP systems




Summary

- Page tables in shared memory
 - Set up by the OS
 - Used by the hardware
- TLB consistency in software by the OS
 - Hardware brings PTE into the TLB from the PT
 - Page replacement algorithm changes the PT and does the TLB shoot-down
- Synchronized atomicity
 - Test-and-set instruction serialized by the shared bus
 - Atomic read-modify-write transaction
- Cache coherence in hardware
 - Invalidation based or update based

Roadmap






- Programmed IO and DMA
 - Chapter 10 (10.1 – 10.7)
- Networking
 - Chapter 13

Beyond the processor and memory



Anything else??


Beyond the processor and memory




Anything else??

A lot of peripheral devices!


Communication: CPU and I/O Devices



Communication: CPU and I/O Devices




Communication: CPU and I/O Devices

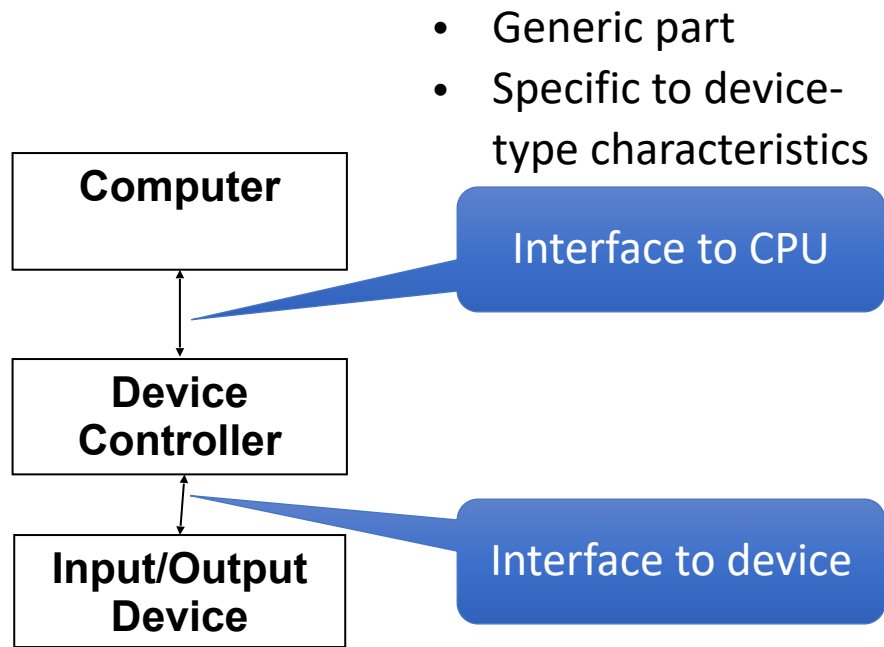


Communication: CPU and I/O Devices


- Generic part



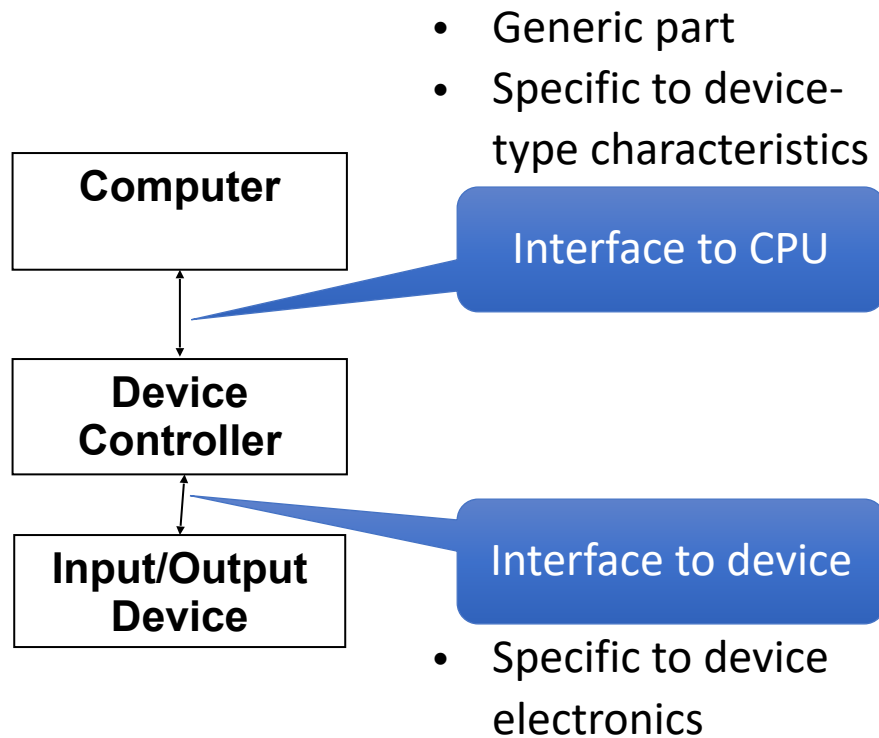
Communication: CPU and I/O Devices



Communication: CPU and I/O Devices




Communication: CPU and I/O Devices



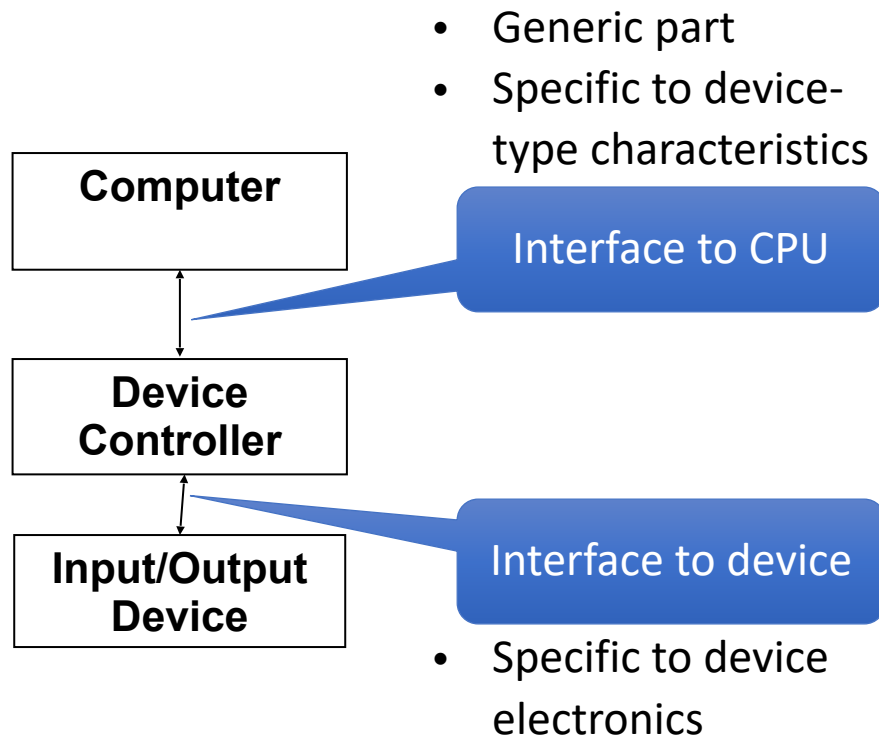
- How does the CPU "talk" to the controller
 - Memory-mapped I/O
 - Load/store instructions
 - (Alternative is special I/O instructions)

Communication: CPU and I/O Devices

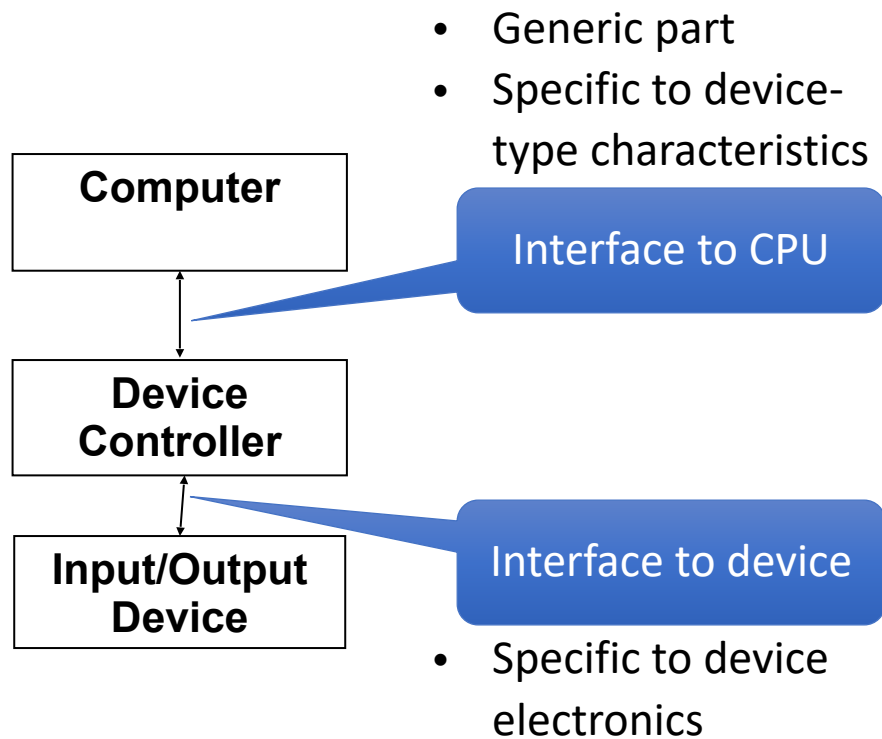


- How does the CPU "talk" to the controller
 - Memory-mapped I/O
 - Load/store instructions
 - (Alternative is special I/O instructions)
- How is data movement effected?
 - Slow speed
 - Programmed I/O
 - High speed
 - DMA & interrupts

Communication: CPU and I/O Devices




Communication: CPU and I/O Devices




- What commands do we issue to devices?
 - Camera
 - Display
 - Audio
 - ...

A keyboard device



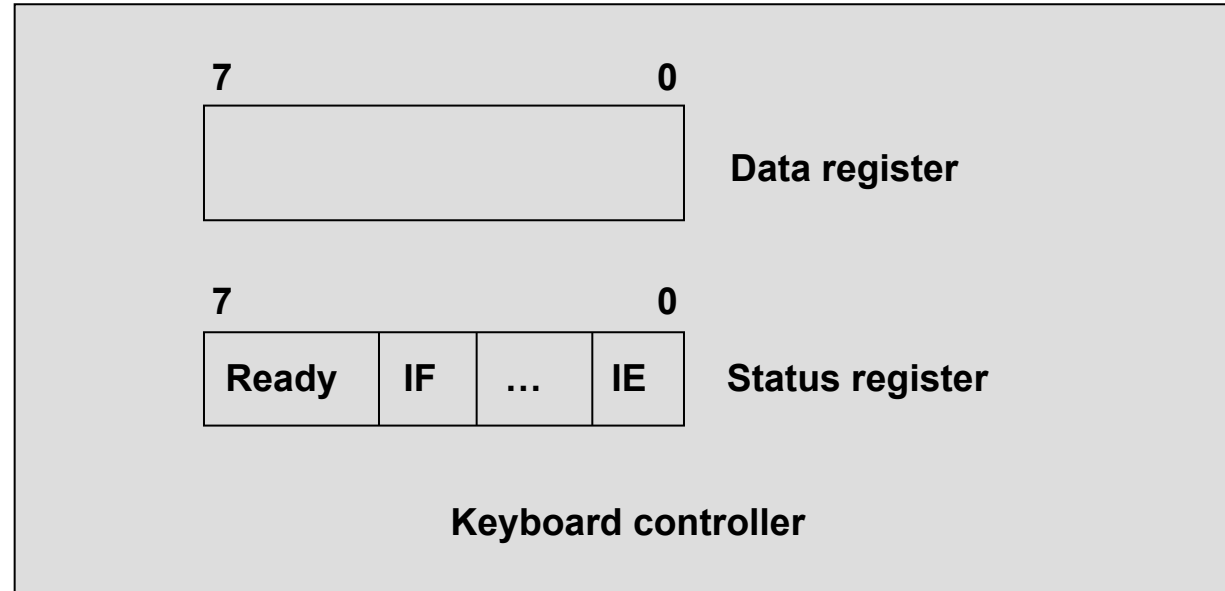
A keyboard device

Controller



A keyboard device

Controller




Device



A keyboard device

Controller



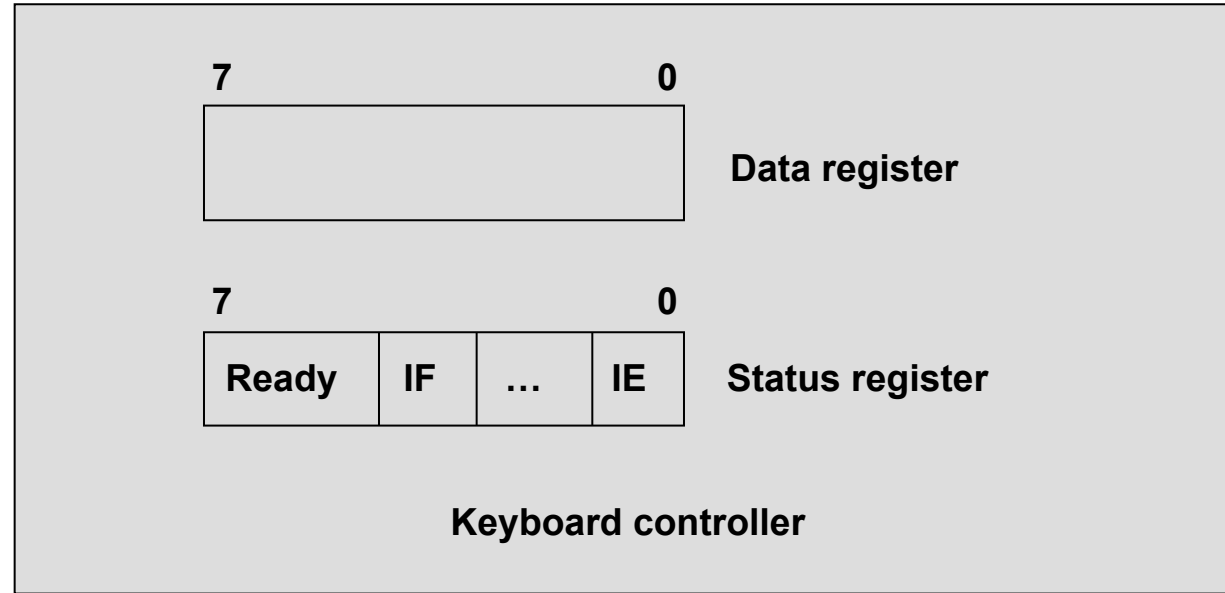
- Slow speed device

Device



A keyboard device

Controller



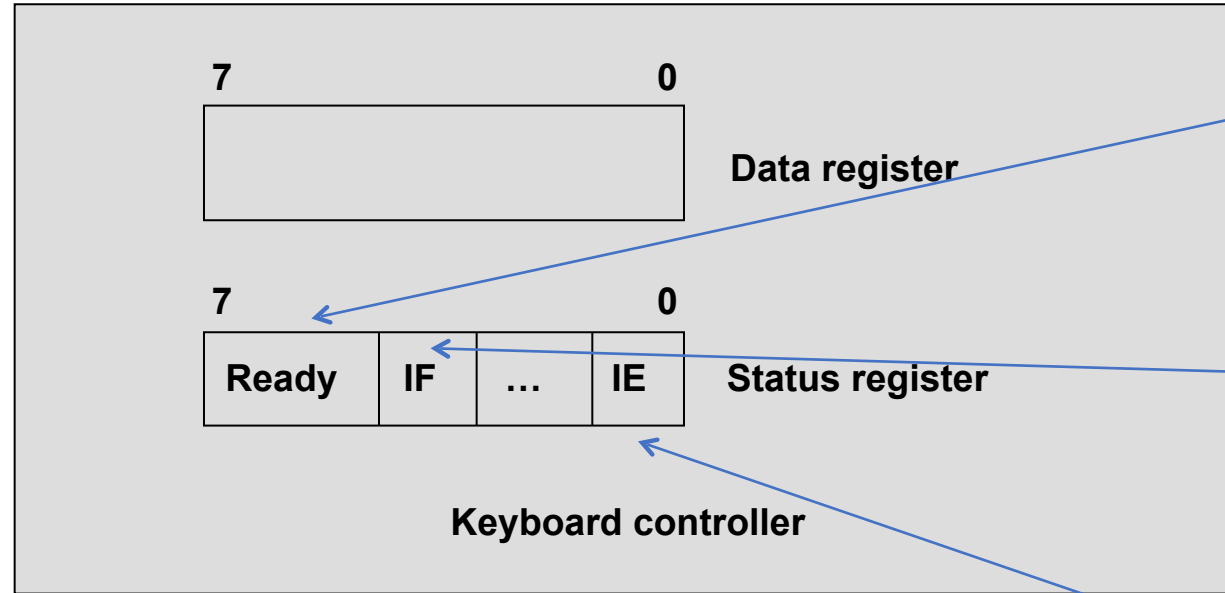
- Slow speed device
- Service it with programmed I/O (PIO)

Device



A keyboard device

Controller



- Ready: there's a byte waiting in the data register

- IF: the device has produced data – controller ready to signal an interrupt

- IE: If 1 interrupt when Ready

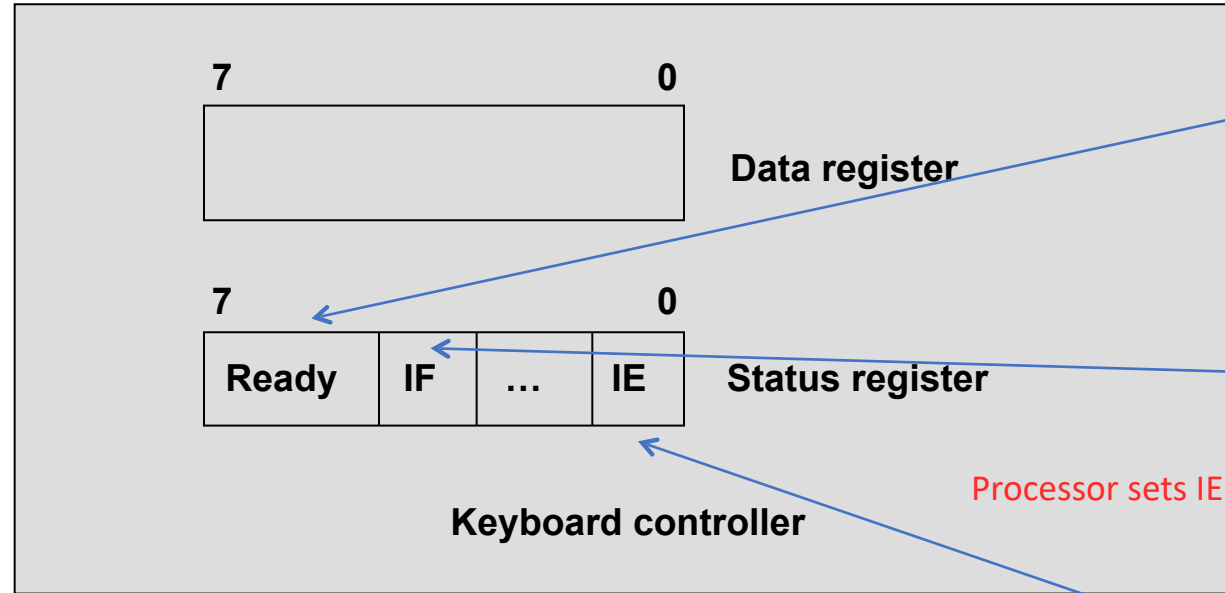
- Slow speed device
- Service it with programmed I/O (PIO)

Device



A keyboard device

Controller



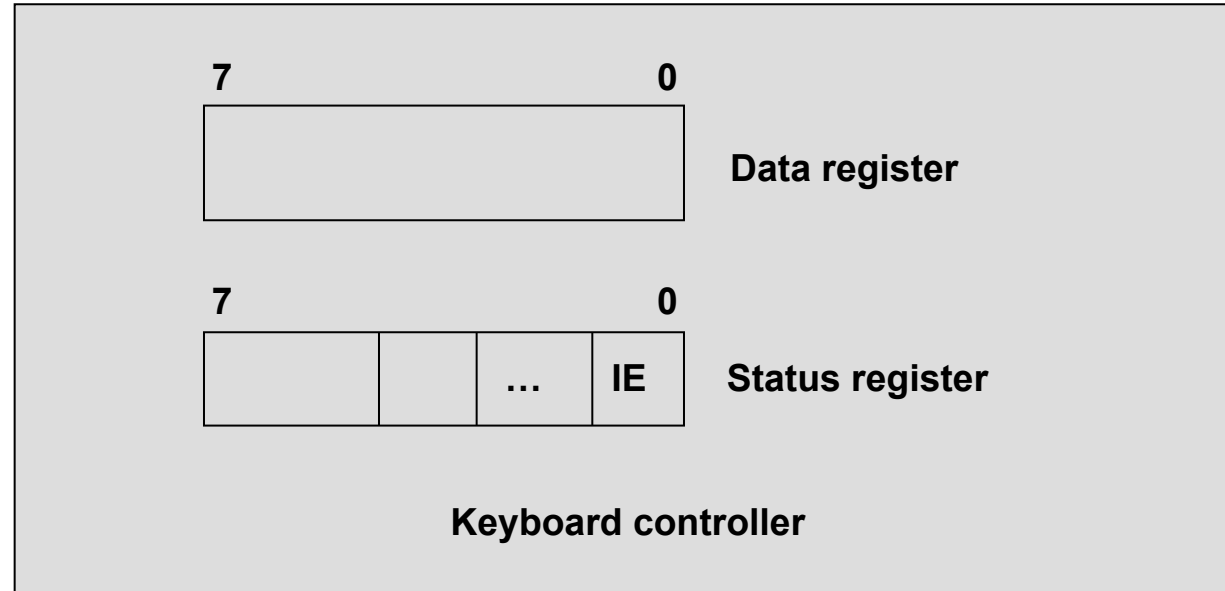
- Ready: there's a byte waiting in the data register
- IF: the device has produced data – controller ready to signal an interrupt
- IE: If 1 interrupt when Ready

Device




- Slow speed device
- Service it with programmed I/O (PIO)

A keyboard device



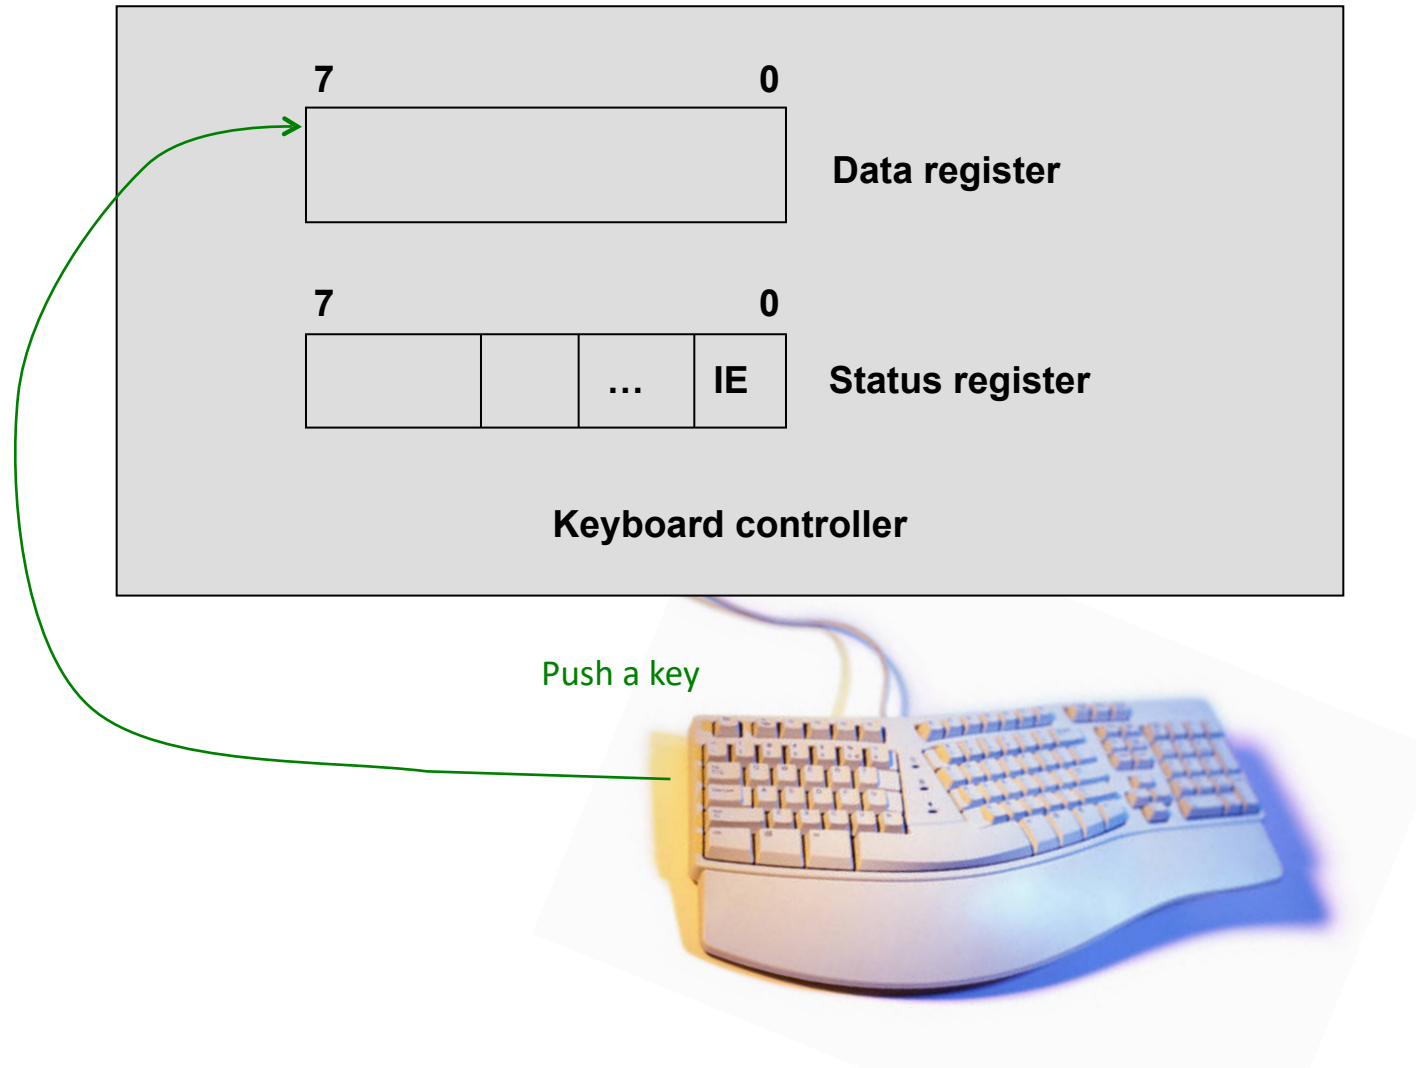
A keyboard device



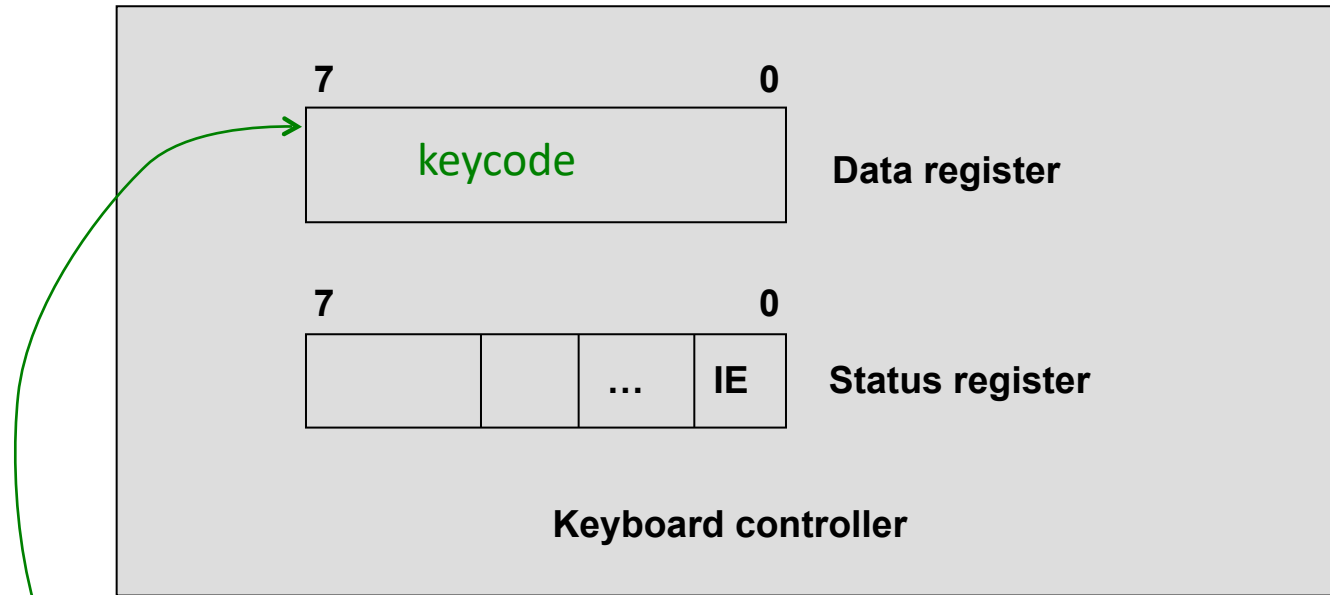
Push a key



A keyboard device



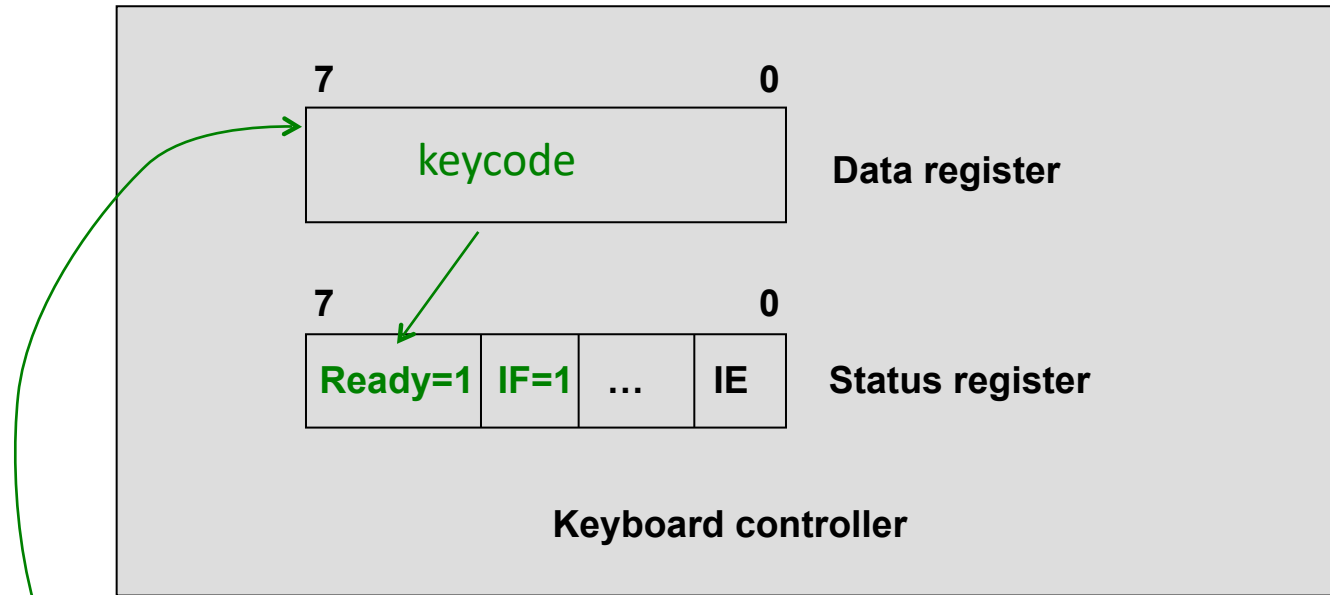
A keyboard device



Push a key




A keyboard device



Push a key



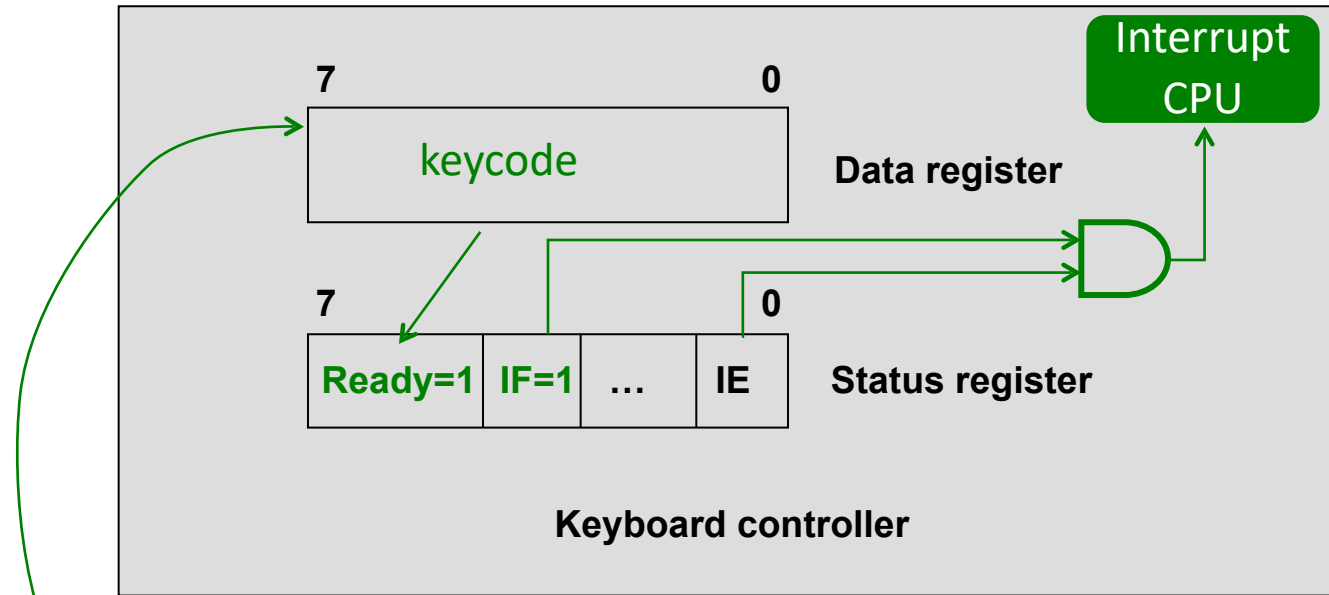
A keyboard device



Push a key



A keyboard device

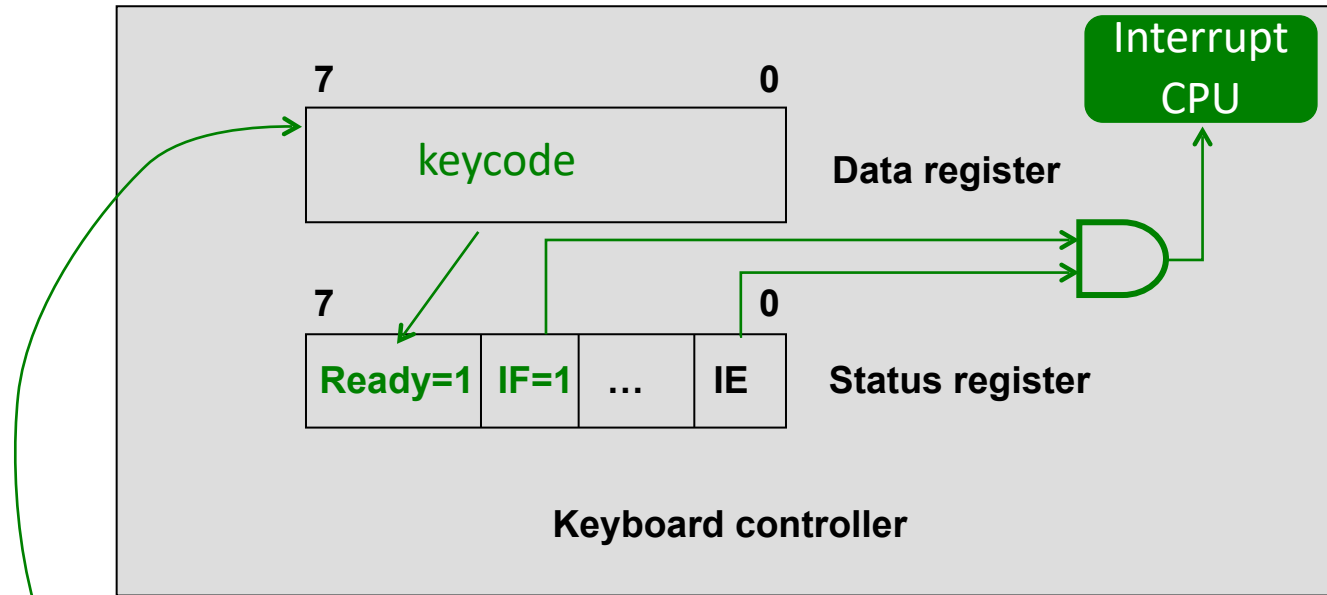


- Commands from the CPU:

Push a key



A keyboard device




- Commands from the CPU:
- Set IE

Push a key



A keyboard device




- Commands from the CPU:
- Set IE
- Check Ready equals 1

Push a key



A keyboard device

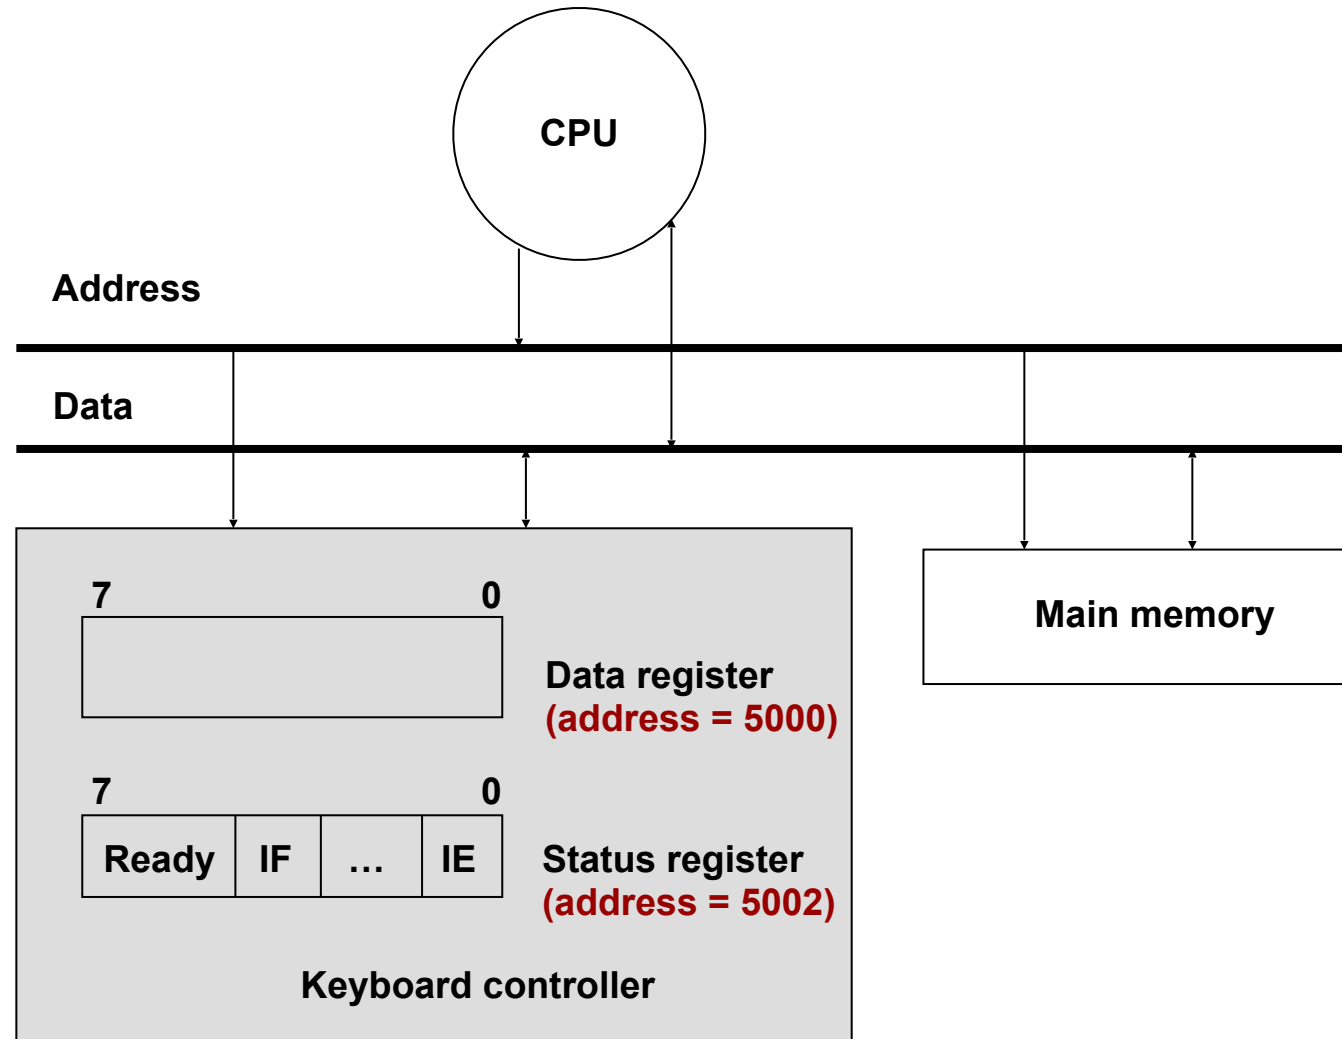


- Commands from the CPU:
- Set IE
- Check Ready equals 1
- Load **keycode** from data register

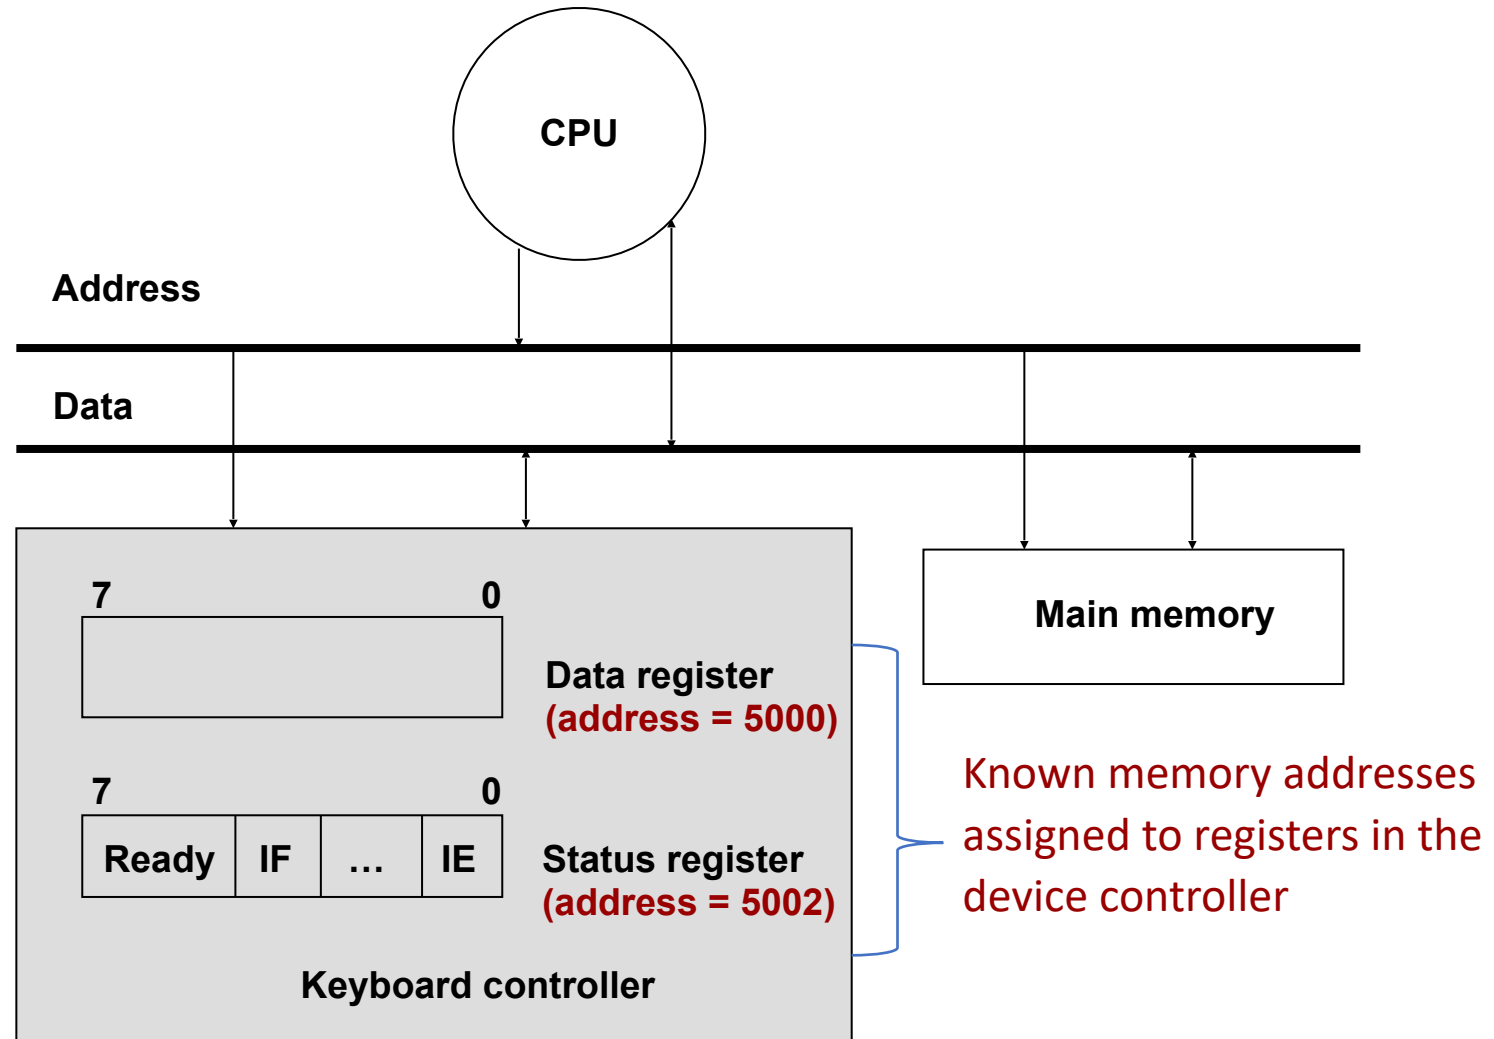
Push a key



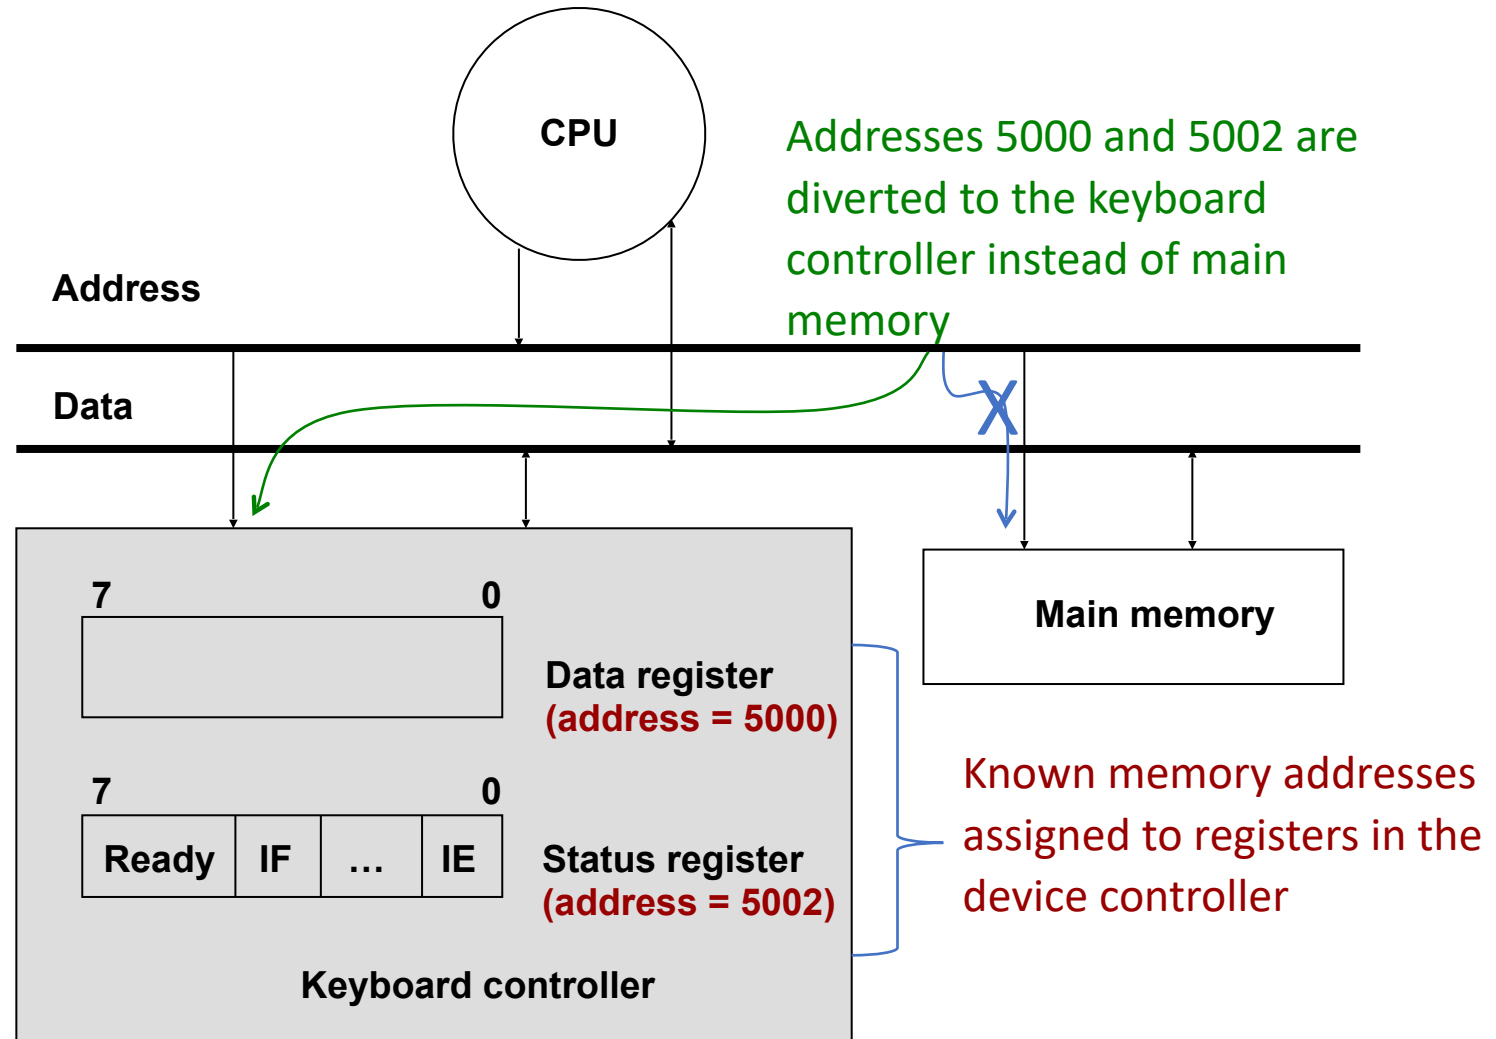
Memory mapped I/O



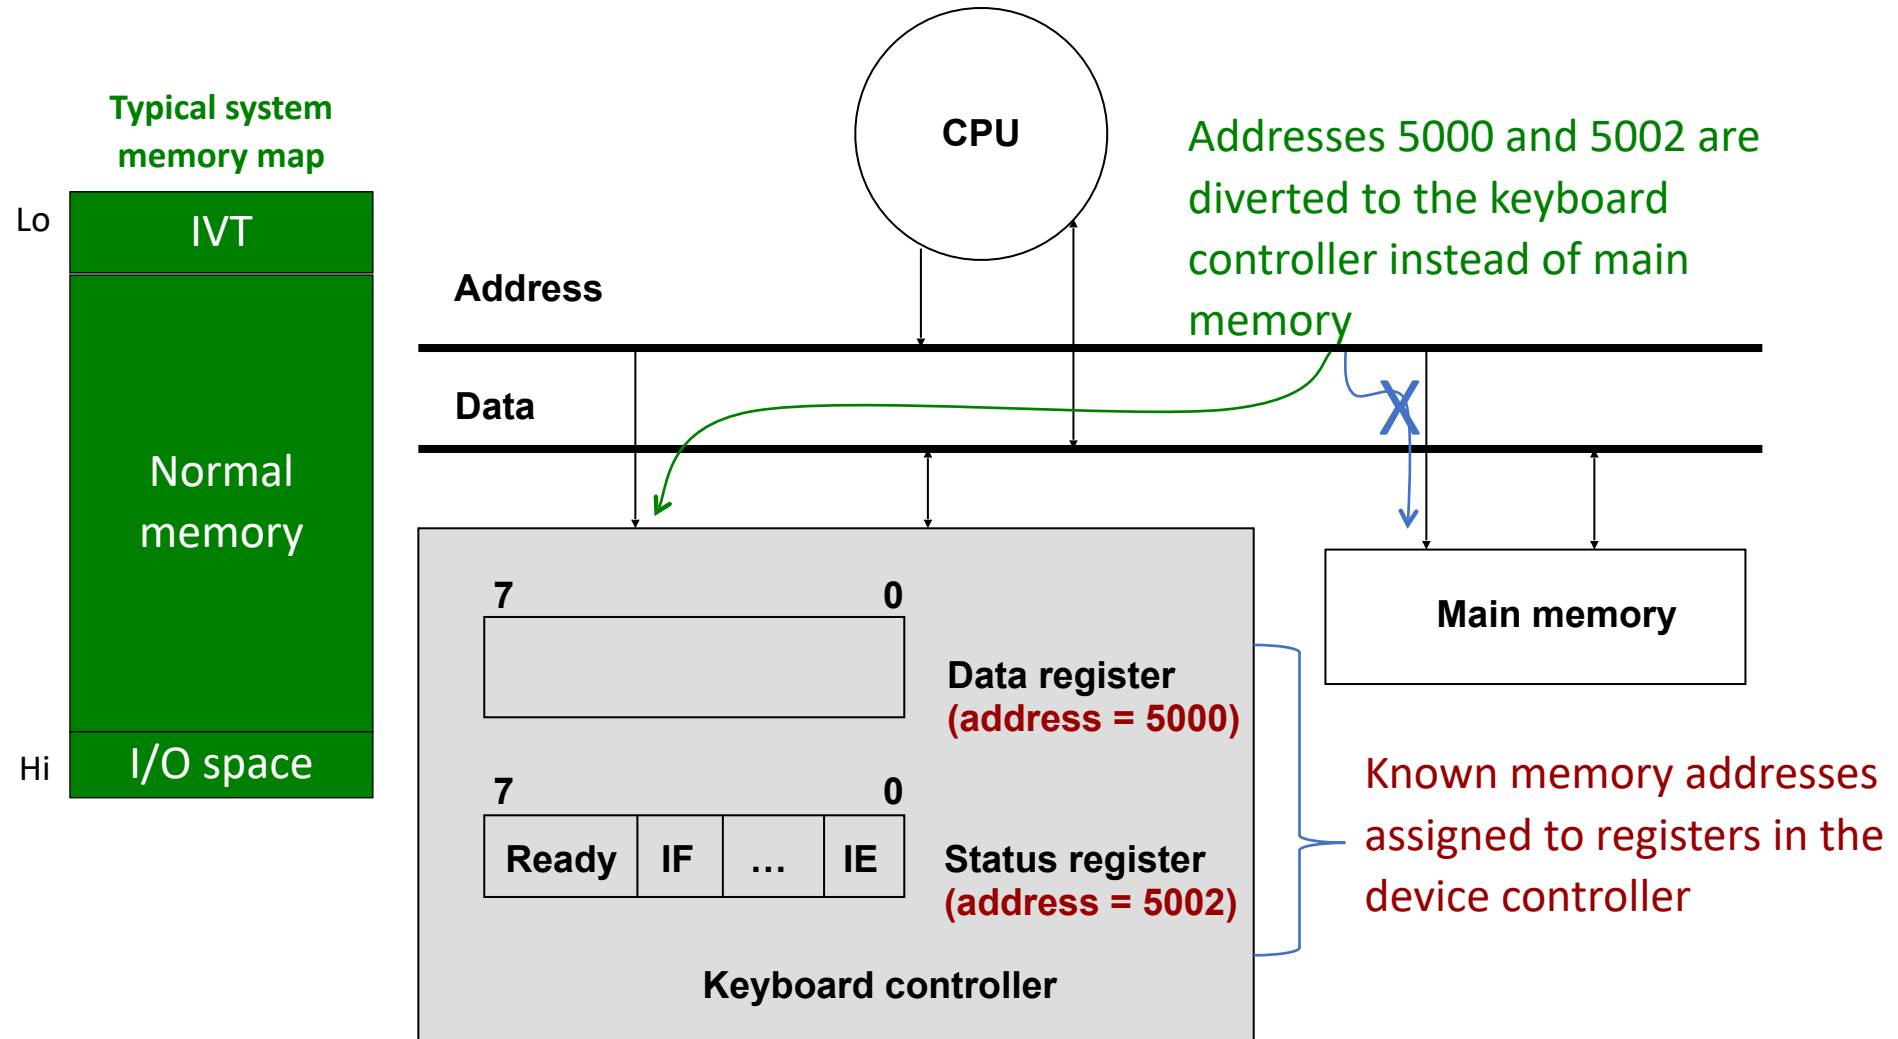
Memory mapped I/O



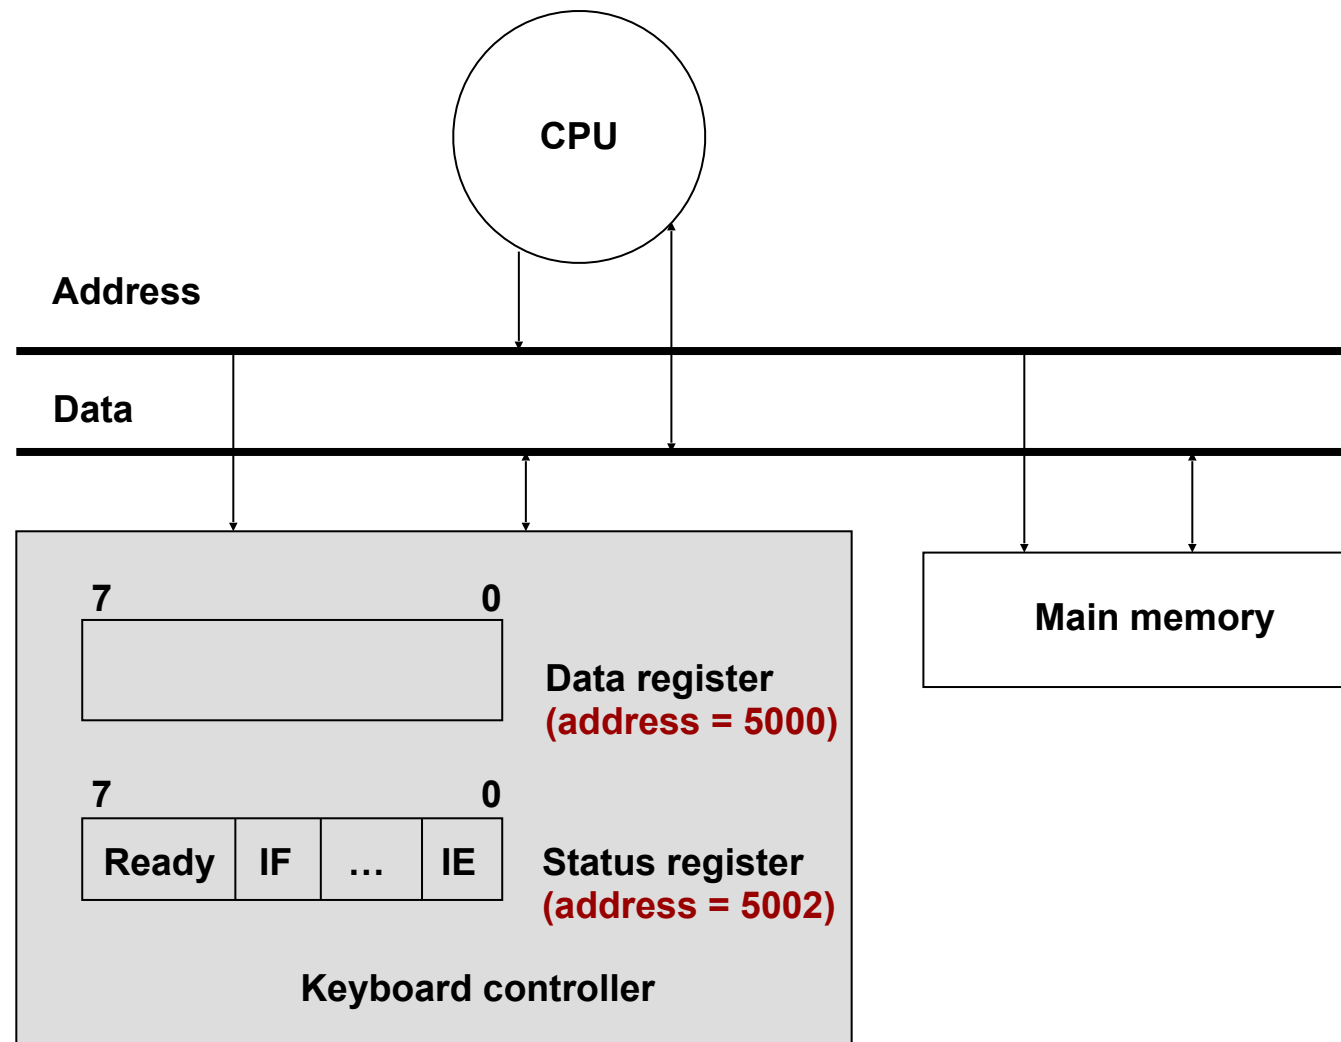
Memory mapped I/O



Memory mapped I/O




Memory mapped I/O



Memory mapped I/O

Commands?




Memory mapped I/O

Commands?

Check for new data

LD RI,mem[5002]



Memory mapped I/O

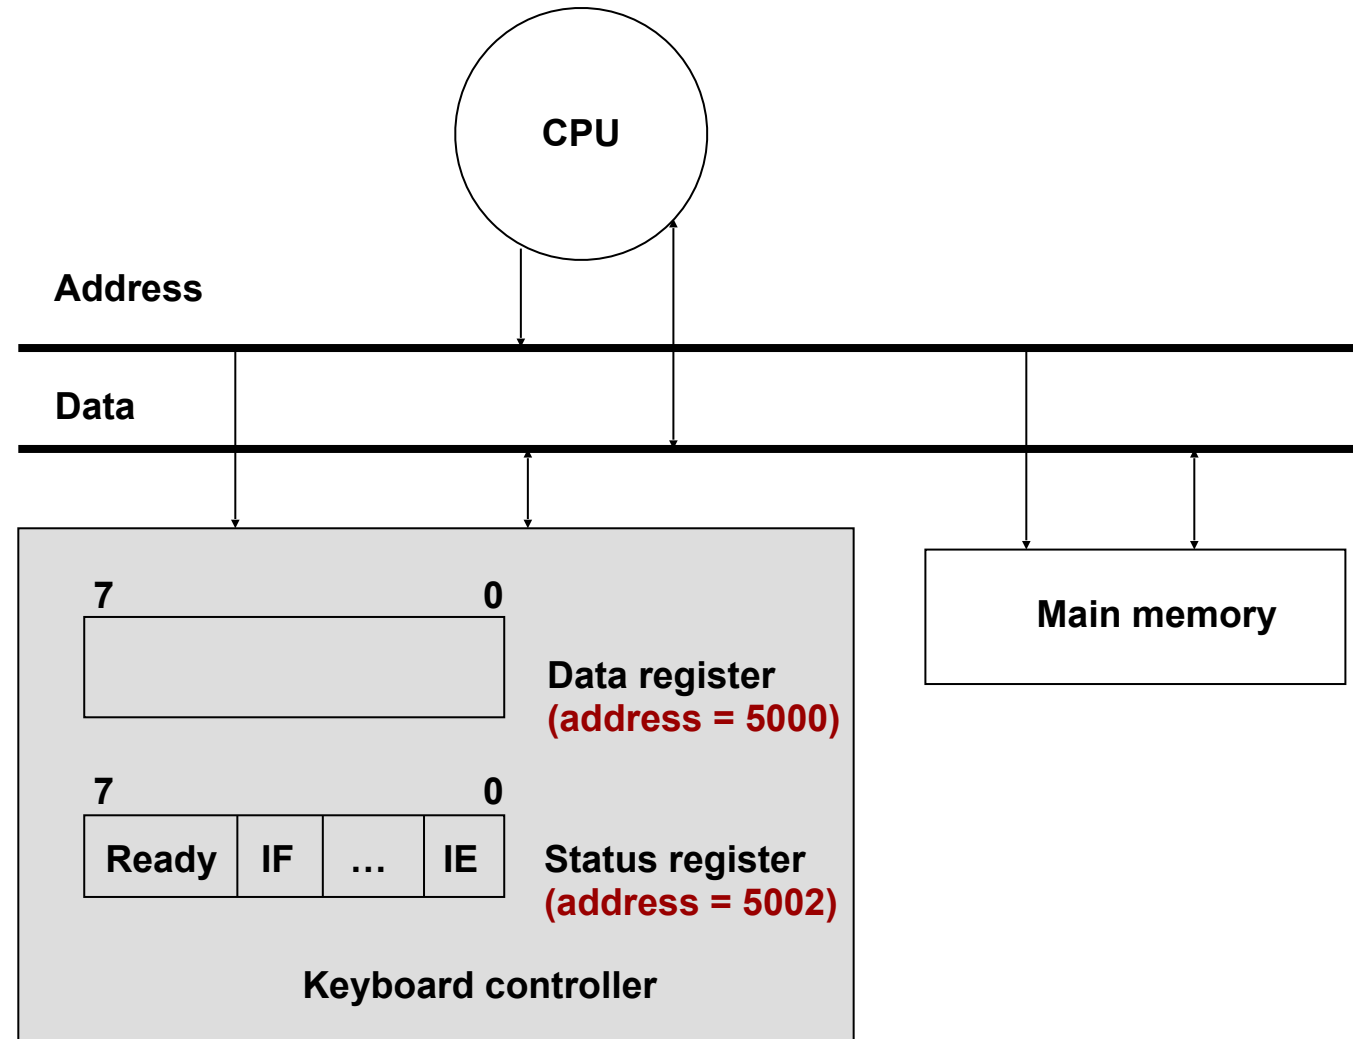
Commands?

Check for new data

LD R1,mem[5002]

Read data

LD R2,mem[5000]



Memory mapped I/O

Commands?

Check for new data


LD R1,mem[5002]

Read data

LD R2,mem[5000]

Set IE

ST #1,mem[5002]



Programmed I/O (PIO)

Programmed I/O (PIO)

- Slow speed devices (Keyboard, mouse, etc.)

Programmed I/O (PIO)

- Slow speed devices (Keyboard, mouse, etc.)
- CPU executes program to move data

Programmed I/O (PIO)

- Slow speed devices (Keyboard, mouse, etc.)
 - CPU executes program to move data
 - Often uses polling
- ```
X: LD R1, statusreg ; Get the status register
 BRZP X ; if high bit is 0, branch back one
 LD R2, datareg ; Load the data into R2
 ST #0, statusreg ; Clear the status register
```

LD usually clears the ready bit as a side effect of reading the data register



# Programmed I/O (PIO)

---

- Slow speed devices (Keyboard, mouse, etc.)
- CPU executes program to move data

- Often uses polling

X: LD        R1, statusreg        ; Get the status register  
     BRZP    X                      ; if high bit is 0, branch back one  
     LD       R2, datareg         ; Load the data into R2  
     ST       #0, statusreg       ; Clear the status register

LD usually clears the ready bit as a side effect of reading the data register

- Busy waiting – wastes processor resources

# Interrupt-driven I/O

---

- Slow speed devices (Keyboard, mouse, etc.)
- CPU executes program for moving data
- Can be interrupt driven

ST    #1, statusreg ; Set the IE bit

Upon interrupt, handler code is executed:

LD    R2, datareg    ; Load the data into R2

# High speed devices

---

# High speed devices

---

- Once a command is given by the CPU to the controller, data comes in continually
  - Streaming

# High speed devices

---

- Once a command is given by the CPU to the controller, data comes in continually
  - Streaming
- PIO has potential for data loss if the CPU doesn't poll or respond to the interrupt quickly enough

# High speed devices

---

- Once a command is given by the CPU to the controller, data comes in continually
  - Streaming
- PIO has potential for data loss if the CPU doesn't poll or respond to the interrupt quickly enough
- Streaming devices move data to/from memory autonomously


# High speed devices

---

- Once a command is given by the CPU to the controller, data comes in continually
  - Streaming
- PIO has potential for data loss if the CPU doesn't poll or respond to the interrupt quickly enough
- Streaming devices move data to/from memory autonomously
- The CPU to controller interface?
  - Convey commands (read/write, IE, etc.)
  - Check status (error, etc.)

# Direct Memory Access - DMA

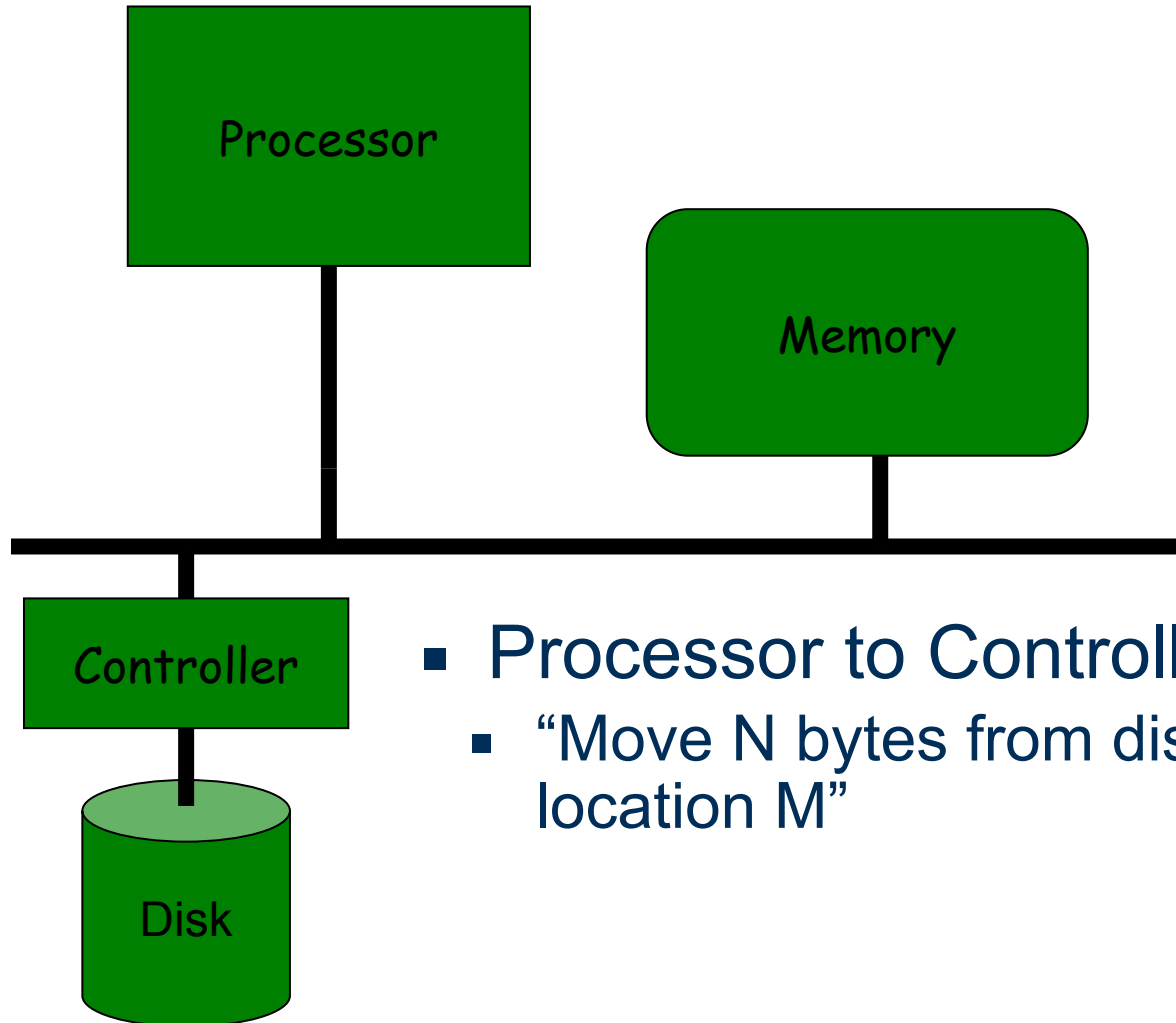
---





# Direct Memory Access - DMA

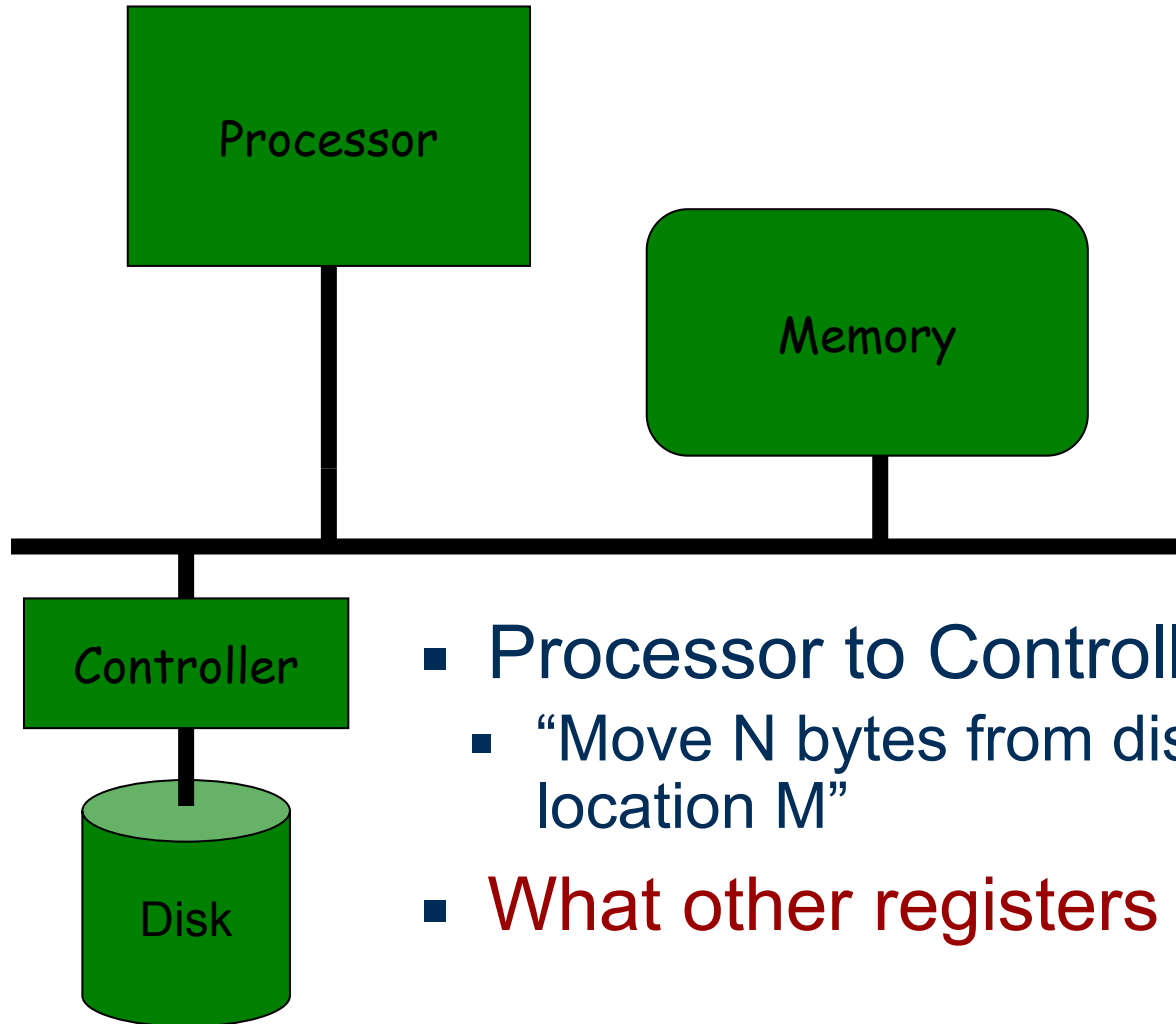
---



- Processor to Controller:
  - “Move N bytes from disk location X to memory location M”

# Direct Memory Access - DMA

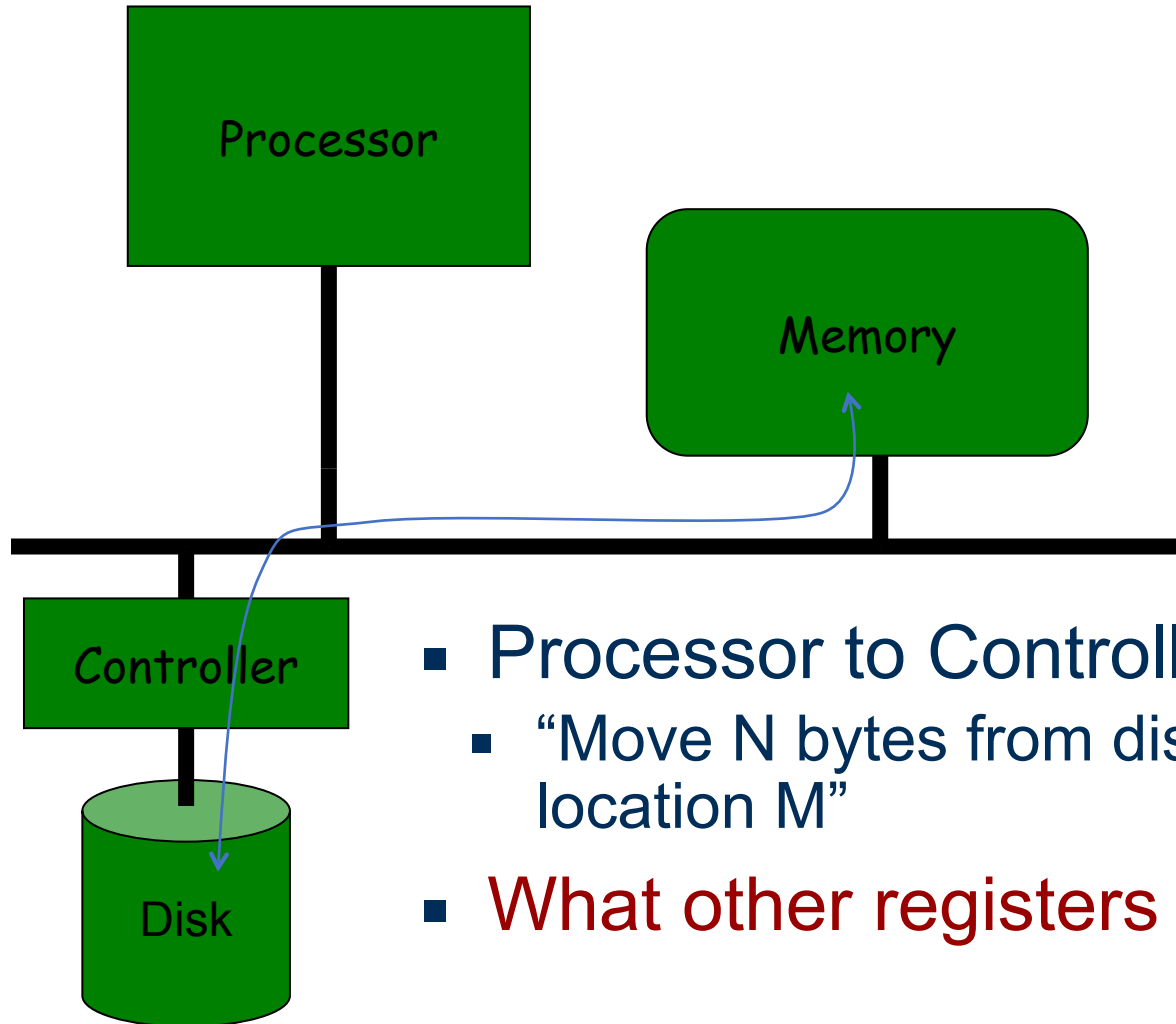
---



- Processor to Controller:
  - “Move N bytes from disk location X to memory location M”
- What other registers do we need?

# Direct Memory Access - DMA


---



- Processor to Controller:
  - “Move N bytes from disk location X to memory location M”
- What other registers do we need?


# DMA registers

---




# DMA registers

---




# DMA registers

---




# DMA registers

---




# DMA registers

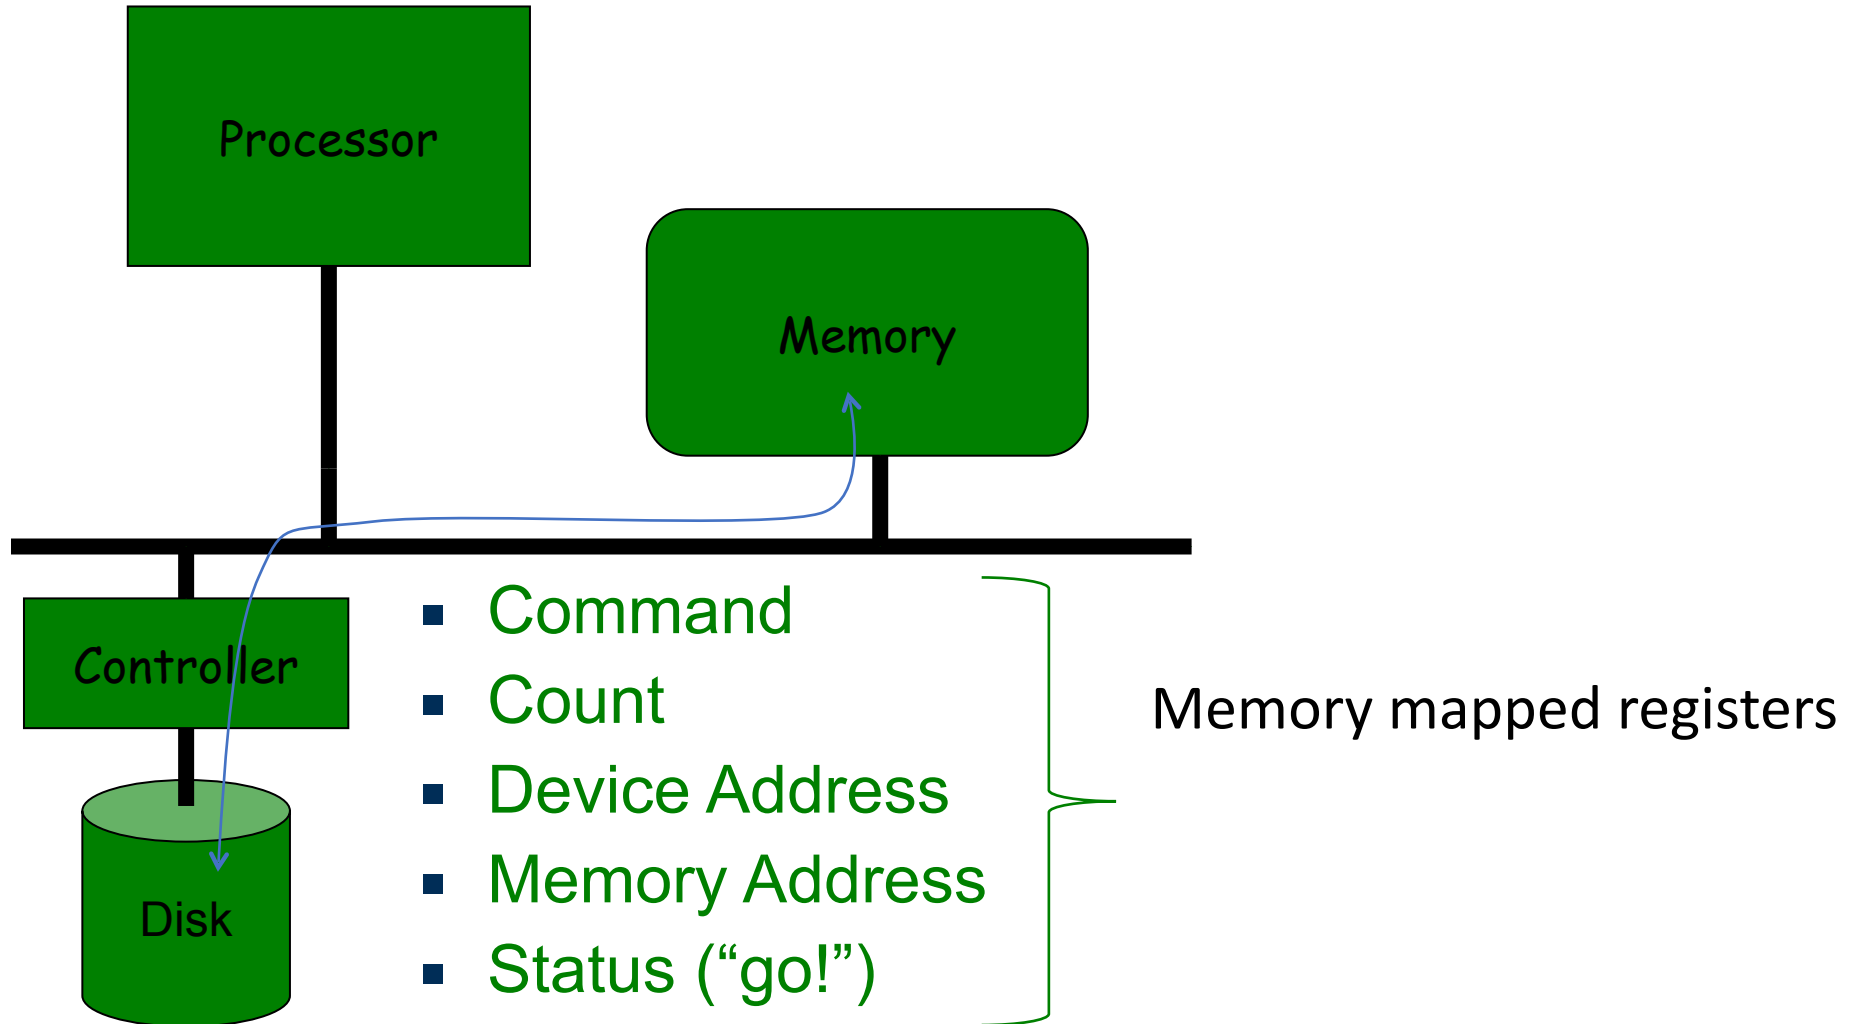




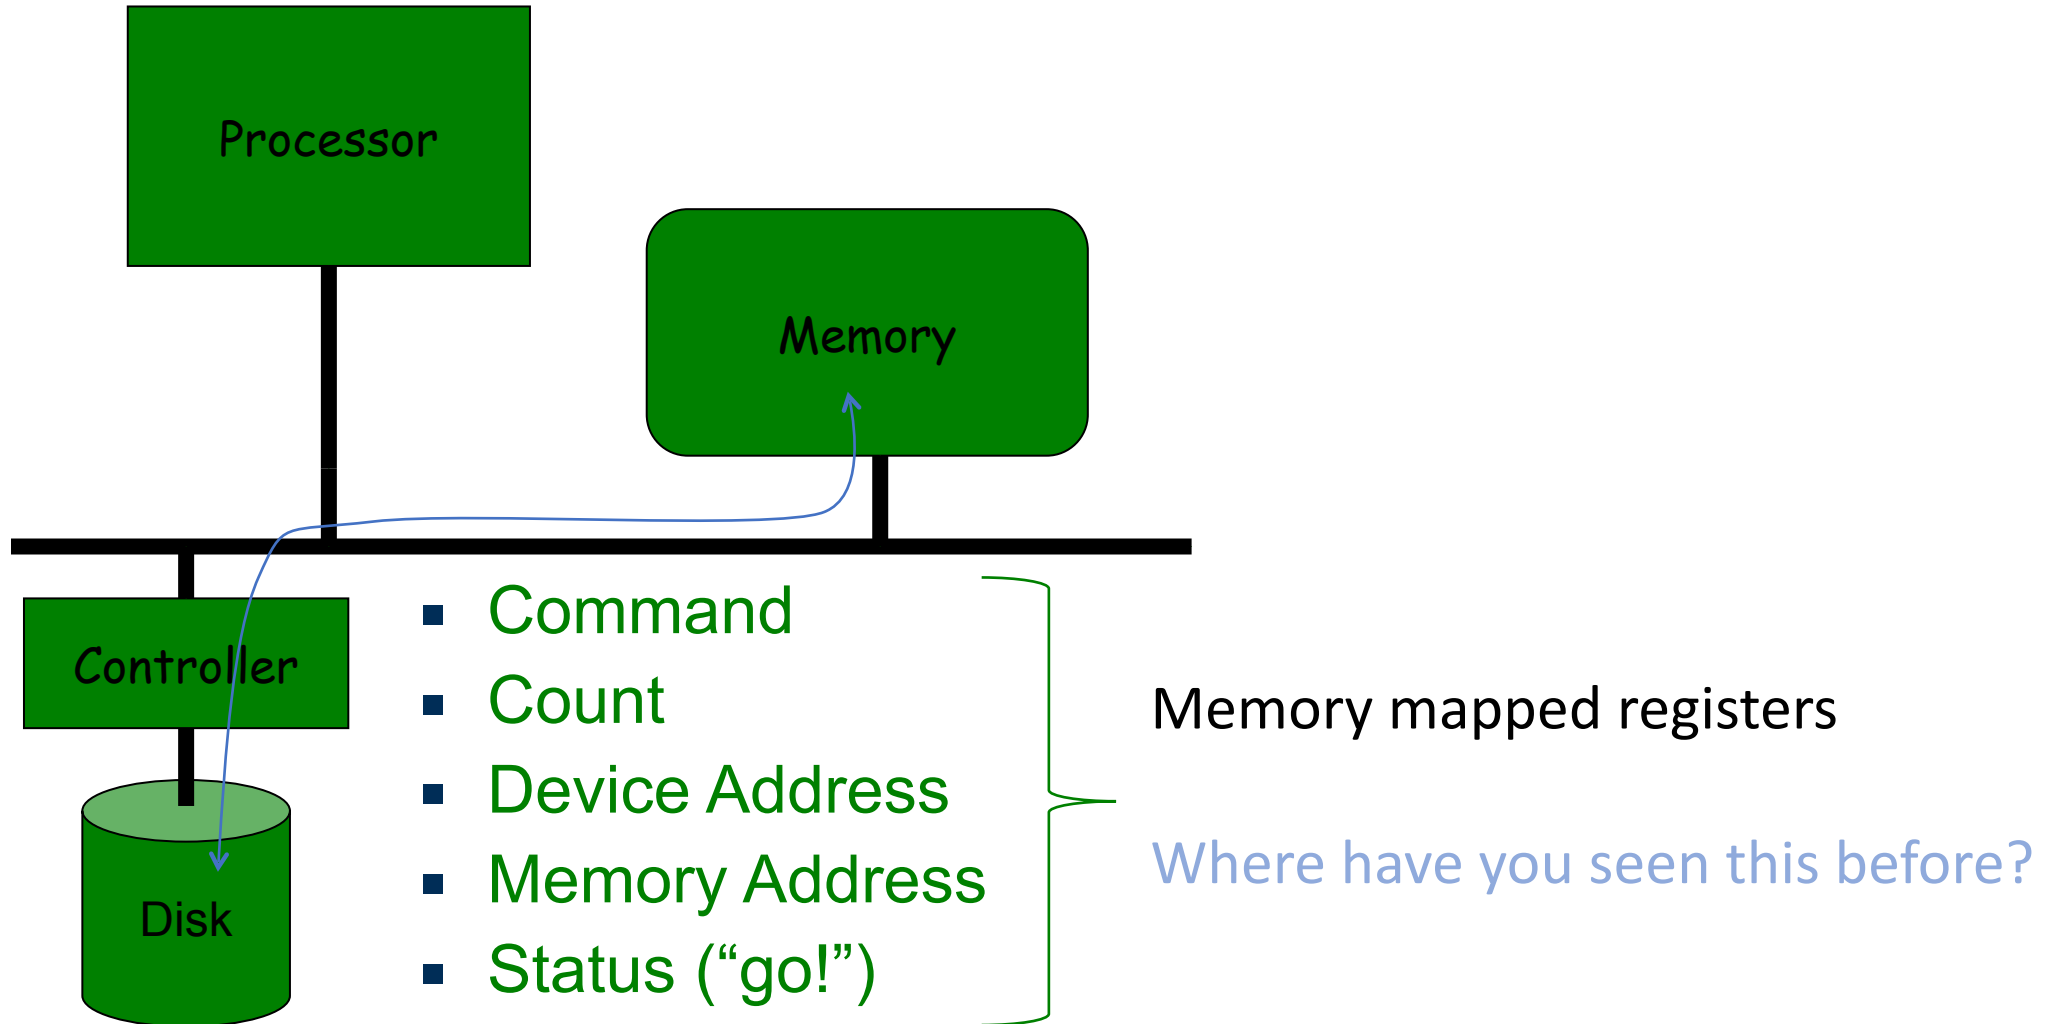
# DMA registers



# DMA registers



# DMA registers



# Program for conveying I/O commands

---

# Program for conveying I/O commands

---

Store N, **count**

Store #block\_num, **device\_addr**

Store #mem\_buf\_addr, **mem\_addr**

Store #write\_command, **command**

Store #l, **status** ; the signal to execute the command

# Program for conveying I/O commands

---

Store N, **count**

Store #block\_num, **device\_addr**

Store #mem\_buf\_addr, **mem\_addr**

Store #write\_command, **command**

Store #I, **status** ; the signal to execute the command

- At this point the CPU's work is done
  - The device controller takes over to do the actual data movement

# Program for conveying I/O commands

---

Store N, **count**

Store #block\_num, **device\_addr**

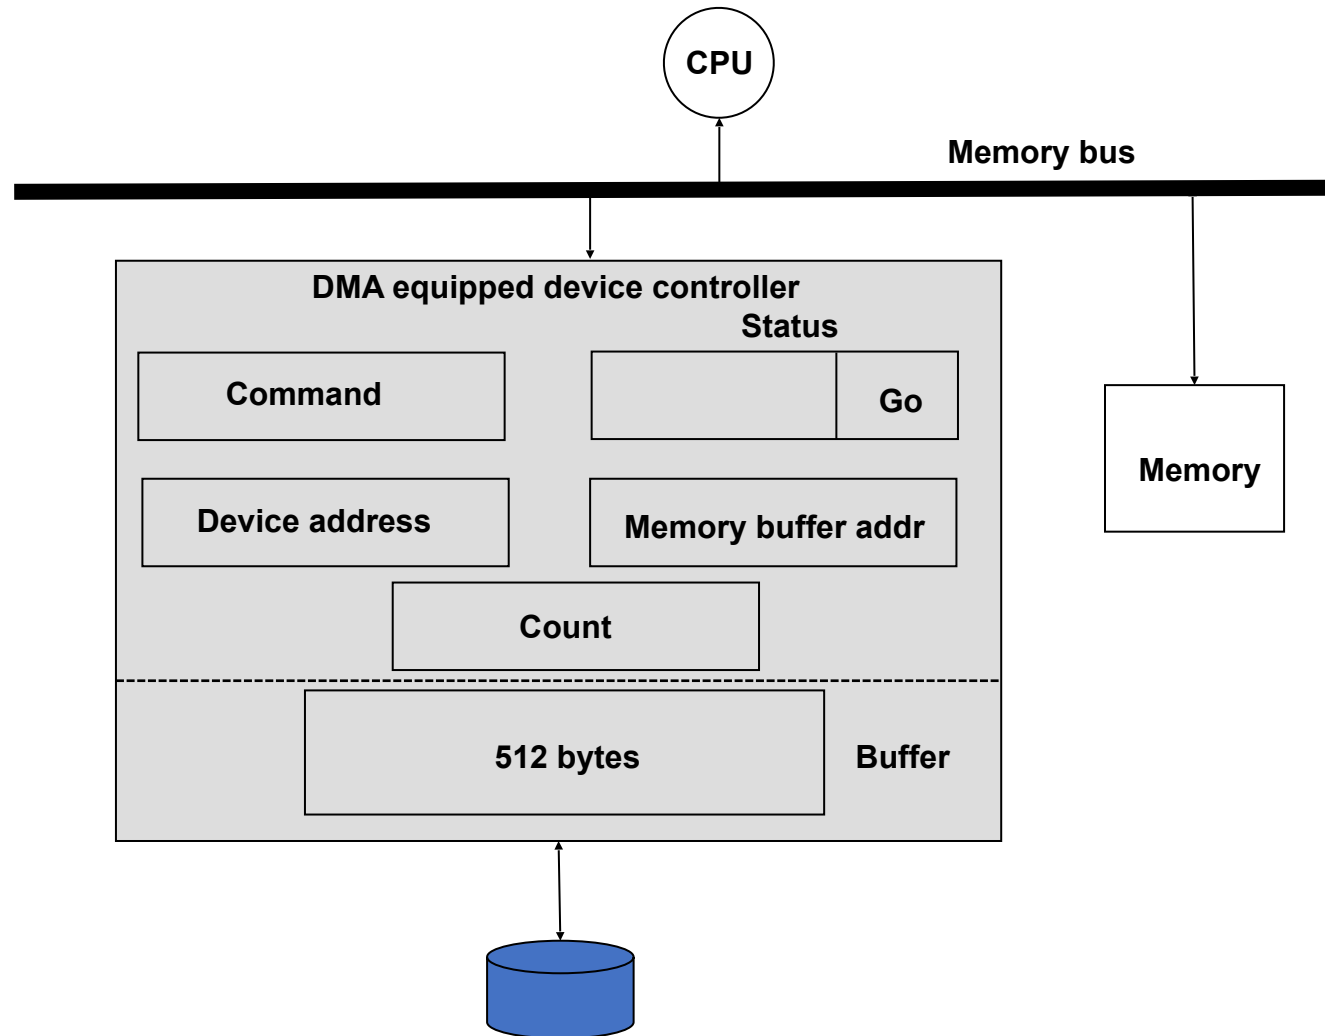
Store #mem\_buf\_addr, **mem\_addr**

Store #write\_command, **command**

Store #I, **status** ; the signal to execute the command

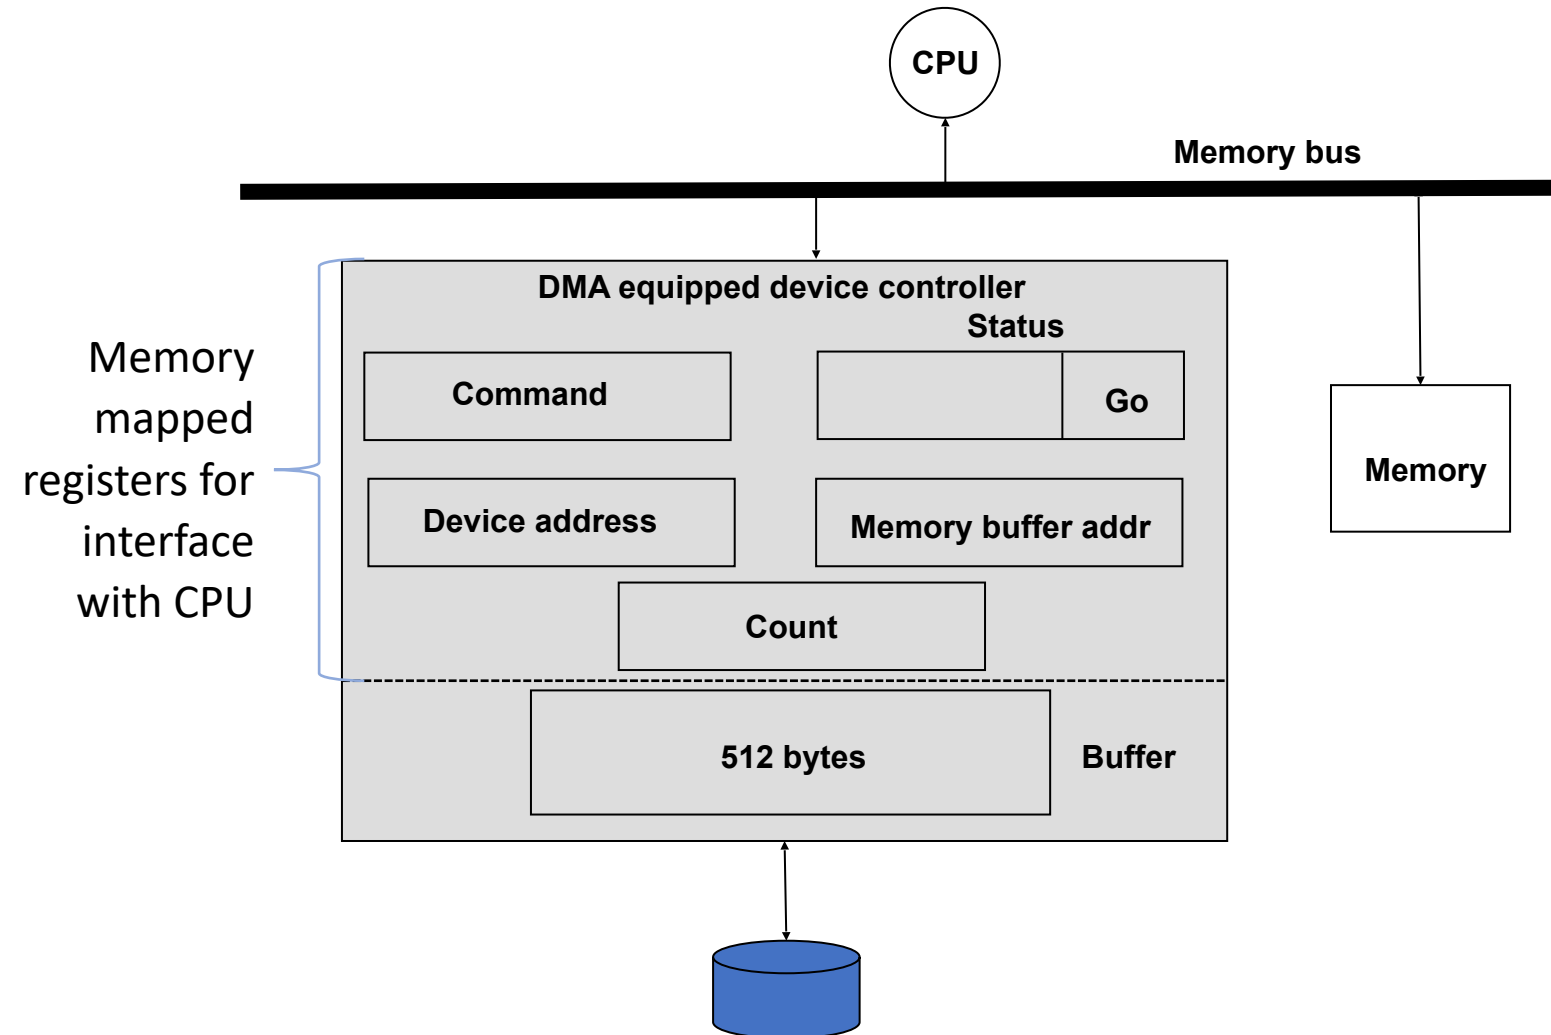
- At this point the CPU's work is done
  - The device controller takes over to do the actual data movement
- Modern CPU chips already contain device controllers for PCI-E devices which share the on-chip memory controllers with the CPU cores

# DMA

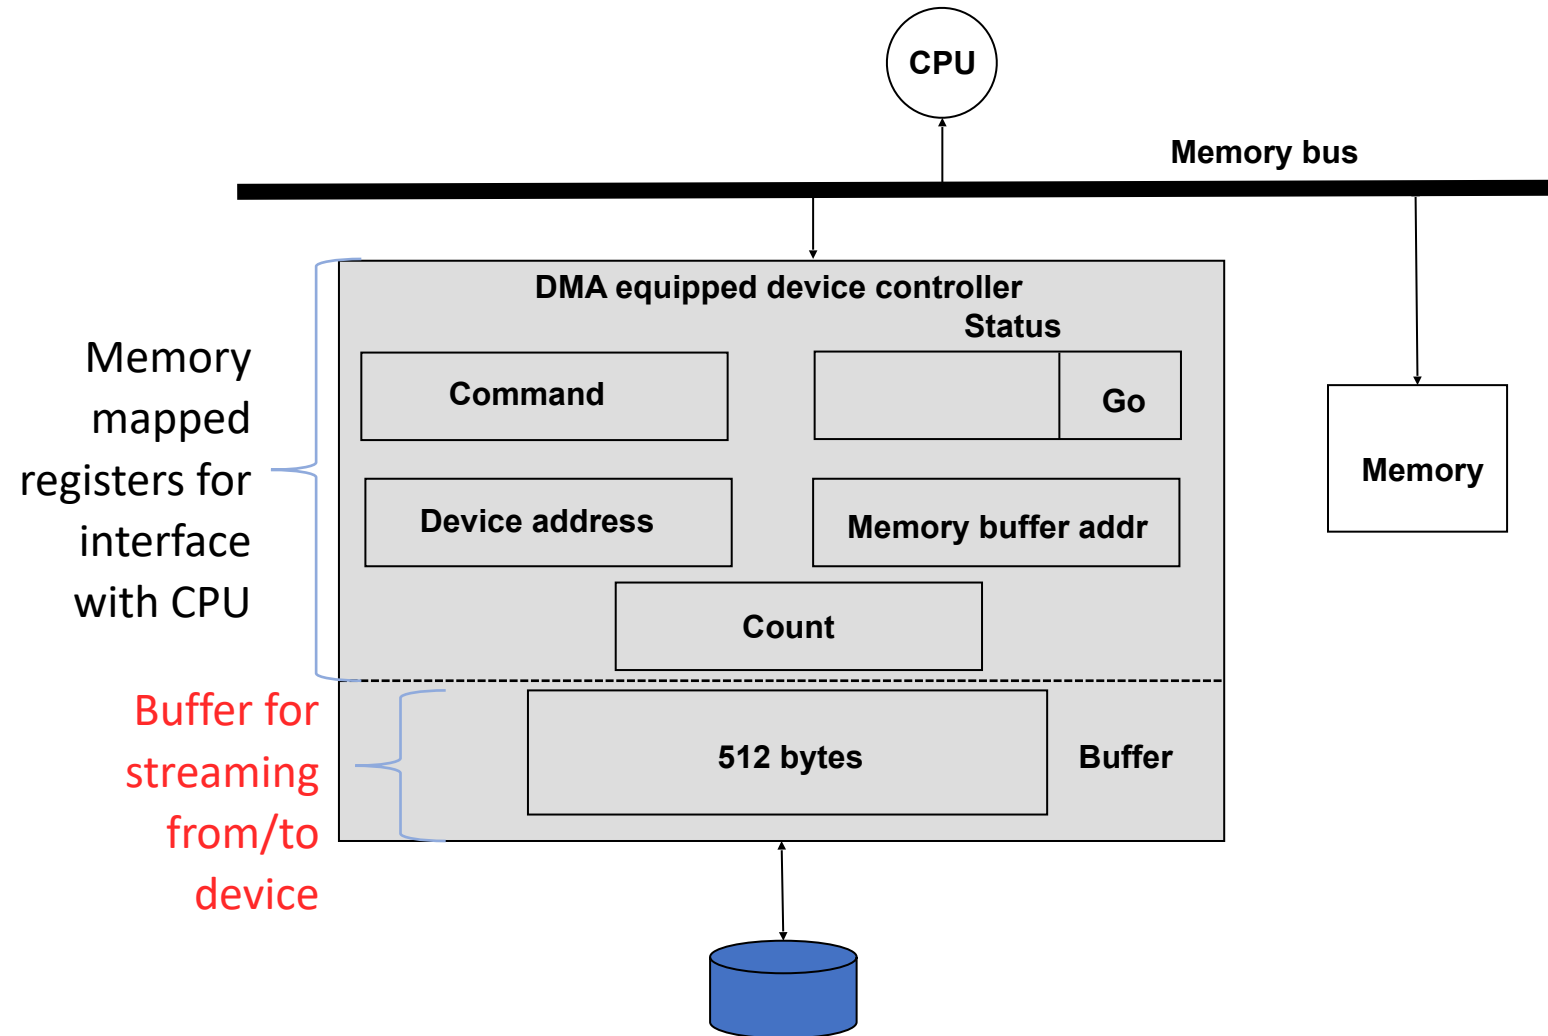




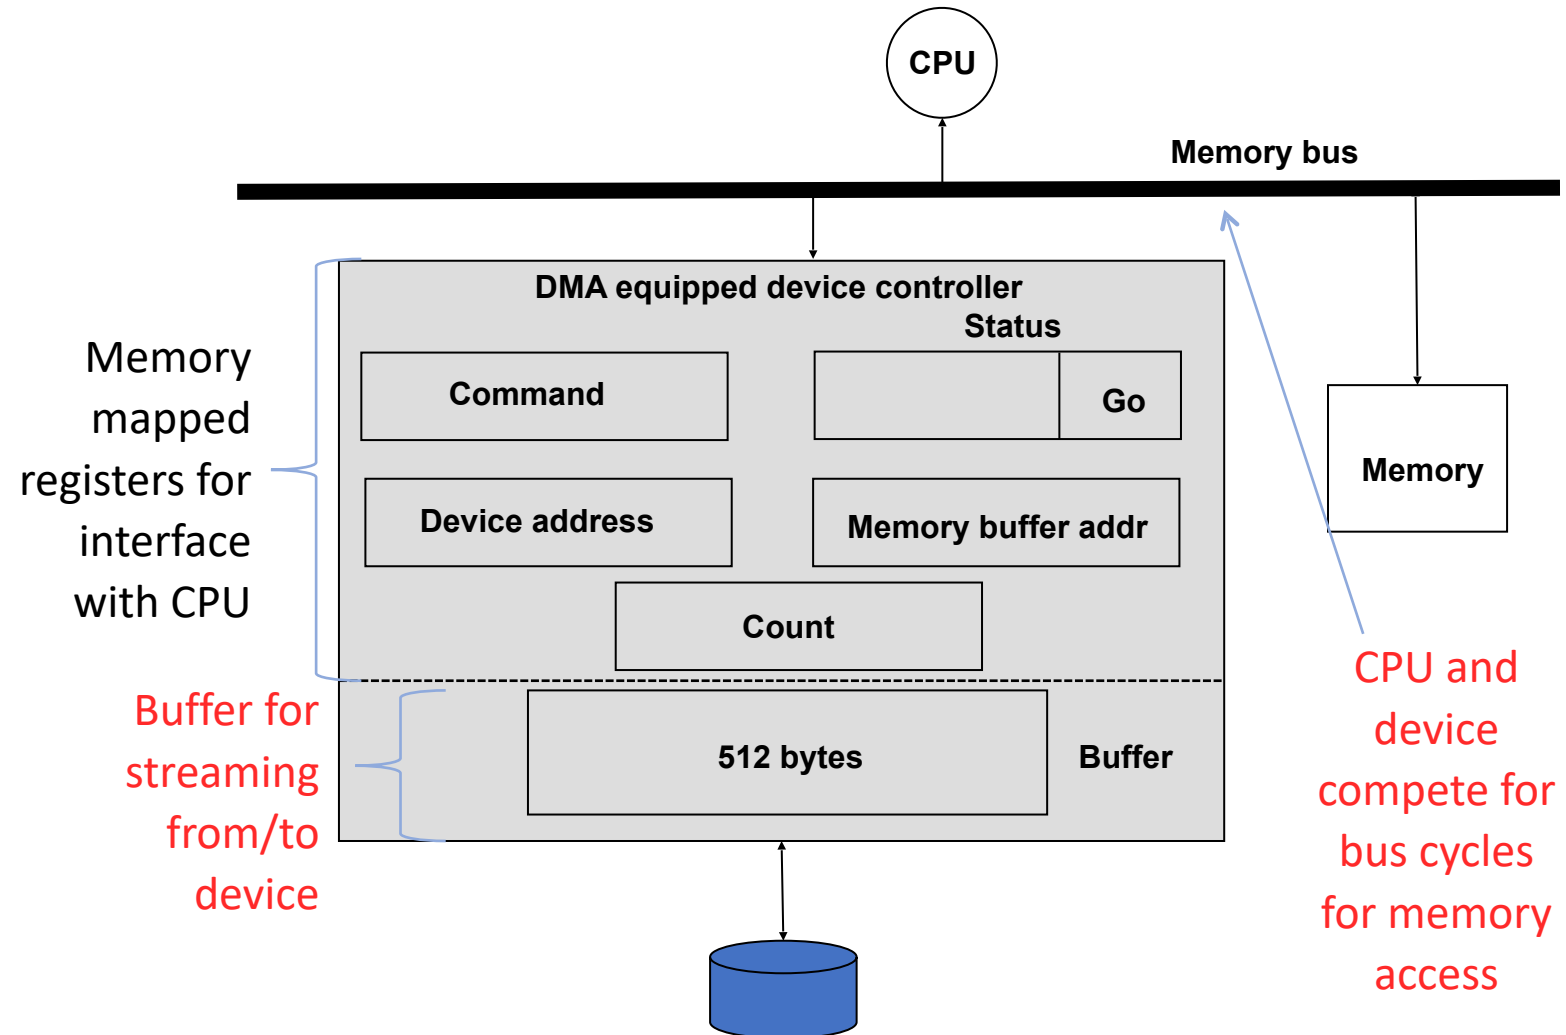
# DMA



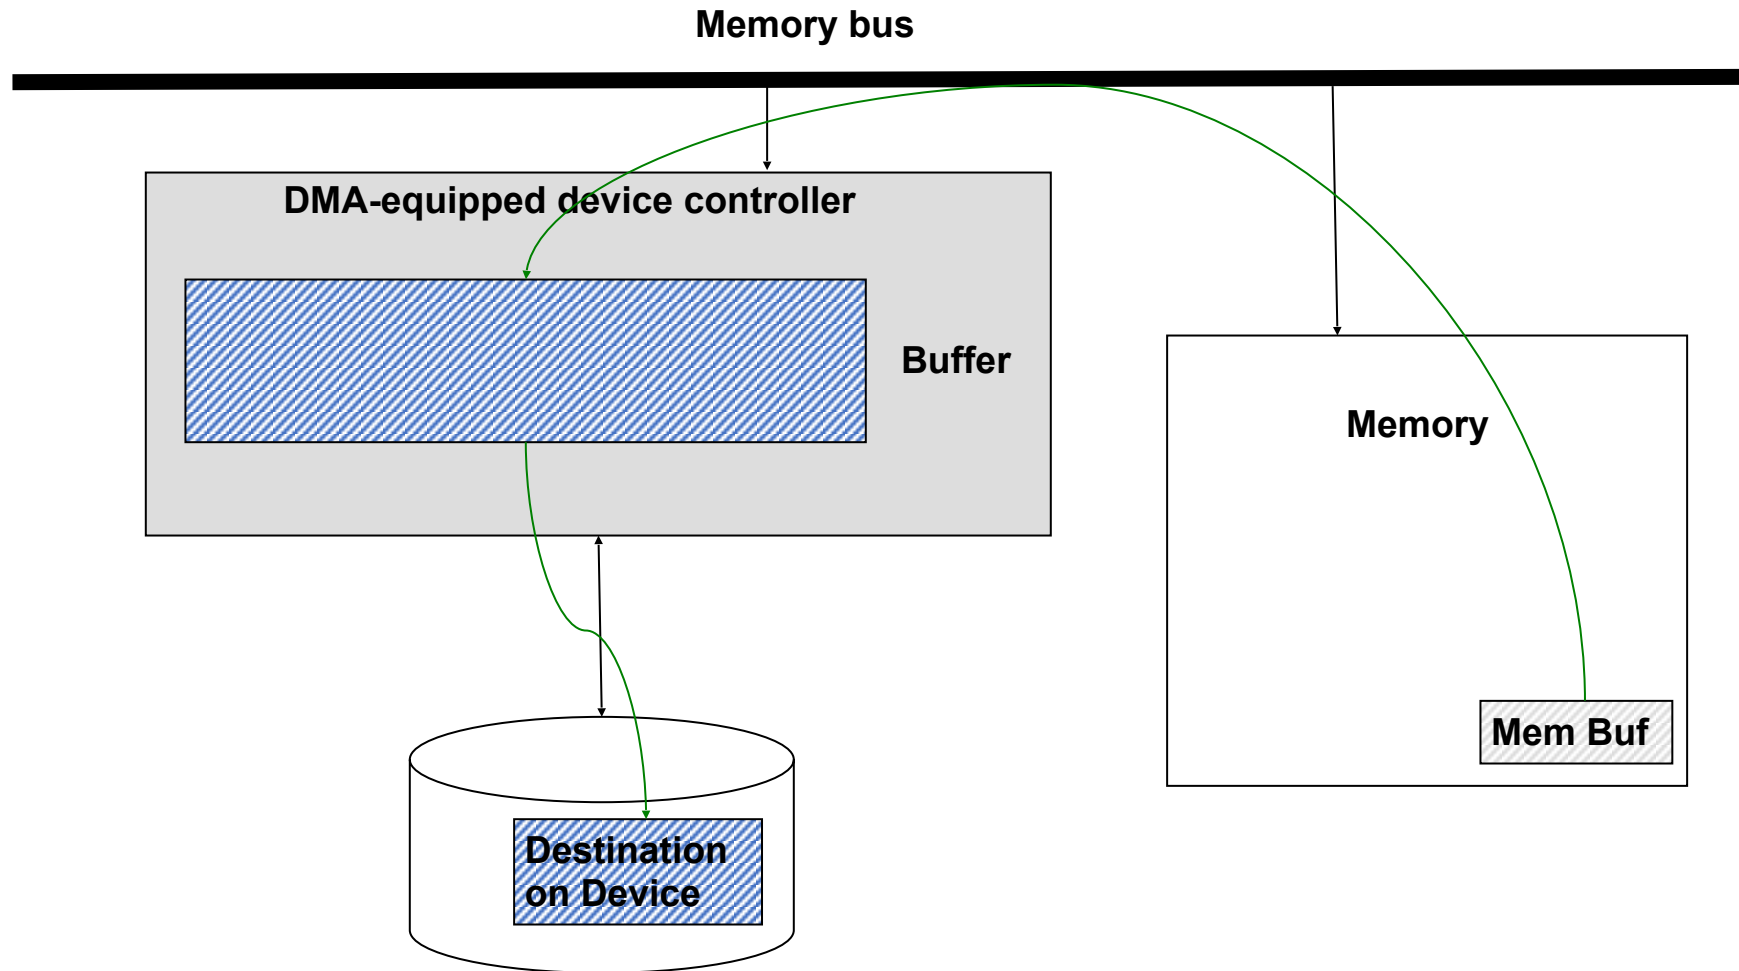
# DMA



# DMA



# DMA Transfer example



# Data transfer speeds

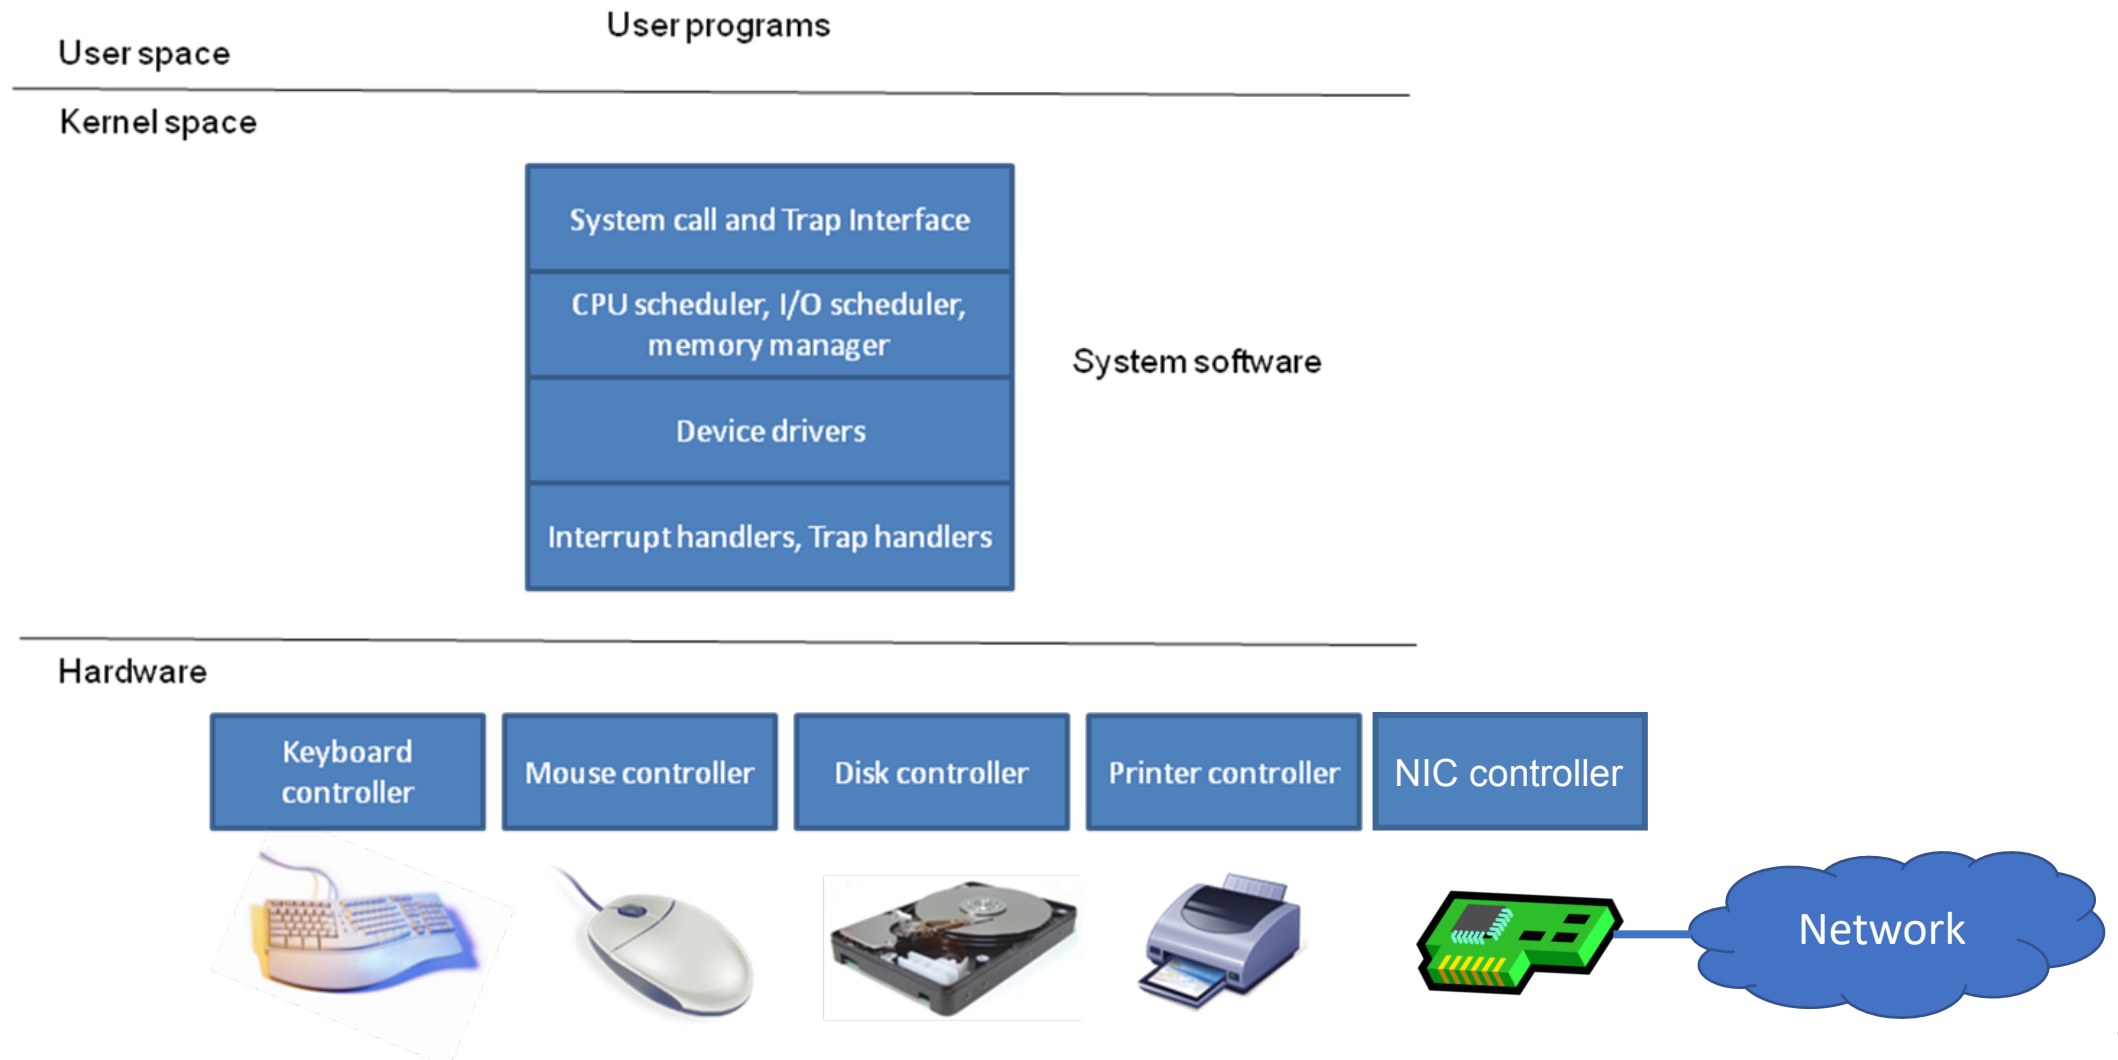
| Device                     | Input/output | Human in the loop | Data rate (circa 2008) | Data rate (circa 2020) | PIO | DMA |
|----------------------------|--------------|-------------------|------------------------|------------------------|-----|-----|
| Keyboard                   | Input        | Yes               | 5-10 bytes/sec         | 5-10 bytes/sec         | X   |     |
| Mouse                      | Input        | Yes               | 80-200 bytes/sec       | 80-200 bytes/sec       | X   |     |
| Graphics display           | Output       | No                | 200-350 MB/sec         | 200-350 MB/sec         |     | X   |
| Disk (hard drive)          | I/O          | No                | 100-200 MB/sec         | 500 MB/sec (each)      |     | X   |
| Network (LAN)              | I/O          | No                | 1 Gbit/sec             | 1-400 Gbit/sec         |     | X   |
| Modem                      | I/O          | No                | 1-8 Mbit/sec           | 1-8 Mbit/sec           |     | X   |
| Inkjet printer             | Output       | No                | 20-40 KB/sec           | 20-40 KB/sec           | X   | X   |
| Laser printer              | Output       | No                | 200-400 KB/sec         | 200-400 KB/sec         |     | X   |
| Voice (microphone/speaker) | I/O          | Yes               | 10 bytes/sec           | 10 bytes/sec           | X   |     |
| Solid State (SSD)          | I/O          | No                | 10-50 MB/sec           | 0.1-7GB/sec            |     | X   |
| CD, DVD, Blu-Ray           | I/O          | No                | 10-20 MB/sec           | 10-20 MB/sec           |     | X   |



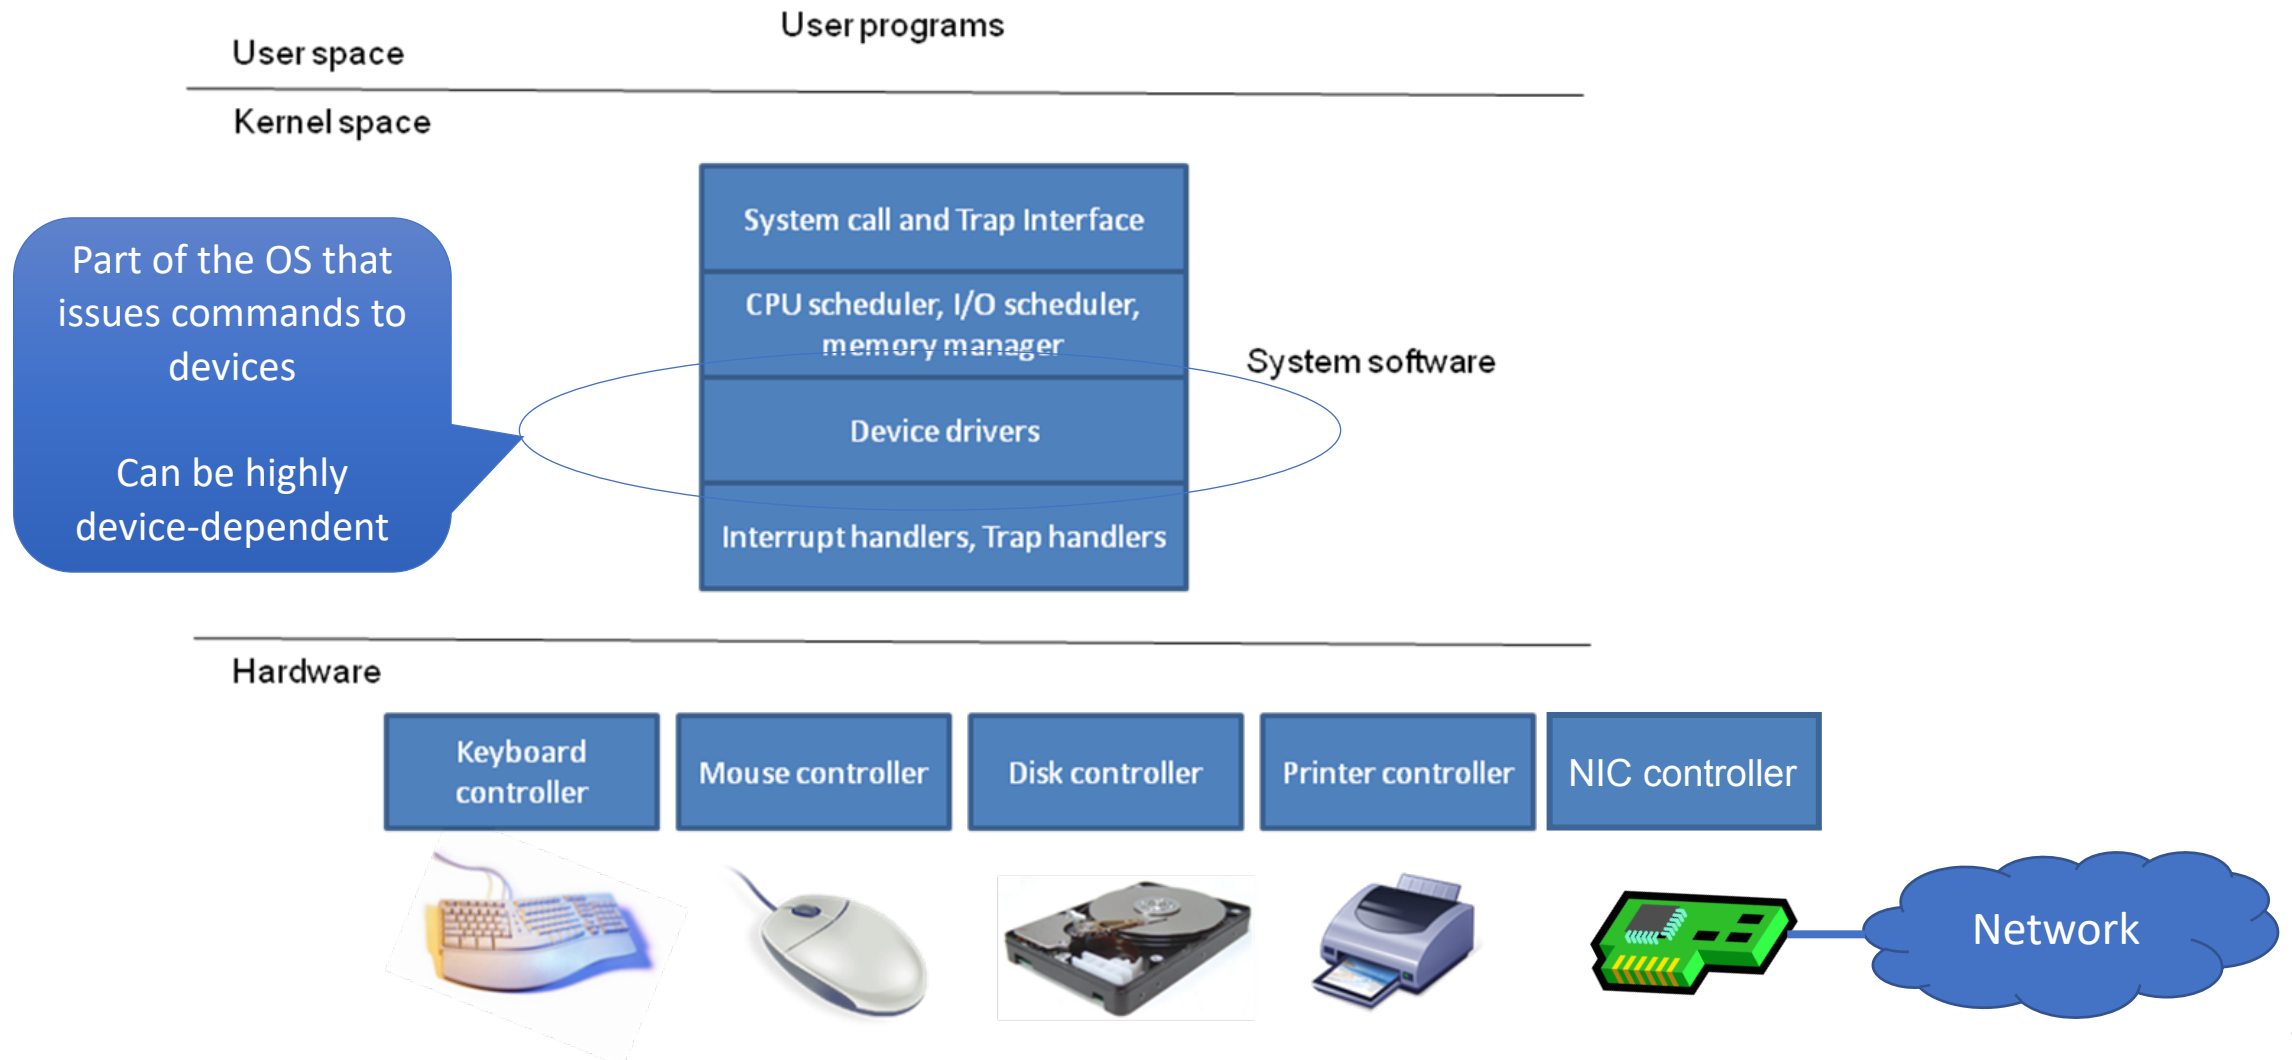
# DMA is used instead of PIO for I/O because

- 33% A. PIO can lose data if the CPU doesn't poll or respond to the interrupt quickly enough
- 10% B. Servicing an interrupt in the CPU is expensive especially if it has to be done for every word of I/O
- 10% C. A DMA controller can transfer data into memory in fewer cycles than a programmed loop in the CPU
- 42% D. All of the above

# Device drivers and OS



# Device drivers and OS





# Device drivers and OS

