



CS2200
Systems and Networks
Spring 2024

Lecture 28: File Systems

Alexandros (Alex) Daglis School of Computer Science Georgia Institute of Technology adaglis@gatech.edu

Lecture slides adapted from Bill Leahy, Charles Lively, Umakishore Ramachandran of Georgia Tech

CIOS is open

- Your feedback helps me improve cs2200!
 - Your future peers will be thankful
 - What worked well? Suggestions for improvement?

- CIOS participation incentive: entire class gets
 - 0.5% bonus to total grade if we reach 85% participation
 - I% bonus to total grade if we reach 95% participation

Agenda

Agenda

- Disk capacity, performance, and scheduling
 - Chapter 10
- File Systems
 - Chapter I I
 - Last cs2200 topic
- It's Pet-a-Disk day!





Disk space allocation policies

- Disk space allocation policies
 - Figures of merit

- Disk space allocation policies
 - Figures of merit
 - Unix file system

- Disk space allocation policies
 - Figures of merit
 - Unix file system

Similarities to memory management

- Disk space allocation policies
 - Figures of merit
 - Unix file system

Similarities to memory management

	Namespace	Physical entity
Virtual memory	Virtual address space	Physical memory
File system	File name	Disk blocks

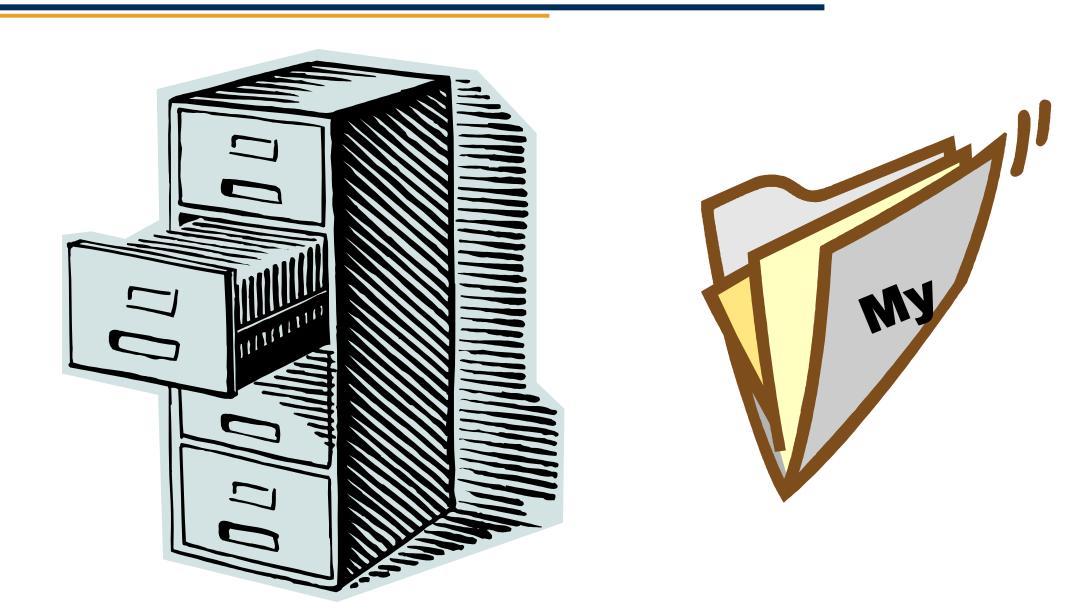
- Disk space allocation policies
 - Figures of merit
 - Unix file system

Similarities to memory management

	Namespace	Physical entity
Virtual memory	Virtual address space	Physical memory
File system	File name	Disk blocks

- Disk block: unit of disk management for the OS
 - May consist of one or more disk sectors

File System Abstraction



Logical block addresses

- We're going to adopt the modern convention of converting cylinder/track/ sector into a logical block address (LBA)
- That is, we calculate an integer 0 to n-1 so that the disk appears to be a simple array of disk blocks
- Modern disk controllers present this abstraction to the CPU to prevent it from having to know the geometry of the physical disk drives

File allocation schemes on the disk

- Contiguous allocation
- Contiguous allocation with overflow
- Linked allocation
- FAT
- Indexed allocation
- Multilevel indexed
- Hybrid indexed

File allocation schemes on the disk

- Contiguous allocation
- Contiguous allocation with overflow
- Linked allocation
- FAT
- Indexed allocation
- Multilevel indexed
- Hybrid indexed

Functionality is similar to memory management

Implemented by "storage manager" or "file system" part of OS

Most OSs support more than one file system.

Fast sequential access

- Fast sequential access
- Fast random access

- Fast sequential access
- Fast random access
- Ability to grow the file

- Fast sequential access
- Fast random access
- Ability to grow the file
- Easy allocation of storage

- Fast sequential access
- Fast random access
- Ability to grow the file
- Easy allocation of storage
- Efficiency of space utilization on the disk

- Fast sequential access
- case?

Use

- Fast random access
- Ability to grow the file
- Easy allocation of storage
- Efficiency of space utilization on the disk

Use

case? Use

case?

- Fast sequential access
- Fast random access
- Ability to grow the file
- Easy allocation of storage
- Efficiency of space utilization on the disk

- Fast sequential access
- Fast random access
- Ability to grow the file
- Easy allocation of storage
- Efficiency of space utilization on the disk

- → Use
- case? Use
- case?
- Flexibility

- Fast sequential access → Use
- Fast random access Case Use
- Ability to grow the file
 Se?
- Easy allocation of storage Flexibility Time
- Efficiency of space utilization on the disk

- Fast sequential access → Use
- Ability to grow the file
 Se?
- Easy allocation of storage Flexibility Time
- Efficiency of space utilization on the disk
 - → Concerns of internal and external fragmentation

■ Fast sequential access → Use

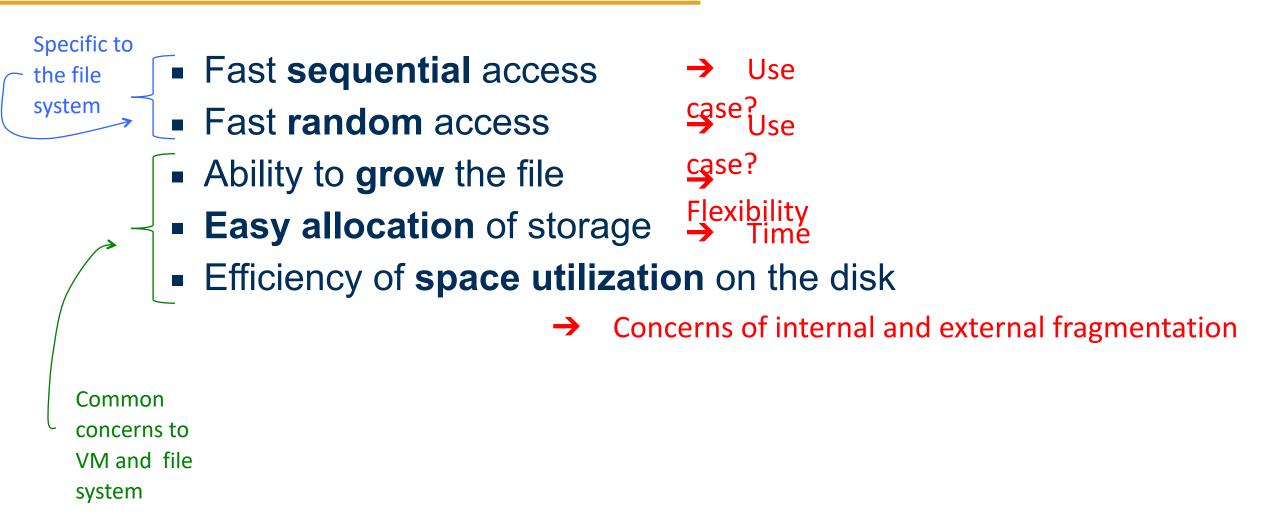
Ability to grow the file
Se?

■ Easy allocation of storage Flexibility Time

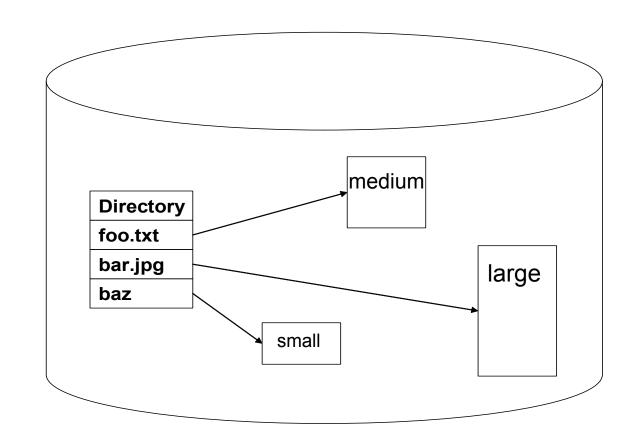
Efficiency of space utilization on the disk

→ Concerns of internal and external fragmentation

Common concerns to VM and file system

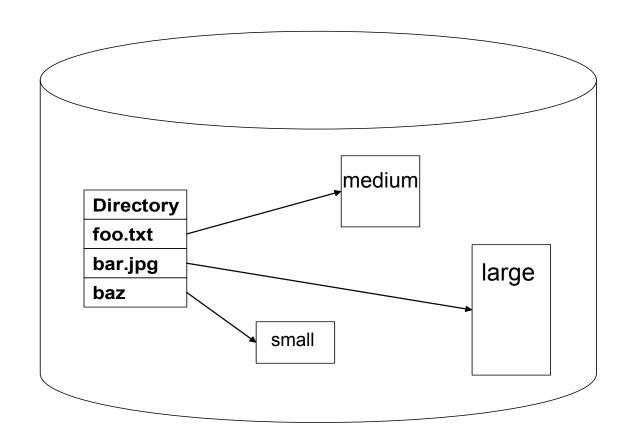


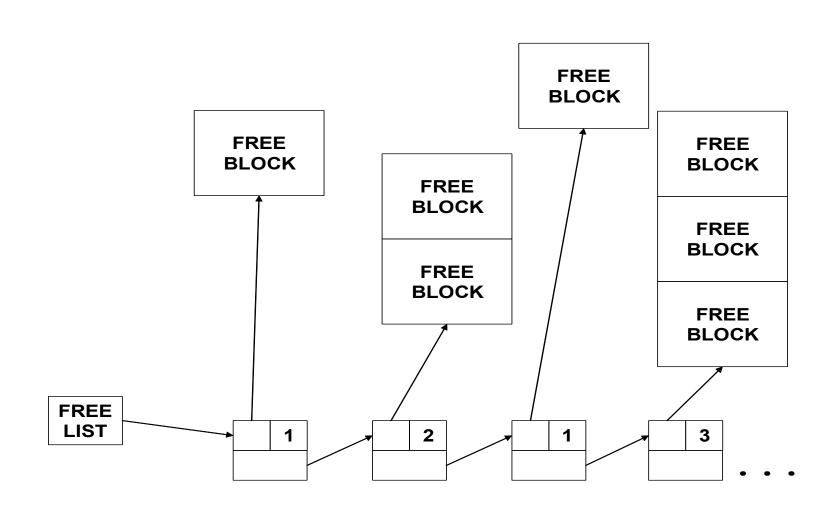
Contiguous allocation

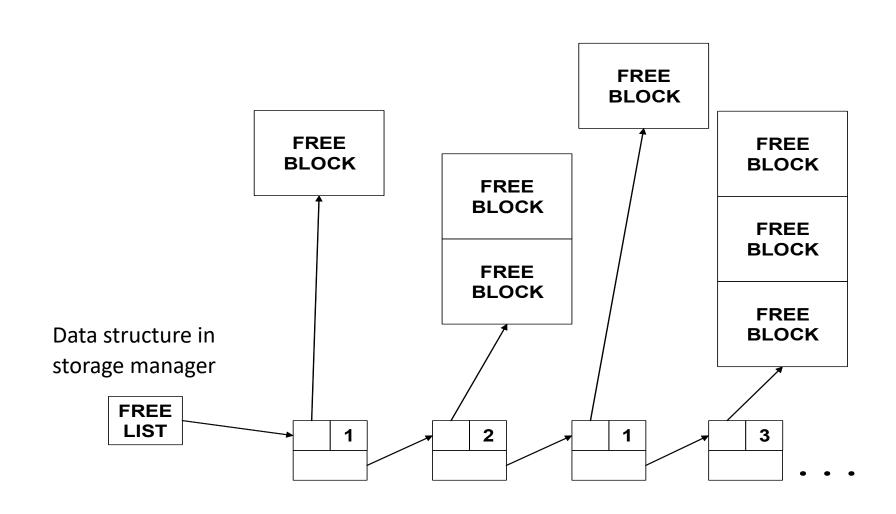


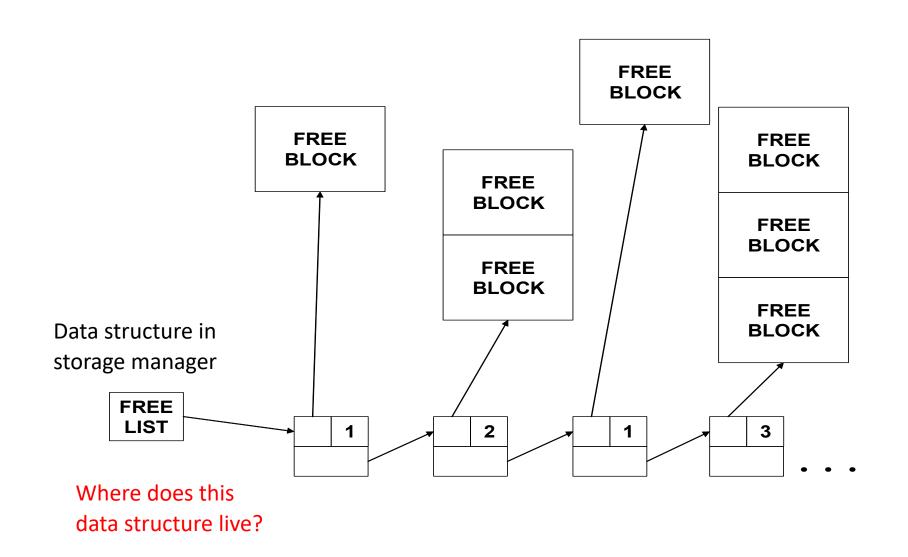
Contiguous allocation

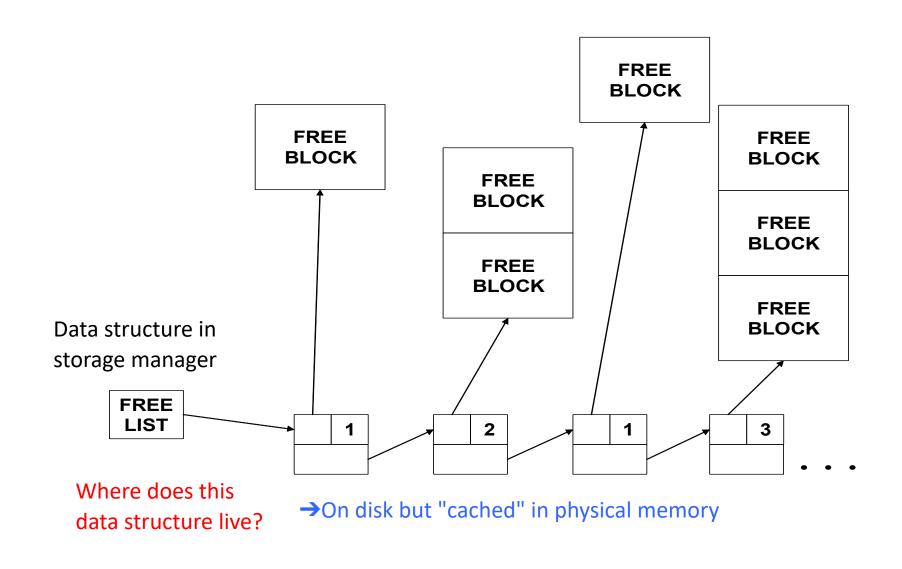
Similar to variable-sized partitioning in memory management



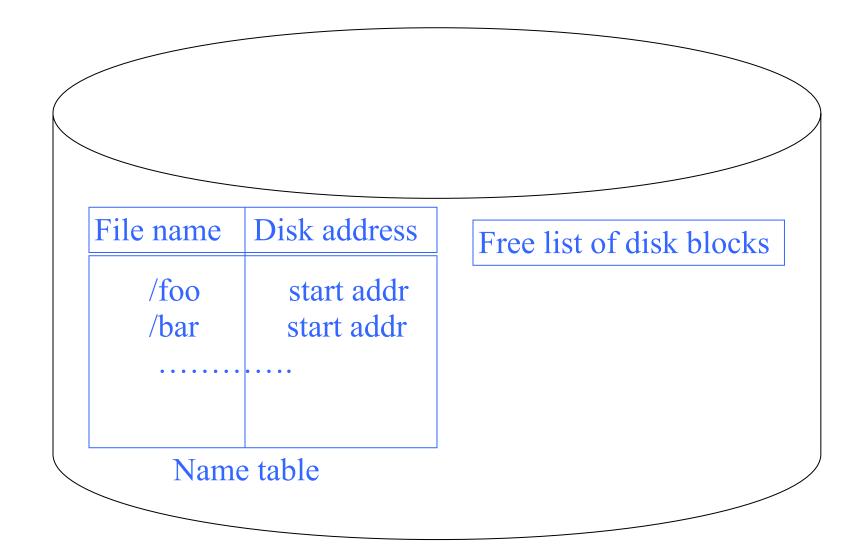




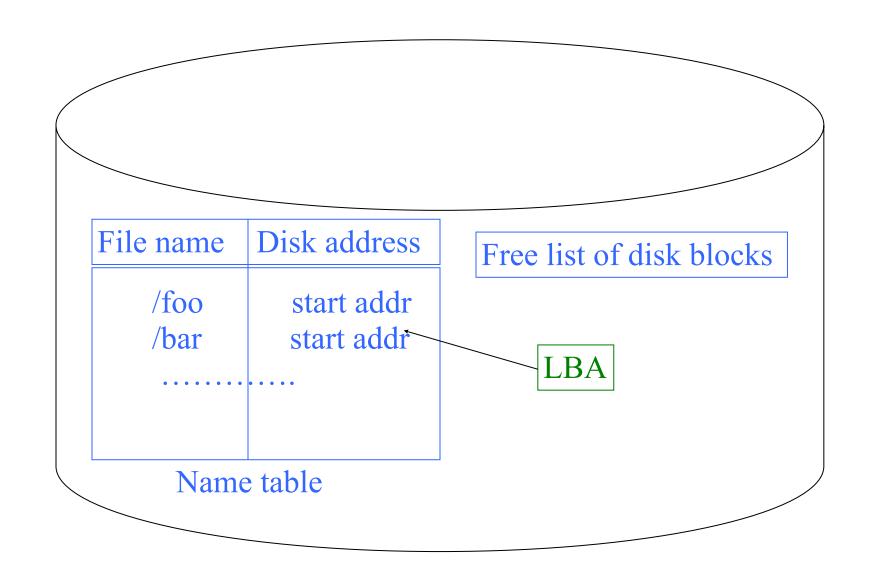




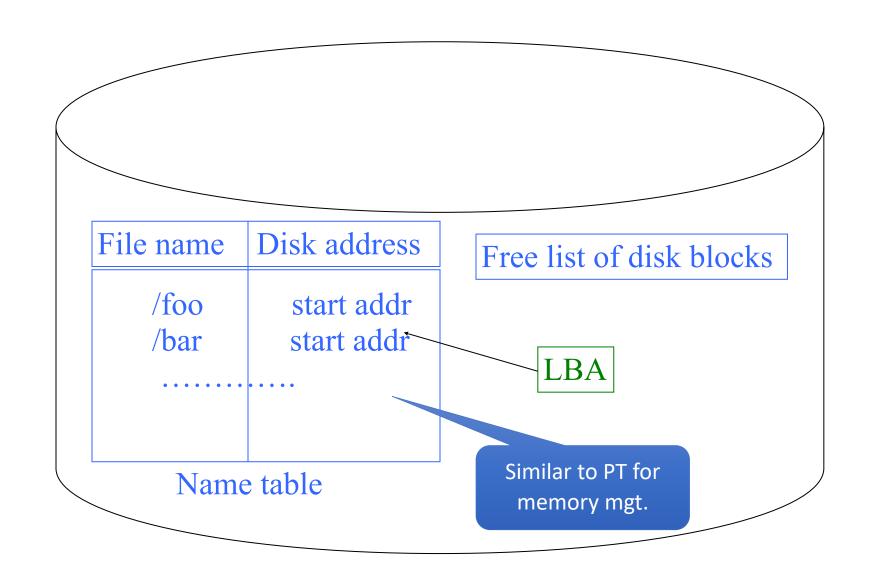
Data structures for storage manager



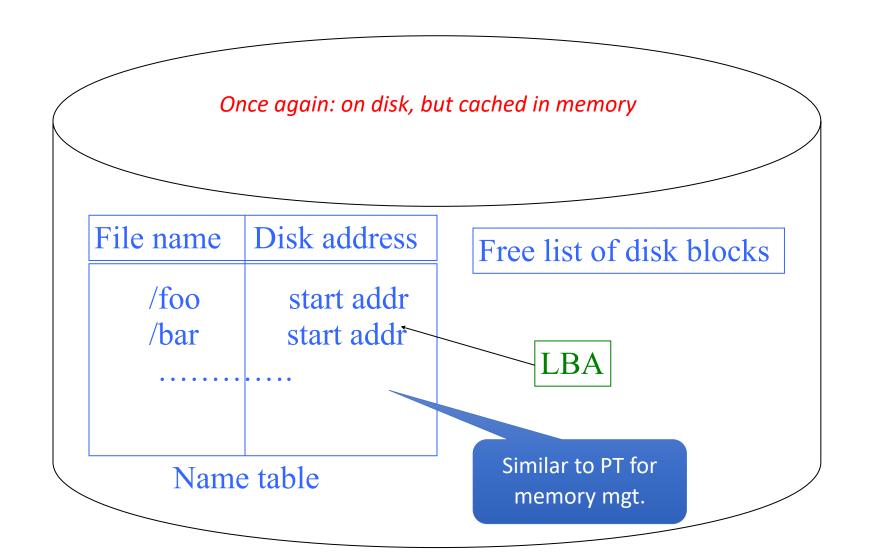
Data structures for storage manager



Data structures for storage manager



Data structures for storage manager



- Growth?
- Quickness of allocation (time)?
- Fragmentation (space)?
- Sequential access?
- Random access?

- X Growth?
 - Quickness of allocation (time)?
 - Fragmentation (space)?
 - Sequential access?
 - Random access?

- X
- Growth?
- 7/
- Quickness of allocation (time)?
- Fragmentation (space)?
- Sequential access?
- Random access?

- First fit
- Best fit
- Compaction

X

Growth?

 $\langle \boldsymbol{\kappa} \rangle$

• Quickness of allocation (time)?

X

- Fragmentation (space)?
 Both internal and external
- Sequential access?
- Random access?

- First fit
- Best fit
- Compaction

- X Growth?
- Quickness of allocation (time)?
- Fragmentation (space)? Both internal and external
- Sequential access?
 - Random access?

- First fit
- Best fit
- Compaction

- X Growth?
- Quickness of allocation (time)?
- Fragmentation (space)? Both internal and external
- ✓ Sequential access?
- ✓ Random access?

- First fit
- Best fit
- Compaction

X Growth?

Can be fixed with "overflow" areas...

 \times

• Quickness of allocation (time)?

X

Fragmentation (space)?
Both internal and external

1

Sequential access?

1

Random access?

- First fit
- Best fit
- Compaction

X Growth?

Can be fixed with "overflow" areas...

→ Contiguous allocation with overflow

• Quickness of allocation (time)?

Fragmentation (space)? Both internal and external

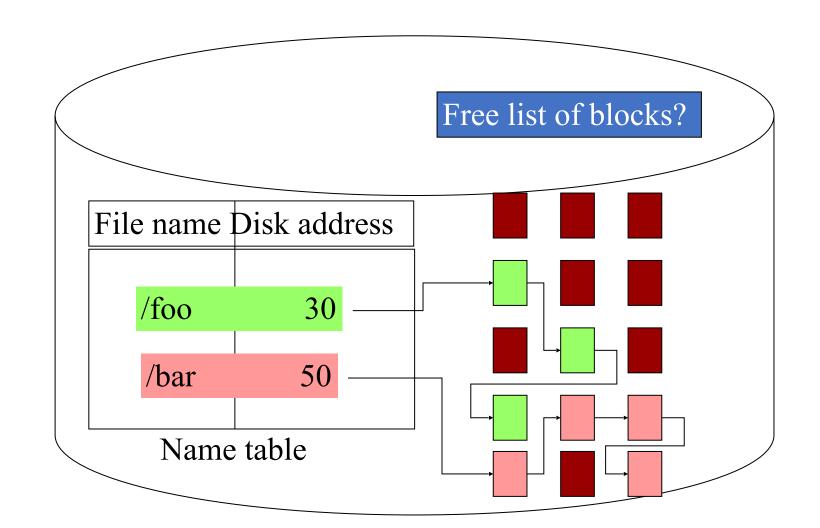
✓ ■ Sequential access?

Random access?

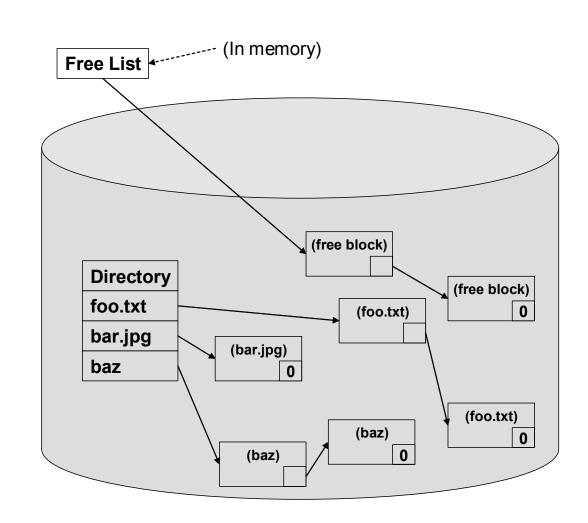
- First fit
- Best fit
- Compaction

File allocation schemes on the disk

- ✓ Contiguous allocation
- Contiguous allocation with overflow
- → Linked allocation
 - FAT
 - Indexed allocation
 - Multilevel indexed
 - Hybrid indexed

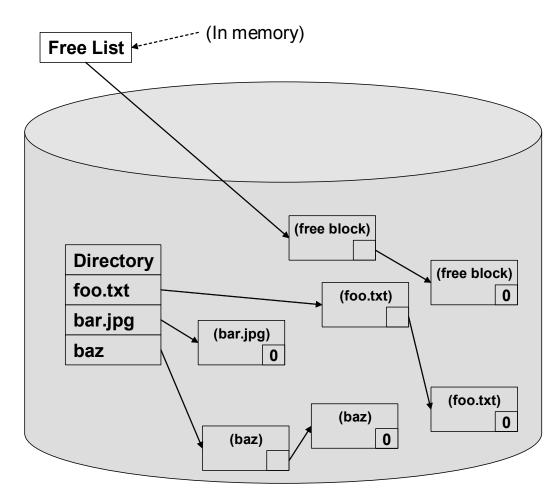


Linked allocation



Linked allocation

As before, data structures are on disk but cached in memory



- Growth?
- Quickness of allocation (time)?
- Fragmentation (space)?
- Sequential access?
- Random access?

- ✓ Growth?
 - Quickness of allocation (time)?
 - Fragmentation (space)?
 - Sequential access?
 - Random access?

- ✓ Growth?
- ✓ Quickness of allocation (time)? One block at a time
 - Fragmentation (space)?
 - Sequential access?
 - Random access?

- ✓ Growth?
- ✓ Quickness of allocation (time)? One
- ✓ Fragmentation (space)?
 - Sequential access?
 - Random access?

One block at a time

No external

- ✓ Growth?
- ✓ Quickness of allocation (time)?
- ✓ Fragmentation (space)?
- ✓ Sequential access?
 - Random access?

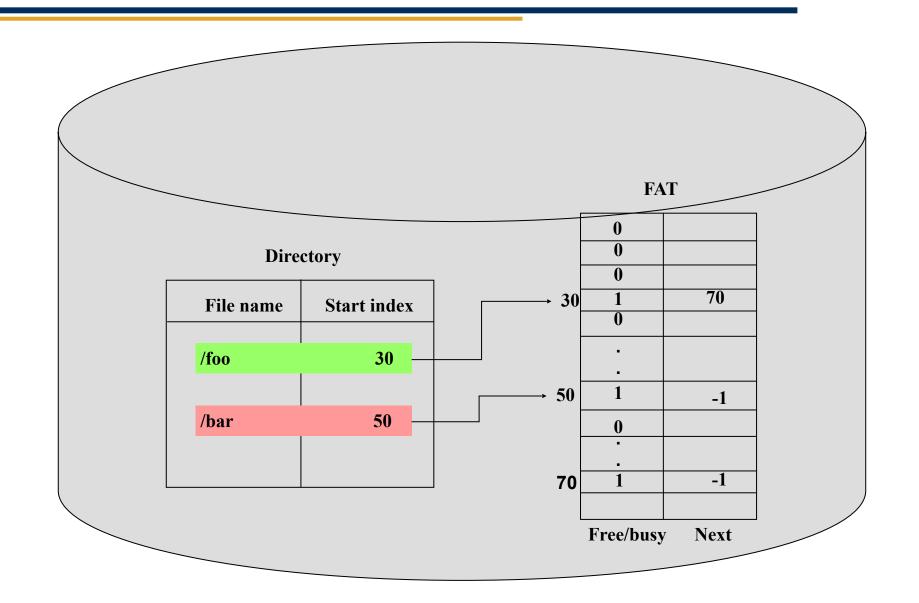
One block at a time

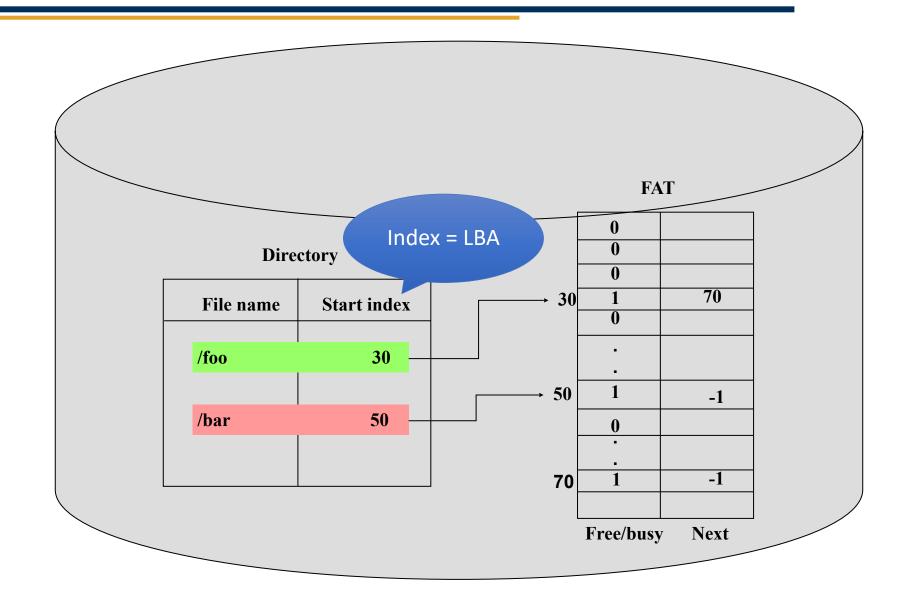
No external

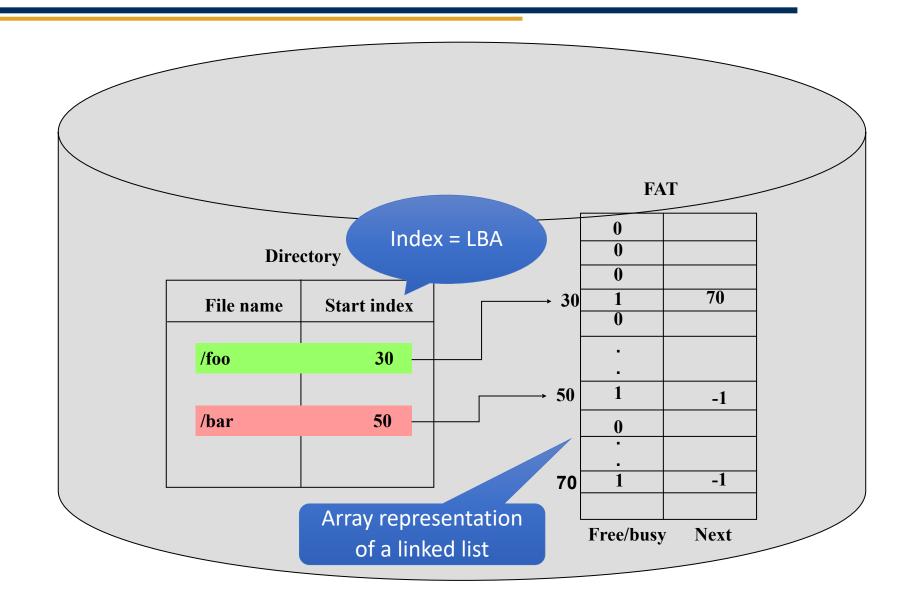
- ✓ Growth?
- ✓ Quickness of allocation (time)? One block at a time
- ✓ Fragmentation (space)? No external
- Sequential access?
- X = Random access? Very dependent on seek time

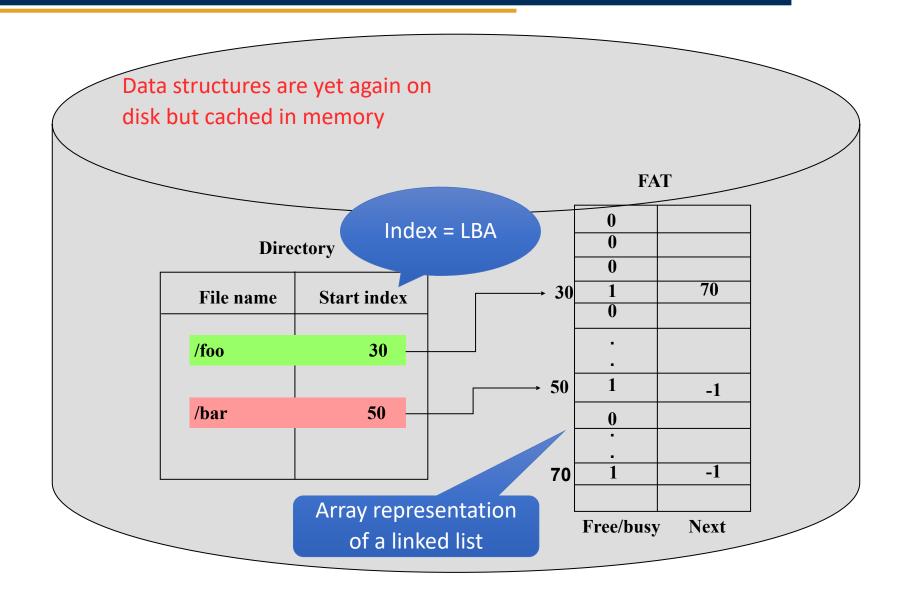
File allocation schemes on the disk

- ✓ Contiguous allocation
- Contiguous allocation with overflow
- ✓ Linked allocation
- → FAT
 - Indexed allocation
 - Multilevel indexed
 - Hybrid indexed









- Growth?
- Quickness of allocation (time)?
- Fragmentation (space)?
- Sequential access?
- Random access?

- ✓ Growth?
 - Quickness of allocation (time)?
 - Fragmentation (space)?
 - Sequential access?
 - Random access?

- ✓ Growth?
- ✓ Quickness of allocation (time)? Better than linked; FAT cached in memory
 - Fragmentation (space)?
 - Sequential access?
 - Random access?

- ✓ Growth?
- ✓ Quickness of allocation (time)?
- Fragmentation (space)?
 - Sequential access?
 - Random access?

Better than linked; FAT cached in memory

Same as linked

- ✓ Growth?
- ✓ Quickness of allocation (time)?
- ✓ Fragmentation (space)?
- Sequential access?
 - Random access?

Better than linked; FAT cached in memory

Same as linked

Similar to linked

- ✓ Growth?
- ✓ Quickness of allocation (time)? Better than linked; FAT cached in memory
- ✓ Fragmentation (space)? Same as linked
- ✓ Sequential access? Similar to linked
- Still depends on seek time, but next-block pointers stored in FAT, so less serialization and

block transfers from disk

✓ ■ Growth?

✓ ■ Quickness of allocation (time)? Better than linked; FAT cached in memory

✓ ■ Fragmentation (space)? Same as linked

✓ ■ Sequential access? Similar to linked

Random access?
Still depends on seek time, but next-block

pointers stored in FAT, so less serialization and

block transfers from disk

Less error prone than linked allocation thanks to data/metadata separation

✓ ■ Growth?

✓ ■ Quickness of allocation (time)? Better than linked; FAT cached in memory

✓ ■ Fragmentation (space)? Same as linked

✓ ■ Sequential access? Similar to linked

Random access?
 Still depends on seek time, but next-block pointers stored in FAT, so less serialization and

block transfers from disk

Less error prone than linked allocation thanks to data/metadata separation

FAT table size limits max supported disk capacity. Must resort to disk partitioning, which becomes a burden to user

✓ ■ Growth?

✓ ■ Quickness of allocation (time)? Better than linked; FAT cached in memory

✓ ■ Fragmentation (space)? Same as linked

✓ ■ Sequential access? Similar to linked

Random access?
Still depends on seek time, but next-block

pointers stored in FAT, so less serialization and

block transfers from disk

Less error prone than linked allocation thanks to data/metadata separation

FAT table size limits max supported disk capacity. Must resort to disk partitioning, which becomes a burden to user

• E.g., FAT-16 limits pointer size to 16 bits; can only address 2^16 blocks

File allocation schemes on the disk

- ✓ Contiguous allocation
- ✓ Contiguous allocation with overflow
- ✓ Linked allocation
- ✓ FAT
- Indexed allocation
 - Multilevel indexed
 - Hybrid indexed

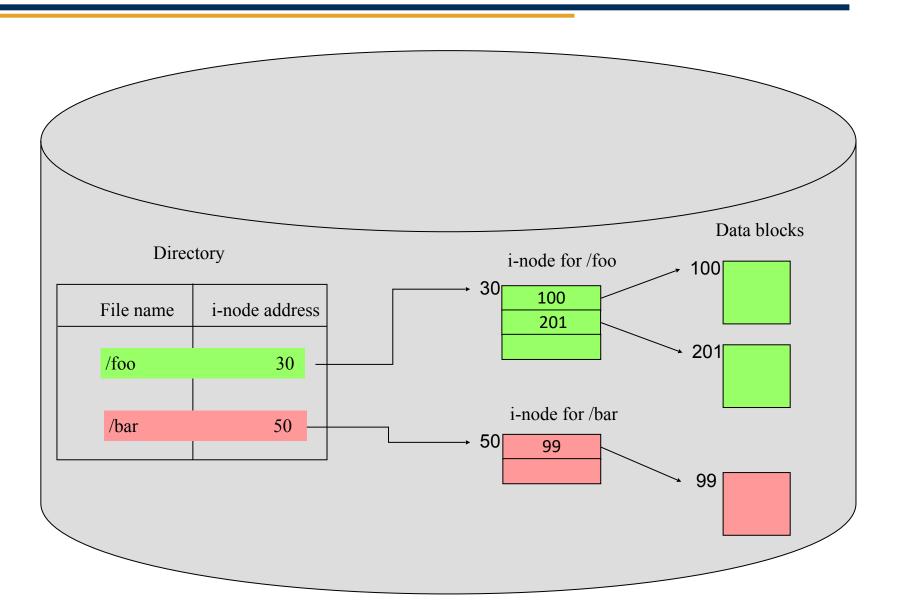
Some Unix terminology

- The Unix file system stored the metadata for a file separately from the directory entry; this was a novel idea at the time
- This technique has been used in a number of subsequent file systems but the Unix terminology is often used to describe it

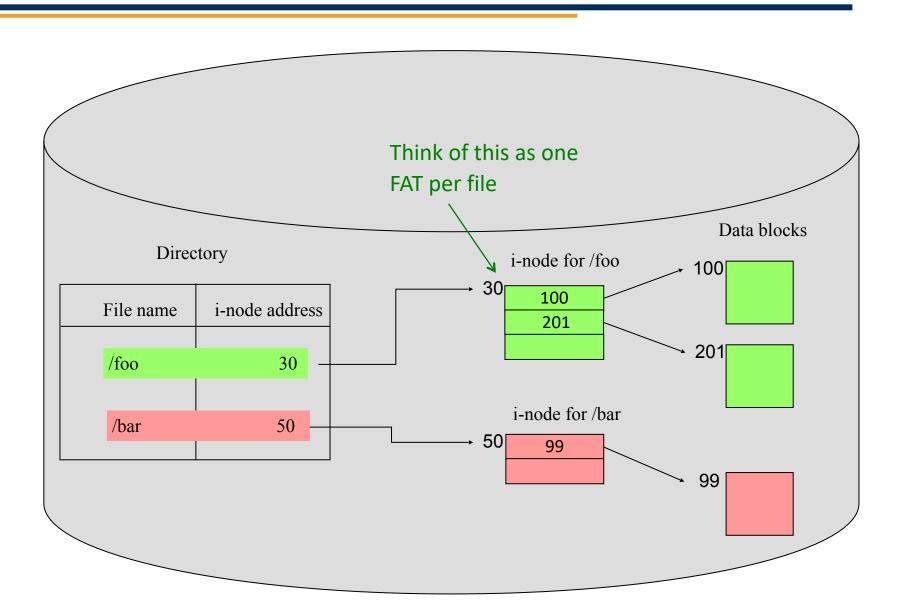
Some Unix terminology

- The Unix file system stored the metadata for a file separately from the directory entry; this was a novel idea at the time
- This technique has been used in a number of subsequent file systems but the Unix terminology is often used to describe it
- i-node (modern usage inode)
 - All of the information about a file except name was stored in a fixed-length entry on disk called an i-node (similar to what a PCB is for a process)
 - The i-nodes were kept in an array on the file system so that the disk location of an i-node could be calculated from its index, often called the i-number
 - i-nodes can be cached by the OS in memory when a file is in use
- Directory
 - Directories contained only file names and the i-number (i-node index number)
 - It is possible for two directory entries to contain the same i-number (more on that later)

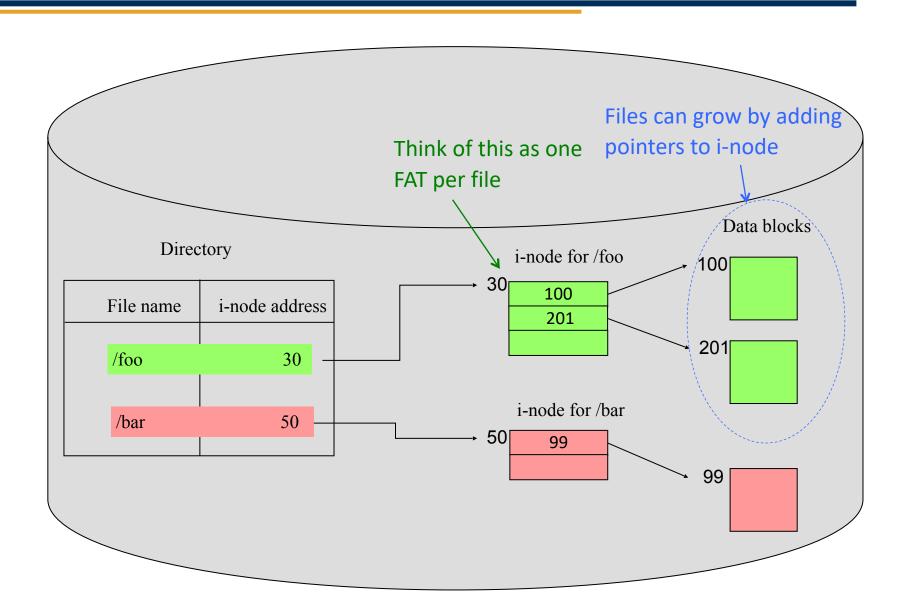




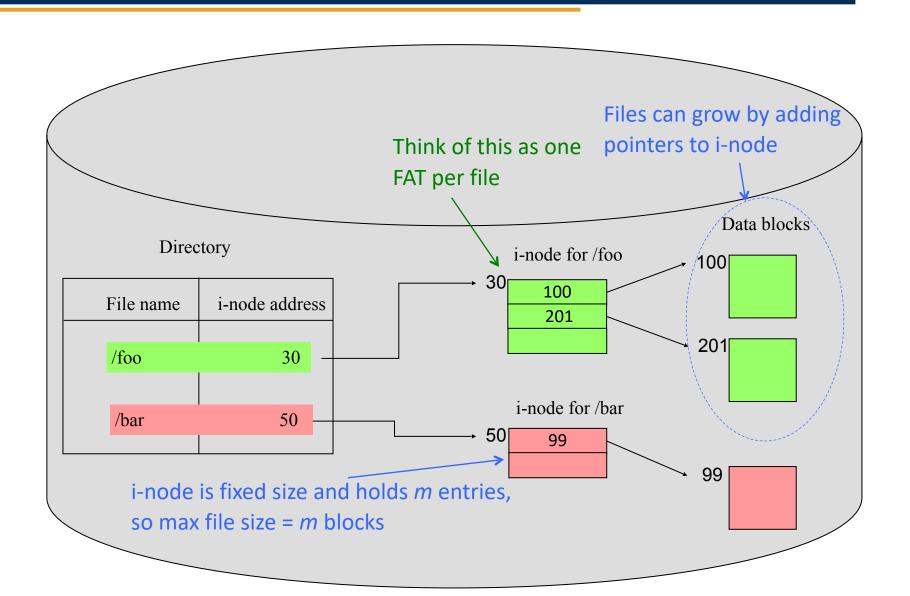














- A. Suffers from external fragmentation
- B. Offers sequential access as fast as FAT
- C. Supports arbitrarily large files

- Growth?
- Quickness of allocation (time)?
- Fragmentation (space)?
- Sequential access?
- Random access?

X Growth?

Big problems

- Quickness of allocation (time)?
- Fragmentation (space)?
- Sequential access?
- Random access?

X Growth?

- Big problems
- ✓ Quickness of allocation (time)?
 - Fragmentation (space)?
 - Sequential access?
 - Random access?

- X Growth?
- ✓ Quickness of allocation (time)?
- ✓ Fragmentation (space)?
 - Sequential access?
 - Random access?

Big problems

- X Growth?
- ✓ Quickness of allocation (time)?
- ✓ Fragmentation (space)?
- ✓ Sequential access?
 - Random access?

Big problems

Same as FAT

X Growth? Big problems

✓ ■ Quickness of allocation (time)?

✓ ■ Fragmentation (space)?

Sequential access?

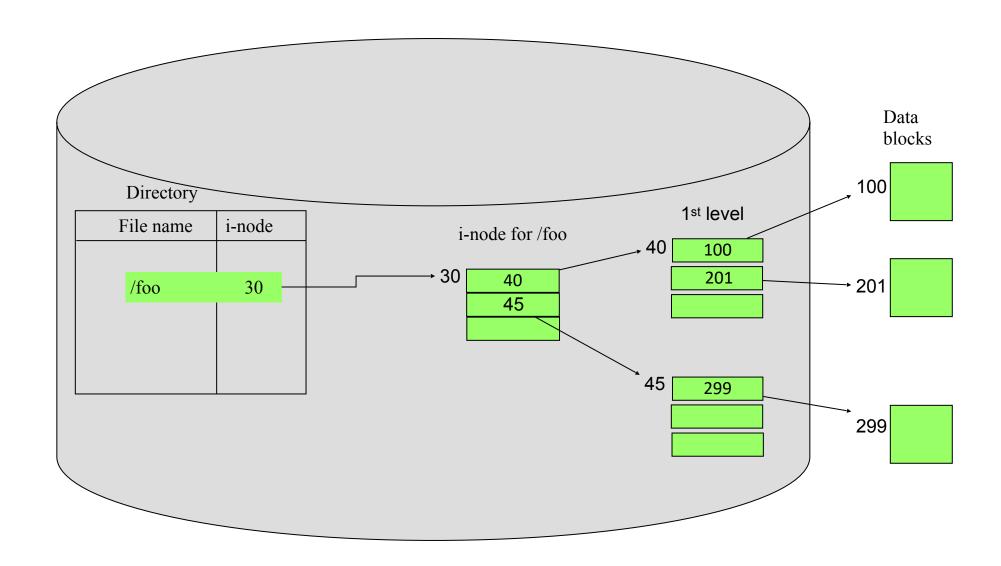
Random access?

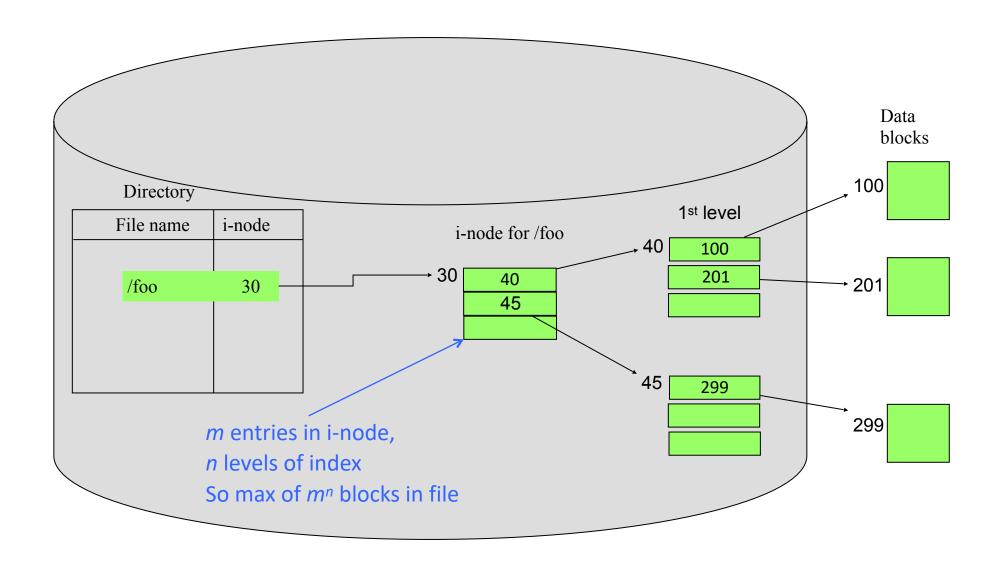
Same as FAT

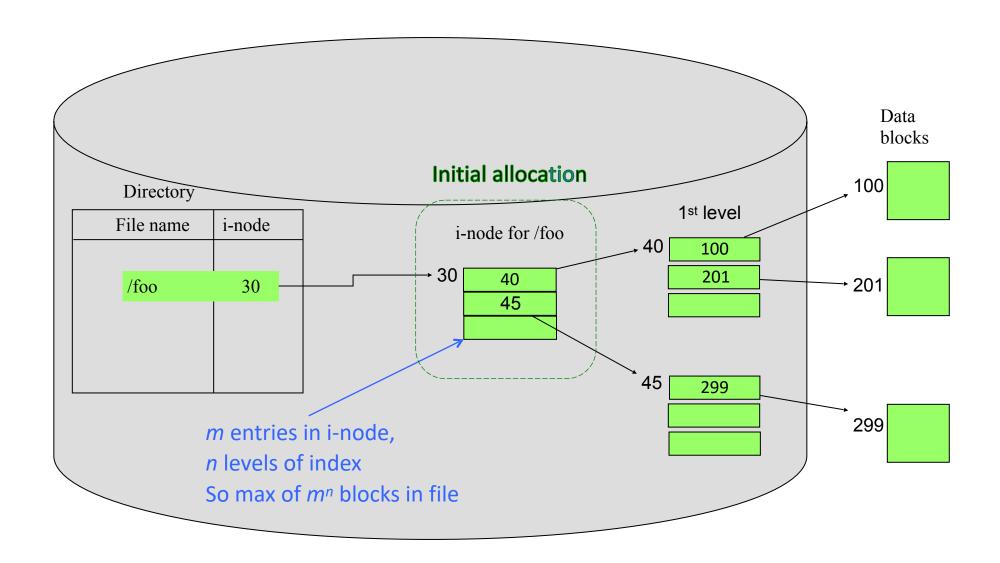
Potentially better than FAT -- Still depends on seek time, but can cache i-node in memory

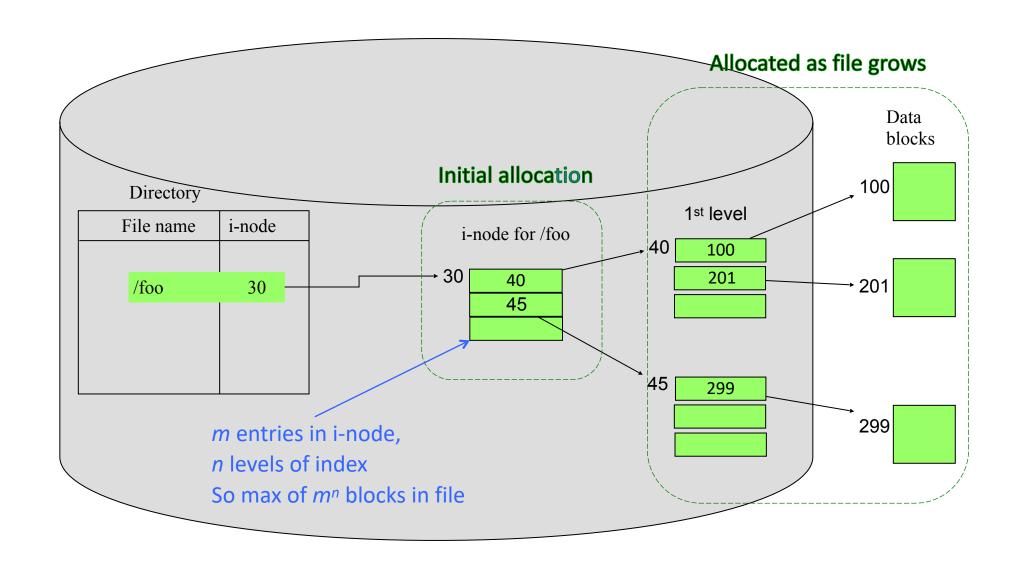
File allocation schemes on the disk

- ✓ Contiguous allocation
- Contiguous allocation with overflow
- ✓ Linked allocation
- ✓ FAT
- Indexed allocation
- Multilevel indexed
- → Hybrid indexed

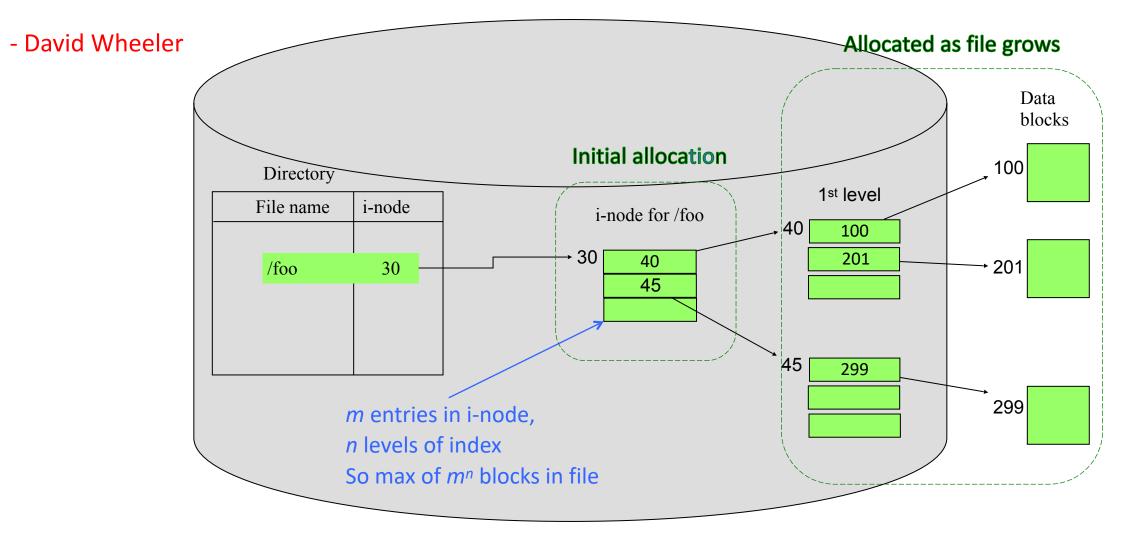




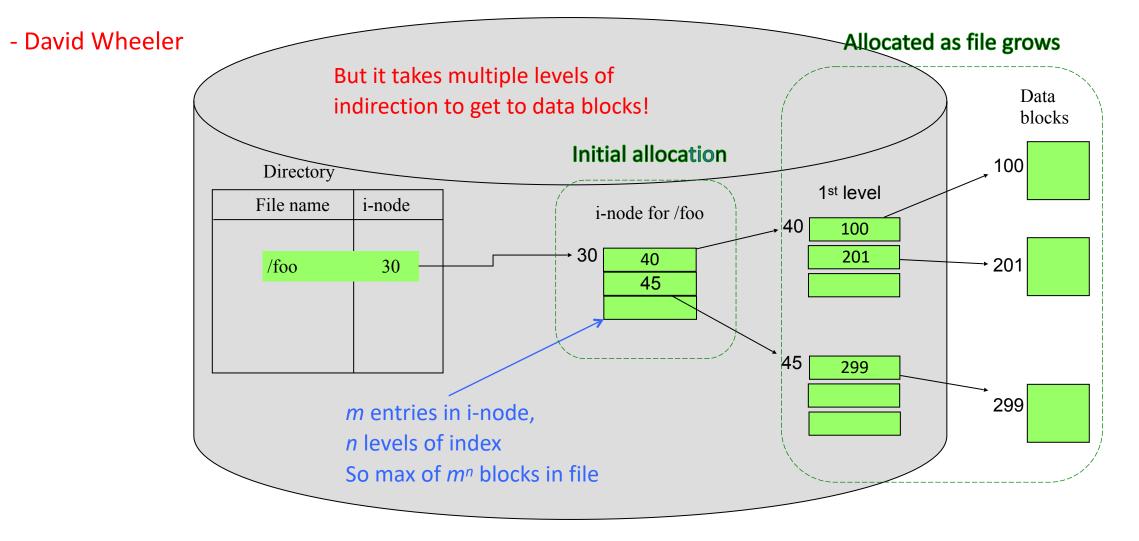




"Every problem in CS can be solved with a layer of indirection..."



"Every problem in CS can be solved with a layer of indirection..."



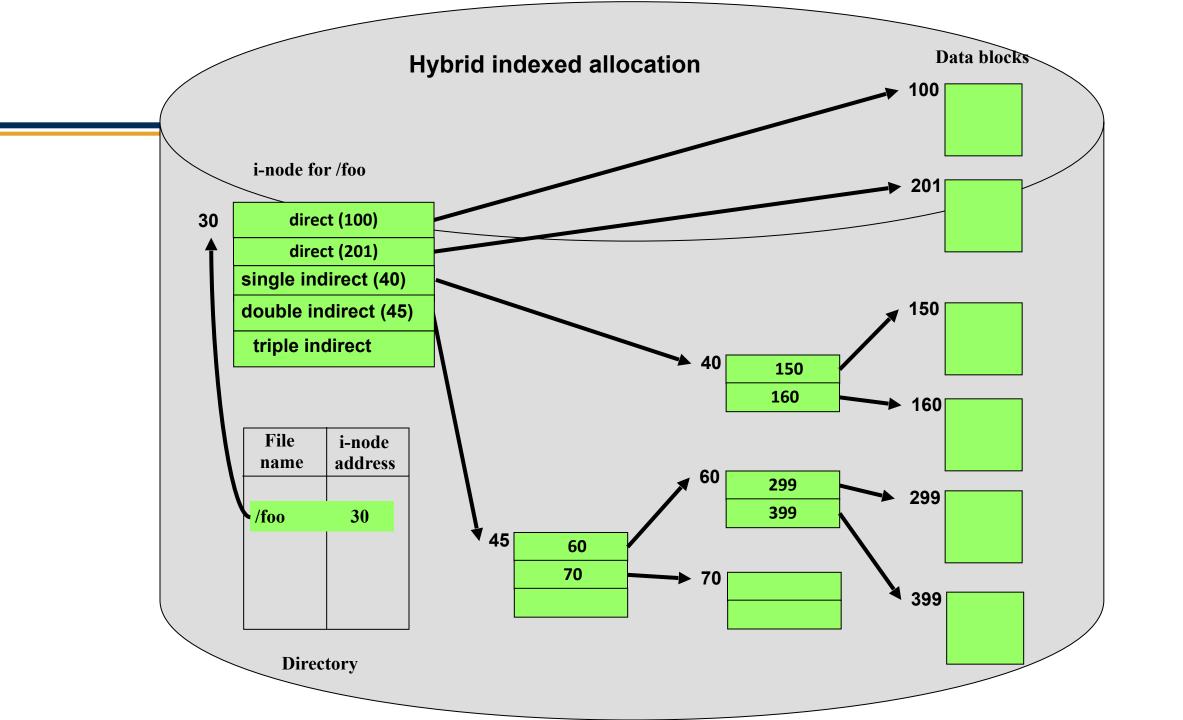
"Every problem in CS can be solved with a layer of indirection..." "...but that usually will create another problem" - David Wheeler Allocated as file grows But it takes multiple levels of Data indirection to get to data blocks! blocks Initial allocation 100 Directory 1st level File name i-node i-node for /foo 40 100 30 201 40 **201** /foo 30 45 45 299 299 *m* entries in i-node, *n* levels of index So max of m^n blocks in file

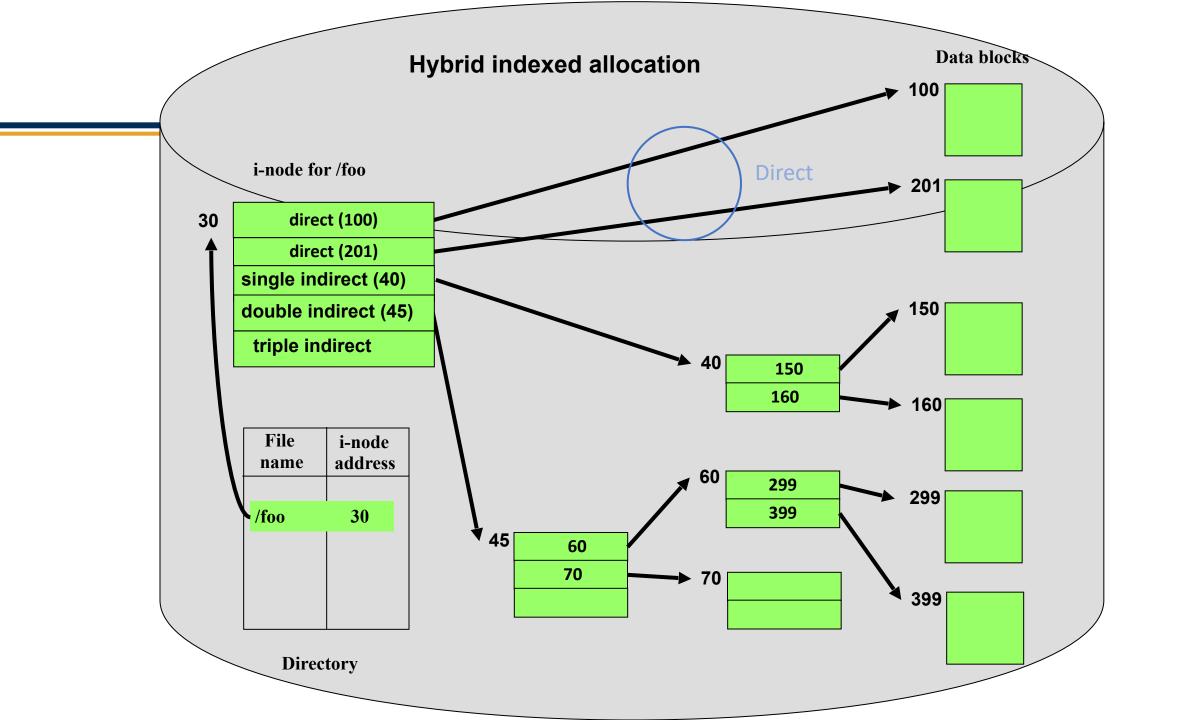
- Small files are a big problem with multilevel indexed allocation
- Why?

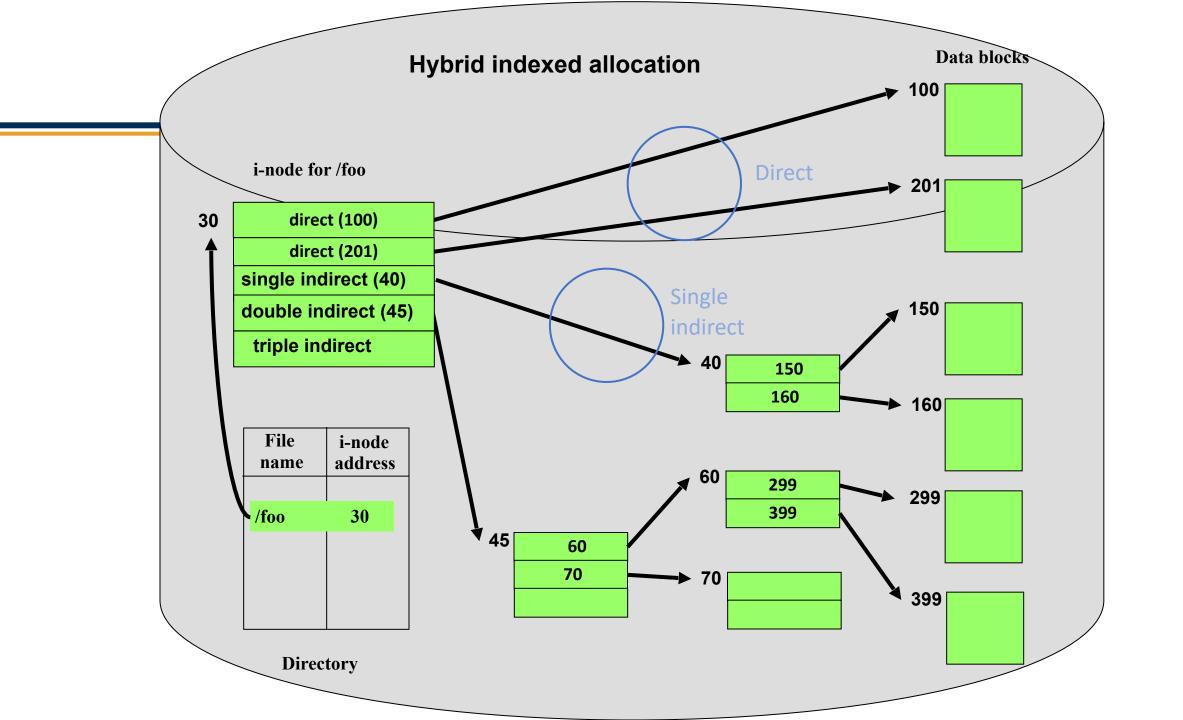
- Small files are a big problem with multilevel indexed allocation
- Why?
- The index block(s) take up as much or more space than the file!

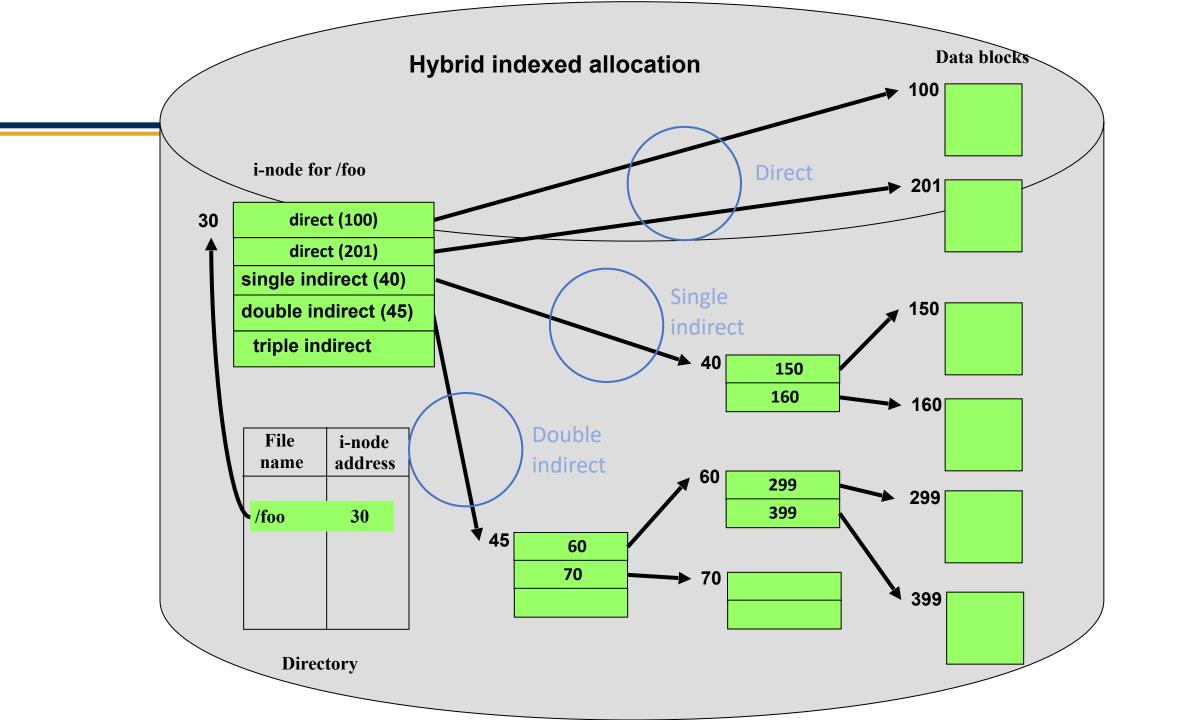
- Small files are a big problem with multilevel indexed allocation
- Why?
- The index block(s) take up as much or more space than the file!
- Too many sequential accesses to get to data

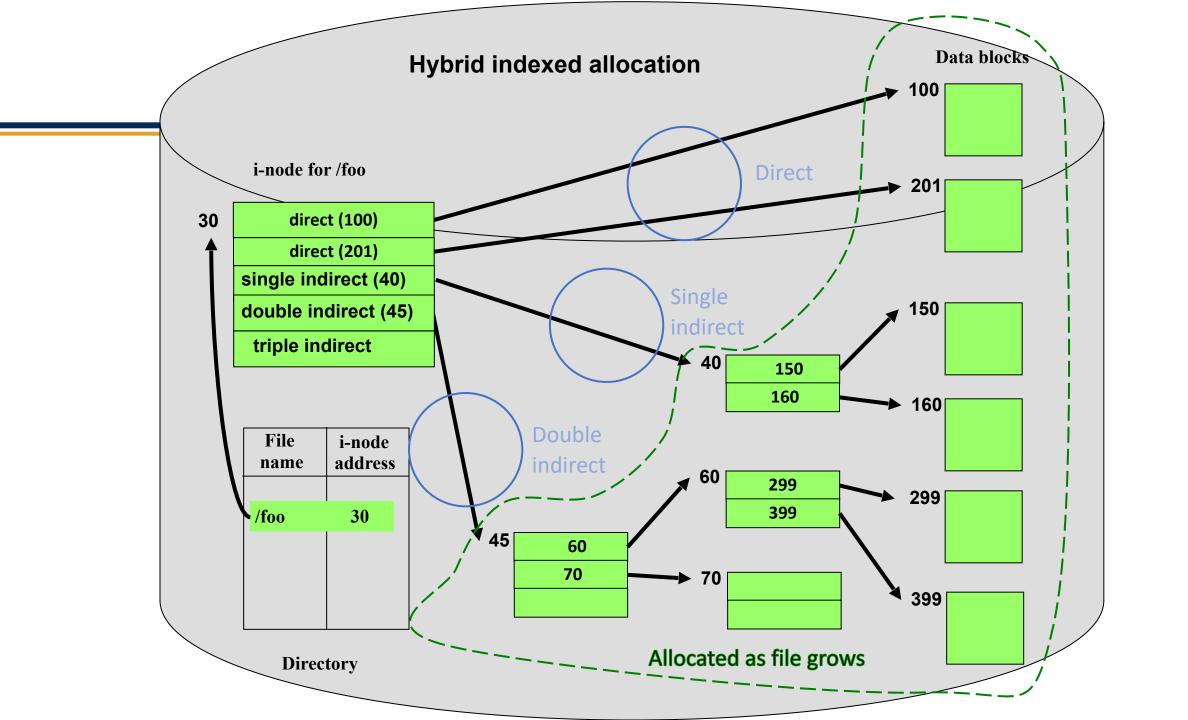
- Small files are a big problem with multilevel indexed allocation
- Why?
- The index block(s) take up as much or more space than the file!
- Too many sequential accesses to get to data
- Enter Hybrid Indexed Allocation











- Growth?
- Quickness of allocation (time)?
- Fragmentation (space)?
- Sequential access?
- Random access?

- ✓ Growth?
 - Quickness of allocation (time)?
 - Fragmentation (space)?
 - Sequential access?
 - Random access?

- ✓ Growth?
- ✓ Quickness of allocation (time)?
 - Fragmentation (space)?
 - Sequential access?
 - Random access?

- ✓ Growth?
- ✓ Quickness of allocation (time)?
- ✓ Fragmentation (space)?
 - Sequential access?
 - Random access?

- ✓ Growth?
- ✓ Quickness of allocation (time)?
- ✓ Fragmentation (space)?
- ✓ Sequential access? Similar to FAT
 - Random access?

- ✓ Growth?
- ✓ Quickness of allocation (time)?
- ✓ Fragmentation (space)?
- ✓ Sequential access? Similar to FAT
- ✓ Random access? Similar to indexed allocation

Best of both worlds: indexed allocation and multilevel indexed allocation

Allocation Strategy	File representation	Free list maintenance	Sequential Access	Random Access	File growth	Allocation Overhead	Space Efficiency
Contiguous	Contiguous blocks	complex	Very good	Very good	messy	Medium to high	Internal and external fragmentati on
Contiguous With Overflow	Contiguous blocks for small files	complex	Very good for small files	Very good for small files	OK	Medium to high	**
Linked List	Non- contiguous blocks	Bit vector	Good but dependent on seek time	Not good	Very good	Small to medium	Excellent
FAT	"	FAT	***	Good but dependent on seek time	11	Small	***
Indexed	"	Bit vector	••	**	limited	"	,,
Multilevel Indexed	11	Bit vector	11	**	good	11	11
Hybrid	***	Bit vector	**	"	"	"	**

Allocation Strategy	File representation	Free list mainte-nance	Sequential Access	Random Access	File growth	Allocation Overhead	Space Efficiency
Contiguous	Contiguous blocks	complex	Very good ✓	Very good	messy	Medium to high	Internal and X external fragmentati on
Contiguous With Overflow	Contiguous blocks for small files	complex	Very good for small files	Very good for small files	ОК	Medium to high	**
Linked List	Non- contiguous blocks	Bit vector	Good but dependent on seek time	Not good	Very good	Small to medium	Excellent
FAT	***	FAT	11	Good but dependent on seek time	"	Small	11
Indexed	"	Bit vector	"	"	limited	"	"
Multilevel Indexed	11	Bit vector	11	11	good	"	"
Hybrid	***	Bit vector	"	"	"	**	**

Allocation Strategy	File representation	Free list maintenance	Sequential Access	Random Access	File growth	Allocation Overhead	Space Efficiency
Contiguous	Contiguous blocks	complex	Very good	Very good ✓	messy X	Medium to high	Internal and X external fragmentati on
Contiguous With Overflow	Contiguous blocks for small files	complex X	Very good for small files	Very good for small files	OK ✓	Medium to high	,, X
Linked List	Non- contiguous blocks	Bit vector	Good but dependent on seek time	Not good	Very good	Small to medium	Excellent
FAT	11	FAT	11	Good but dependent on seek time	11	Small	***
Indexed	"	Bit vector	"	**	limited	"	"
Multilevel Indexed	"	Bit vector	11	***	good	"	11
Hybrid	***	Bit vector	**	"	**	"	"

Allocation Strategy	File representation	Free list maintenance	Sequential Access	Random Access	File growth	Allocation Overhead	Space Efficiency
Contiguous	Contiguous blocks	complex	Very good ✓	Very good	messy	Medium to high	Internal and X external fragmentati on
Contiguous With Overflow	Contiguous blocks for small files	complex X	Very good for small files	Very good for small files	OK ✓	Medium to high	,, X
Linked List	Non- contiguous blocks	Bit vector	Good but dependent on seek time	Not good	Very good	Small to medium	Excellent
FAT	"	FAT	"	Good but dependent on seek time	"	Small	"
Indexed	11	Bit vector	"	**	limited	"	"
Multilevel Indexed	"	Bit vector	"	***	good	"	11
Hybrid	"	Bit vector	***	***	***	***	**

Allocation Strategy	File representation	Free list mainte- nance	Sequential Access	Random Access	File growth	Allocation Overhead	Space Efficiency
Contiguous	Contiguous blocks	complex X	Very good ✓	Very good	messy	Medium to high	Internal and X external fragmentati on
Contiguous With Overflow	Contiguous blocks for small files	complex X	Very good for small files	Very good for small files	OK	Medium to high	,, X
Linked List	Non- contiguous blocks	Bit vector	Good but dependent on seek time	Not good X	Very good	Small to medium	Excellent
FAT	"	FAT ✓	" ✓	Good but dependent on seek time	" ✓	Small 🗸	" *
Indexed	11	Bit vector	"	**	limited	***	11
Multilevel Indexed	**	Bit vector	"	11	good	"	"
Hybrid	11	Bit vector	***	***	***	***	11

Allocation Strategy	File representation	Free list mainte- nance	Sequential Access	Random Access	File growth	Allocation Overhead	Space Efficiency
Contiguous	Contiguous blocks	complex	Very good ✓	Very good	messy	Medium to high	Internal and X external fragmentati on
Contiguous With Overflow	Contiguous blocks for small files	complex X	Very good for small files	Very good for small files	OK ✓	Medium to high	X
Linked List	Non- contiguous blocks	Bit vector	Good but dependent on seek time	Not good	Very good	Small to medium	Excellent
FAT Partitioning not good for user	***	FAT 🗸	" ✓	Good but dependent on seek time	" *	Small 🗸	" √
Indexed	***	Bit vector	***	11	limited	**	"
Multilevel Indexed	***	Bit vector	11	11	good	11	"
Hybrid	"	Bit vector	**	"	***	"	**

Allocation Strategy	File representation	Free list mainte- nance	Sequential Access	Random Access	File growth	Allocation Overhead	Space Efficiency
Contiguous	Contiguous blocks	complex	Very good ✓	Very good	messy X	Medium to high	Internal and X external fragmentati on
Contiguous With Overflow	Contiguous blocks for small files	complex X	Very good for small files	Very good for small files	OK ✓	Medium to high	 X
Linked List	Non- contiguous blocks	Bit vector	Good but dependent on seek time	Not good X	Very good	Small to medium	Excellent
FAT Partitioning not good for user	***	FAT 🗸	" *	Good but dependent on seek time	" V	Small 🗸	" ✓
Indexed	11	Bit vector	" /	" 🗸	limited X	" 🗸	" 🗸
Multilevel Indexed	"	Bit vector	11	11	good	11	"
Hybrid	"	Bit vector	"	"	"	**	**

Allocation Strategy	File representation	Free list maintenance	Sequential Access	Random Access	File growth	Allocation Overhead	Space Efficiency
Contiguous	Contiguous blocks	complex	Very good ✓	Very good	messy X	Medium to high	Internal and X external fragmentati on
Contiguous With Overflow	Contiguous blocks for small files	complex X	Very good for small files	Very good for small files	OK ✓	Medium to high	 X
Linked List	Non- contiguous blocks	Bit vector	Good but dependent on seek time	Not good X	Very good	Small to medium	Excellent
FAT Partitioning not good for user	***	FAT 🗸	" *	Good but dependent on seek time	" **	Small 🗸	" *
Indexed	"	Bit vector	" /	" 🗸	limited X	" 🗸	" 🗸
Multilevel Indexed	***	Bit vector	" •	"	good	" 🗸	" 🗸
Hybrid	"	Bit vector	"	"	"	**	"

Allocation Strategy	File representation	Free list maintenance	Sequential Access	Random Access	File growth	Allocation Overhead	Space Efficiency
Contiguous	Contiguous blocks	complex	Very good ✓	Very good ✓	messy X	Medium to high	Internal and X external fragmentati on
Contiguous With Overflow	Contiguous blocks for small files	complex X	Very good for small files	Very good for small files	OK ✓	Medium to high	" X
Linked List	Non- contiguous blocks	Bit vector	Good but dependent on seek time	Not good X	Very good	Small to medium	Excellent
FAT Partitioning not good for user	11	FAT 🗸	" ✓	Good but dependent on seek time	" ✓	Small	" ~
Indexed	"	Bit vector	" /	" 🗸	limited X	" 🗸	" 🗸
Multilevel Small Indexed Files 🖰	"	Bit vector	" •	"	good	" 🗸	" 🗸
Hybrid	"	Bit vector	"	"	11	**	***

Allocation Strategy	File representation	Free list mainte- nance	Sequential Access	Random Access	File growth	Allocation Overhead	Space Efficiency
Contiguous	Contiguous blocks	complex	Very good	Very good	messy X	Medium to high	Internal and X external fragmentati on
Contiguous With Overflow	Contiguous blocks for small files	complex X	Very good for small files	Very good for small files	OK _	Medium to high	,, X
Linked List	Non- contiguous blocks	Bit vector	Good but dependent on seek time	Not good X	Very good	Small to medium	Excellent
FAT Partitioning not good for user	"	FAT 🗸	" ✓	Good but dependent on seek time	" √	Small 🗸	" ✓
Indexed	"	Bit vector	" /	" 🗸	limited X	" /	" 🗸
Multilevel Small Indexed Files 🖰	"	Bit vector	" •	"	good	" 🗸	" 🗸
Hybrid	**	Bit vector	**	**	11	**	11
		✓	✓	√	√	✓	√

Hybrid FS Example

Given the following:

```
Size of index block = 512 bytes

Size of data block = 2048 bytes

Size of pointer = 8 bytes (to index or data blocks)
```

The i-node consists of

2 direct data block pointers, I single indirect pointer, and I double indirect pointer.

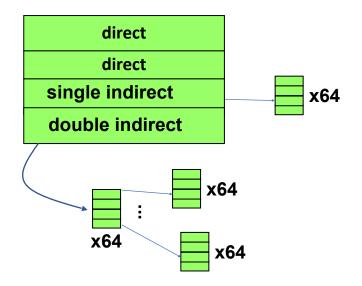
An index block is used for the i-node as well as for the index blocks that store pointers to other index blocks and data blocks. Note that the index blocks and data blocks are allocated on a need basis.

- (a) What is the maximum size (in bytes) of a file that can be stored in this file system?
- (b) How many data blocks are needed for storing the same data file of 266 KB?
- (c) How many index blocks are needed for storing a data file of size 266 KB?

```
Size of index block = 512 bytes
Size of data block = 2048 bytes
Size of pointer = 8 bytes (to index or data blocks)

The i-node consists of

2 direct data block pointers,
I single indirect pointer, and
I double indirect pointer.
```



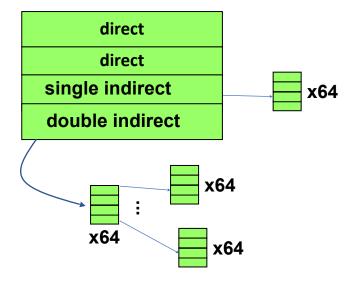
(a) Maximum file size:

- An index block can hold 512/8 = 64 entries
- 2 direct data block pointers in i-node
- 64 entries in first-level index
- 64 * 64 entries in second-level index
- Total: 4, I 62 blocks * 2,048 bytes = 8,523,776 bytes

Size of index block = 512 bytes
Size of data block = 2048 bytes
Size of pointer = 8 bytes (to index or data blocks)

The i-node consists of

2 direct data block pointers,
I single indirect pointer, and
I double indirect pointer.

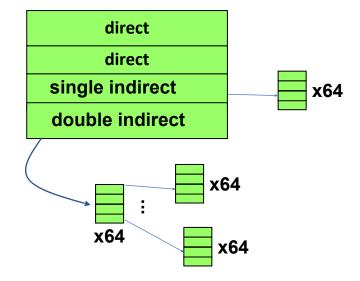


- (a) Maximum file size:
 - An index block can hold 512/8 = 64 entries
 - 2 direct data block pointers in i-node
 - 64 entries in first-level index
 - 64 * 64 entries in second-level index
 - Total: 4, I 62 blocks * 2,048 bytes = 8,523,776 bytes
- (b) Number of data blocks to hold 266KB:
 - \bullet 266 * 210 / 2048 = 266 / 2 = 133 blocks

Size of index block = 512 bytes
Size of data block = 2048 bytes
Size of pointer = 8 bytes (to index or data blocks)

The i-node consists of

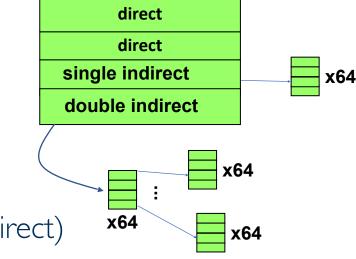
2 direct data block pointers,
I single indirect pointer, and
I double indirect pointer.

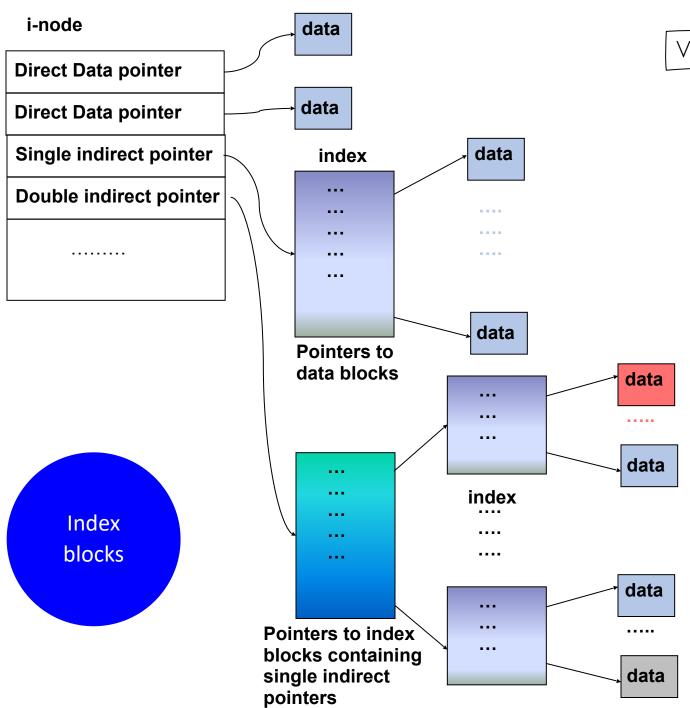


- (a) Maximum file size:
 - An index block can hold 512/8 = 64 entries
 - 2 direct data block pointers in i-node
 - 64 entries in first-level index
 - 64 * 64 entries in second-level index
 - Total: 4, I 62 blocks * 2,048 bytes = 8,523,776 bytes
- (b) Number of data blocks to hold 266KB:
 - \bullet 266 * 210 / 2048 = 266 / 2 = 133 blocks
- (c) How many index blocks to hold 266KB (133 blocks):
 - 2 direct + 64 first-level blocks + (64 + 3) second-level blocks
 - I i-node + I first-level index (under single indirect)
 + I first-level index + 2 second-level index (under double indirect)

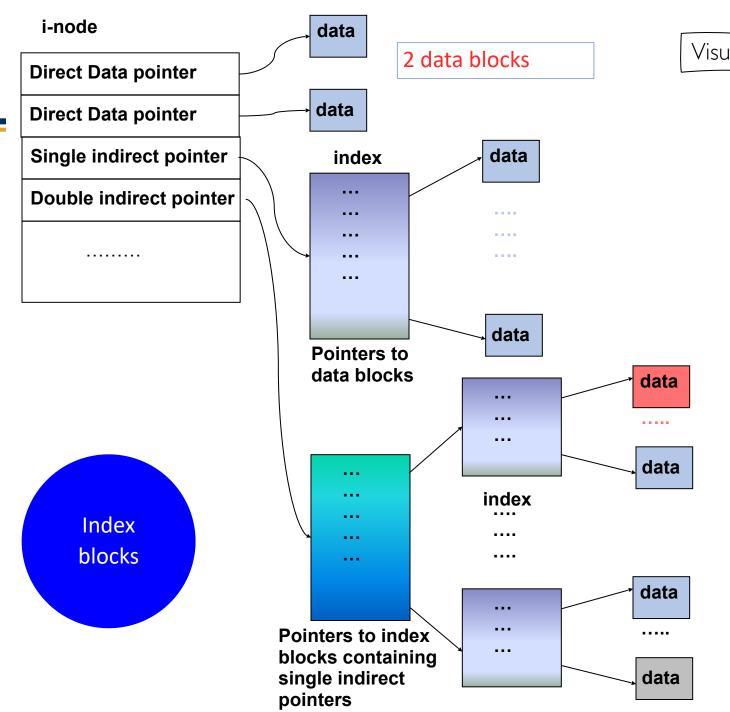
Size of index block = 512 bytes
Size of data block = 2048 bytes
Size of pointer = 8 bytes (to index or data blocks)

The i-node consists of
2 direct data block pointers,
I single indirect pointer, and
I double indirect pointer.

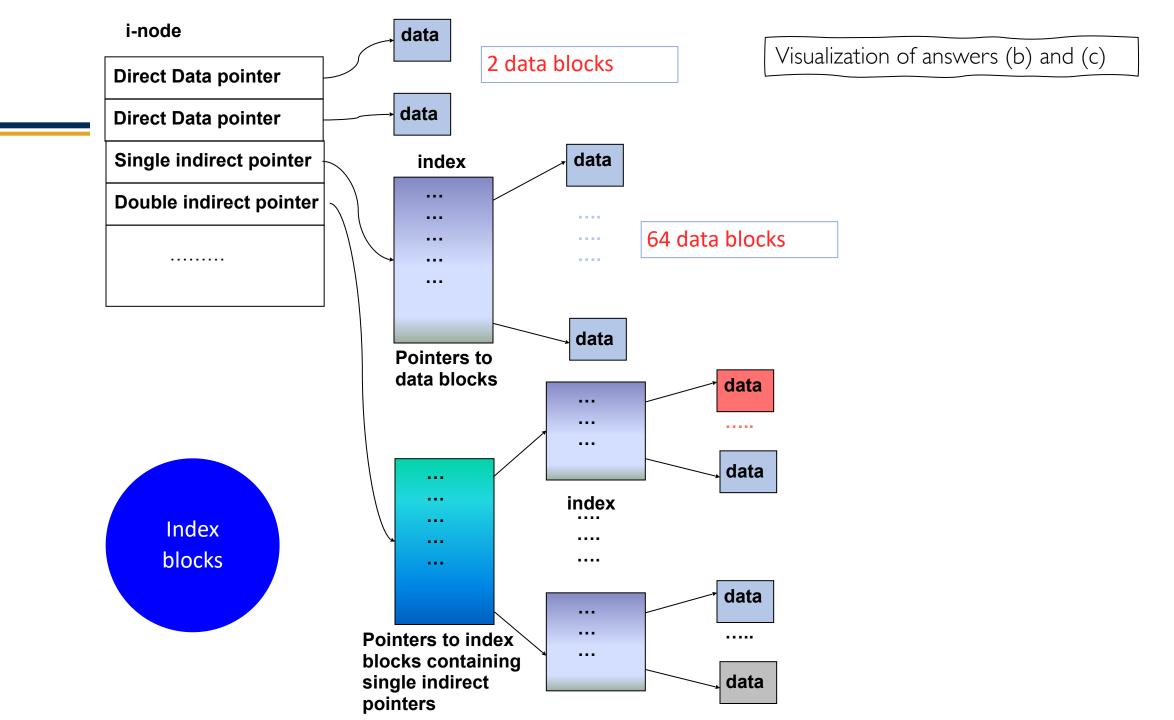


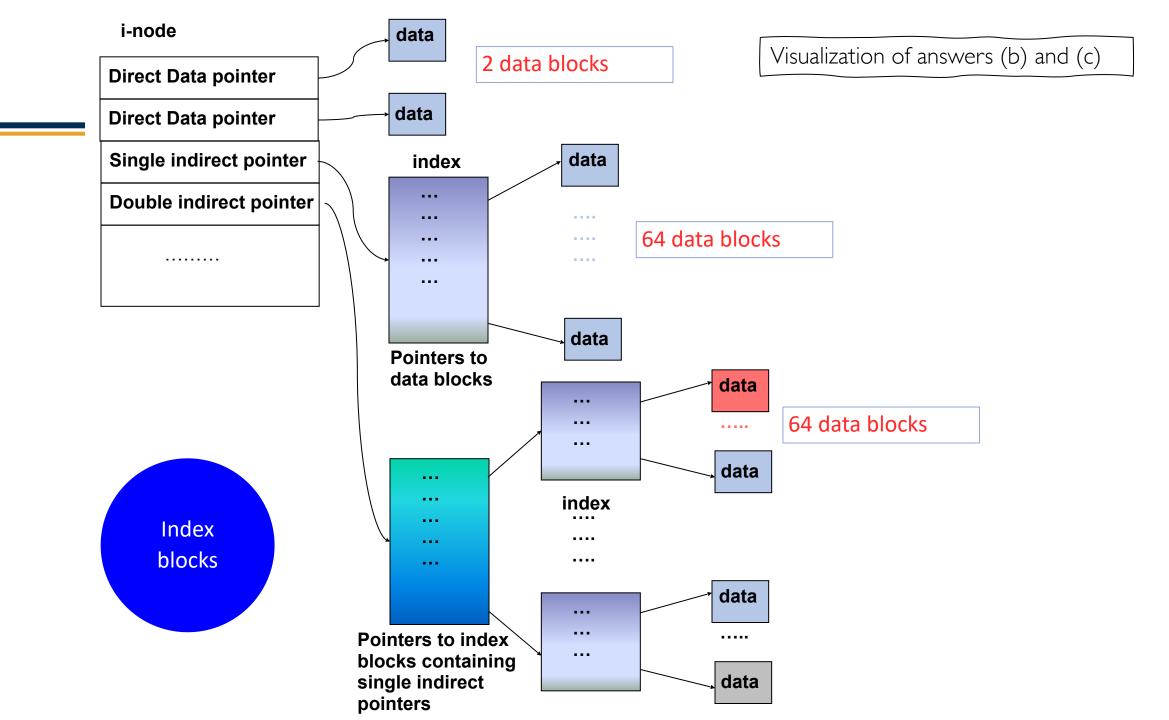


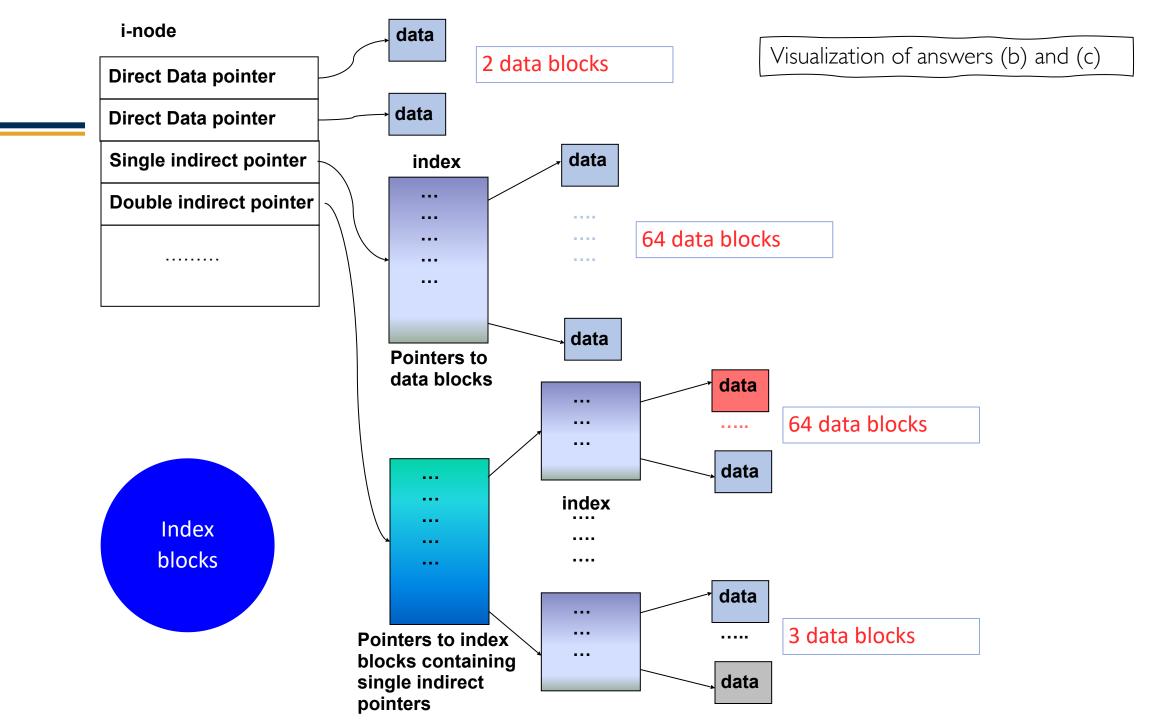
Visualization of answers (b) and (c)

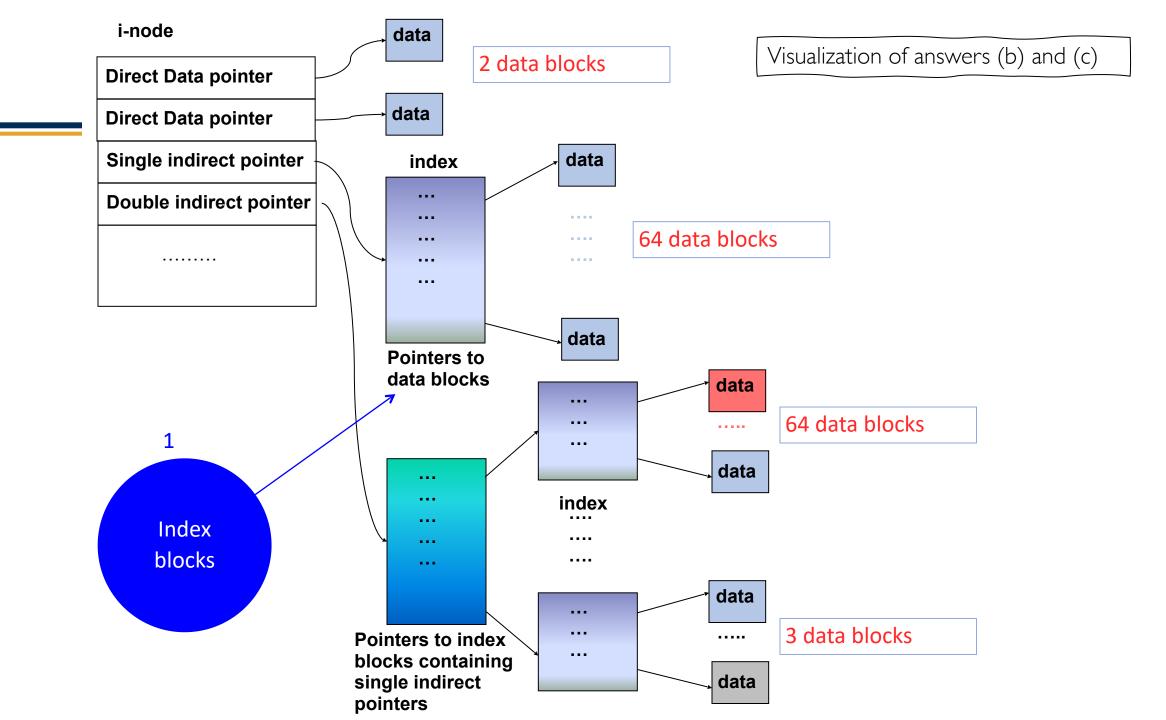


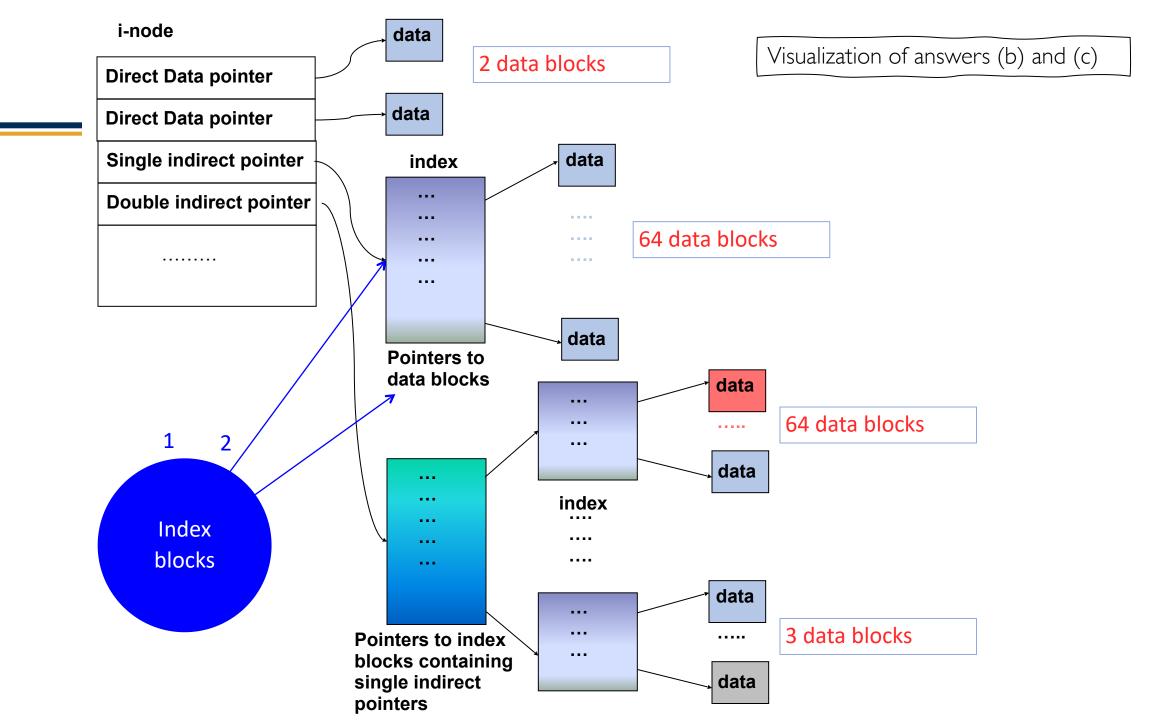
Visualization of answers (b) and (c)

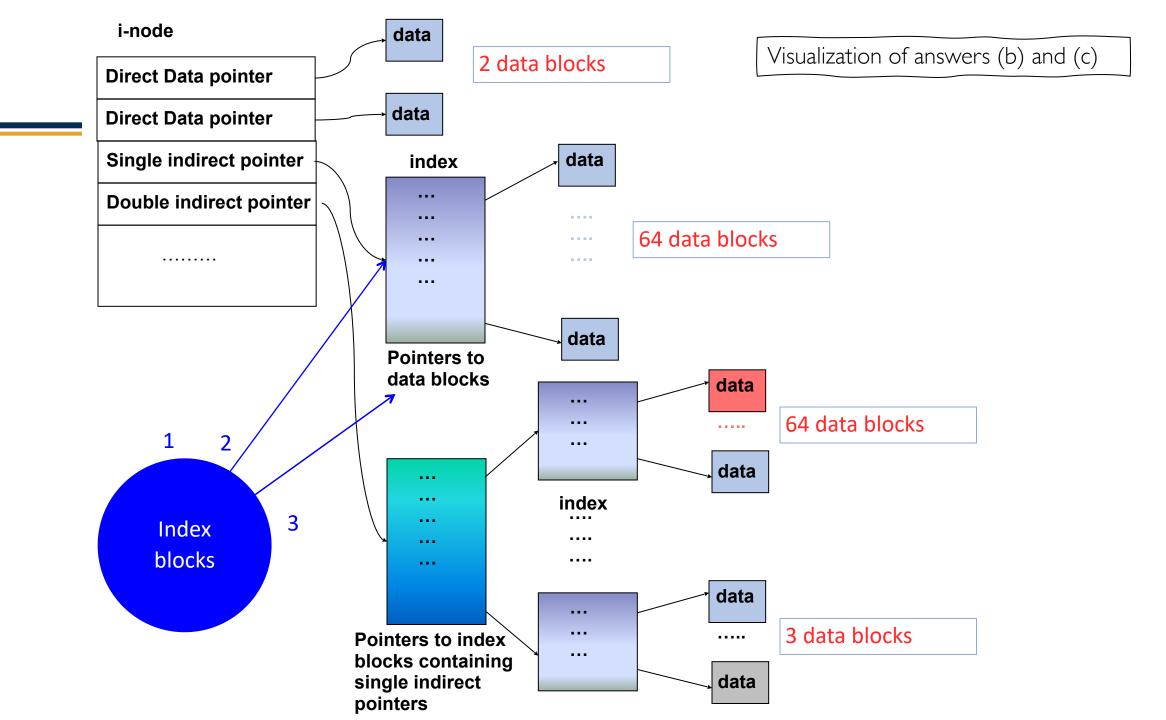


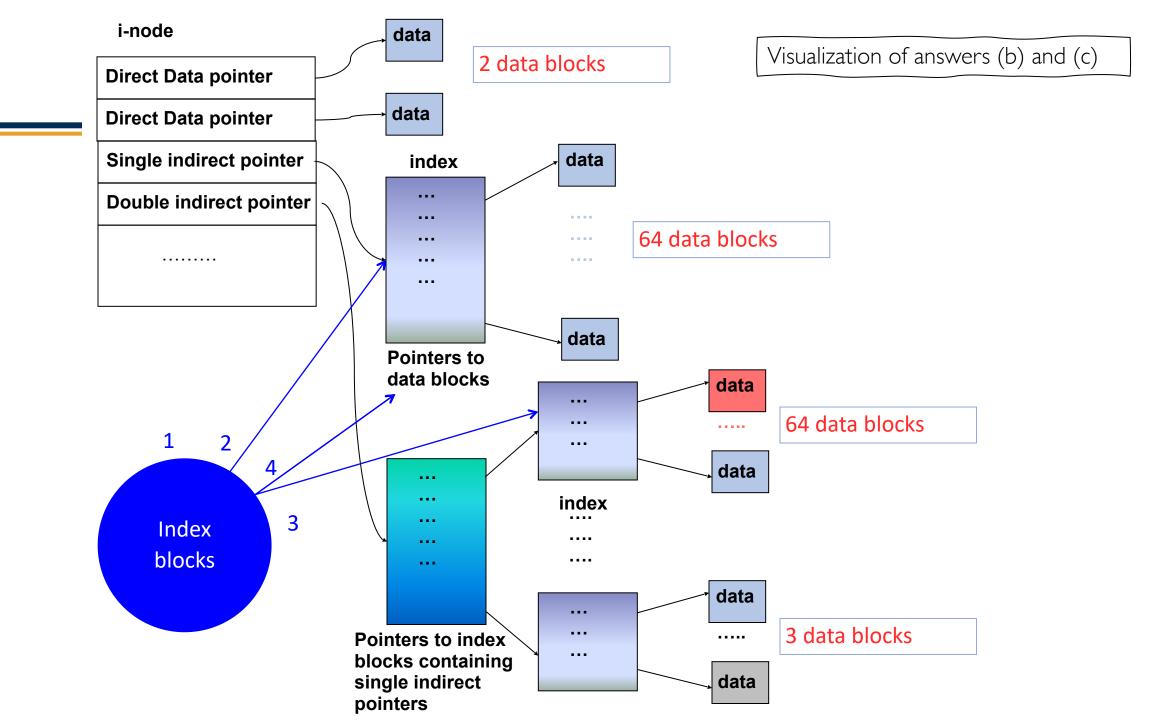


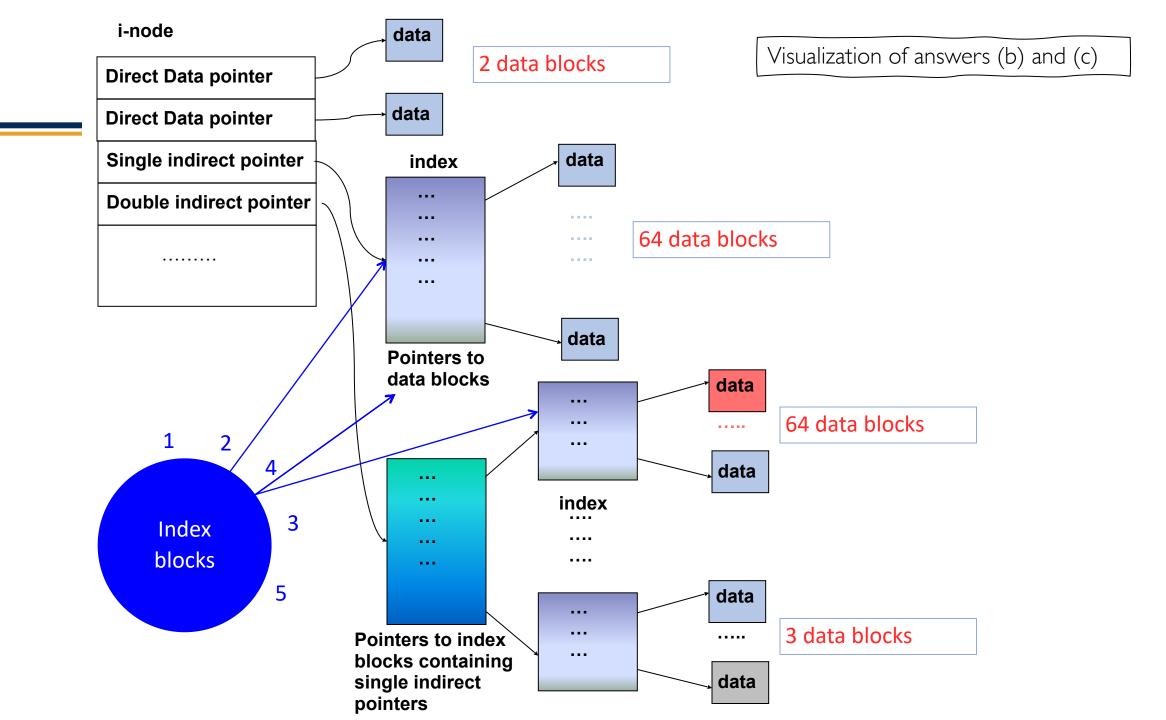


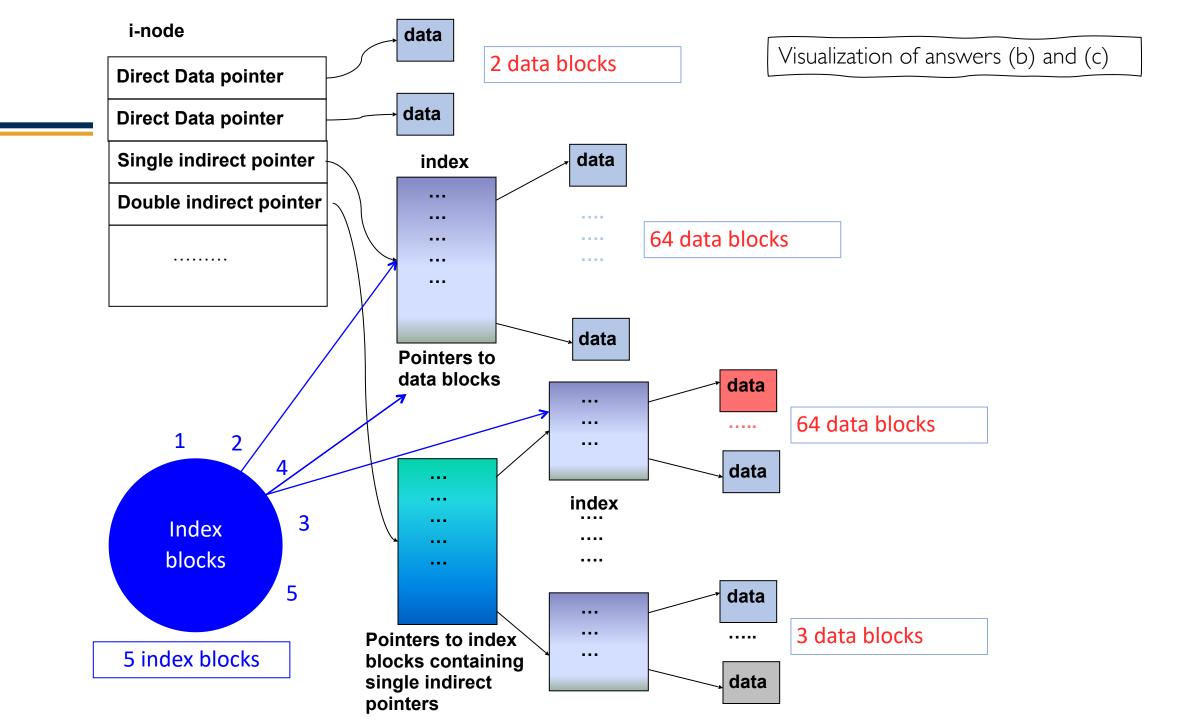






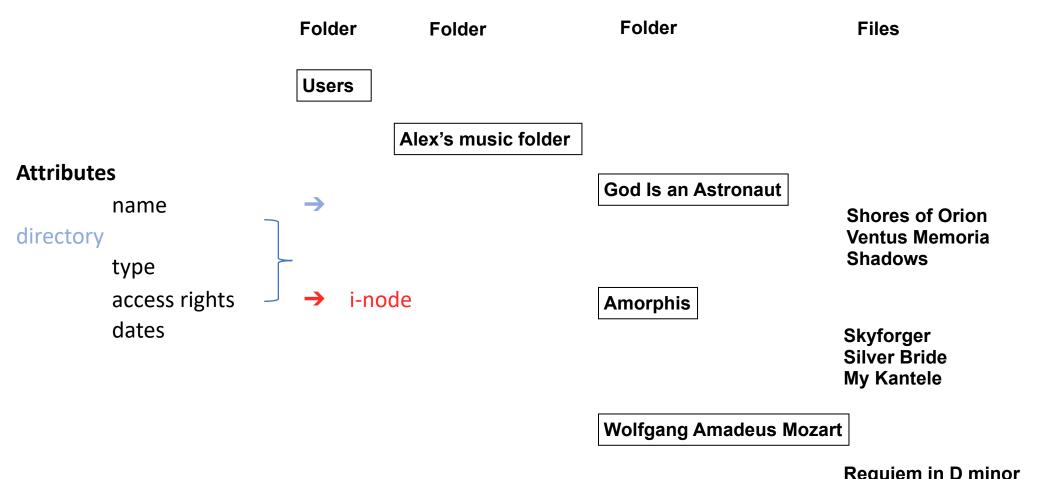




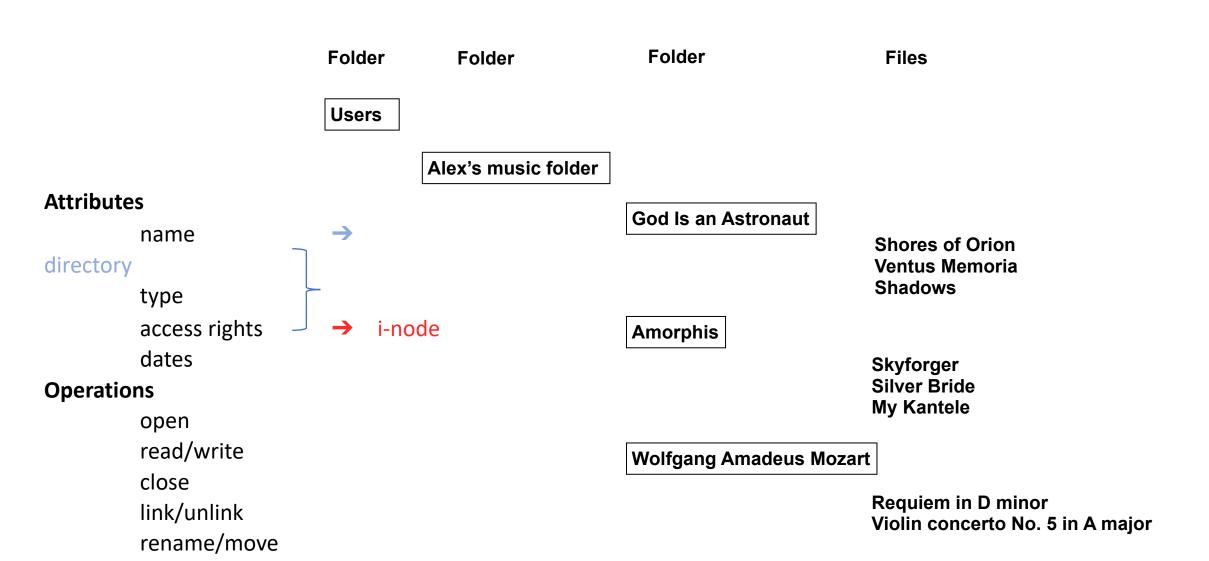


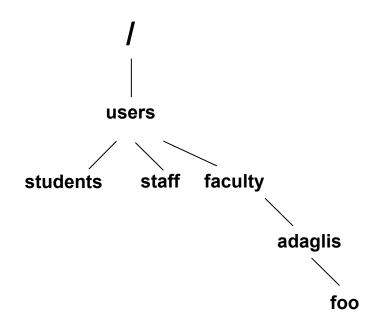
Folder Folder **Files** Folder Users Alex's music folder **God Is an Astronaut Shores of Orion Ventus Memoria Shadows Amorphis** Skyforger Silver Bride My Kantele **Wolfgang Amadeus Mozart**

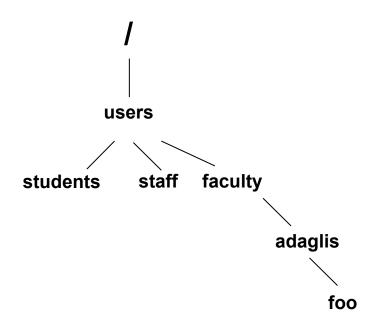
Requiem in D minor Violin concerto No. 5 in A major



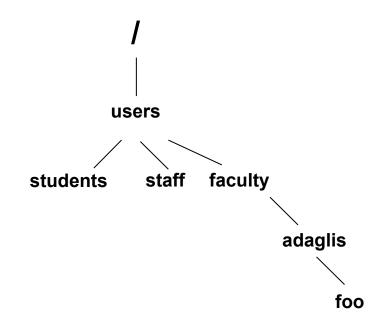
Requiem in D minor Violin concerto No. 5 in A major

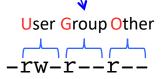




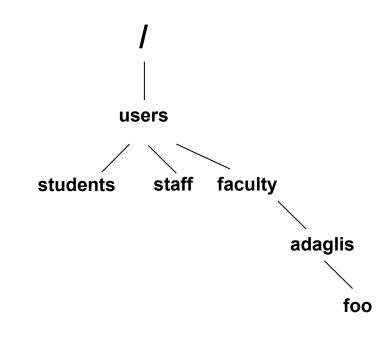


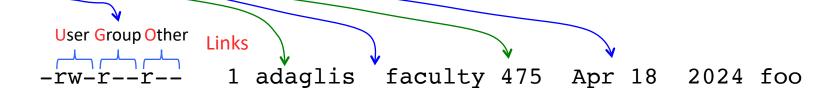
- Metadata in the i-node
 - Access rights: U G O
 - User name (uid)
 - Group name (gid)
 - Size
 - Modification date
 - Data/index block addrs



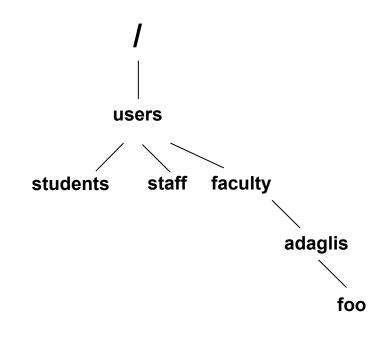


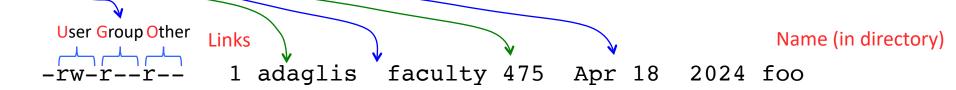
- Metadata in the i-node
 - Access rights: U G O
 - User name (uid)
 - **—** Group name (gid)
 - Size
 - Modification date
 - Data/index block addrs

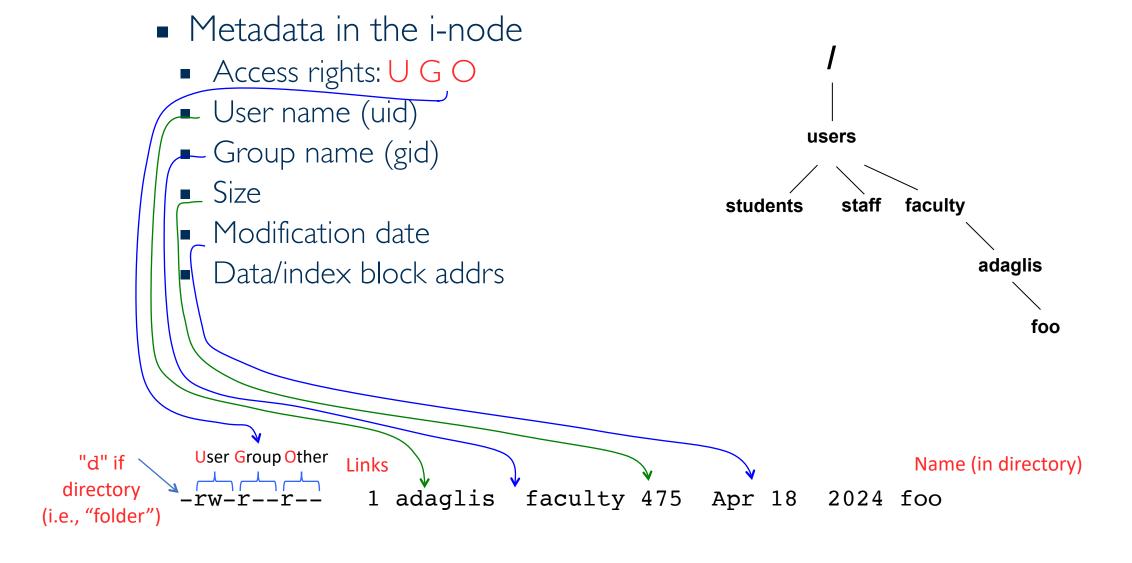




- Metadata in the i-node
 - Access rights: U G O
 - User name (uid)
 - **—** Group name (gid)
 - Size
 - Modification date
 - Data/index block addrs







Links

- Remember that Unix-style file system separates the directory entry from the i-node; this means more than one directory entry can index the same i-node.
- When you create a file, you initialize an i-node and create a directory entry; this sets the link count in the i-node to I
- You can also create additional links to the file with the link() system call or In command; this
 increments the link count in the i-node for each link; this is called a "hard link"
- Curiously, there is no such thing as a primary or secondary link; both directory entries are names for the file and neither link has precedence over the other

Links

- Remember that Unix-style file system separates the directory entry from the i-node; this means more than one directory entry can index the same i-node.
- When you create a file, you initialize an i-node and create a directory entry; this sets the link count in the i-node to I
- You can also create additional links to the file with the link() system call or In command; this
 increments the link count in the i-node for each link; this is called a "hard link"
- Curiously, there is no such thing as a primary or secondary link; both directory entries are names for the file and neither link has precedence over the other
- There is no system call to "remove" a file; you can only unlink a directory entry
- The file is deleted from the file system only when its link count reaches zero; you remove a link with the unlink() system call or the *rm* command

Links

- Remember that Unix-style file system separates the directory entry from the i-node; this means more than one directory entry can index the same i-node.
- When you create a file, you initialize an i-node and create a directory entry; this sets the link count in the i-node to I
- You can also create additional links to the file with the link() system call or In command; this
 increments the link count in the i-node for each link; this is called a "hard link"
- Curiously, there is no such thing as a primary or secondary link; both directory entries are names for the file and neither link has precedence over the other
- There is no system call to "remove" a file; you can only unlink a directory entry
- The file is deleted from the file system only when its link count reaches zero; you remove a link with the unlink() system call or the *rm* command
- Even better, each i-node also has an in-use count which counts the number of processes that have the file open; the i-node is not deallocated until the link count AND the in-use count both reach zero, so you can have an open file that has an i-node but no directory entry and hence no name in the file system

Symbolic links

- Unix also provides a way to create file aliases; these are called symbolic links
- A symbolic link occupies an i-node (contrast with a hard link doesn't take an additional i-node)
- The i-node is marked as a symbolic link and the data blocks contain a path name instead of user data

Symbolic links

- Unix also provides a way to create file aliases; these are called symbolic links
- A symbolic link occupies an i-node (contrast with a hard link doesn't take an additional i-node)
- The i-node is marked as a symbolic link and the data blocks contain a path name instead of user data
- When a symbolic link is encountered while following a path name, the path is replaced by the contents of the symbolic link and the search continues by following the path in the symbolic link
- The file system doesn't promise that symbolic links point to anything; the presence of a symbolic link pointing to a file name doesn't prevent that file from being unlinked

Symbolic links

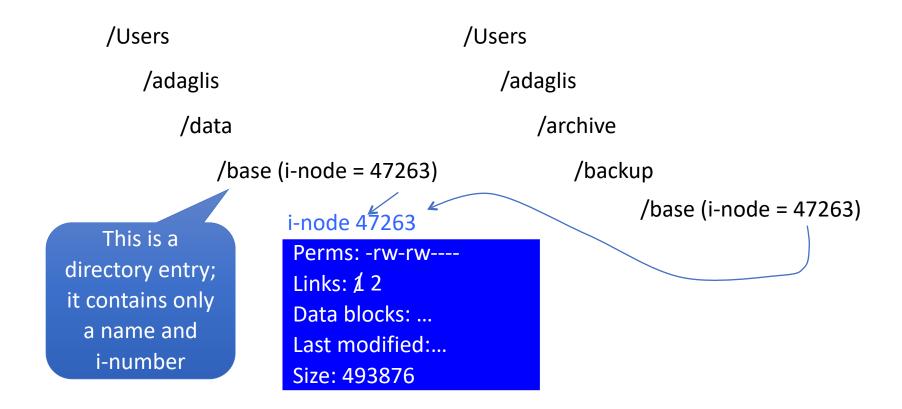
- Unix also provides a way to create file aliases; these are called symbolic links
- A symbolic link occupies an i-node (contrast with a hard link doesn't take an additional i-node)
- The i-node is marked as a symbolic link and the data blocks contain a path name instead of user data
- When a symbolic link is encountered while following a path name, the path is replaced by the contents of the symbolic link and the search continues by following the path in the symbolic link
- The file system doesn't promise that symbolic links point to anything; the presence of a symbolic link pointing to a file name doesn't prevent that file from being unlinked
- Symlinks work across file systems contrasted with hard links that do not

```
/Users
   /adaglis
       /data
           /base (i-node = 47263)
                  i-node 47263
                   Perms: -rw-rw----
                  Links: 1
                   Data blocks: ...
                   Last modified:...
                   Size: 493876
```

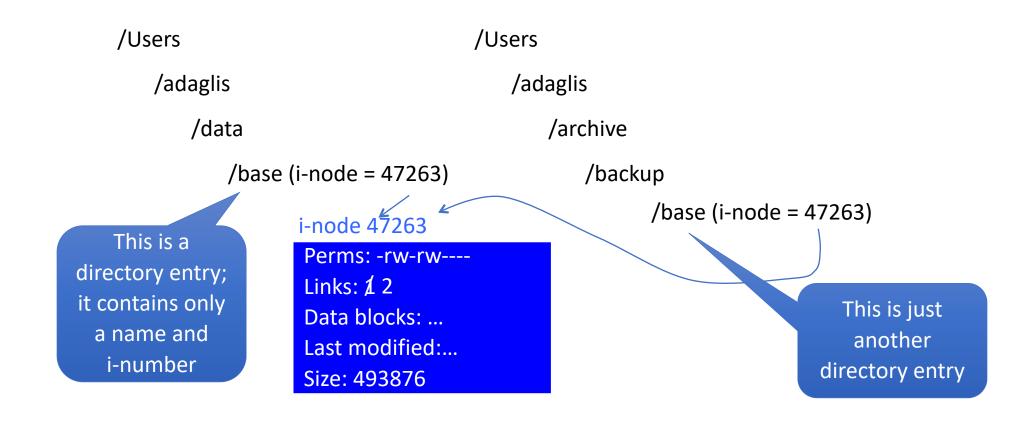
```
/Users
       /adaglis
           /data
               /base (i-node = 47263)
                      i-node 47263
   This is a
                       Perms: -rw-rw----
directory entry;
                       Links: 1
it contains only
                       Data blocks: ...
 a name and
                       Last modified:...
  i-number
                       Size: 493876
```

```
/Users
       /adaglis
           /data
               /base (i-node = 47263)
                      i-node 47263
   This is a
                       Perms: -rw-rw----
directory entry;
                       Links: 1
it contains only
                       Data blocks: ...
 a name and
                       Last modified:...
  i-number
                       Size: 493876
```

ln /Users/adaglis/data/base /Users/adaglis/archive/backup/base



ln /Users/adaglis/data/base /Users/adaglis/archive/backup/base



ln /Users/adaglis/data/base /Users/adaglis/archive/backup/base

```
/Users
                     /adaglis
                         /archive
                             /backup
                                    /base (i-node = 47263)
i-node 47263
Perms: -rw-rw----
Links: 121
                                                       This is just
Data blocks: ...
                                                         another
Last modified:...
                                                     directory entry
Size: 493876
```

ln /Users/adaglis/data/base /Users/adaglis/archive/backup/base
rm /Users/adaglis/data/base # remember this calls unlink()

Attribute	Meaning	Elaboration
Name	Name of the file	Attribute set at the time of creation or renaming
Alias	Other names that exist for the same physical file	Attribute gets set when an alias is created; system such as Unix provide explicit commands for creating aliases for a given file; Unix supports aliasing at two different levels (physical or hard, and symbolic or soft)
Owner	Usually the user who created the file	Attribute gets set at the time of creation of a file; systems such as Unix provide mechanism for the file's ownership to be changed by the superuser
Creation time	Time when the file was created first	Attribute gets set at the time a file is created or copied from some other place
Last write time	Time when the file was last written to	Attribute gets set at the time the file is written to or copied; in most file systems the creation time attribute is the same as the last write time attribute; Note that moving a file from one location to another preserves the creation time of the file
Privileges •Read •Write •Execute	The permissions or access rights to the file specifies who can do what to the file	Attribute gets set to default values at the time of creation of the file; usually, file systems provide commands to modify the privileges by the owner of the file; modern Linux and Windows file systems such as ext4 and NTFS also provide an access control list (ACL) to give more granular access to different users
Size	Total space occupied on the file system	Attribute gets set every time the size changes due to modification to the file

Unix command	Semantics	Elaboration
touch <name></name>	Create a file with the name < name >	Creates a zero byte file with the name <name> and a creation time equal to the current wall clock time</name>
mkdir <sub-dir></sub-dir>	Create a sub-directory <sub-dir></sub-dir>	The user must have write privilege to the current working directory (if <sub-dir> is a relative name) to be able to successfully execute this command</sub-dir>
rm <name></name>	Remove (or delete) the file named <name></name>	Only the owner of the file (and/or superuser) can delete a file
rmdir <sub-dir></sub-dir>	Remove (or delete) the sub-directory named <sub-dir></sub-dir>	Only the owner of the <sub-dir> (and/or the superuser) can remove the named sub-directory</sub-dir>
ln -s <orig> <new></new></orig>	Create a name <new> and make it symbolically equivalent to the file <orig></orig></new>	This is name equivalence only; so if the file <orig> is deleted, the storage associated with <orig> is reclaimed, and hence <new> will be a dangling reference to a non-existent file</new></orig></orig>
In <orig> <new></new></orig>	Create a name <new> and make it physically equivalent to the file <orig></orig></new>	Even if the file <orig> is deleted, the physical file remains accessible via the name <new></new></orig>
chmod <rights> <name></name></rights>	Change the access rights for the file <name> as specified in the mask <rights></rights></name>	Only the owner of the file (and/or the superuser) can change the access rights
chown <user> <name></name></user>	Change the owner of the file <name> to be <user></user></name>	Only superuser can change the ownership of a file
chgrp <group> <name></name></group>	Change the group associated with the file <name> to be <group></group></name>	Only the owner of the file (and/or the superuser) can change the group associated with a file
cp <orig> <new></new></orig>	Create a new file <new> that is a copy of the file <orig></orig></new>	The copy is created in the same directory if <new> is a file name; if <new> is a directory name, then a copy with the same name <orig> is created in the directory <new></new></orig></new></new>
mv <orig> <new></new></orig>	Renames the file <orig> with the name <new></new></orig>	Renaming happens in the same directory if <new> is a file name; if <new> is a directory name, then the file <orig> is moved into the directory <new> preserving its name <orig></orig></new></orig></new></new>
cat/more/less <name></name>	View the file contents	

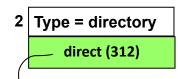
Let's study the path /users/faculty/adaglis/foo i-node for /

2 Type = directory

direct (312)

Let's study the path /users/faculty/adaglis/foo

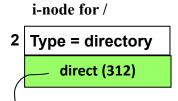
i-node for /



Block 312

	
File	i-node
name	address
users	5

Let's study the path /users/faculty/adaglis/foo

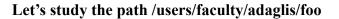


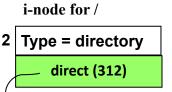
Block 312

File	i-node
name	address
users	5

i-node for /users

Type = directory
direct (291)





Block 312

File	i-node
name	address
users	5

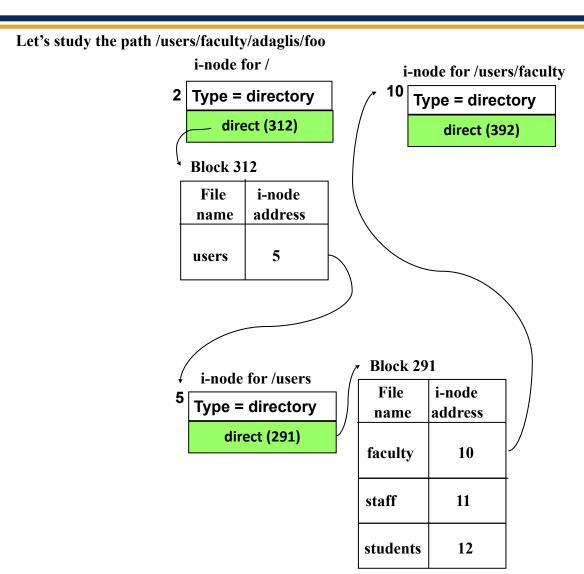
i-node for /users

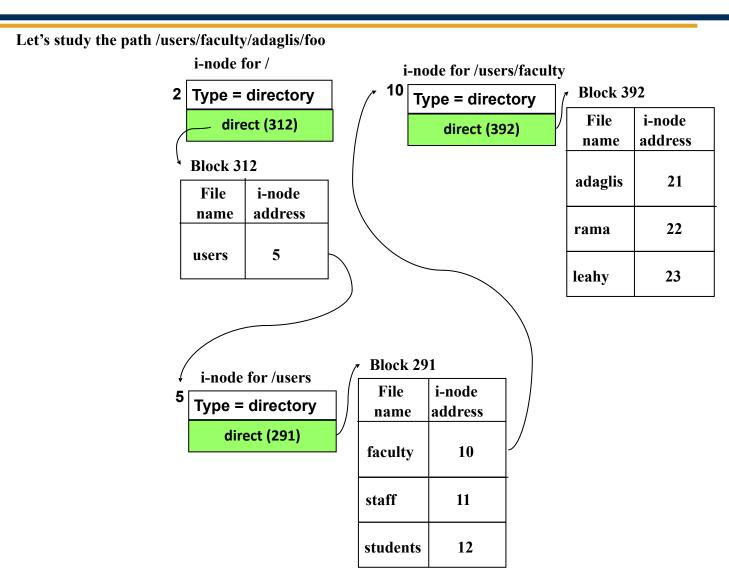
Type = directory

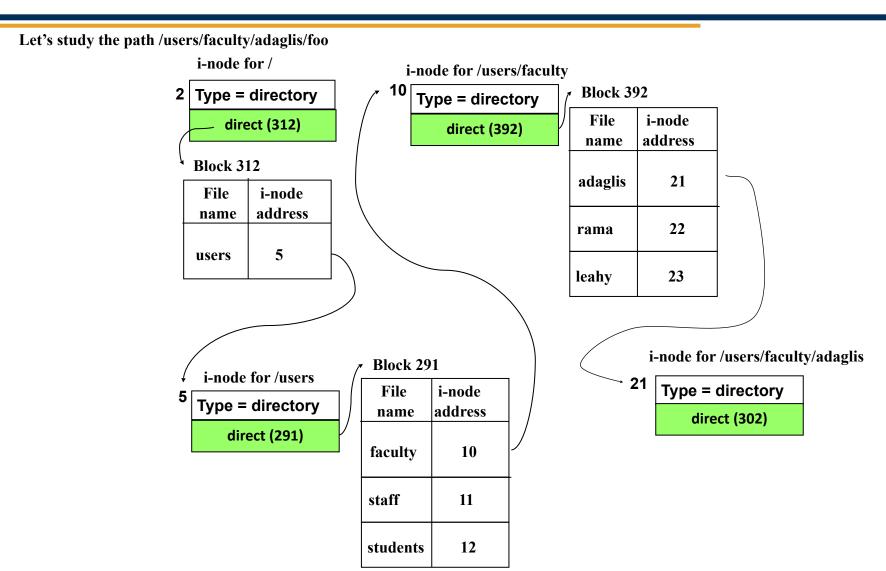
direct (291)

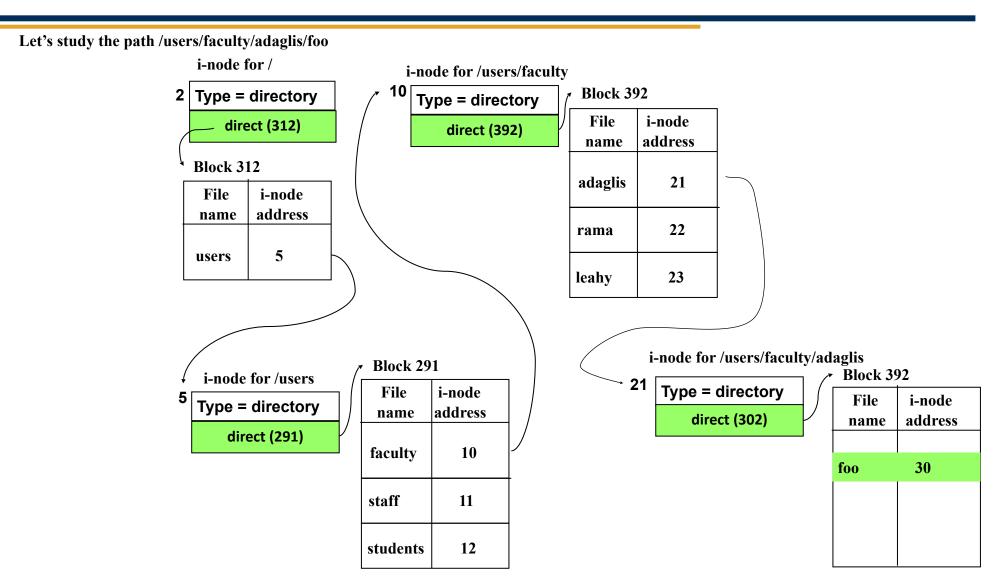
Block 291

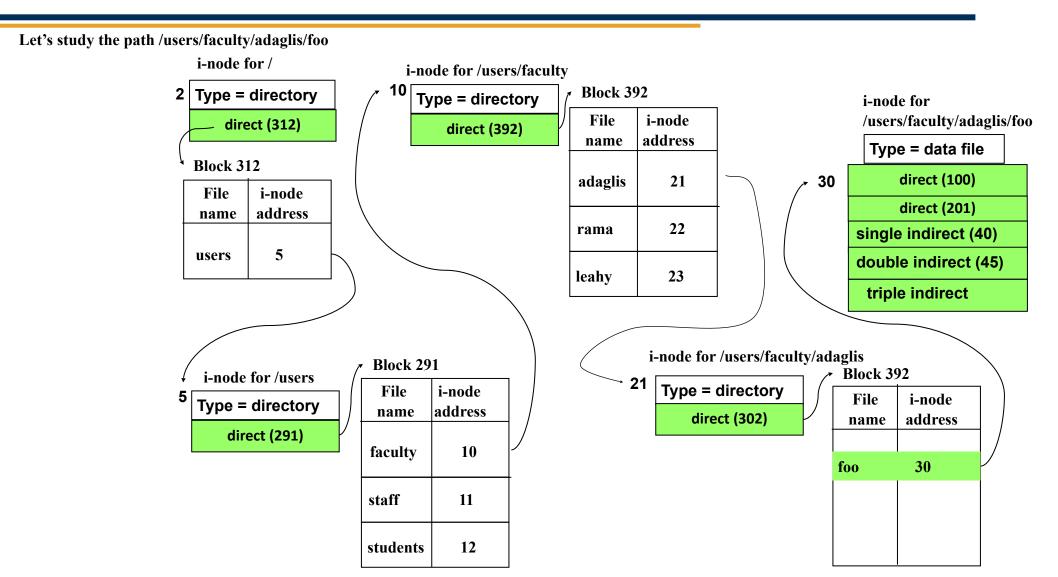
File name	i-node address
faculty	10
staff	11
students	12

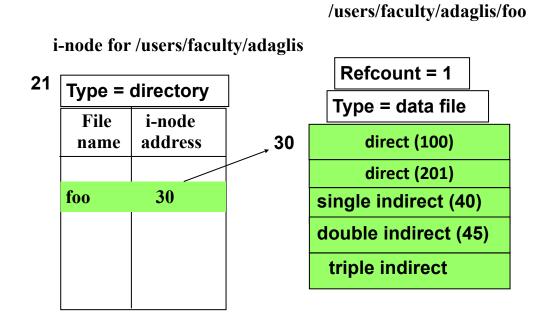












Reference count on files

i-node for

i-node for /users/faculty/adaglis

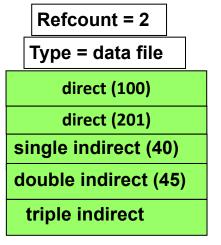
Type = directory

File i-node name address

foo 30

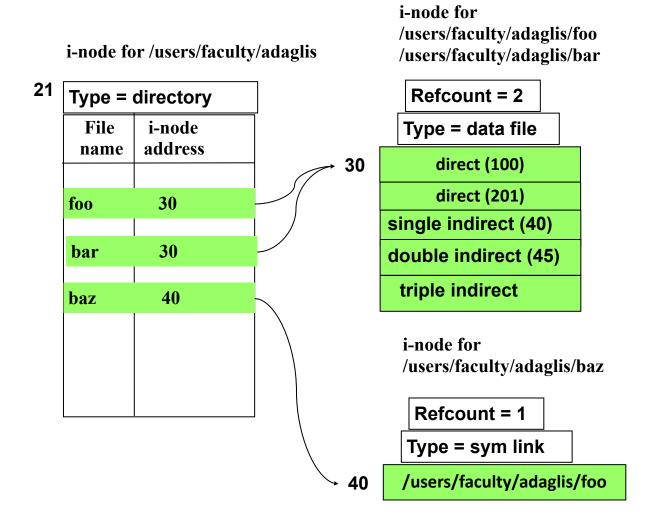
bar 30

i-node for /users/faculty/adaglis/foo /users/faculty/adaglis/bar



Hard links

ln foo bar

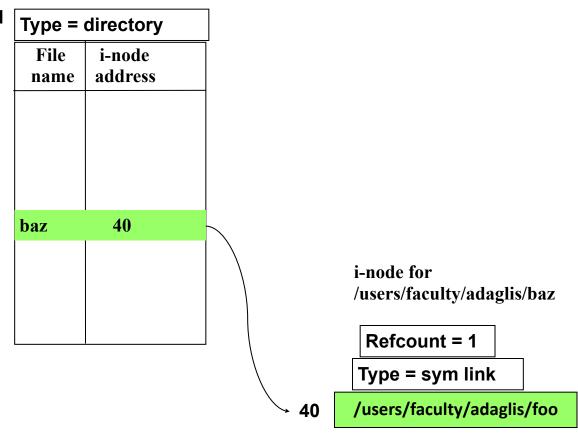


Sym links

ln -s foo baz

i-node for /users/faculty/adaglis

2′



rm foo bar

Examples in book

Work out examples for i-node creation and deletion in book on your Own:

Examples 5, 6, and 7